

Taskonomy: Disentangling Task Transfer Learning

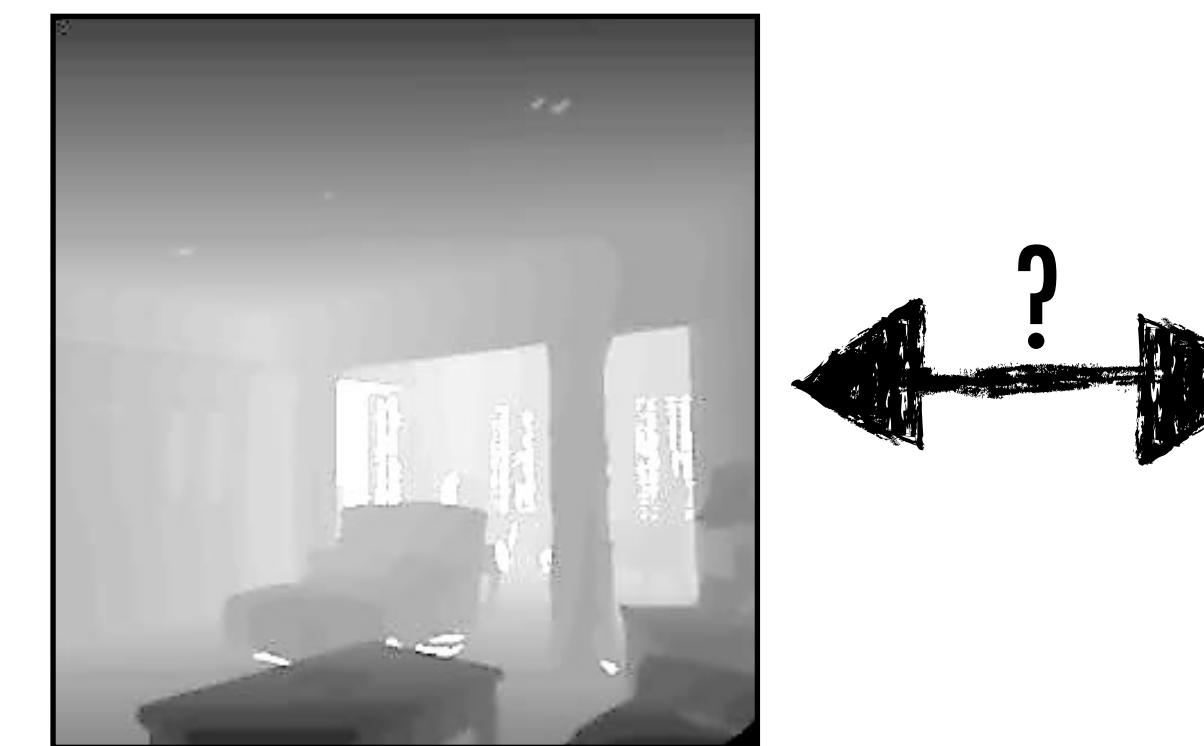
**Amir Zamir, Alexander Sax, William Shen, Leonidas Guibas,
Jitendra Malik, Silvio Savarese**

Stanford, UC Berkeley

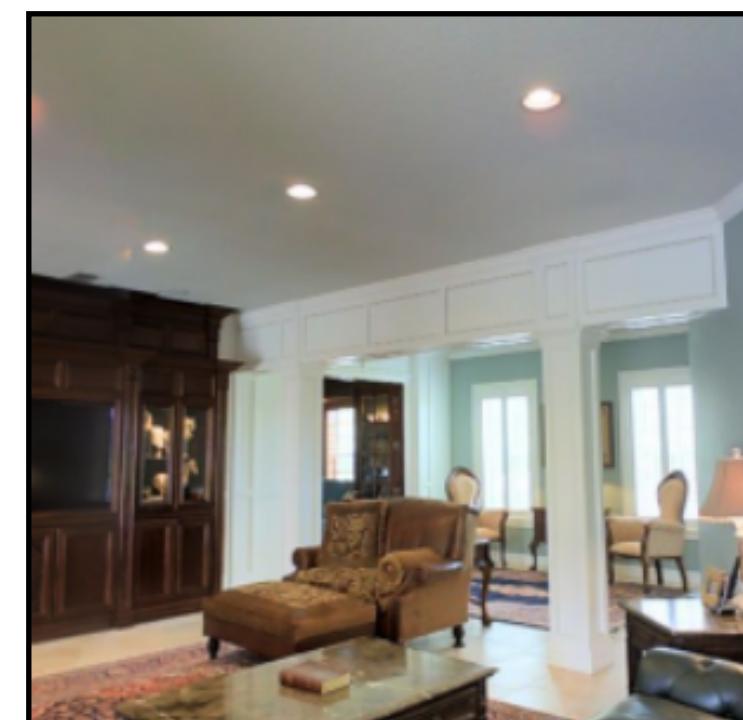
CVPR18 [Best Paper Award]



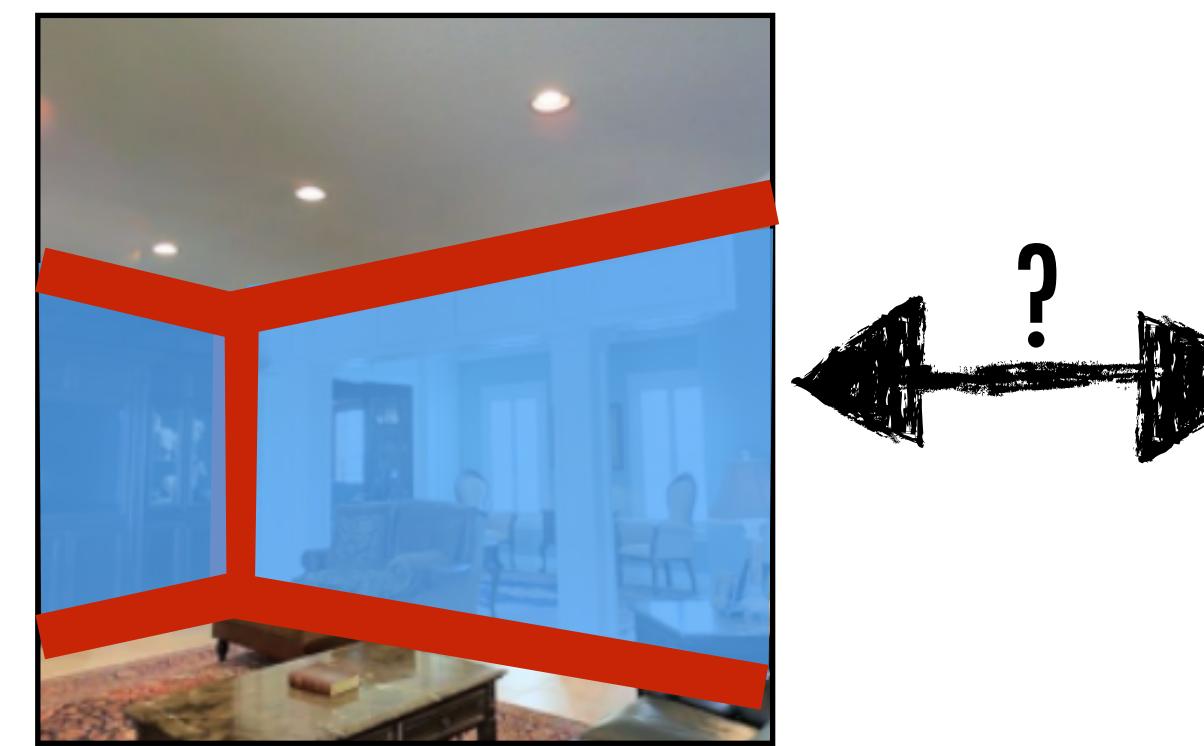
Question: Vision problems - related or independent?



Depth



Normals



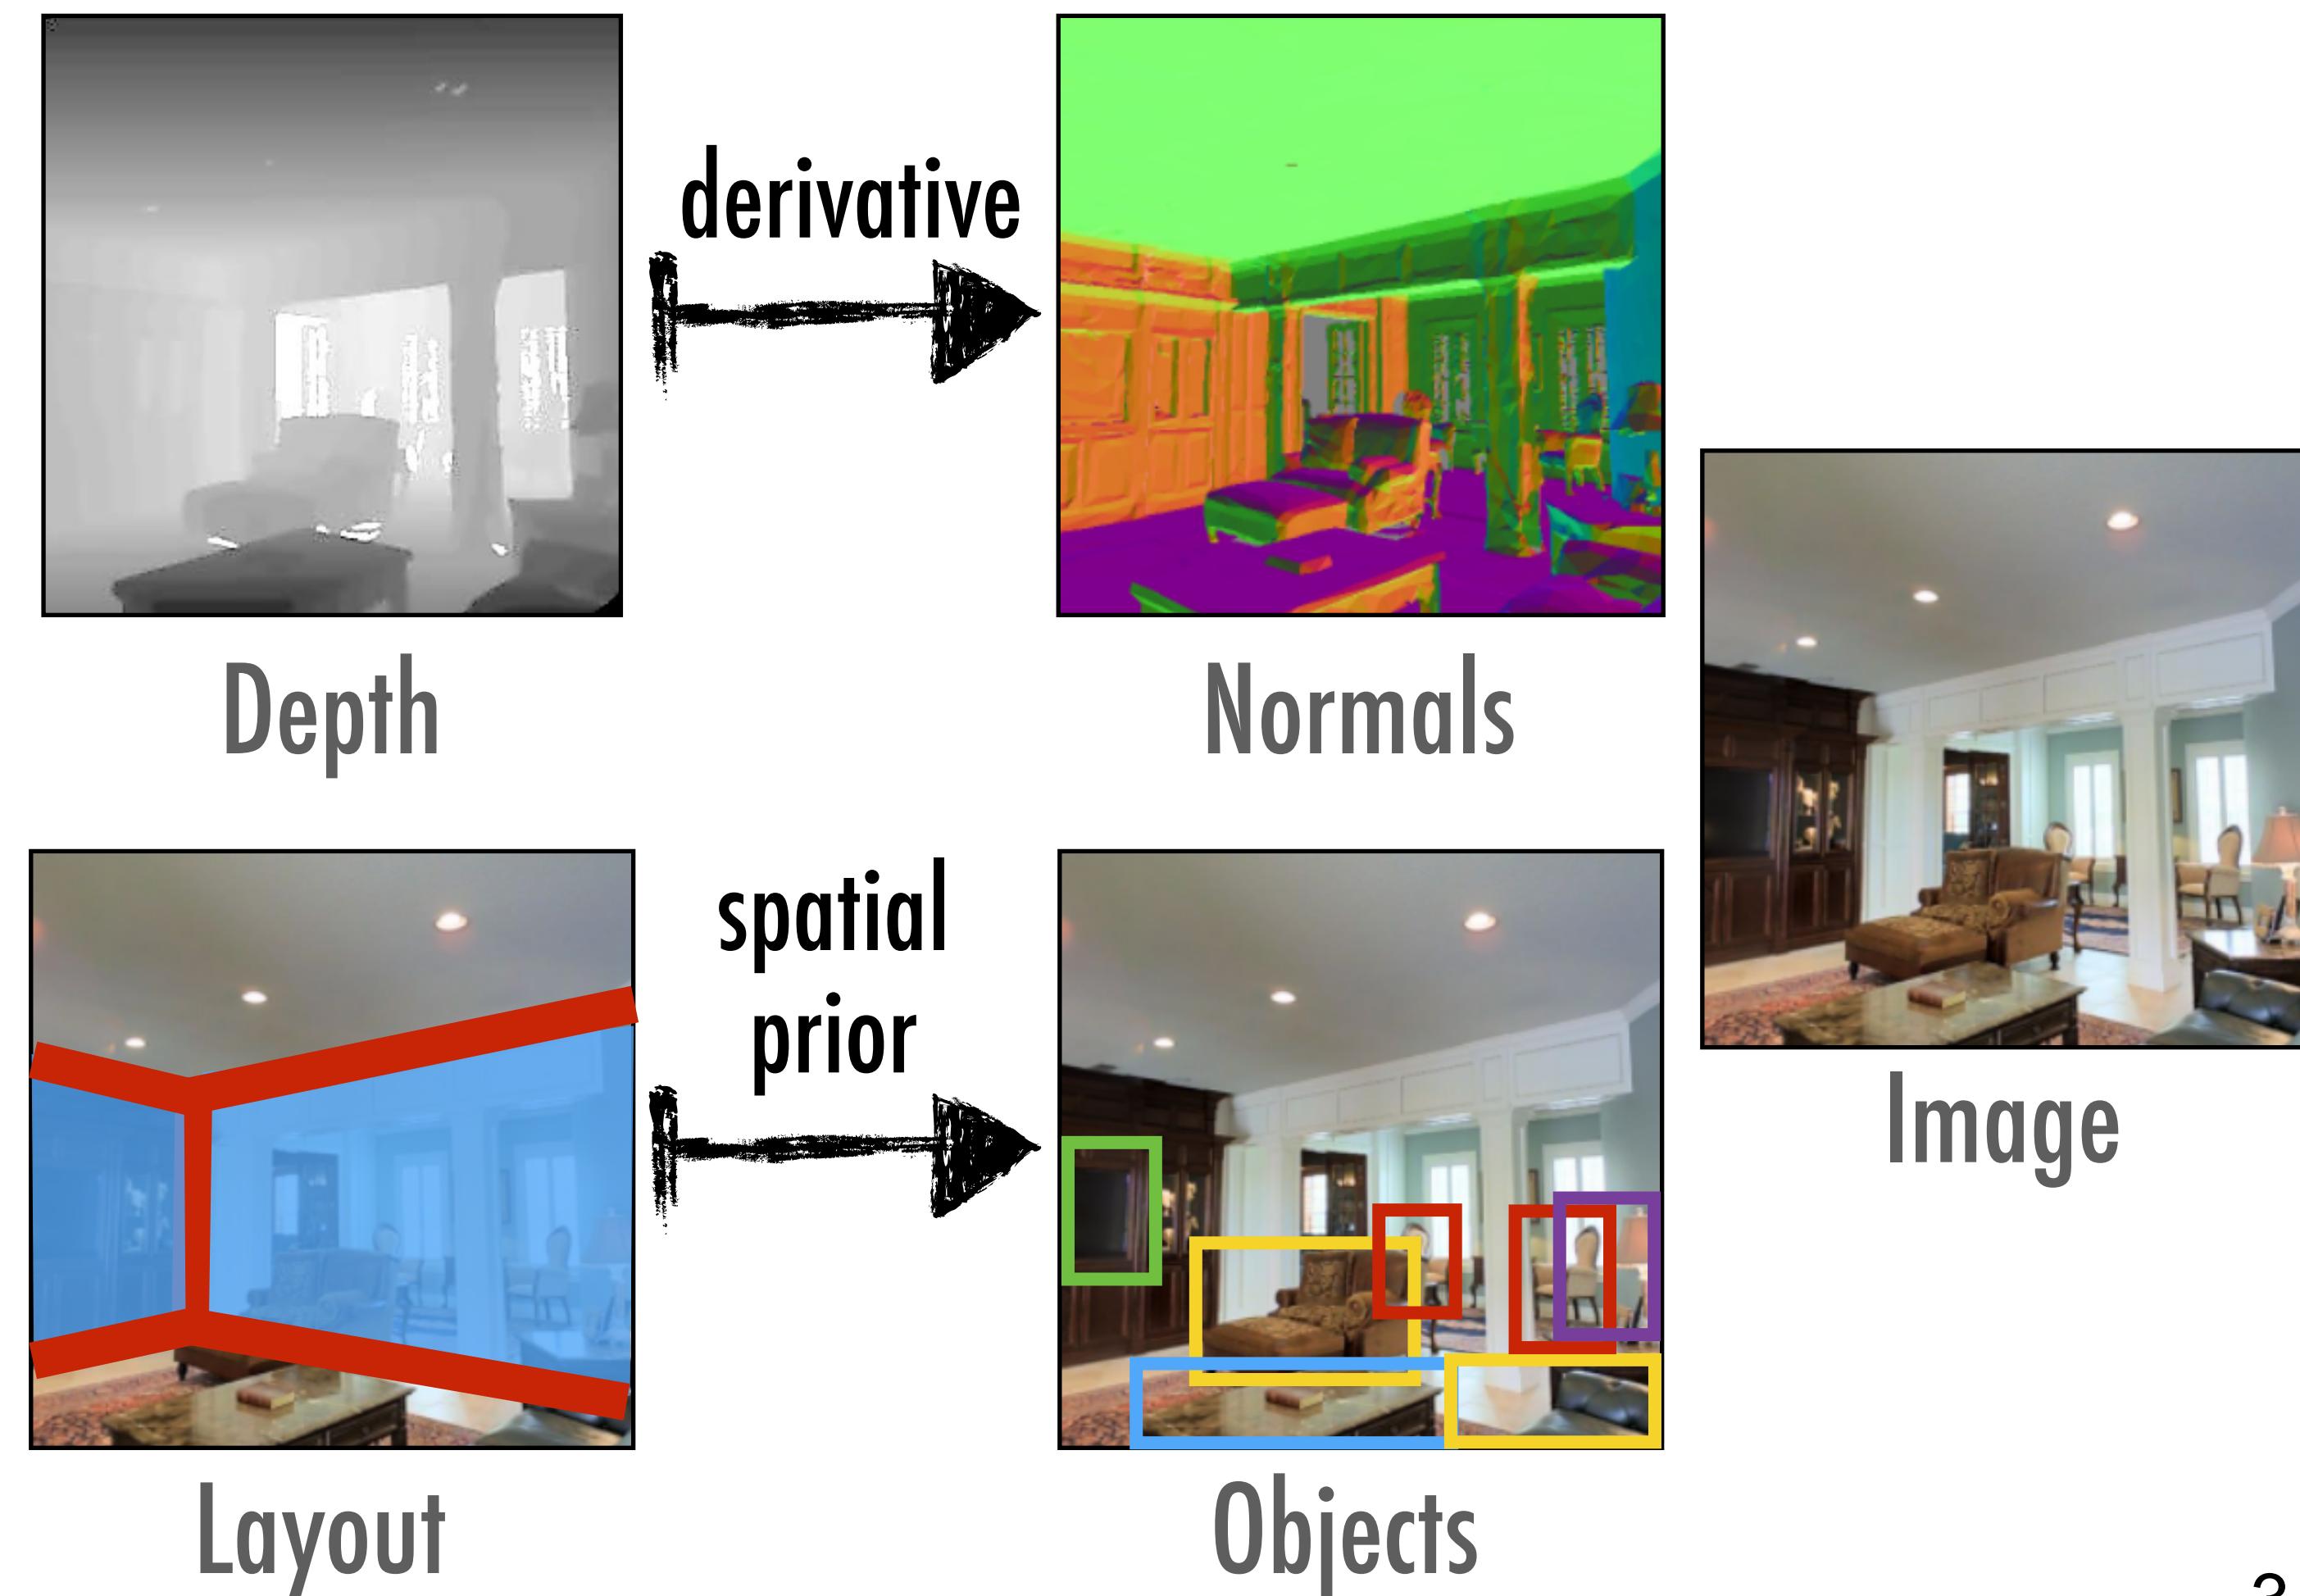
Layout



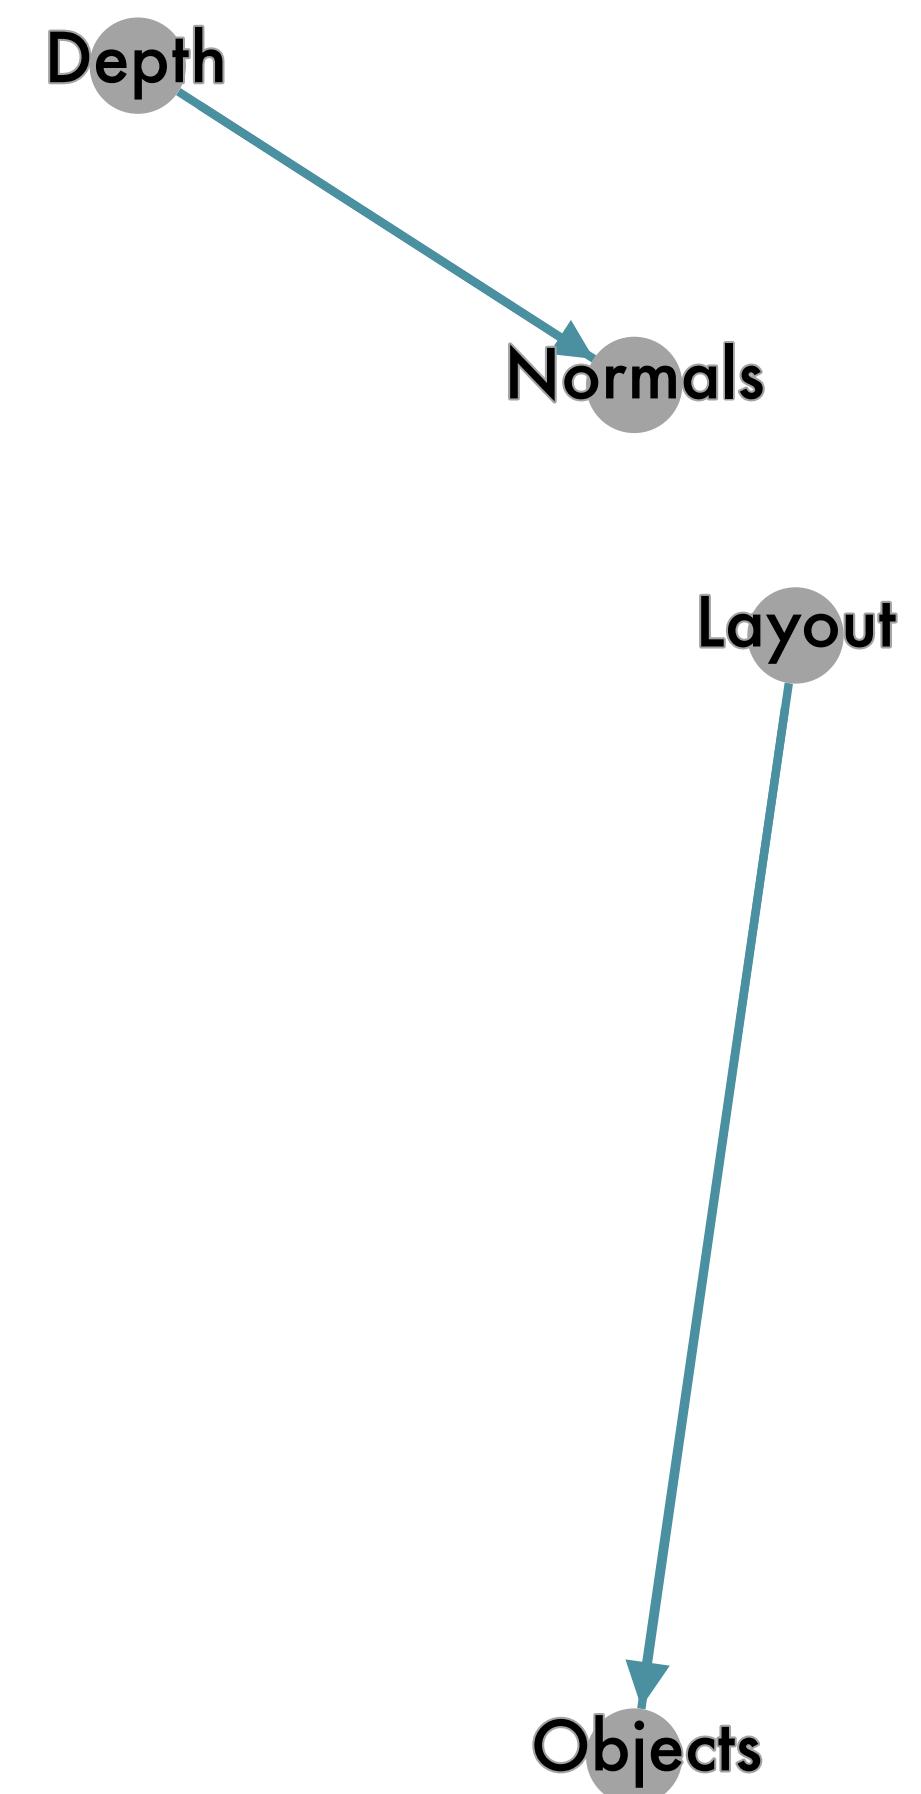
Objects

Question: Vision problems - related or independent?

- Task relationships exist
- Can be computationally measured
- Tasks belonging to a structured space
- Unified model for transfer learning

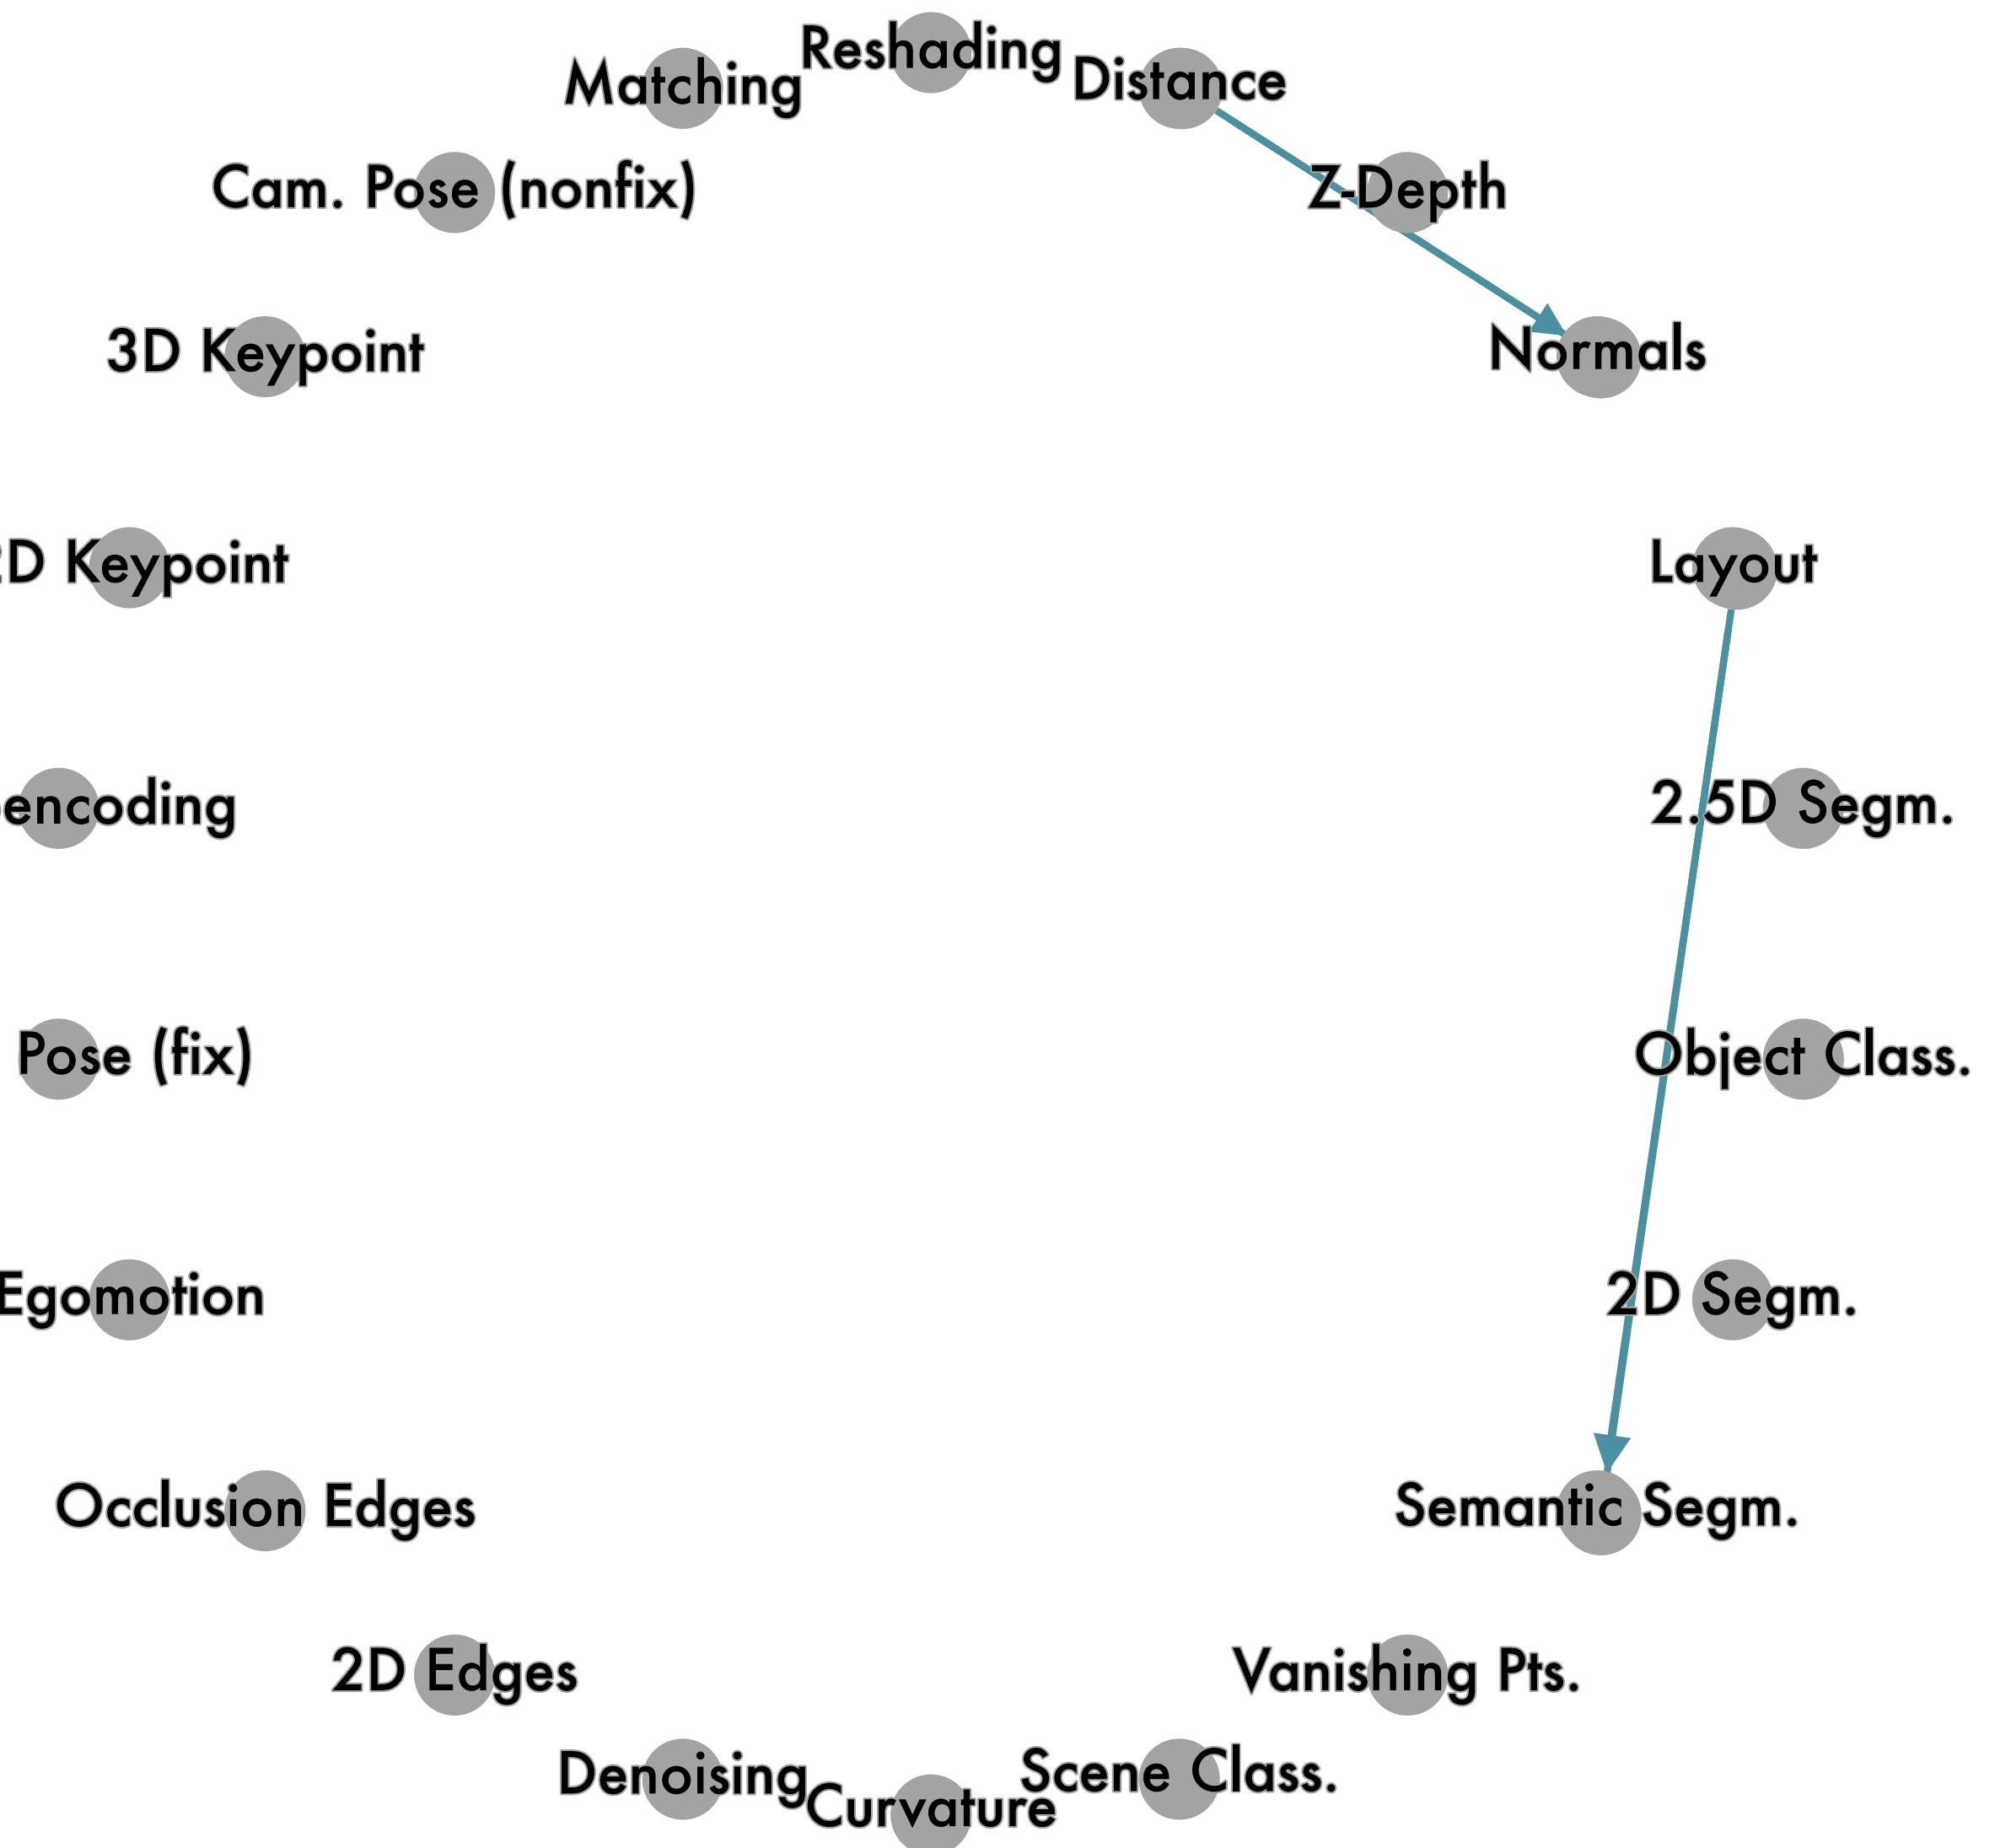


Task Relationships

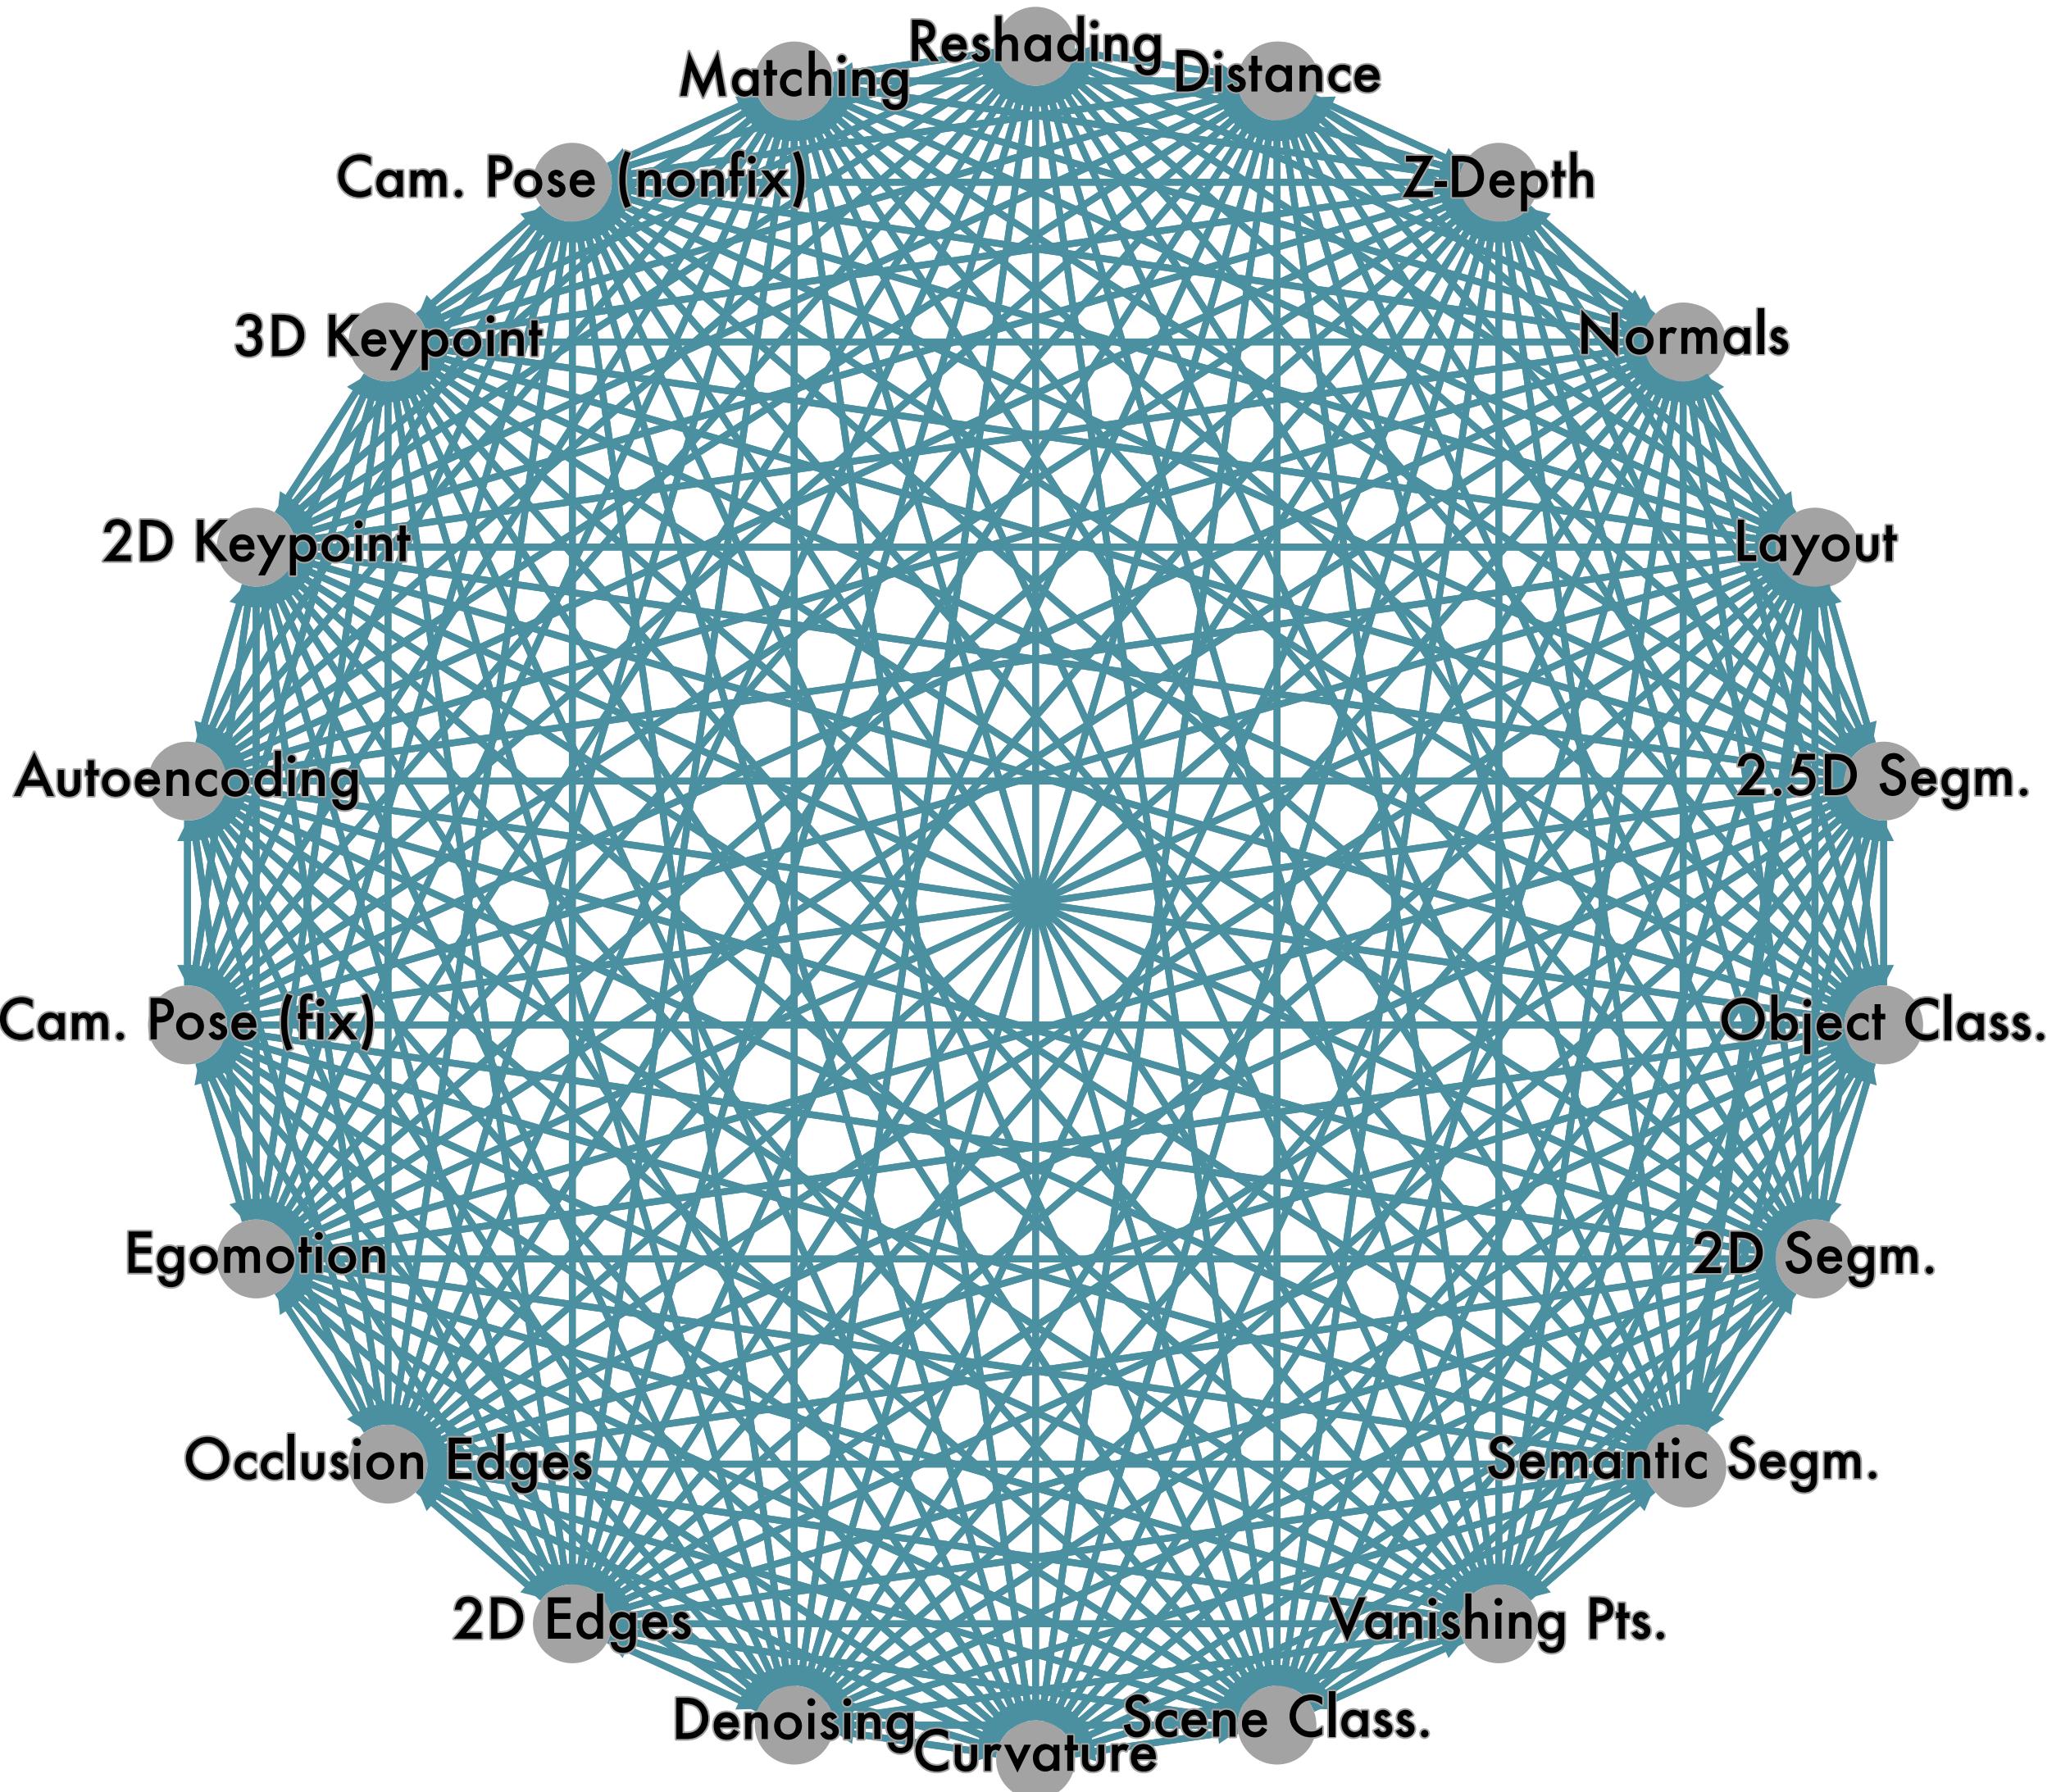
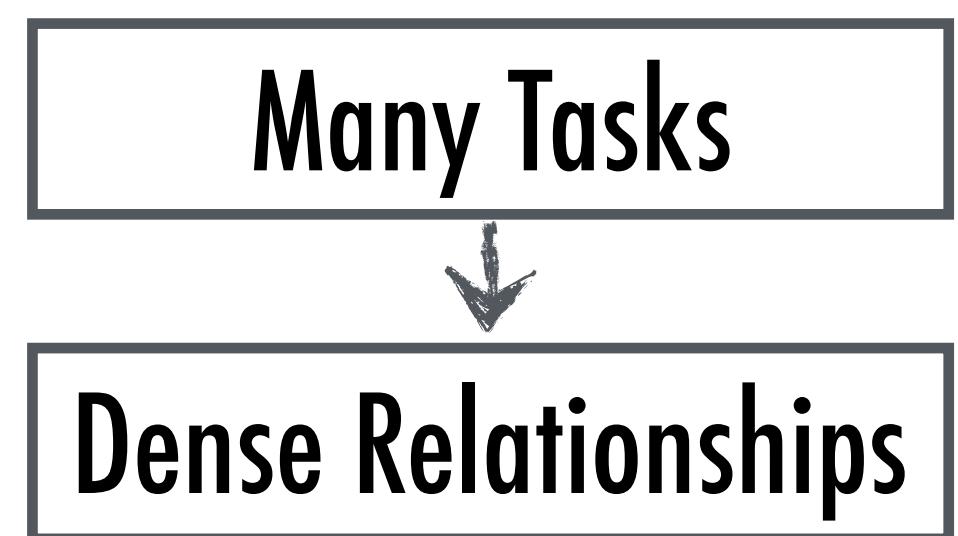


Task Relationships

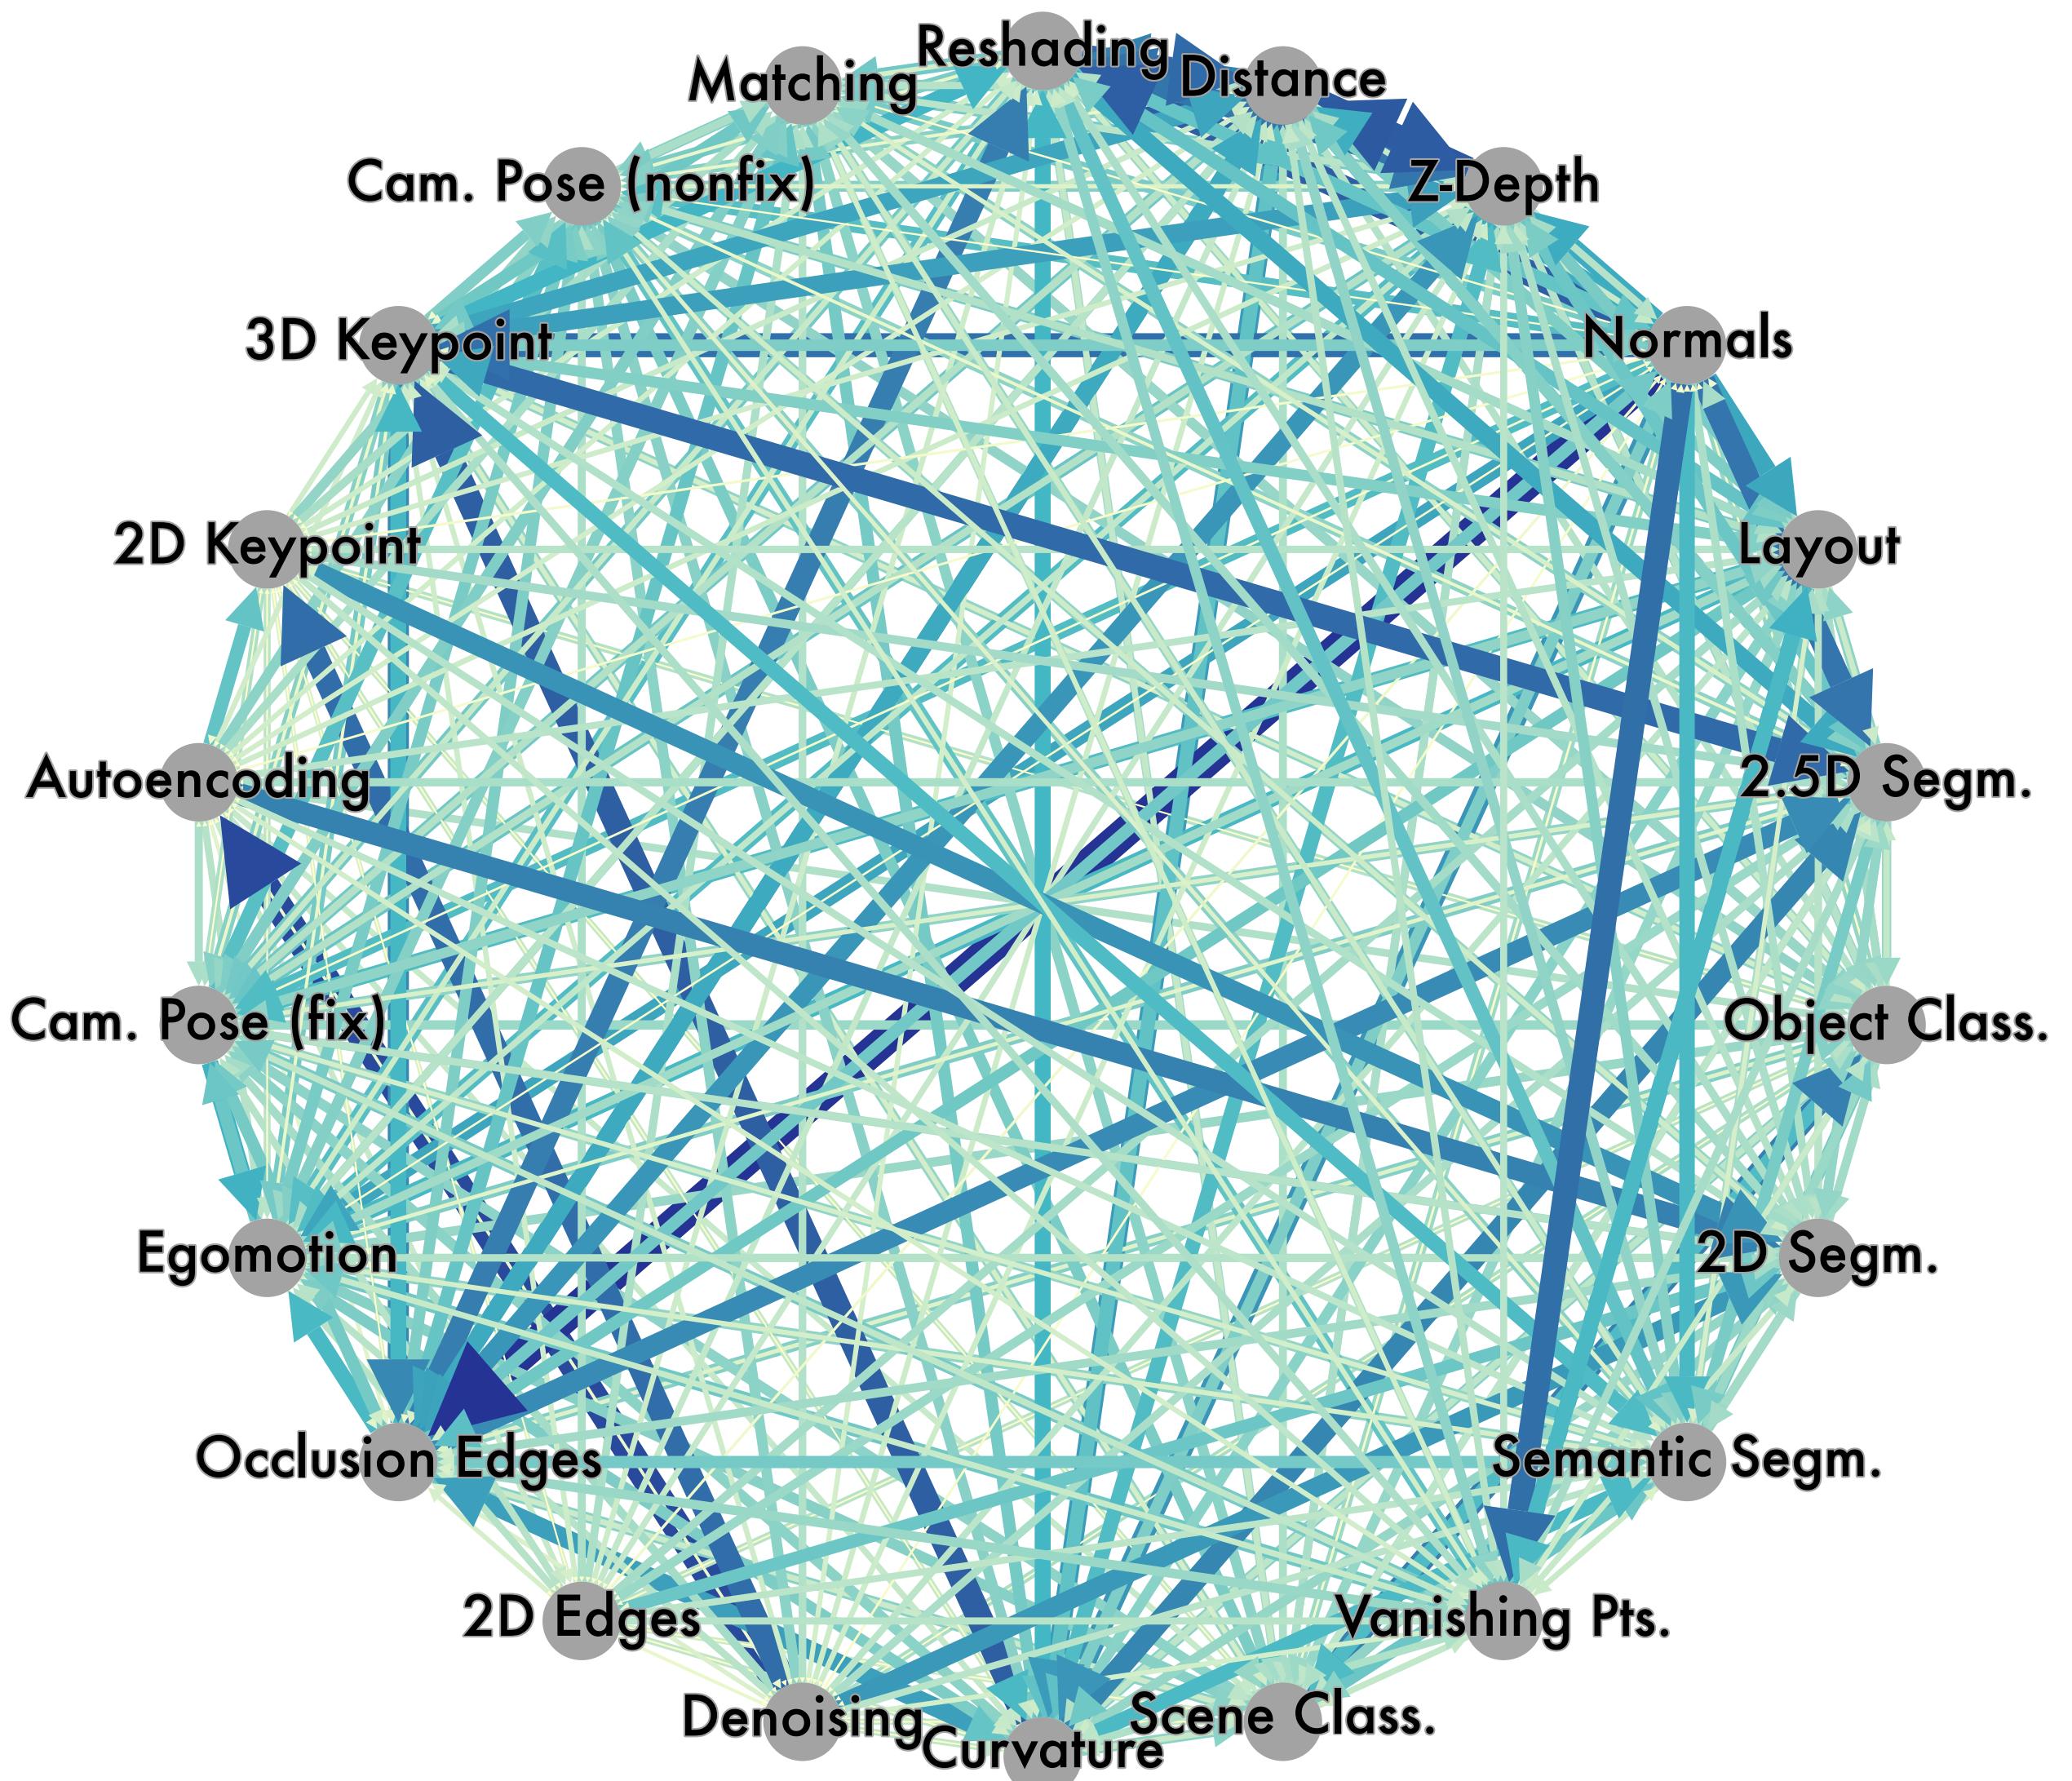
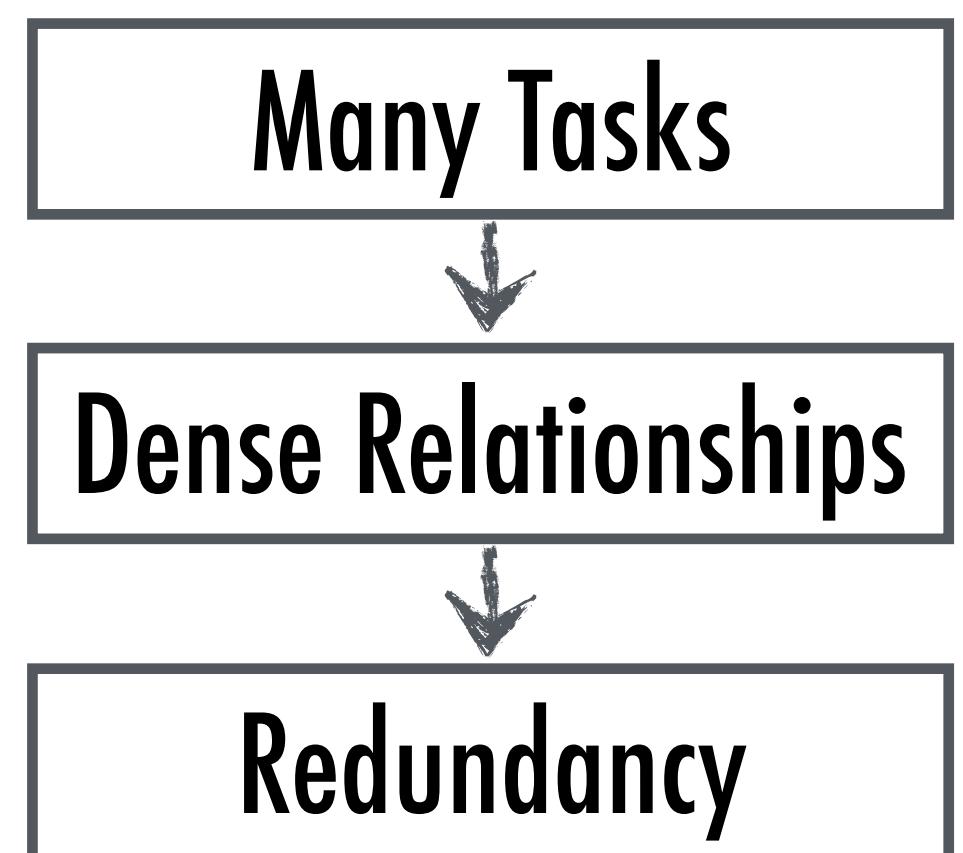
Many Tasks



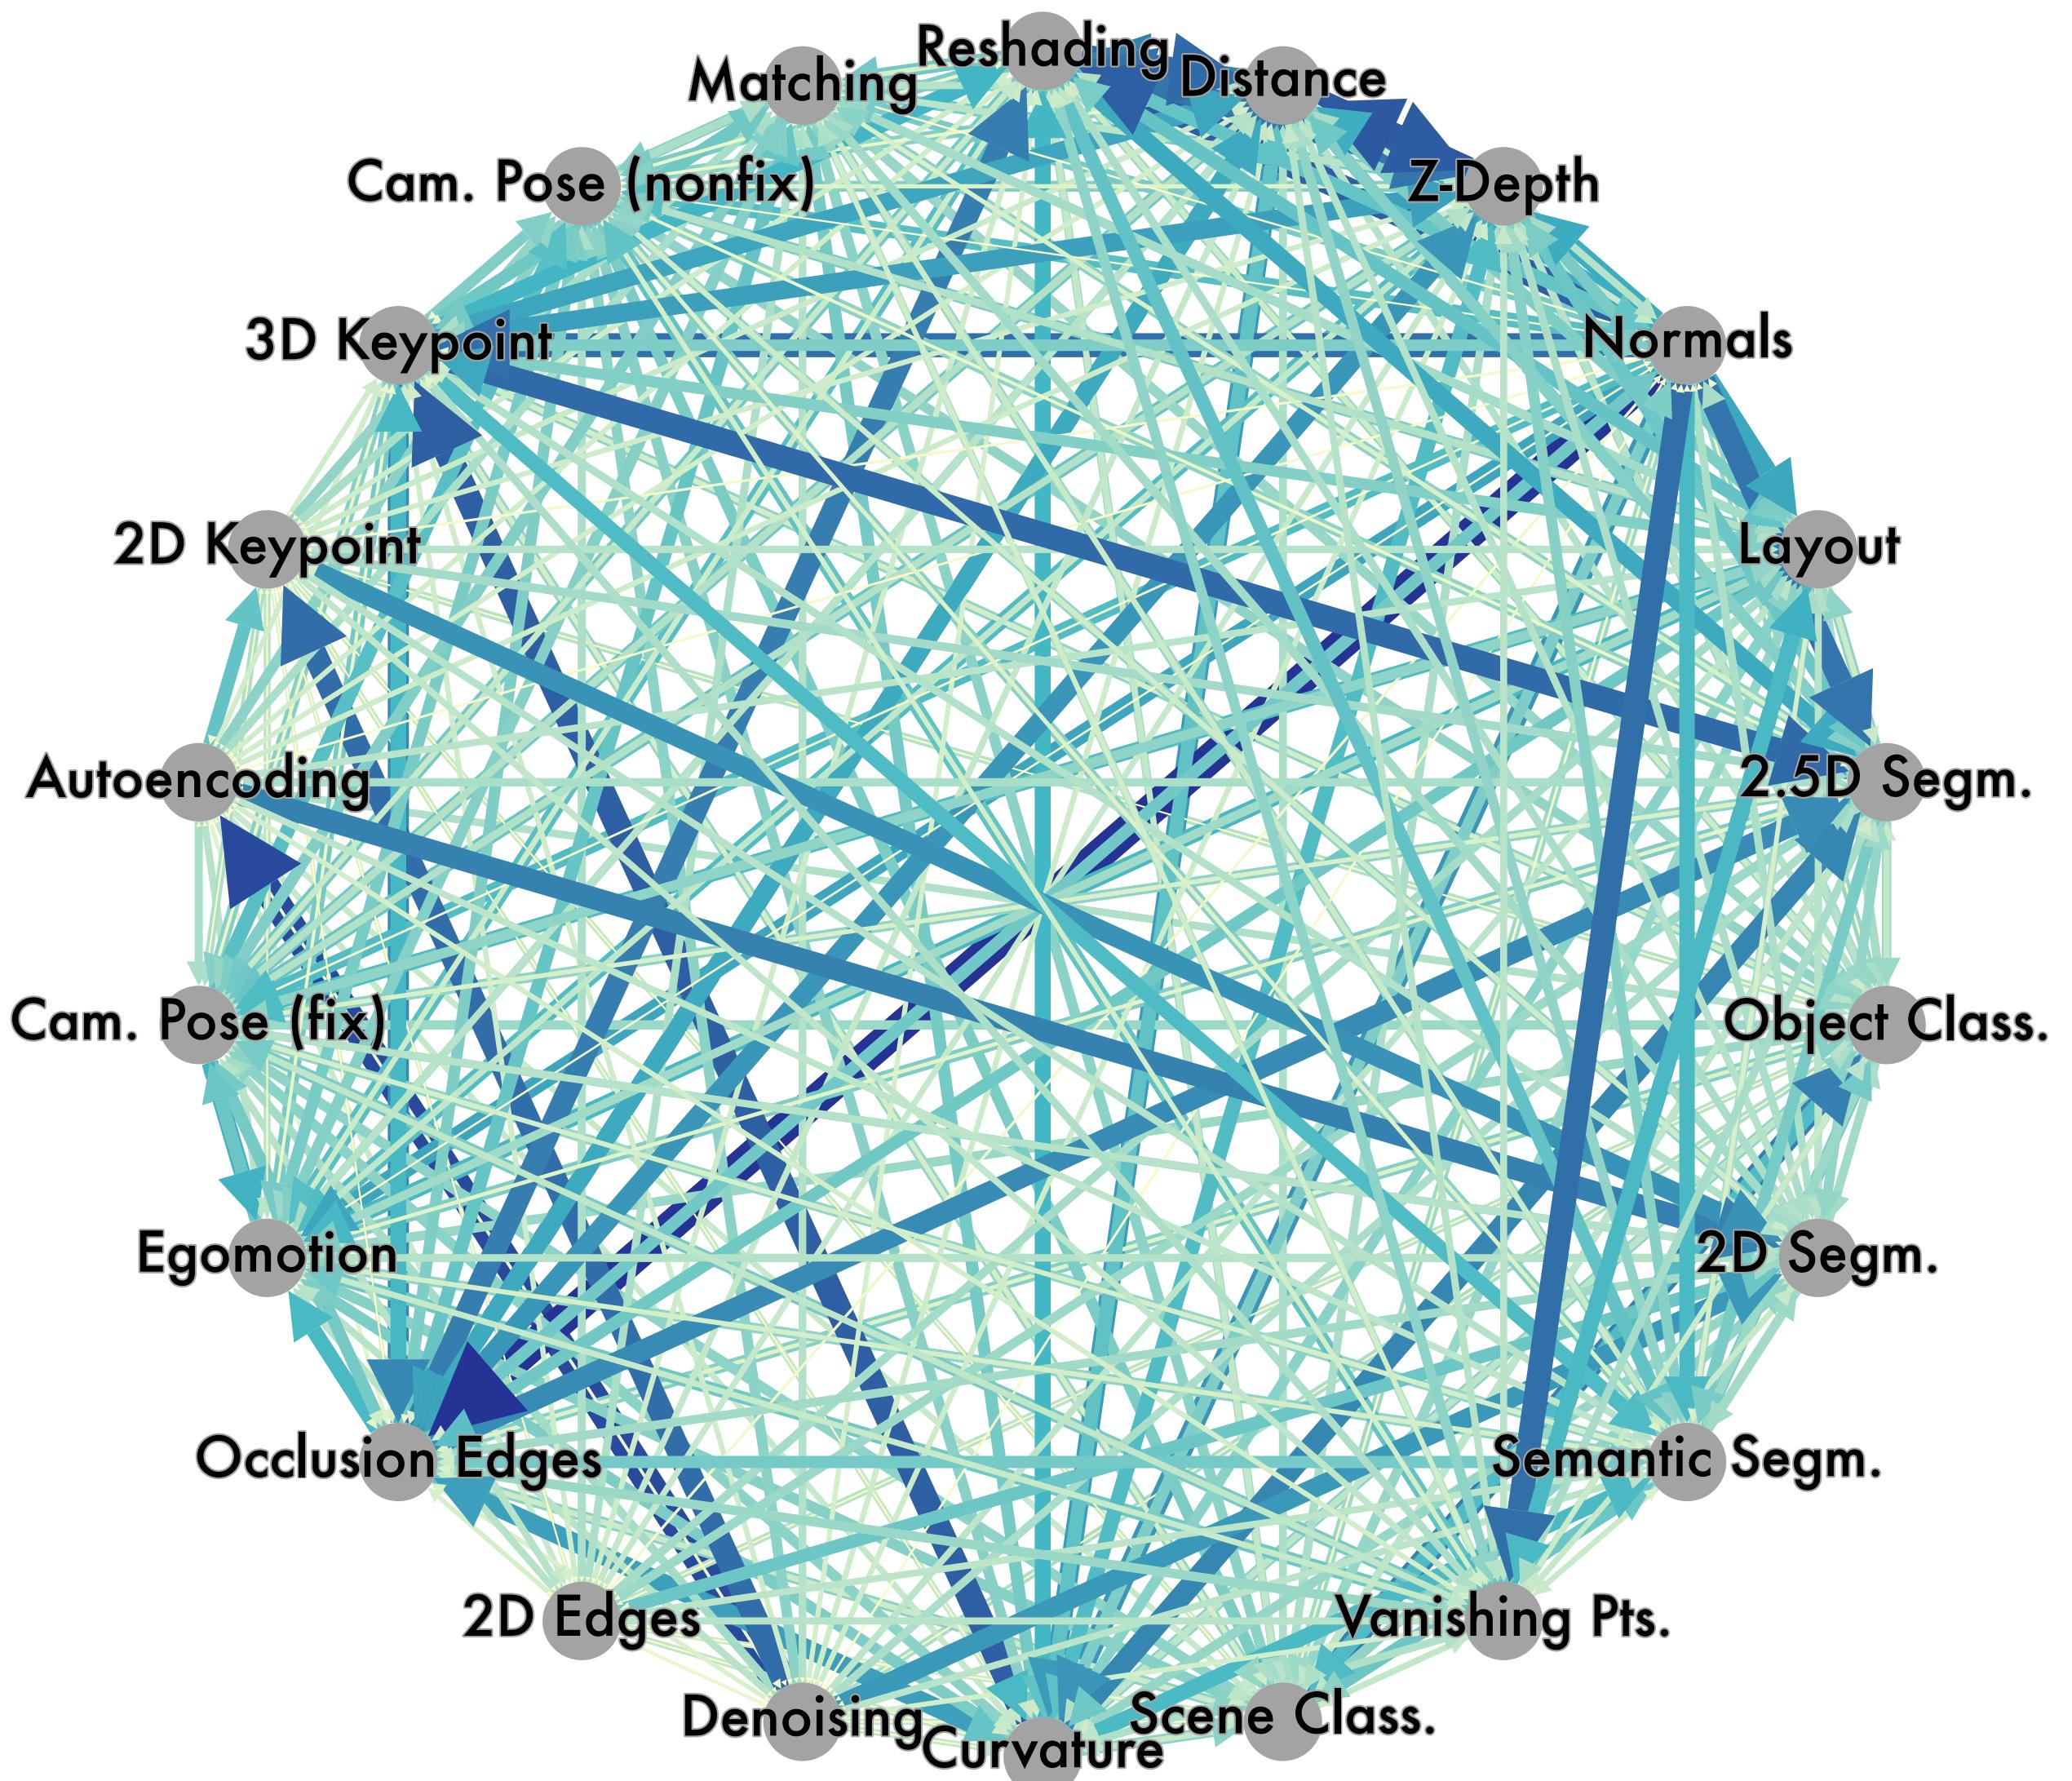
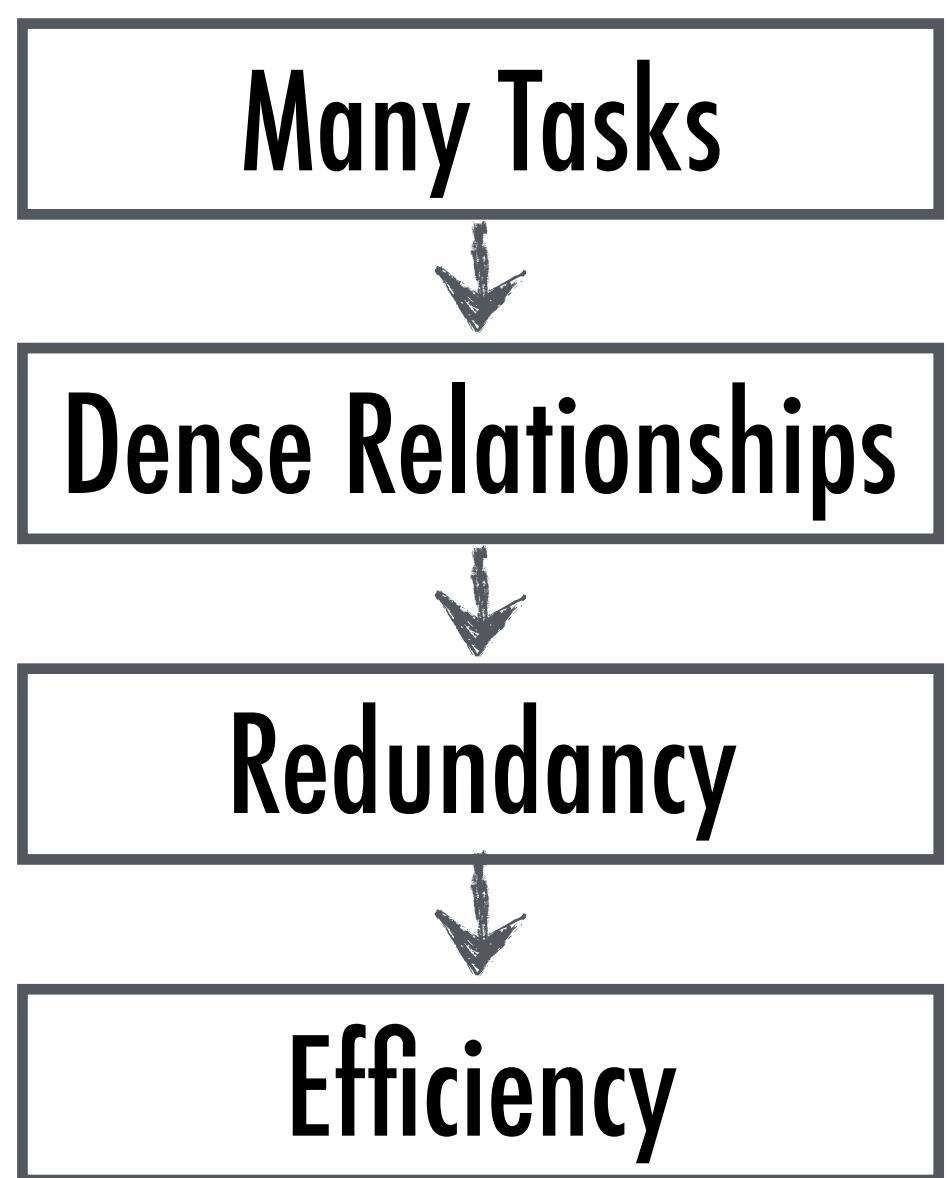
Task Relationships



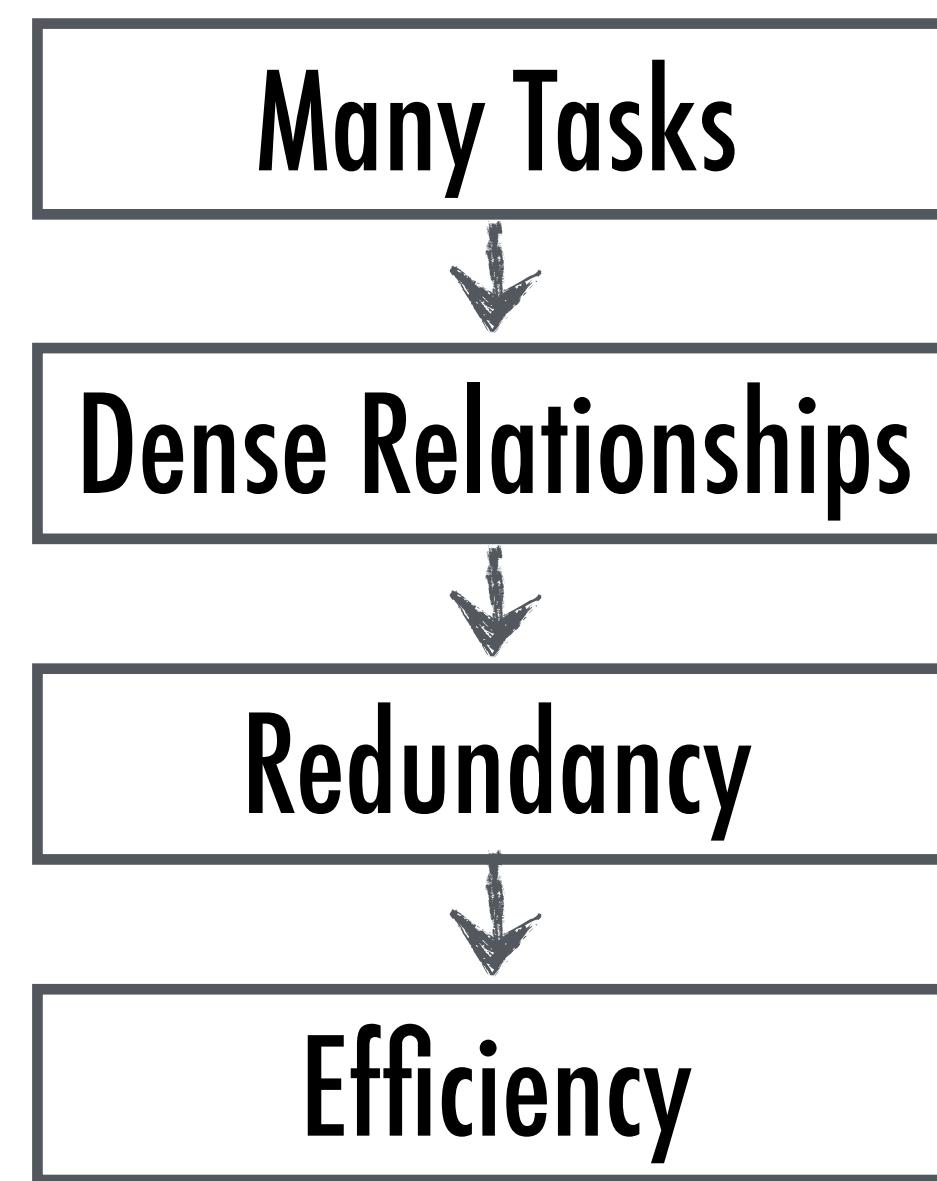
Task Relationships



Task Relationships

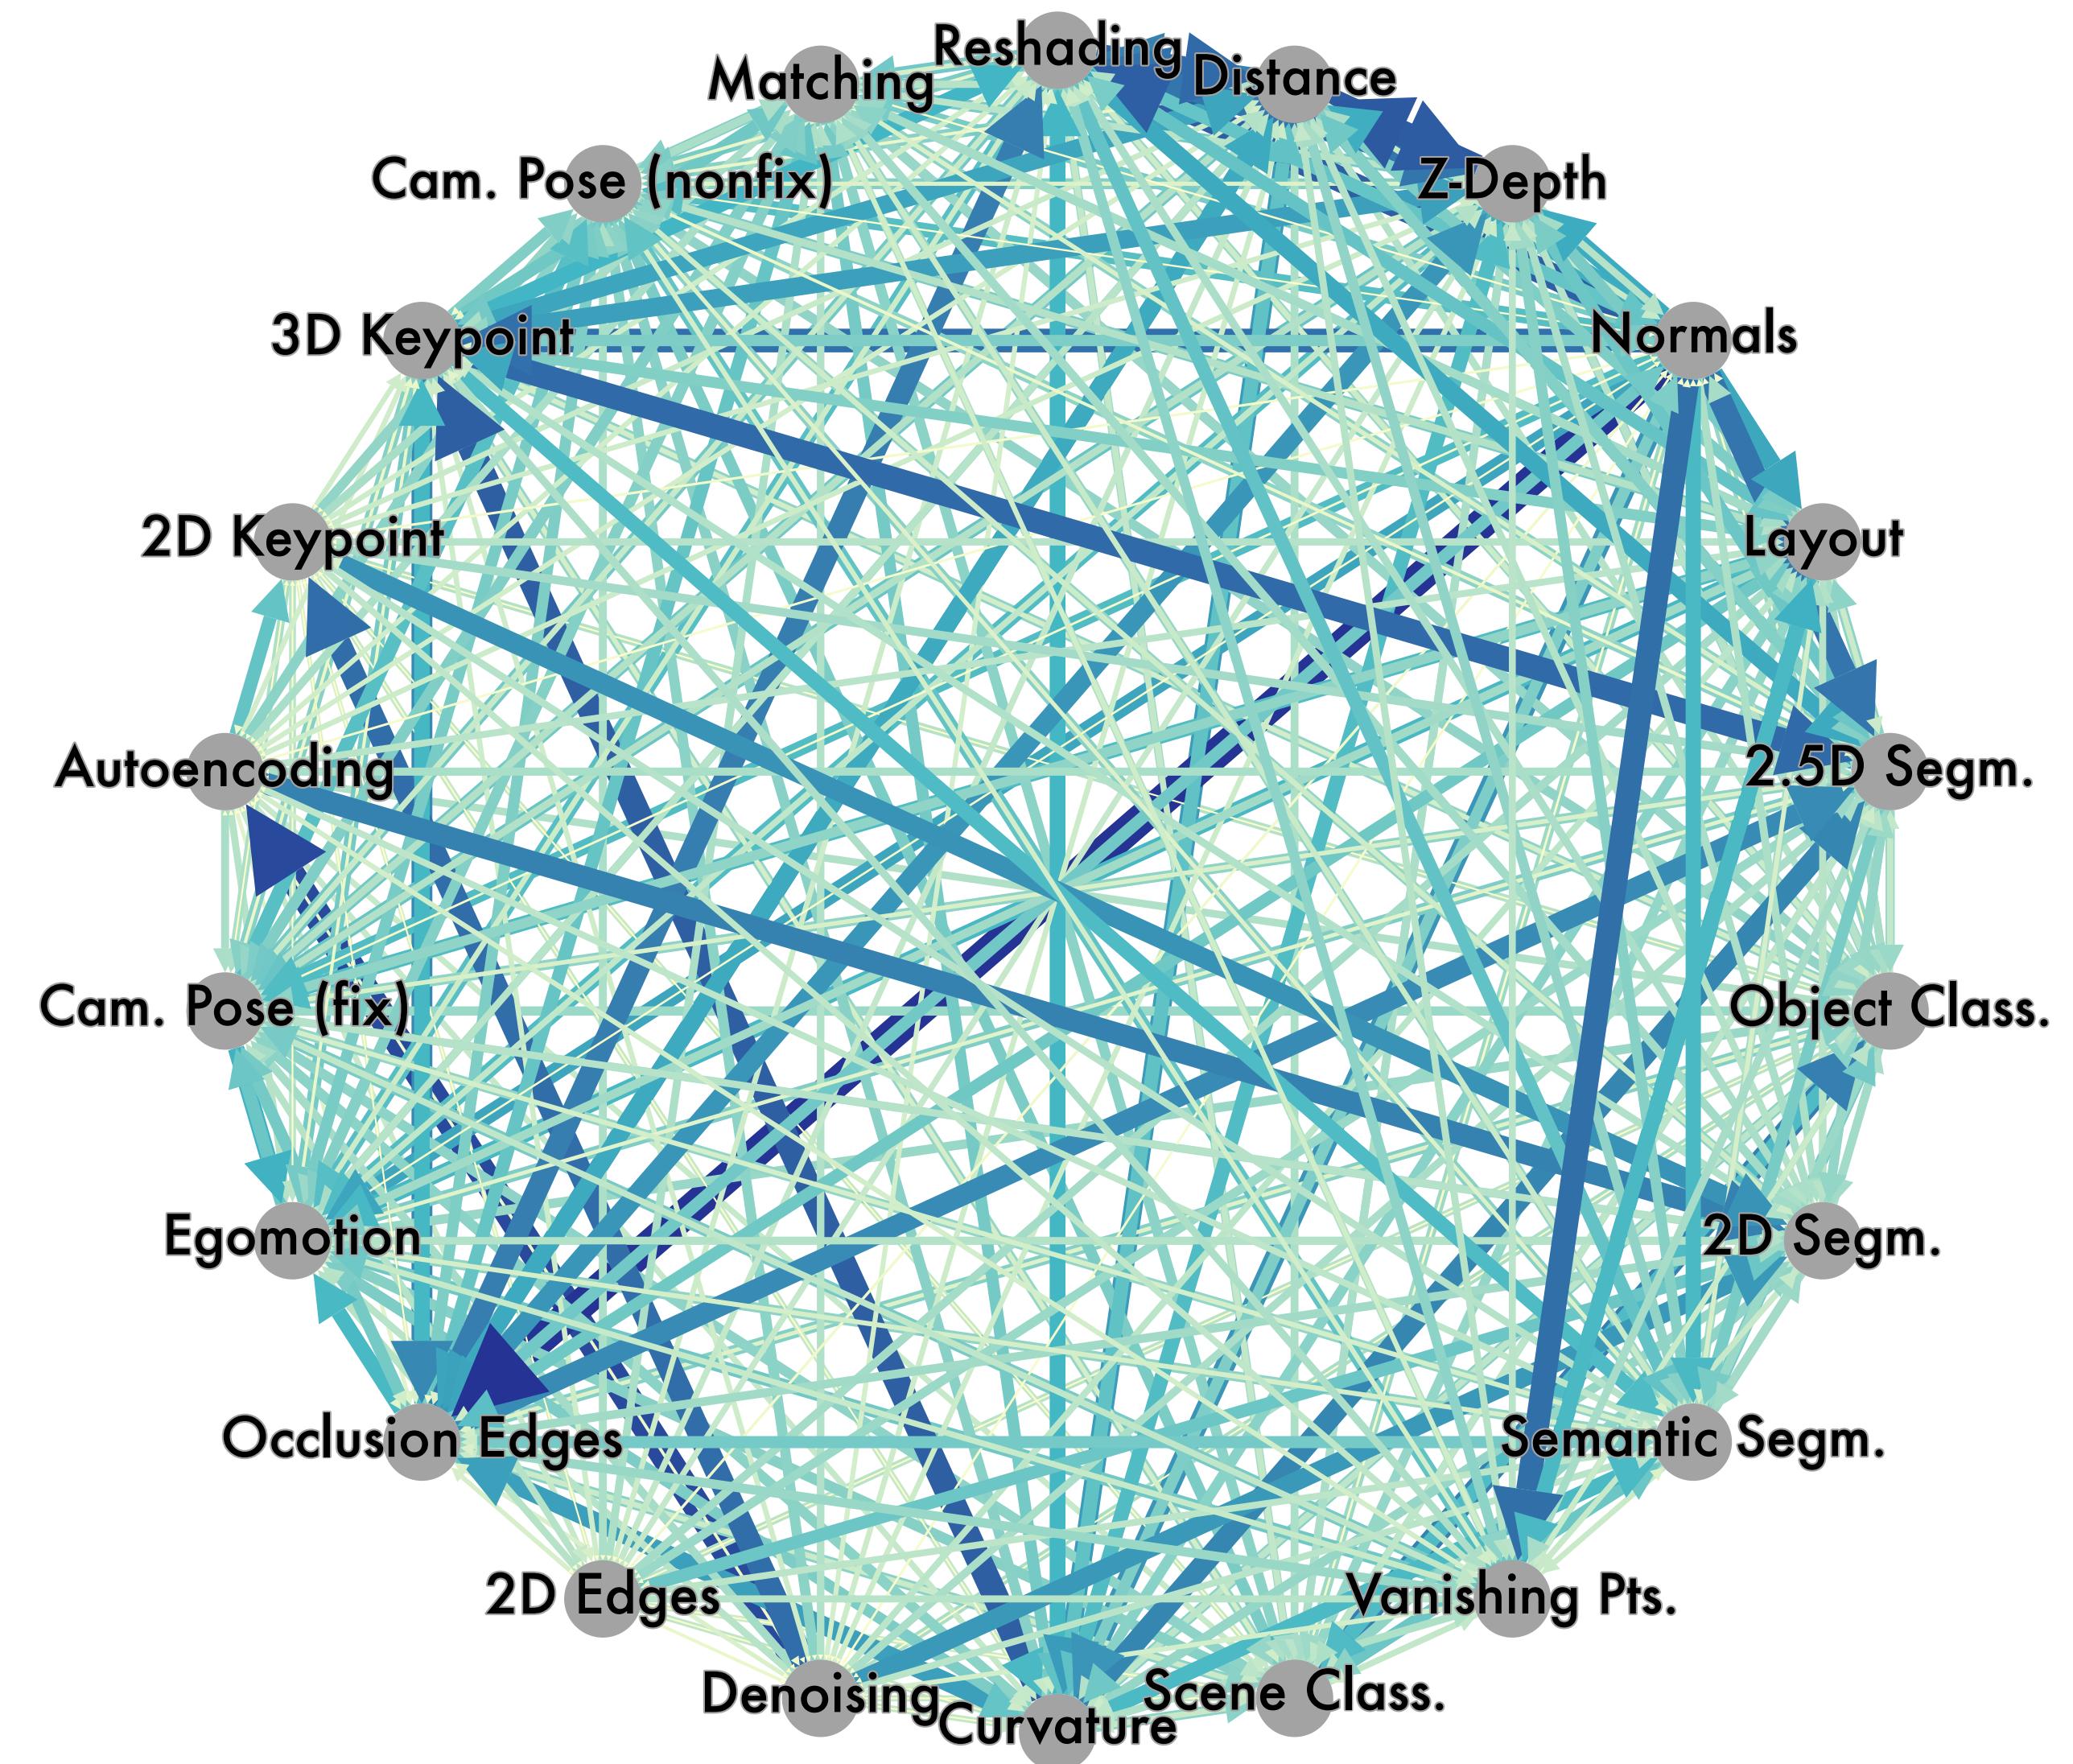


Task Relationships

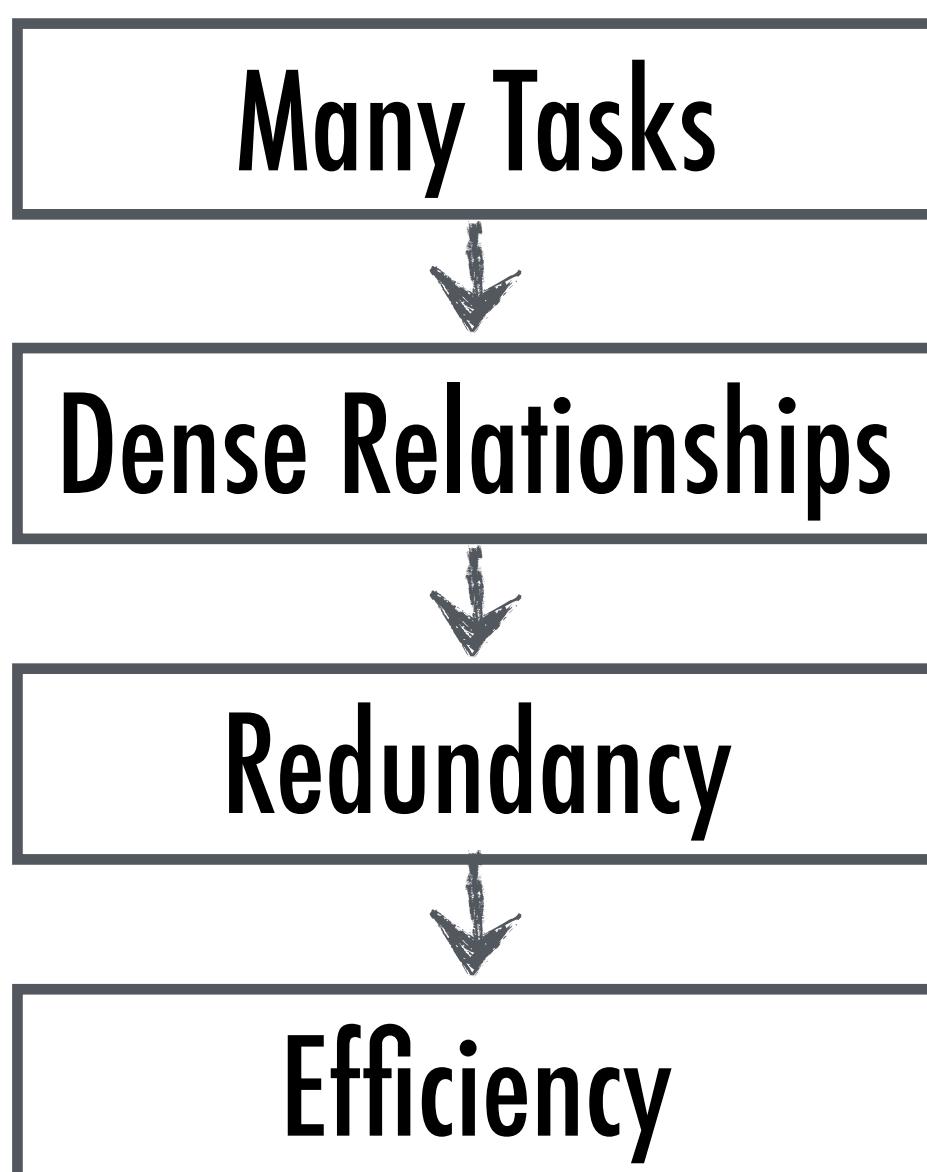


Supervision Efficiency

- Self/Un Supervised Learning [1,2,3]
- ImageNet features. “Fine-Tuning” [4,5,6]
- Meta Learning [7,8,9]
- Domain Adaptation [10,11,12]



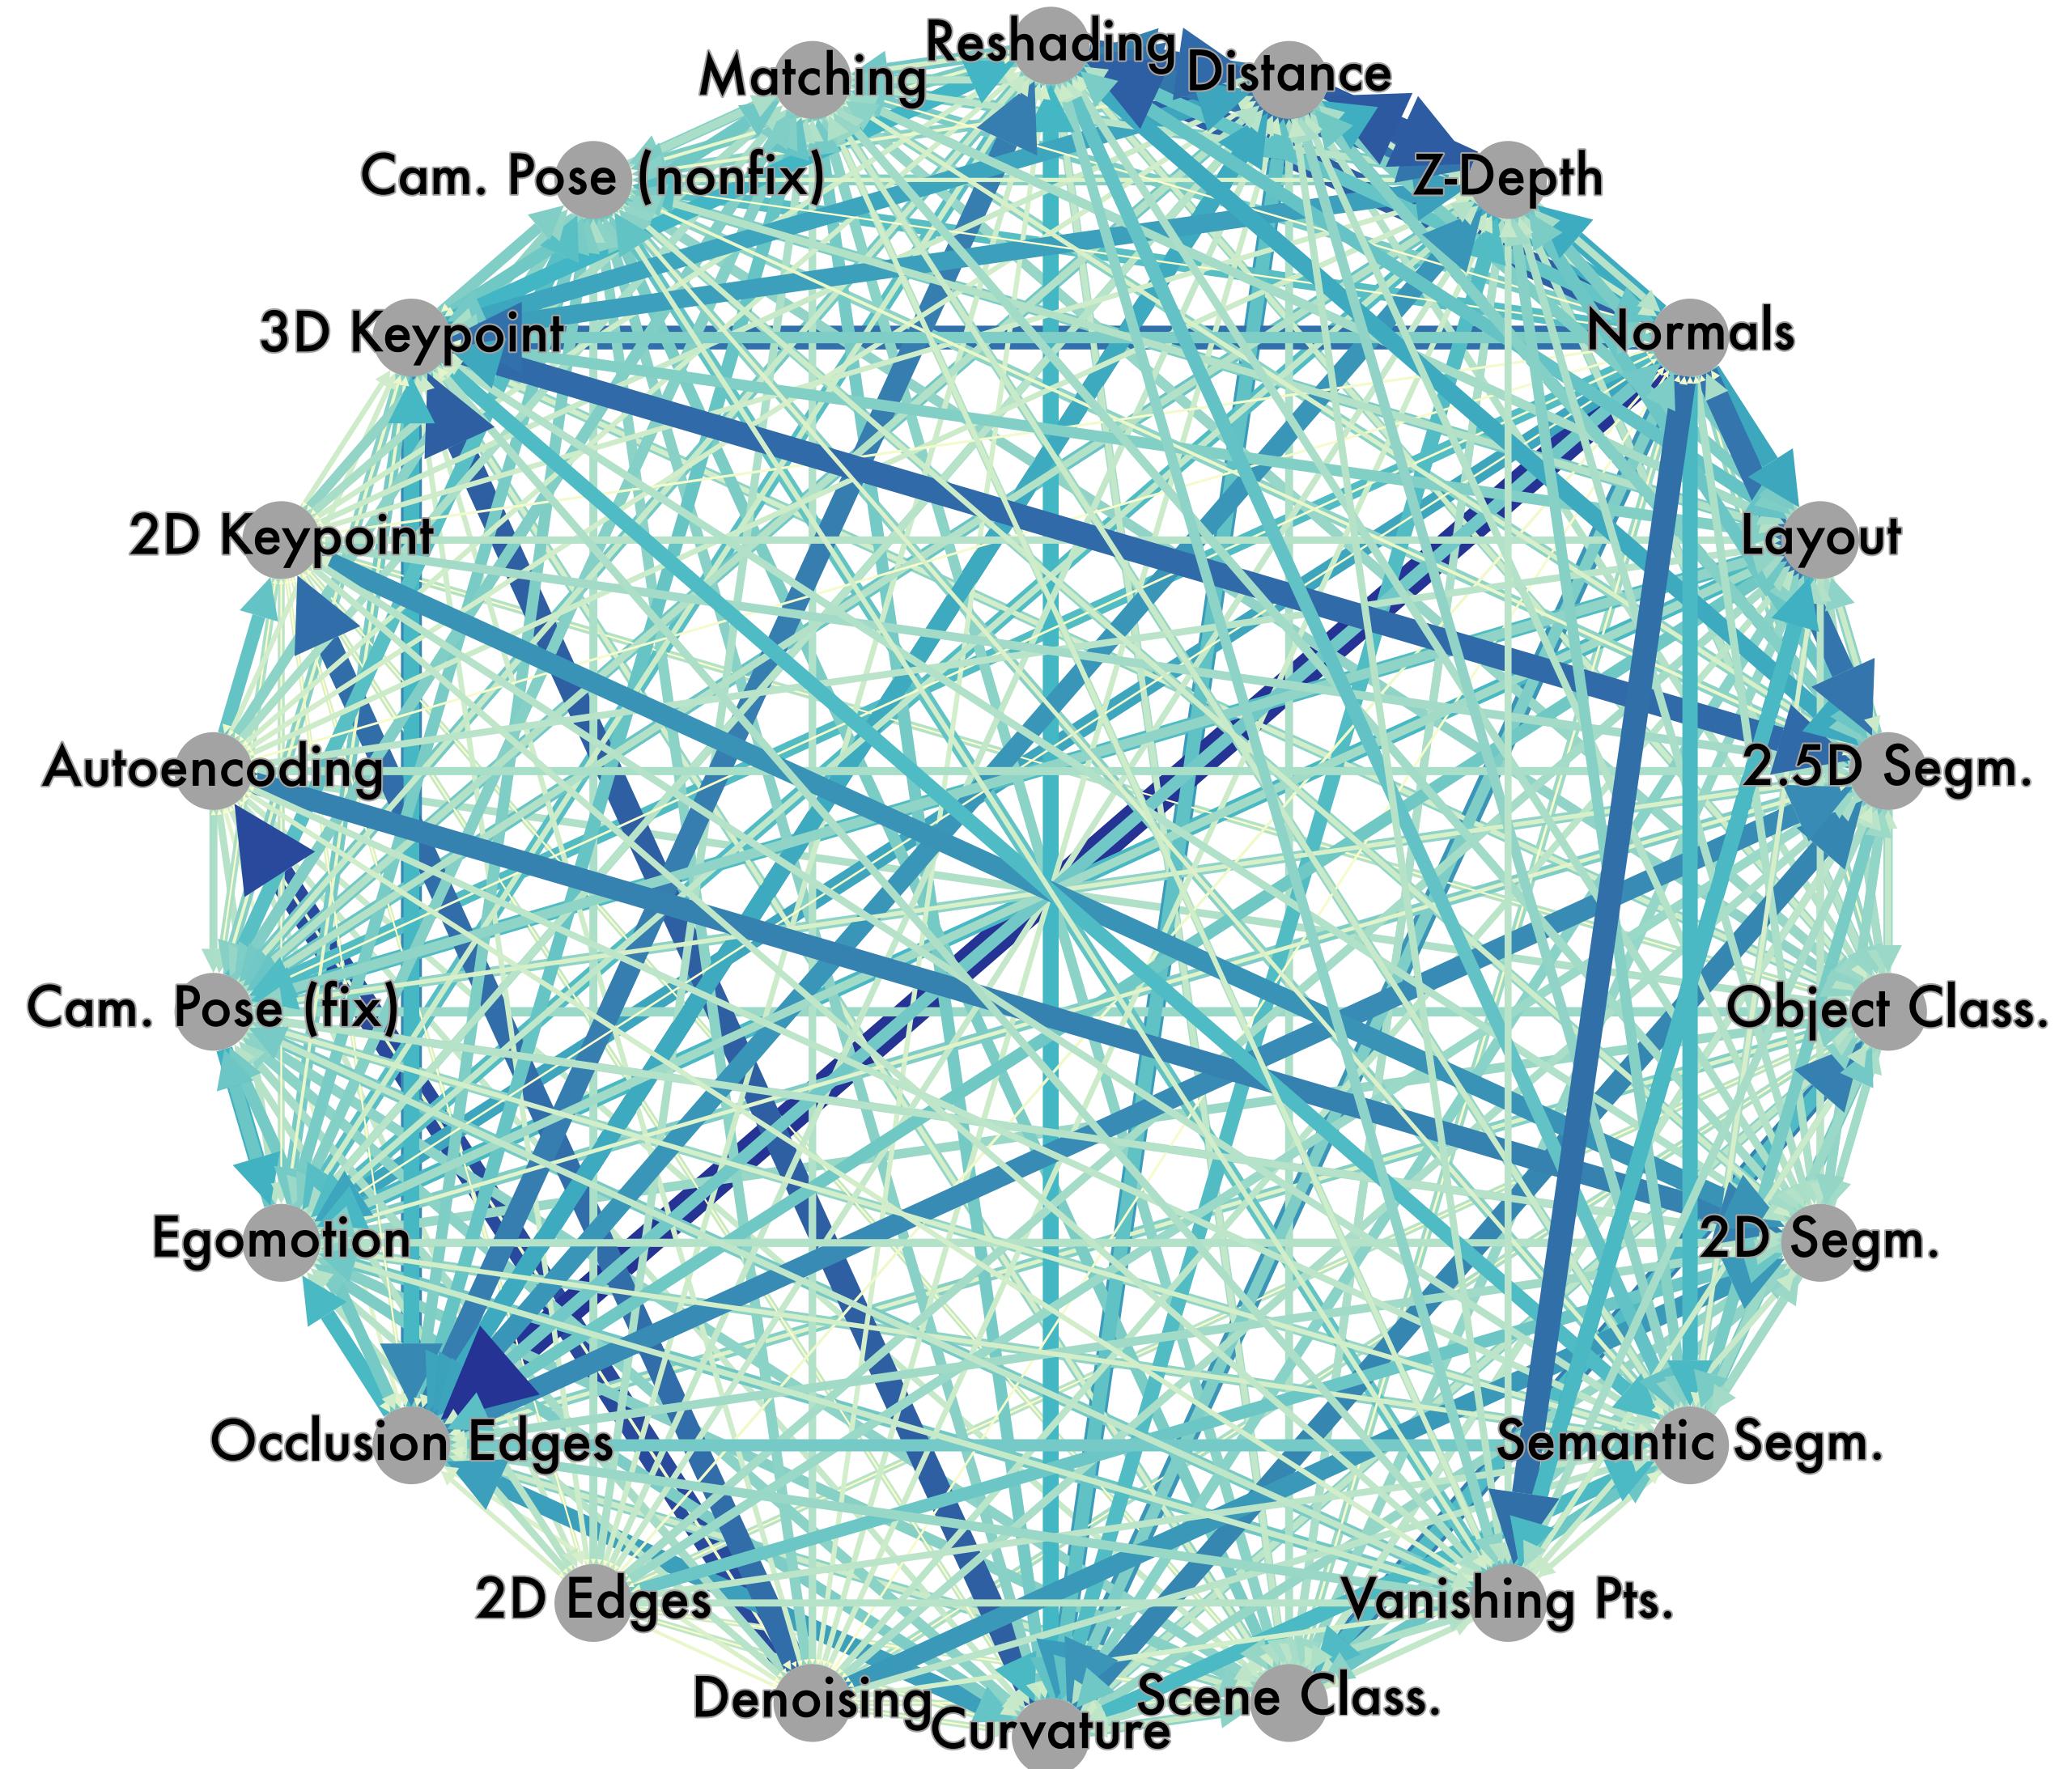
Task Relationships



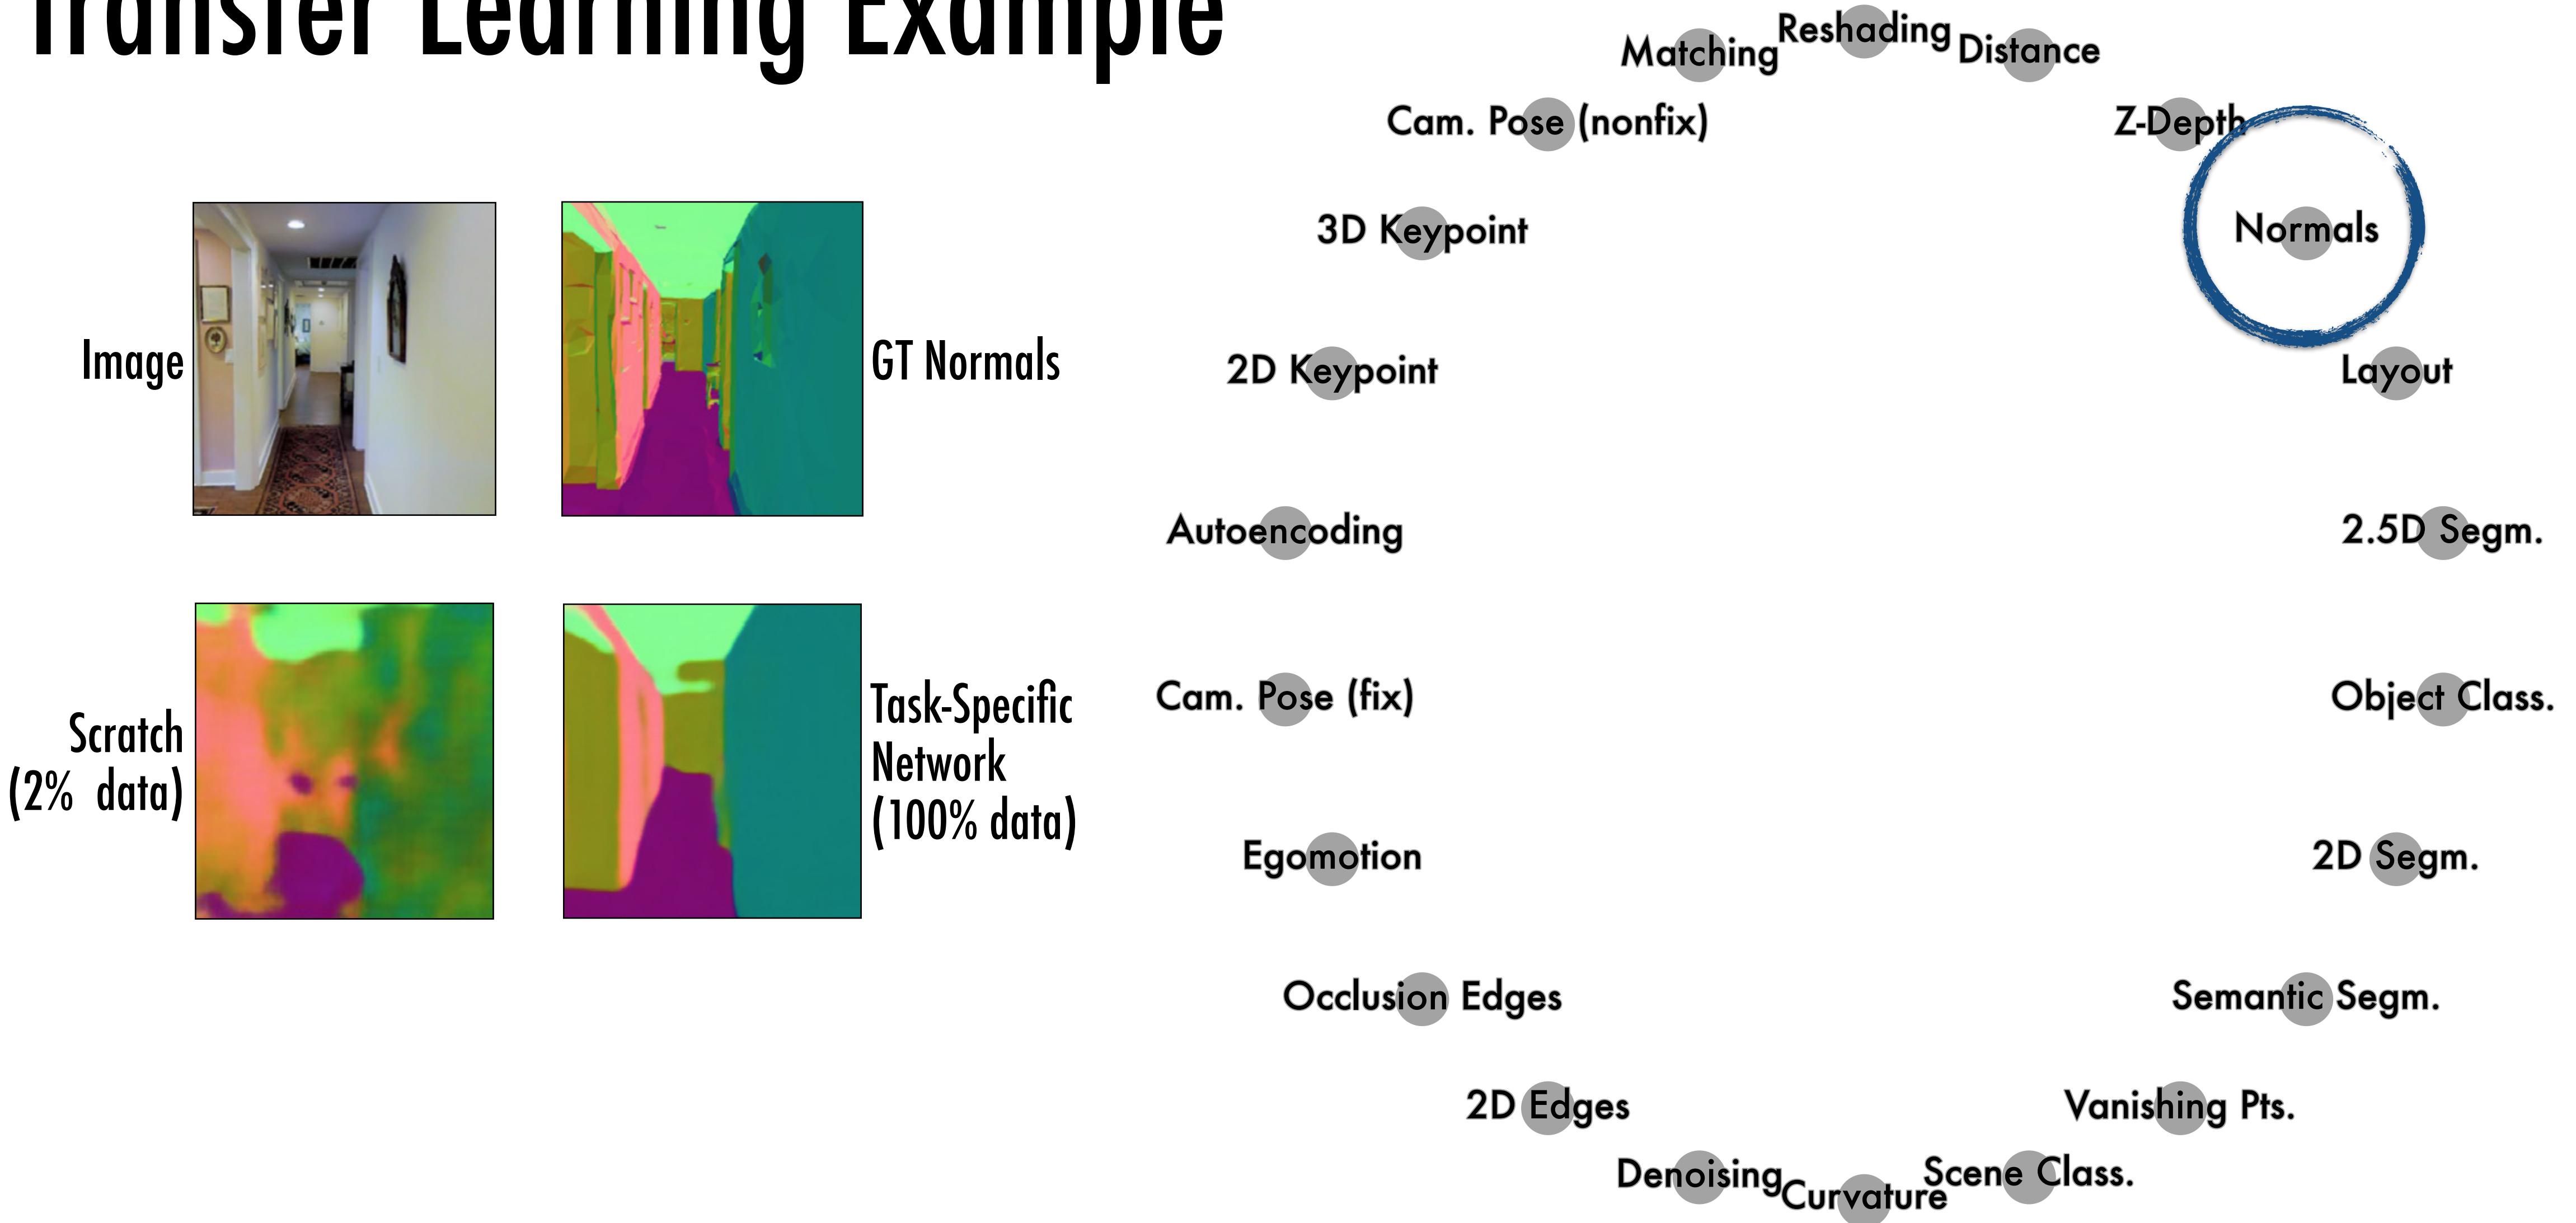
Supervision
Efficiency

Transfer Learning:

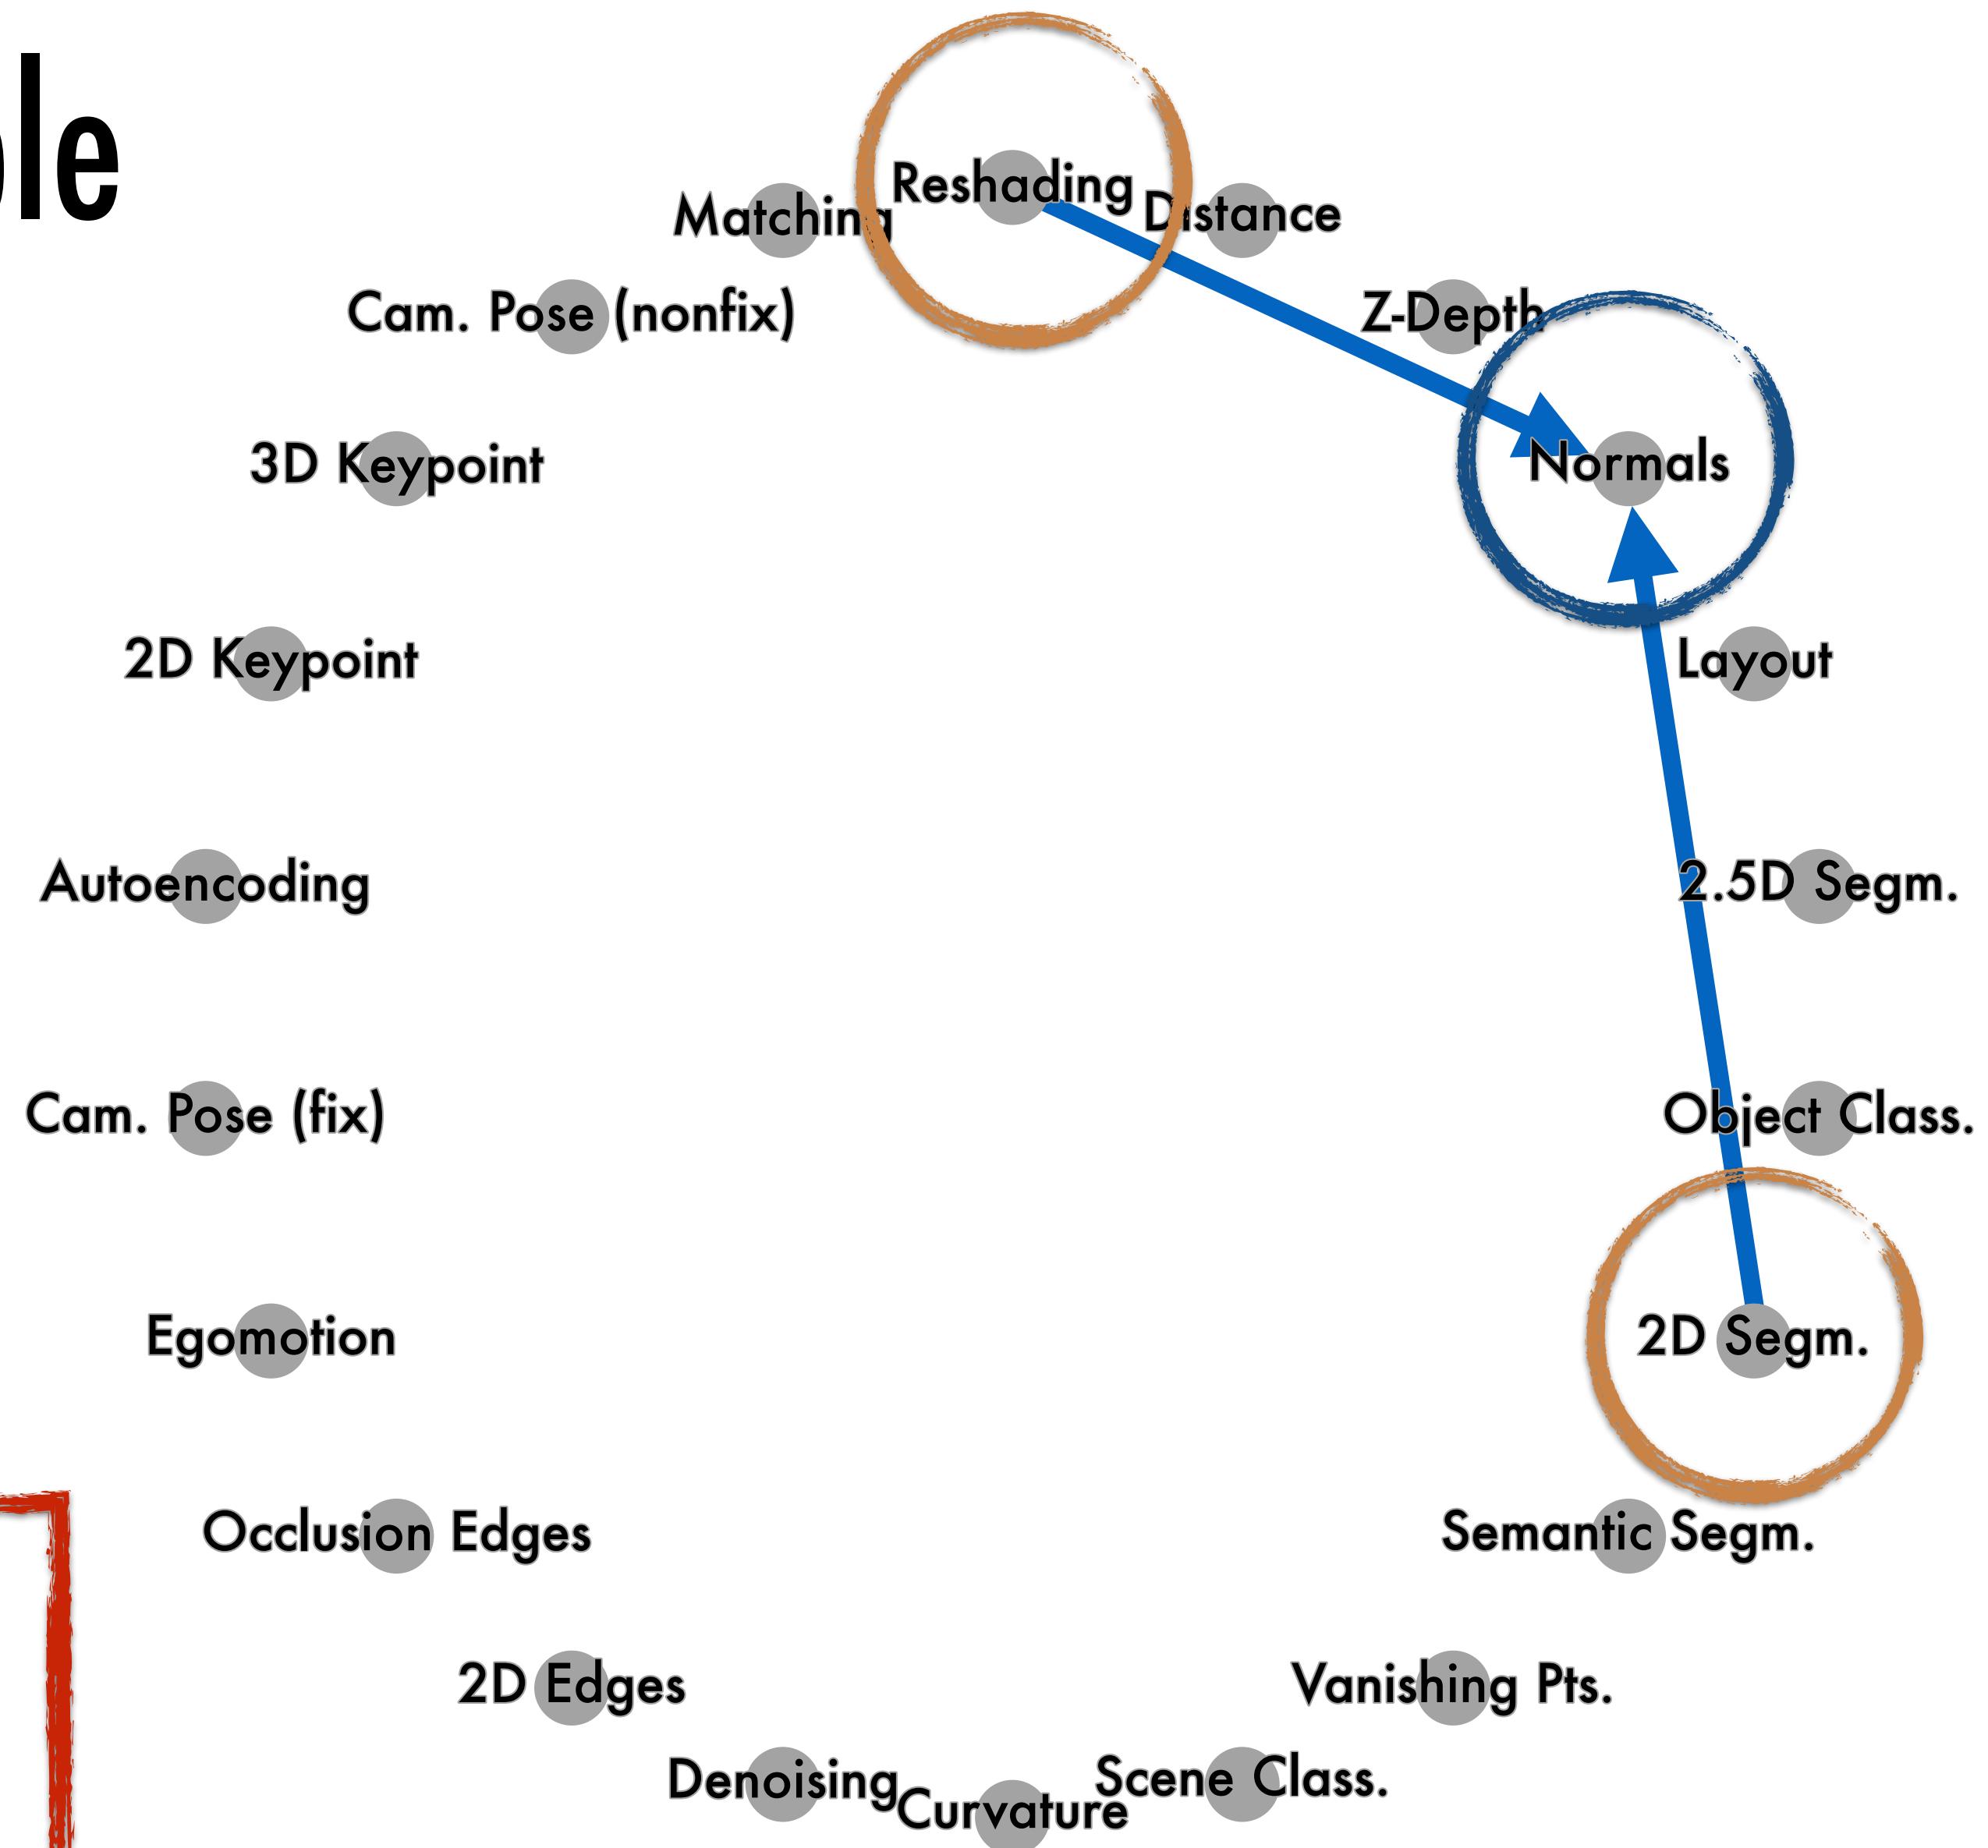
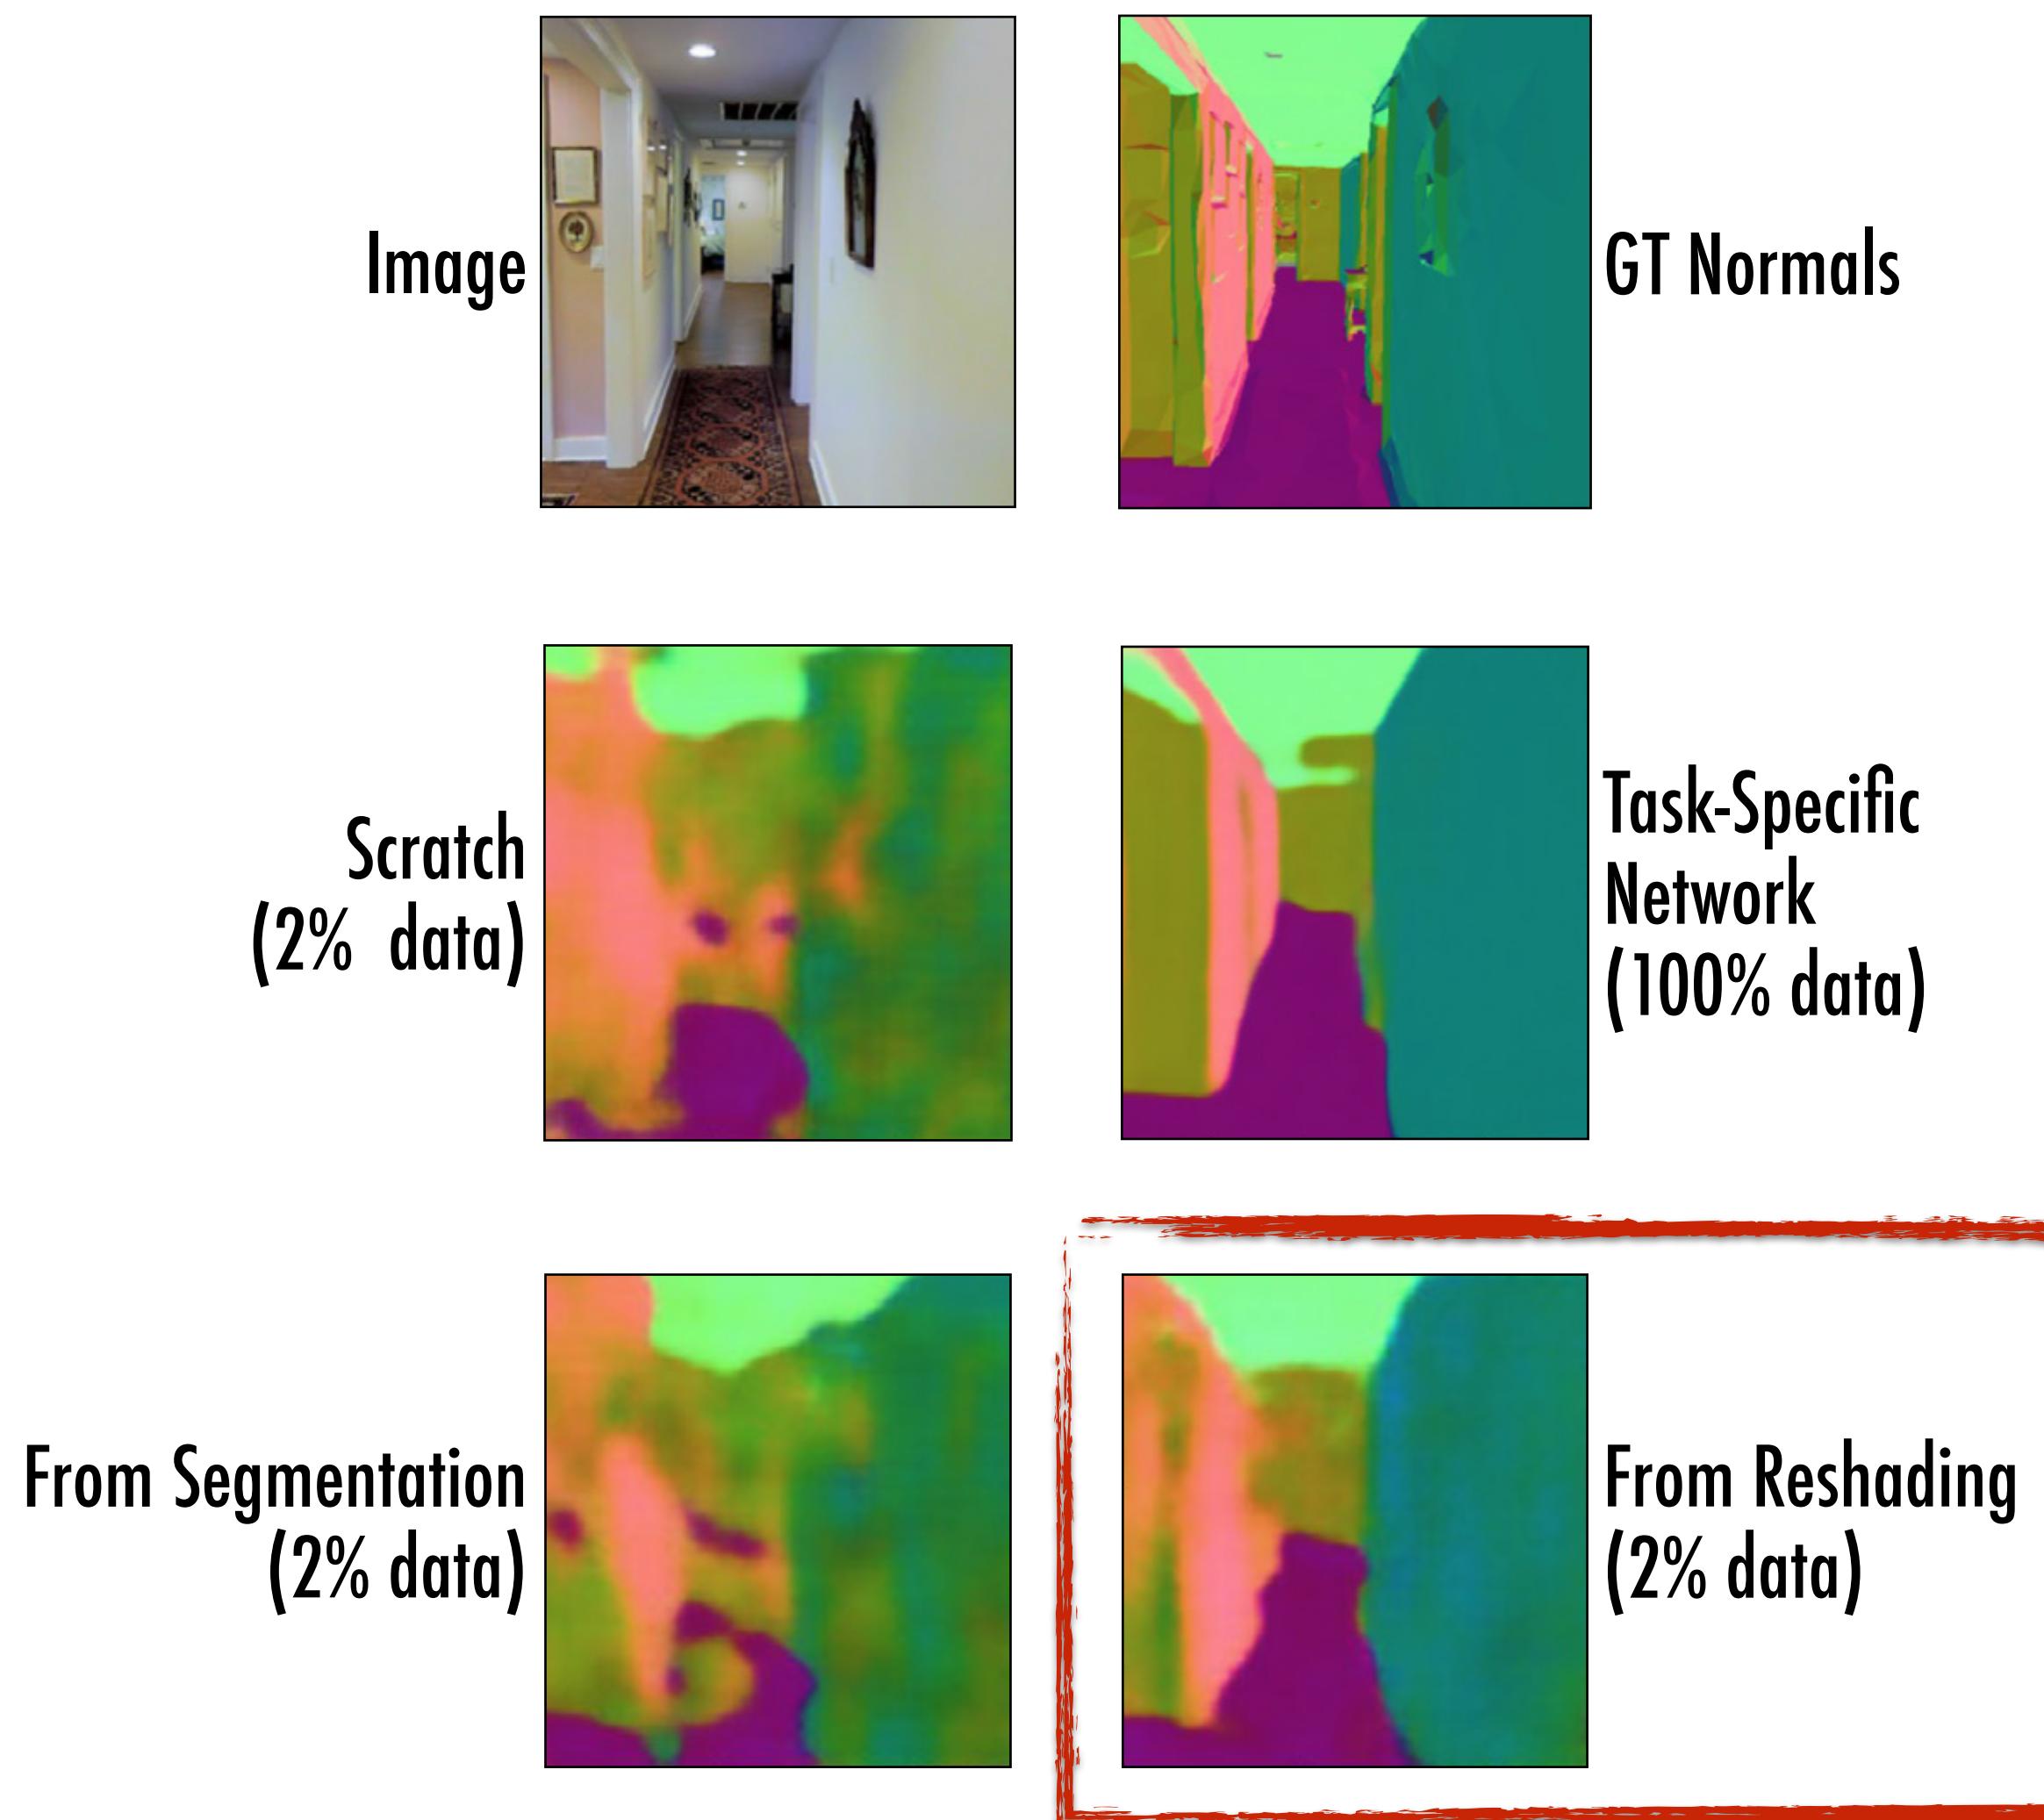
Model developed for task X may be useful
for solving task Y, if X and Y related. [1]



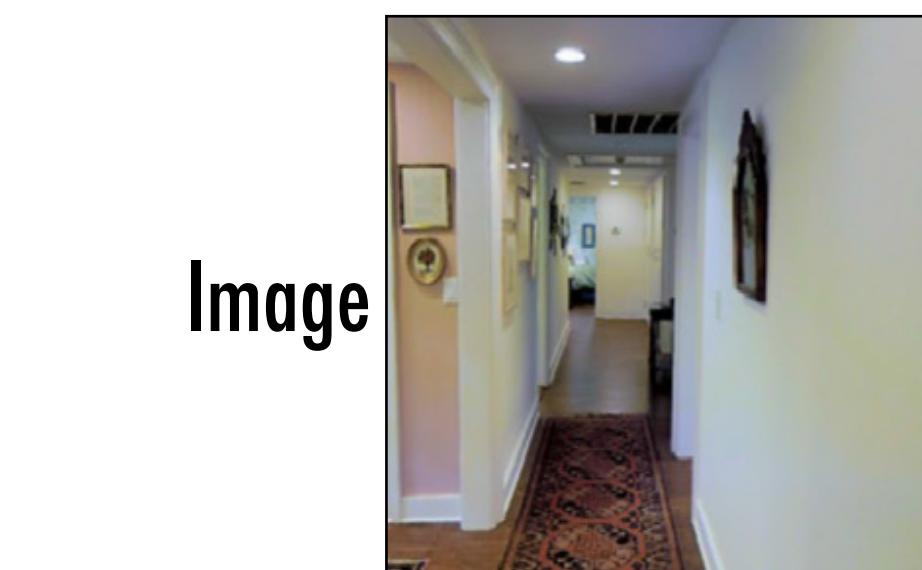
Transfer Learning Example



Transfer Learning Example



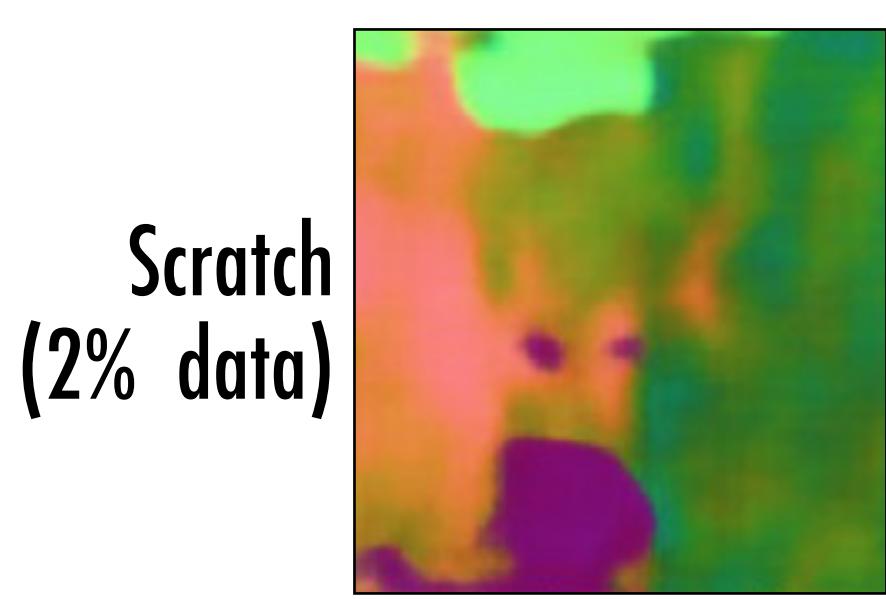
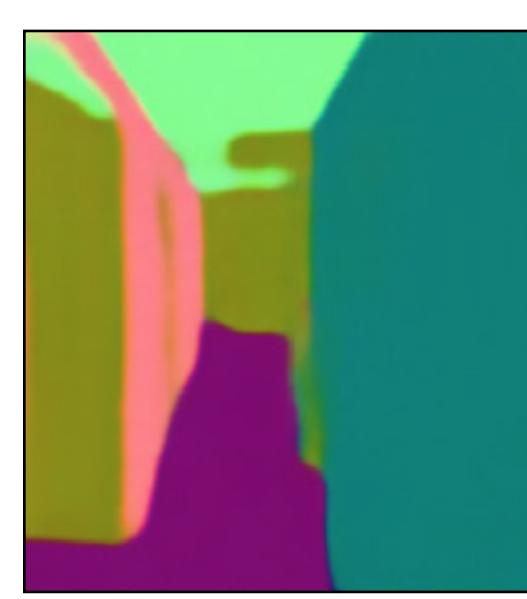
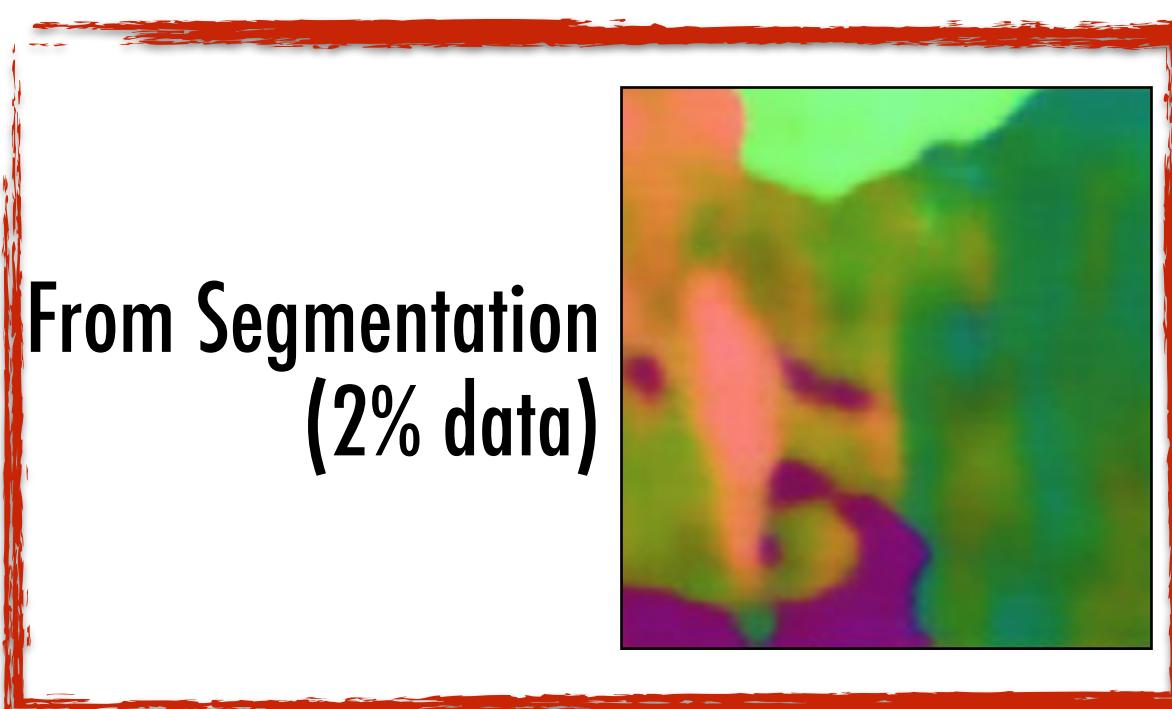
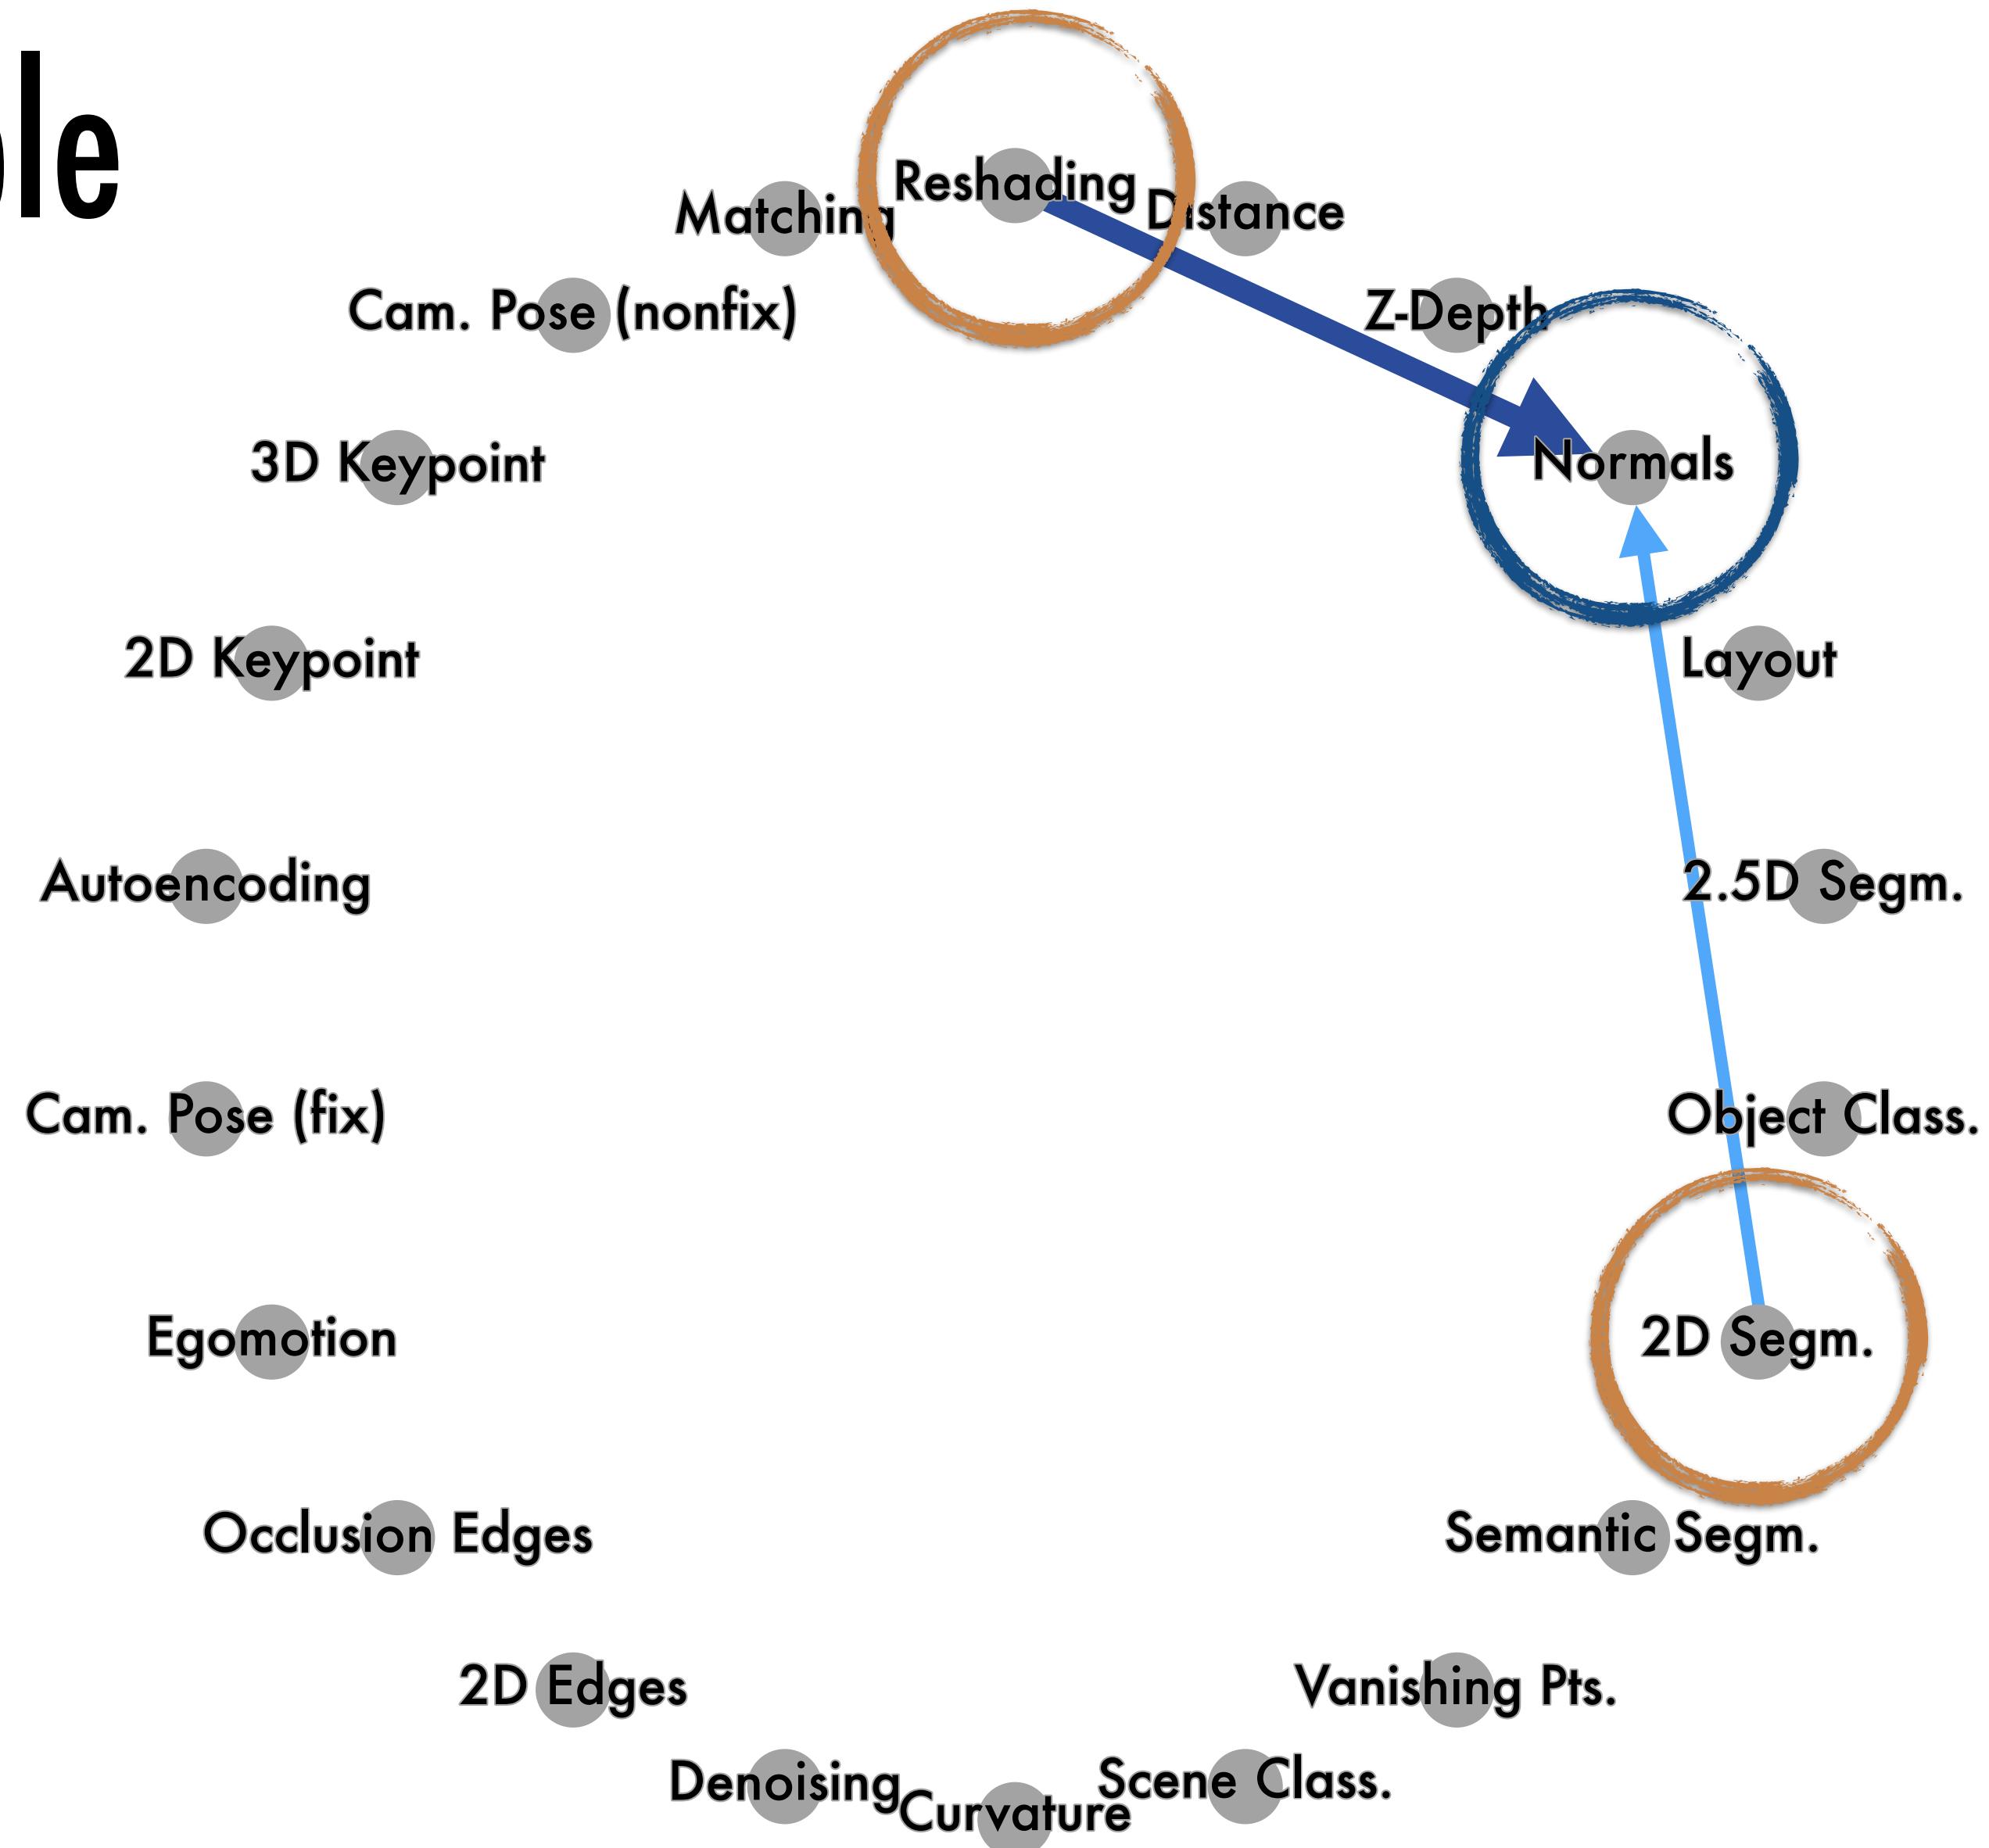
Transfer Learning Example

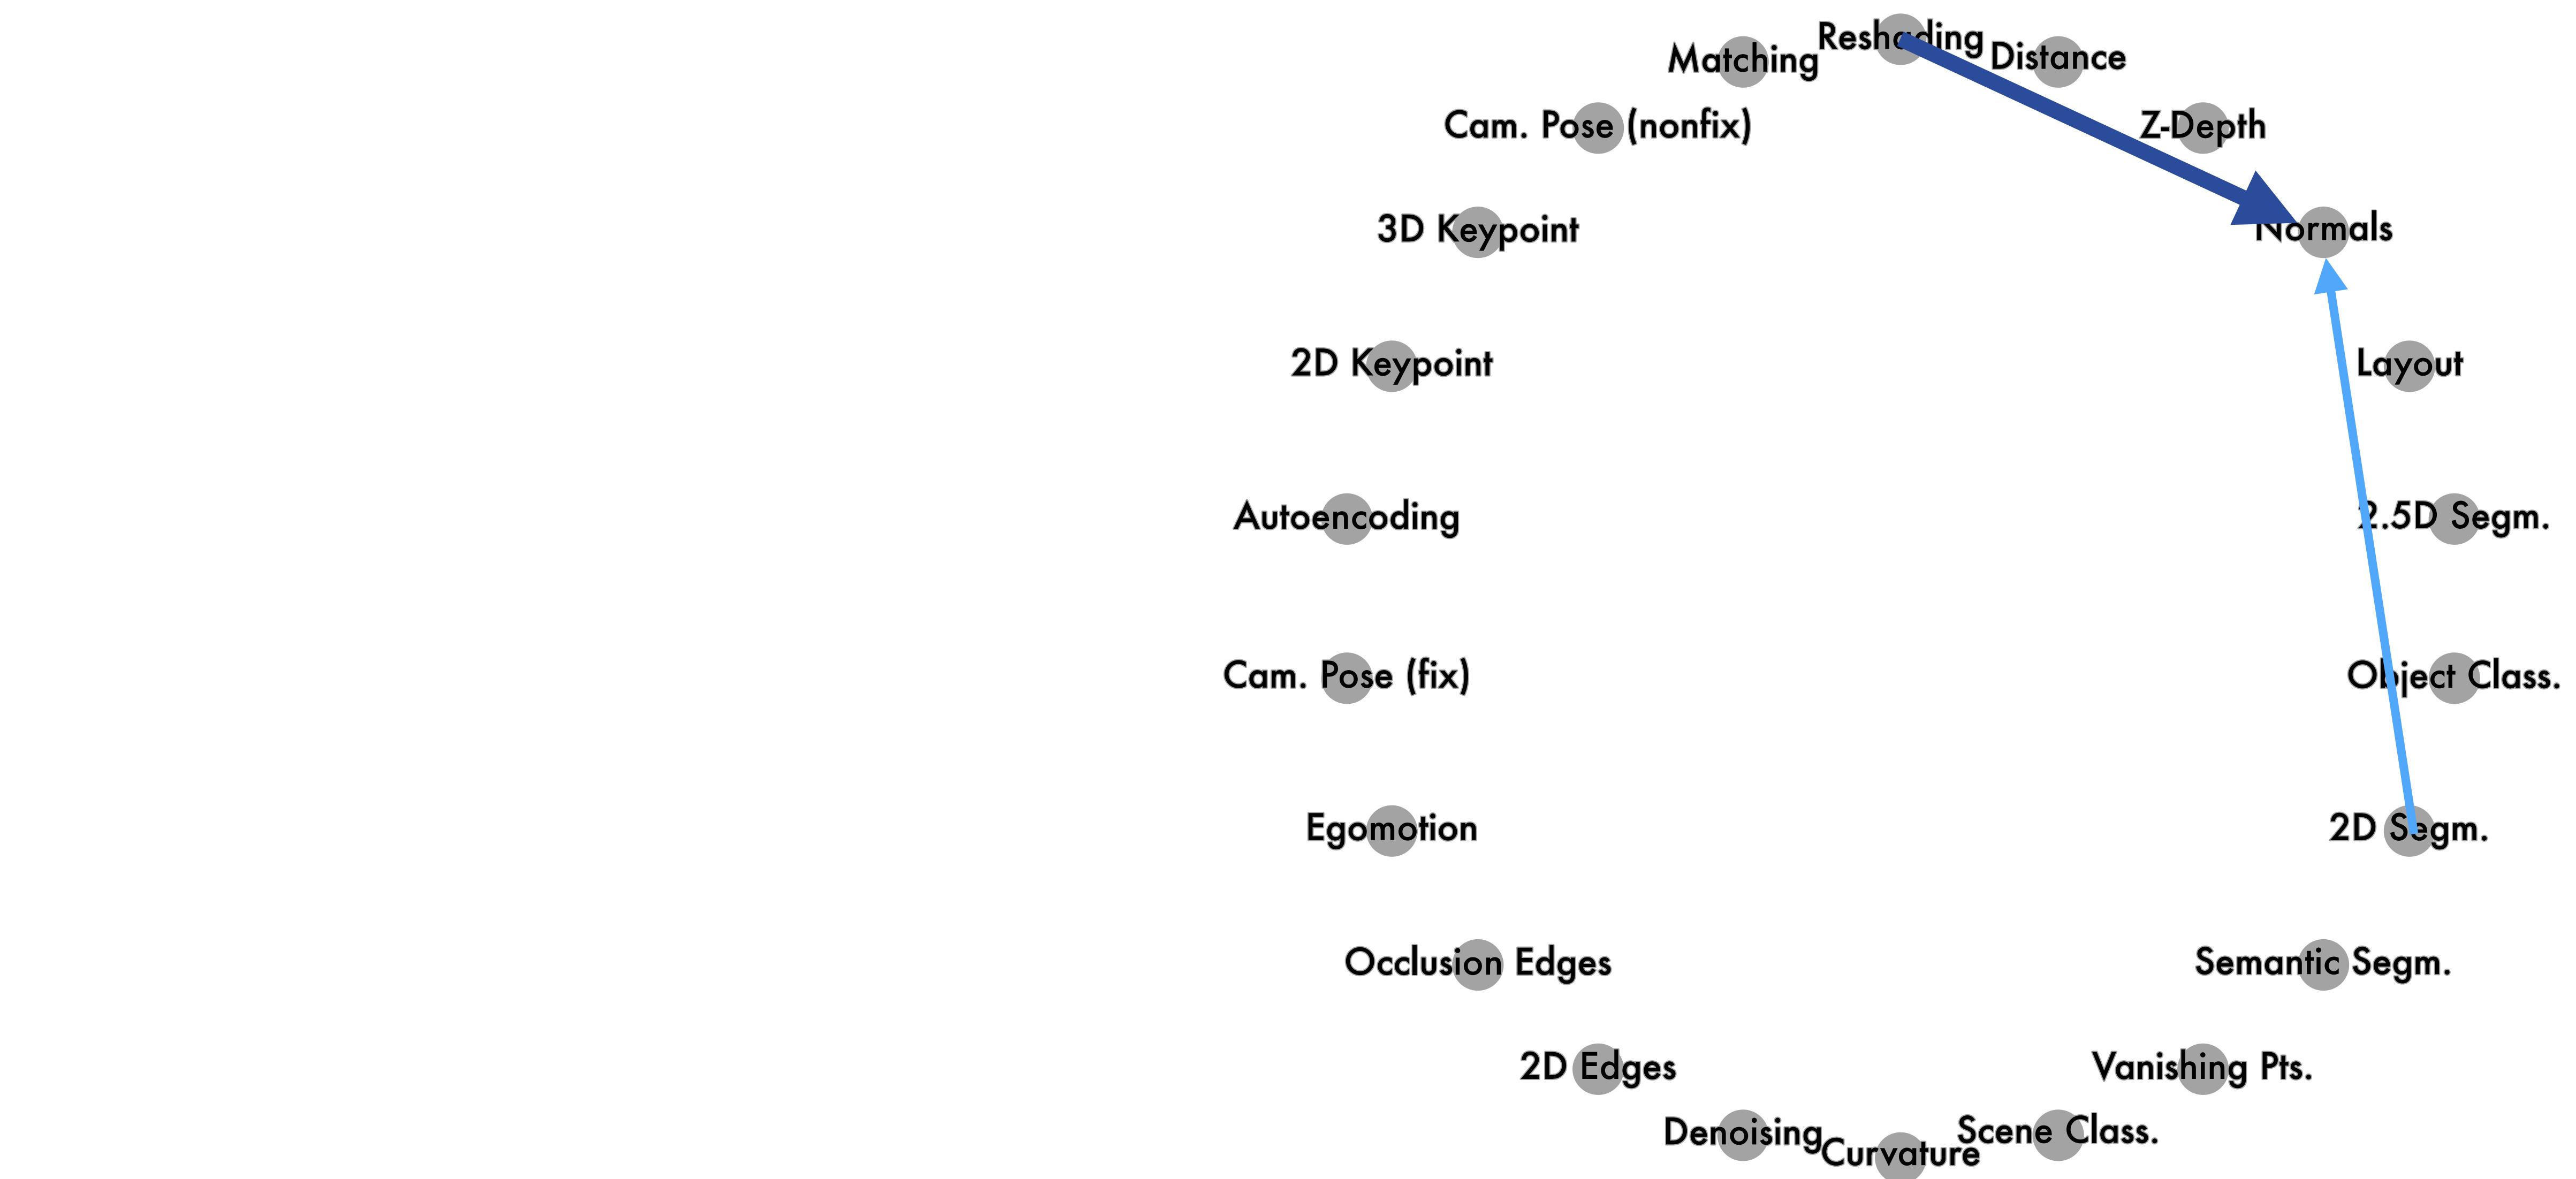


Image



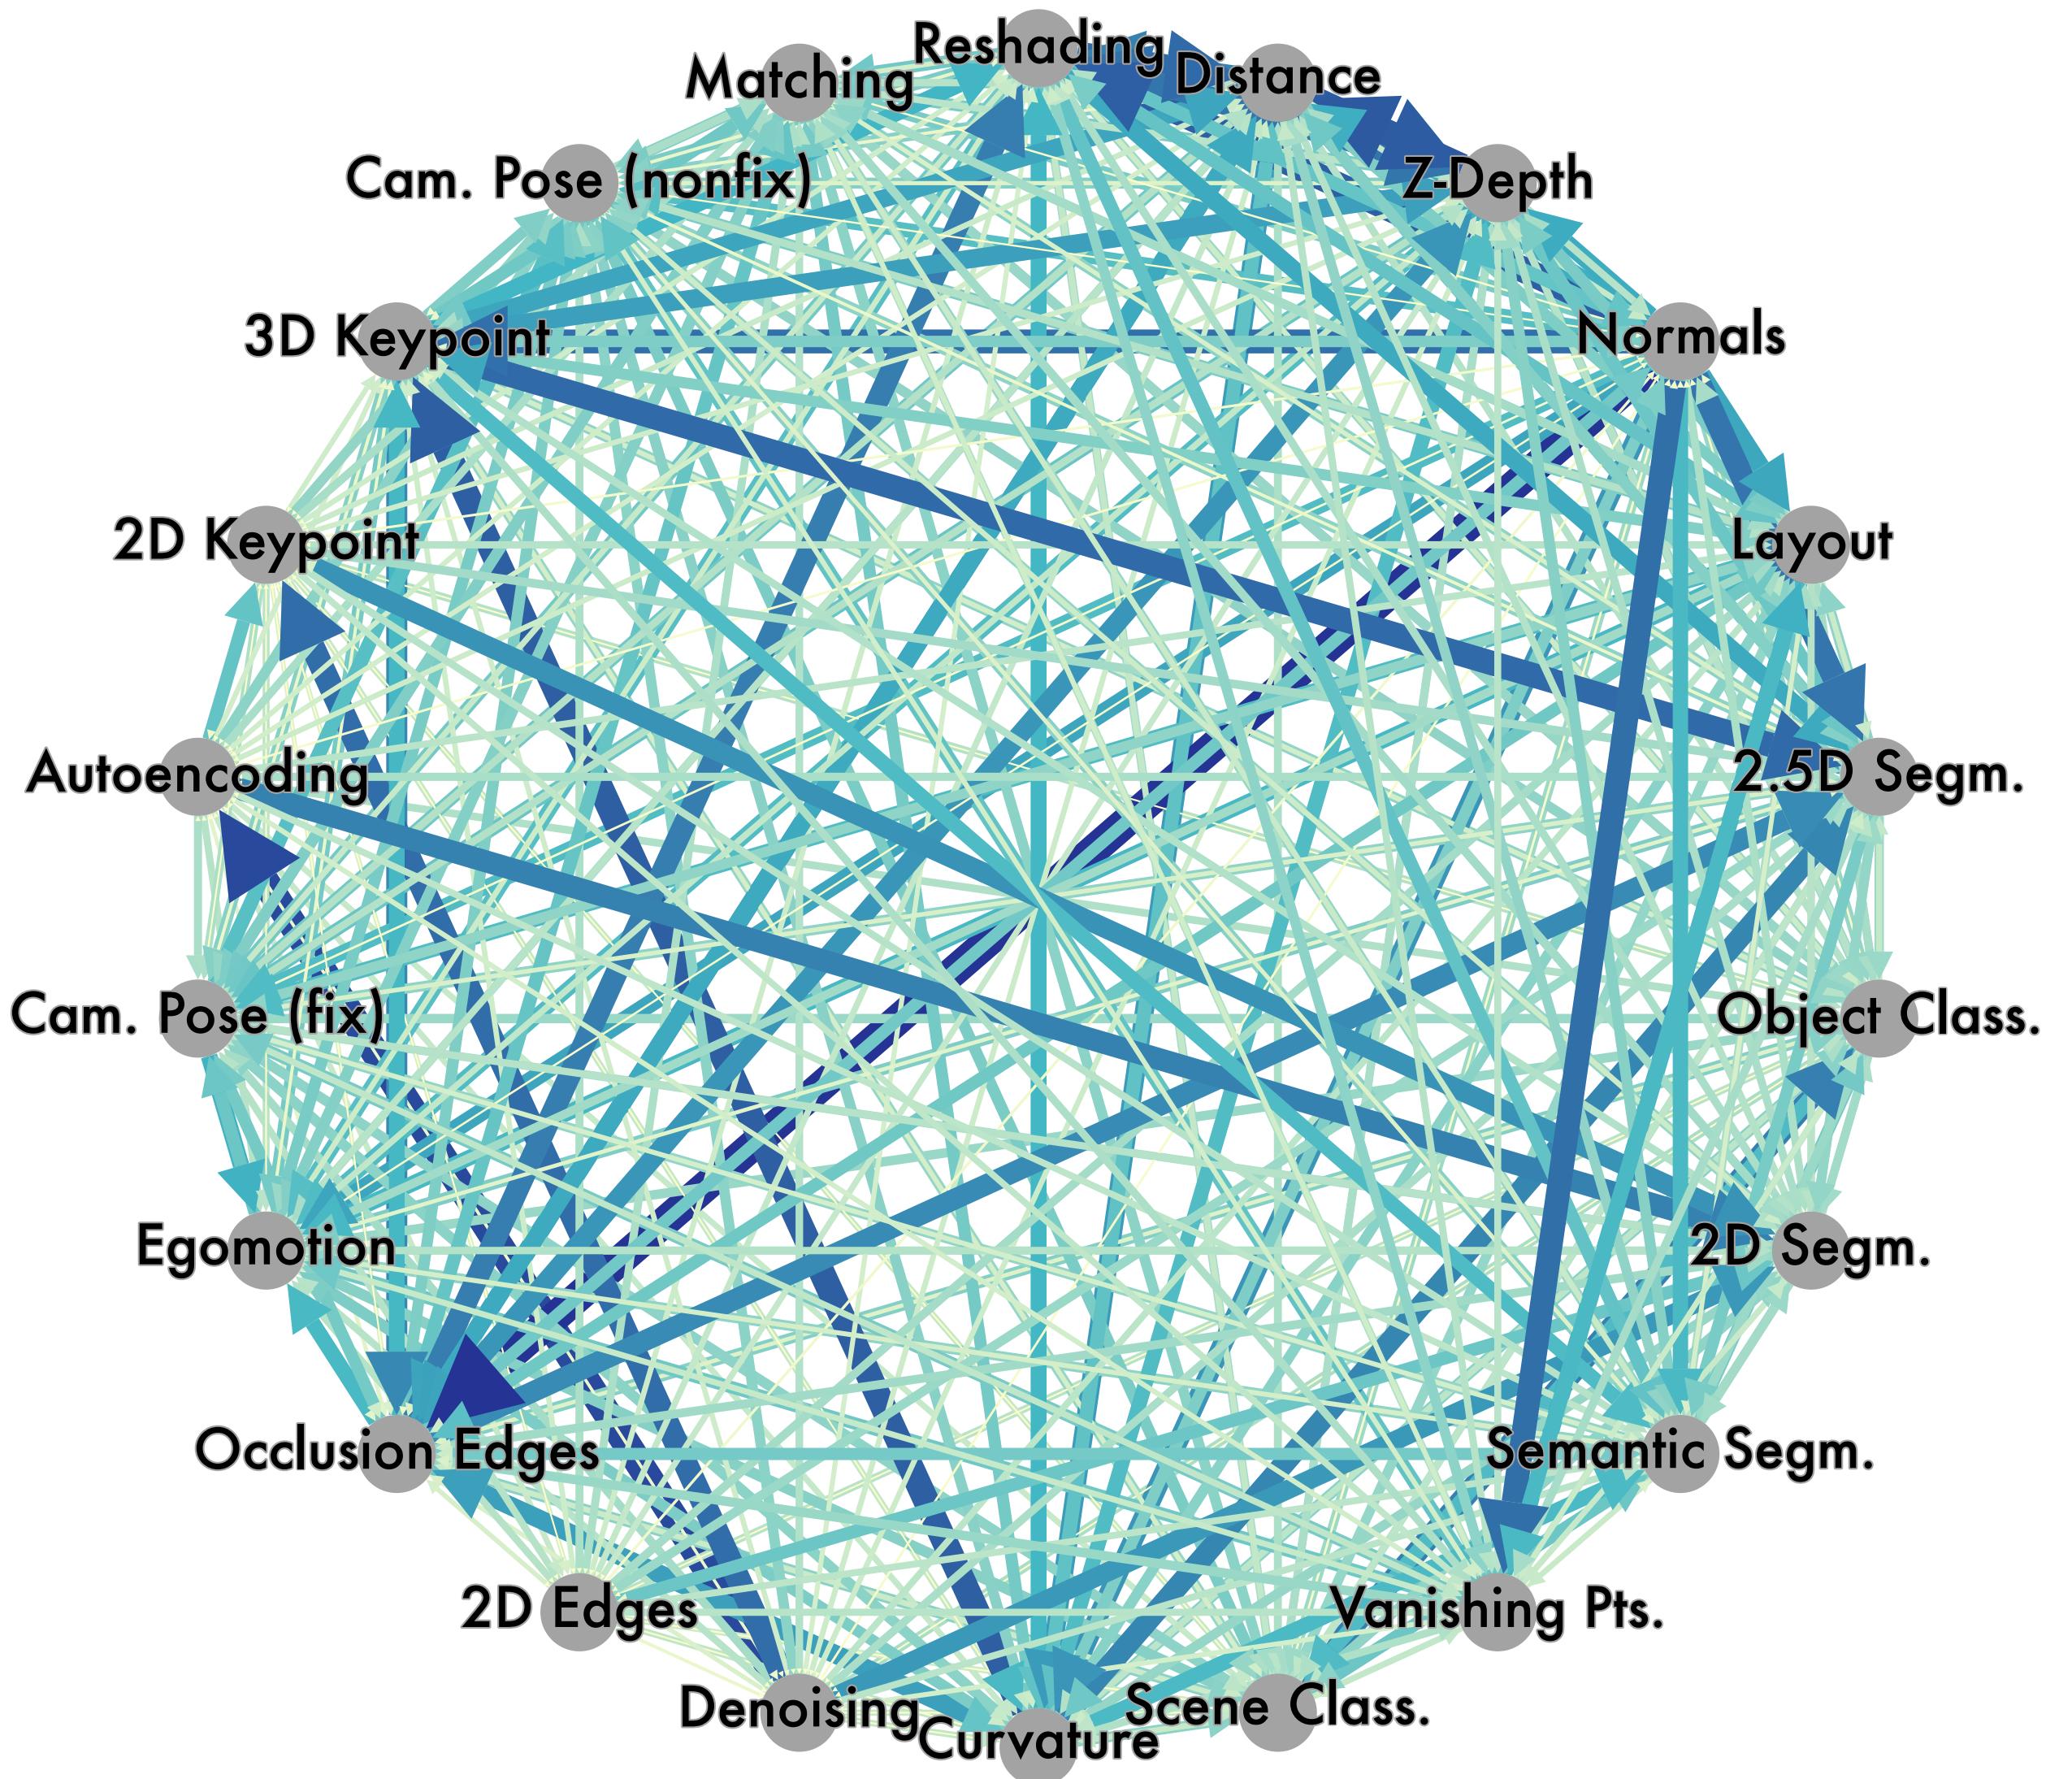
GT Normals

Scratch
(2% data)Task-Specific
Network
(100% data)From Segmentation
(2% data)From Reshading
(2% data)



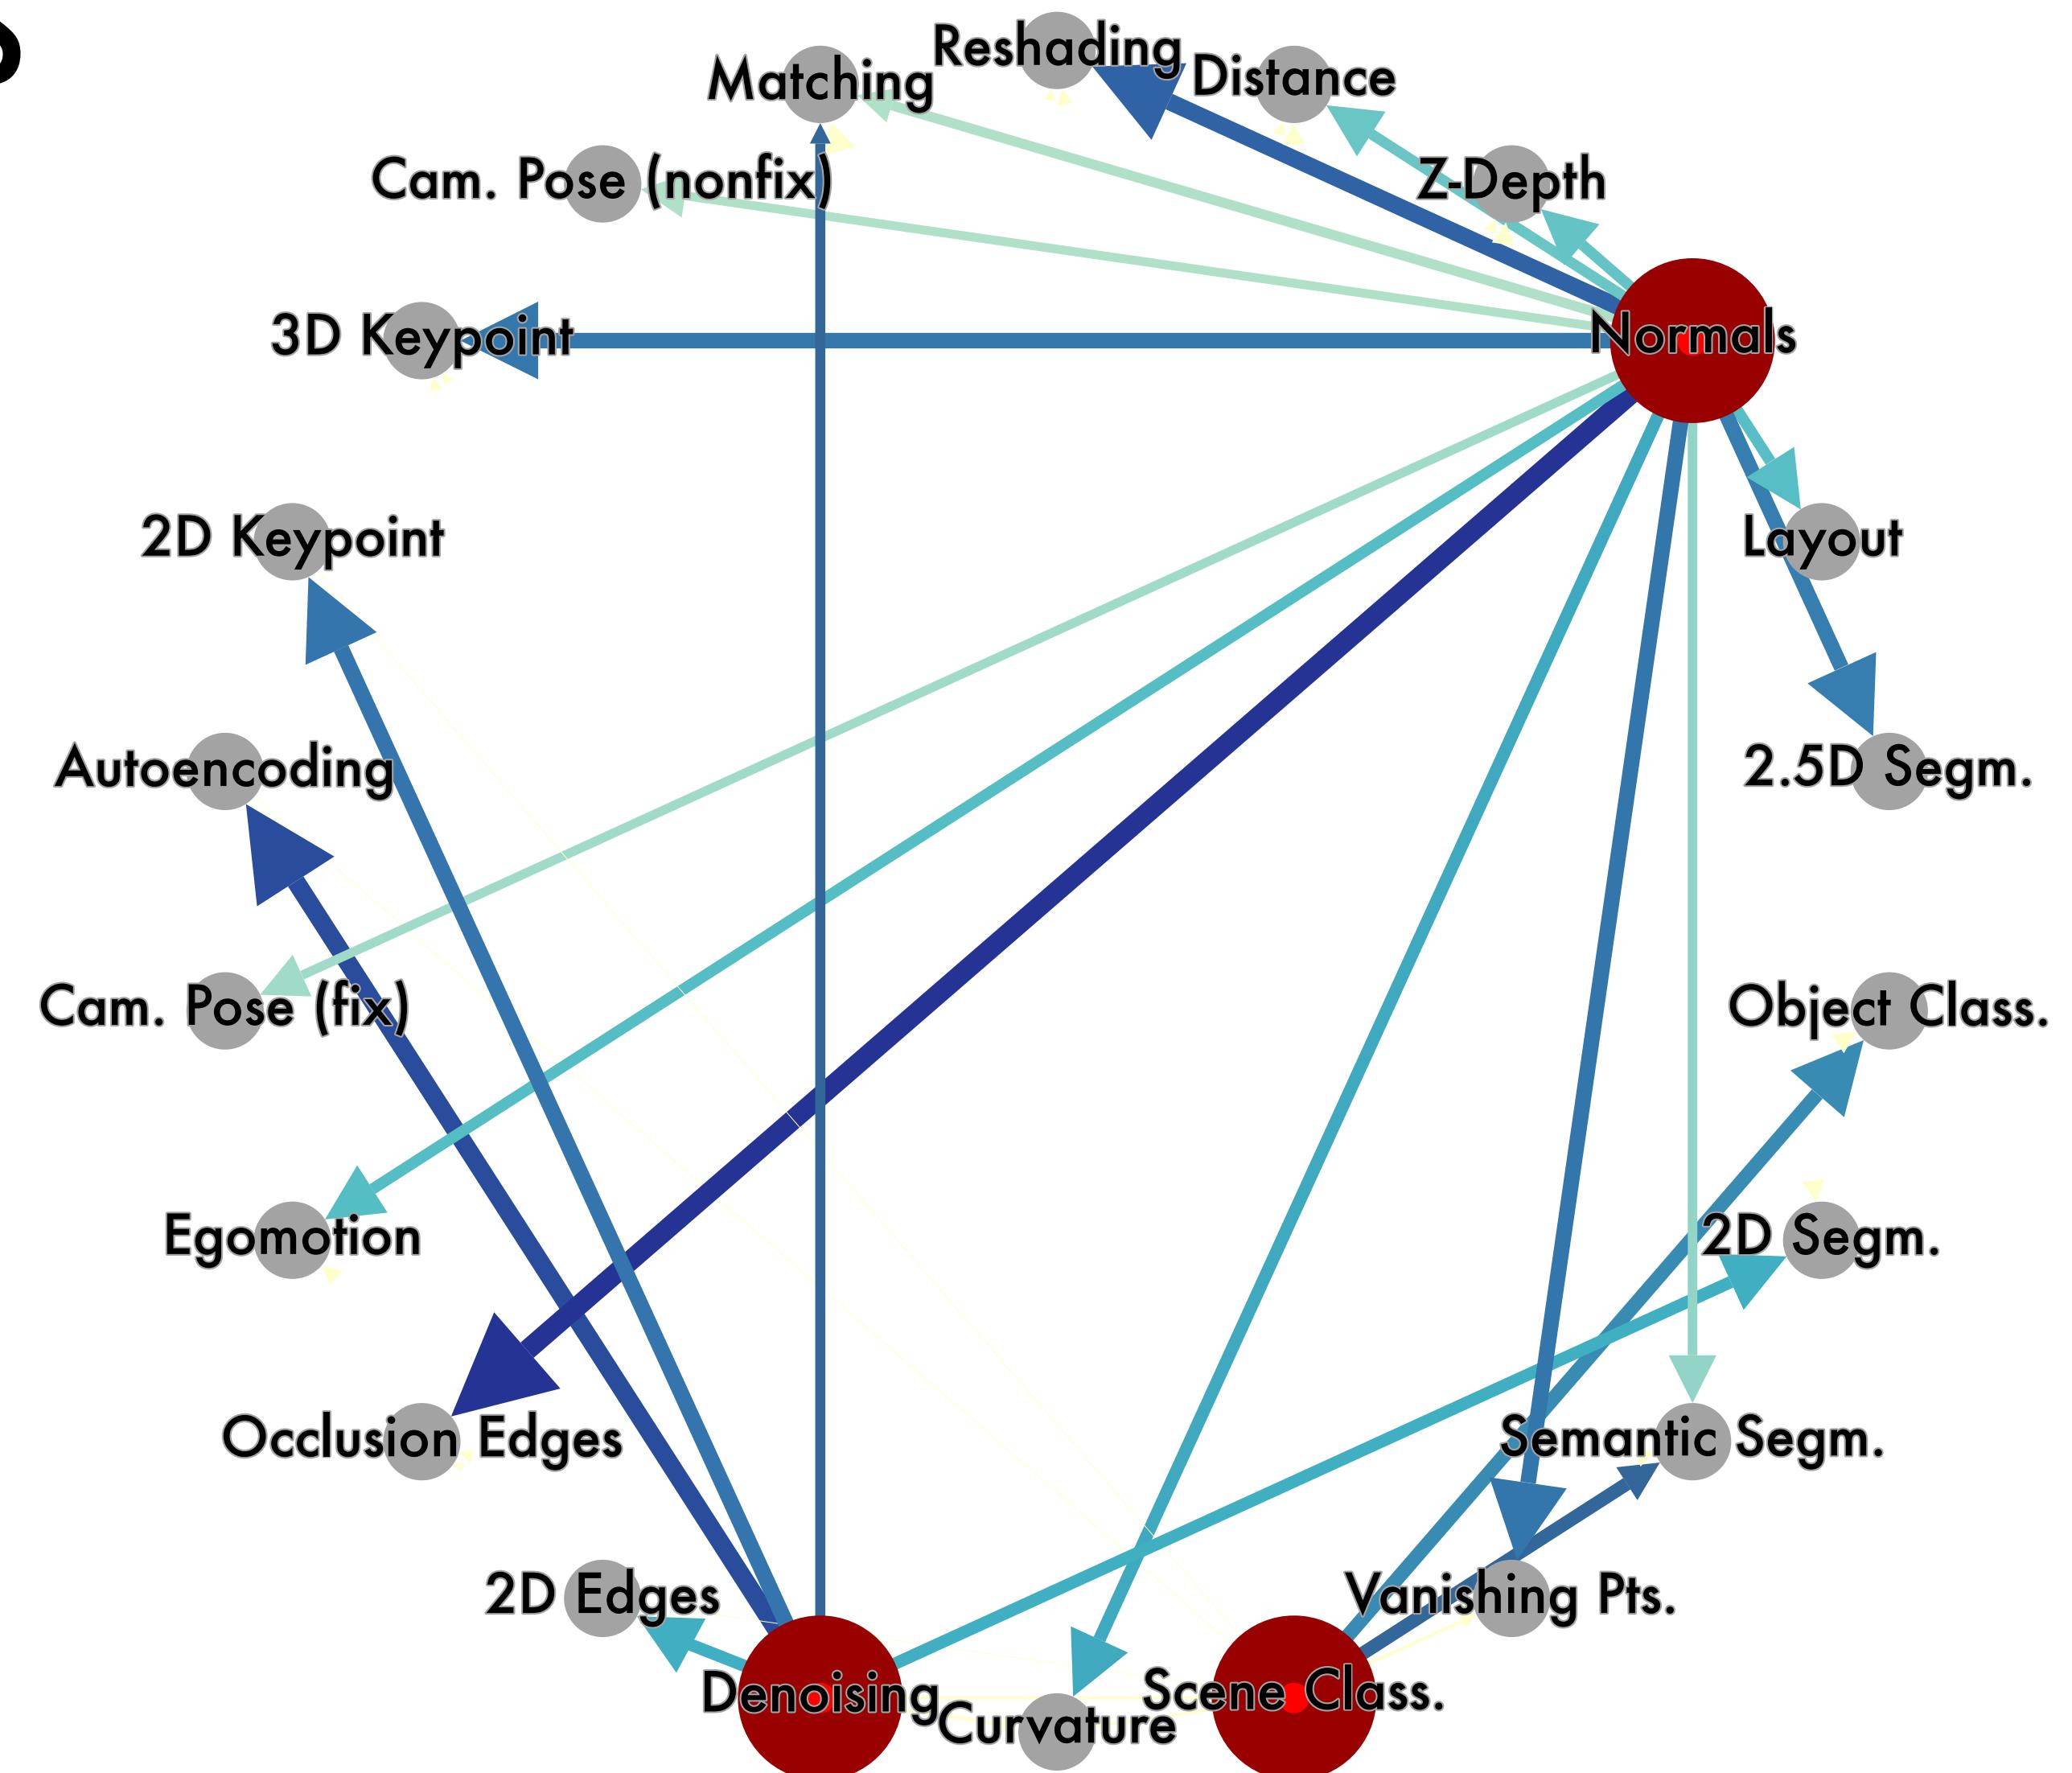
Dense Task Relationships

- Global window to redundancies



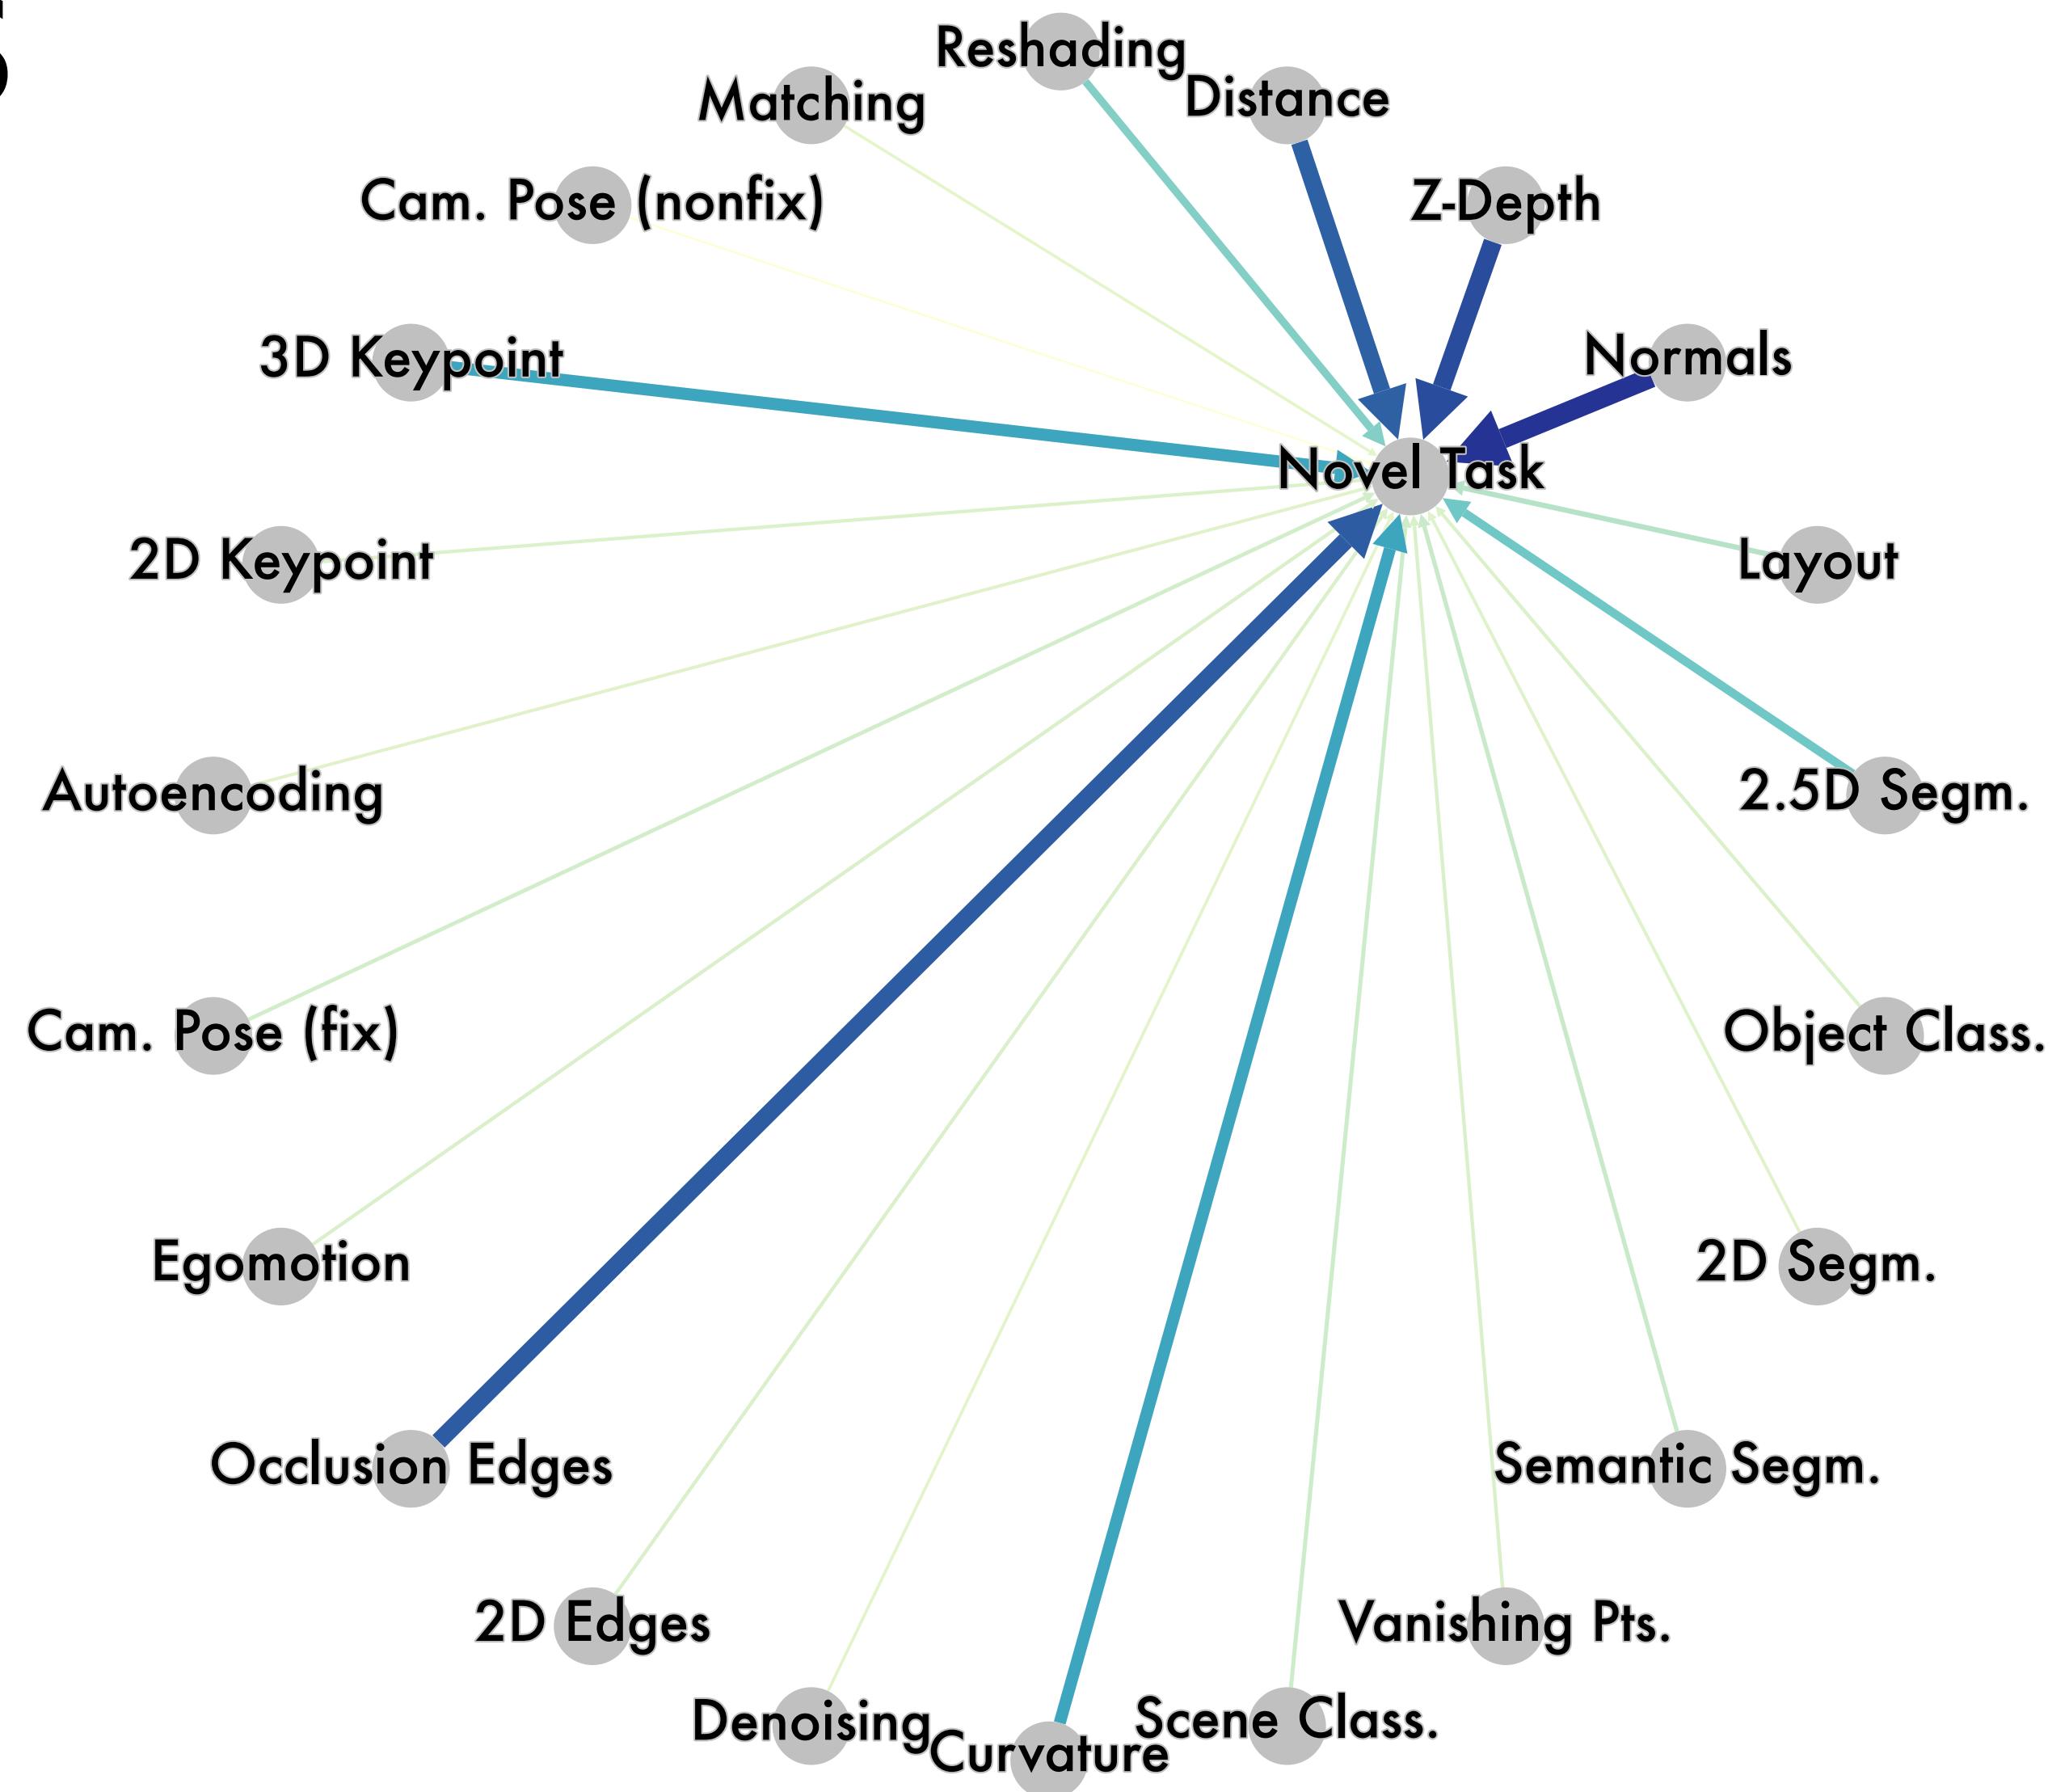
Dense Task Relationships

- Global window to redundancies
- Devise a policy to solve a set of tasks
 - in concert
 - recycle supervision



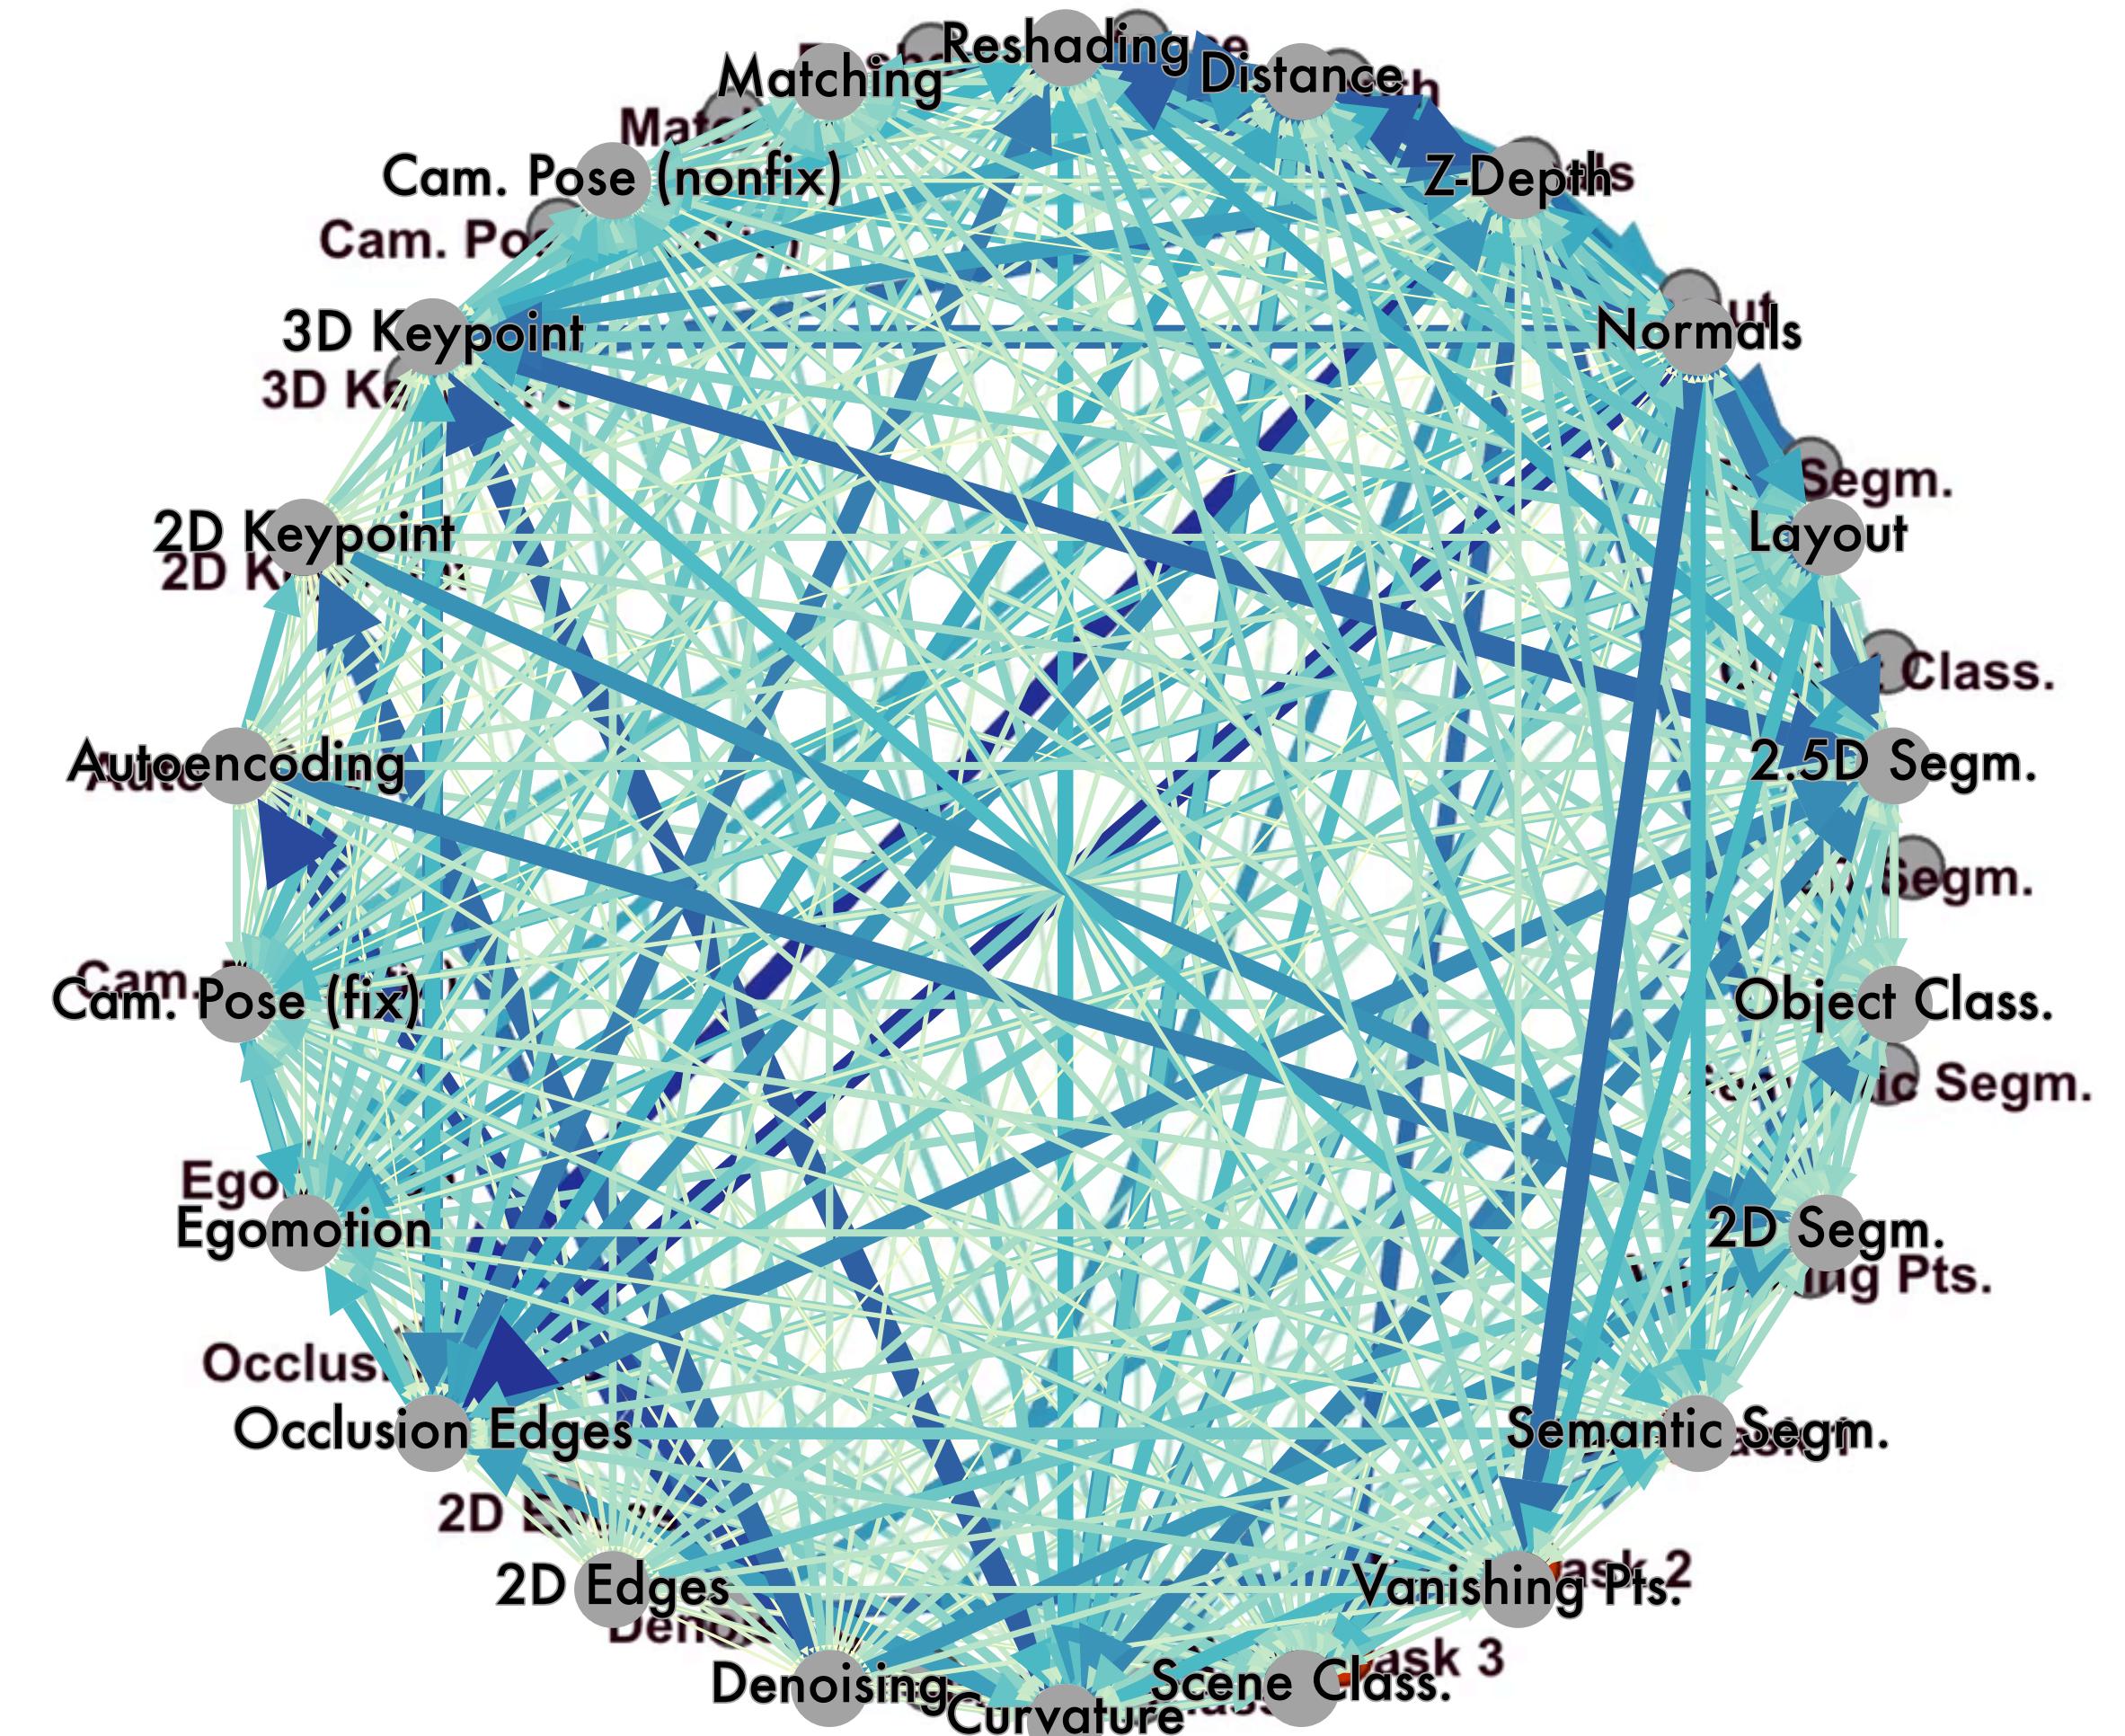
Dense Task Relationships

- Global window to redundancies
- Devise a policy to solve a set of tasks
 - in concert
 - recycle supervision
- Solve a novel task with little data
 - by “insertion” into this structure

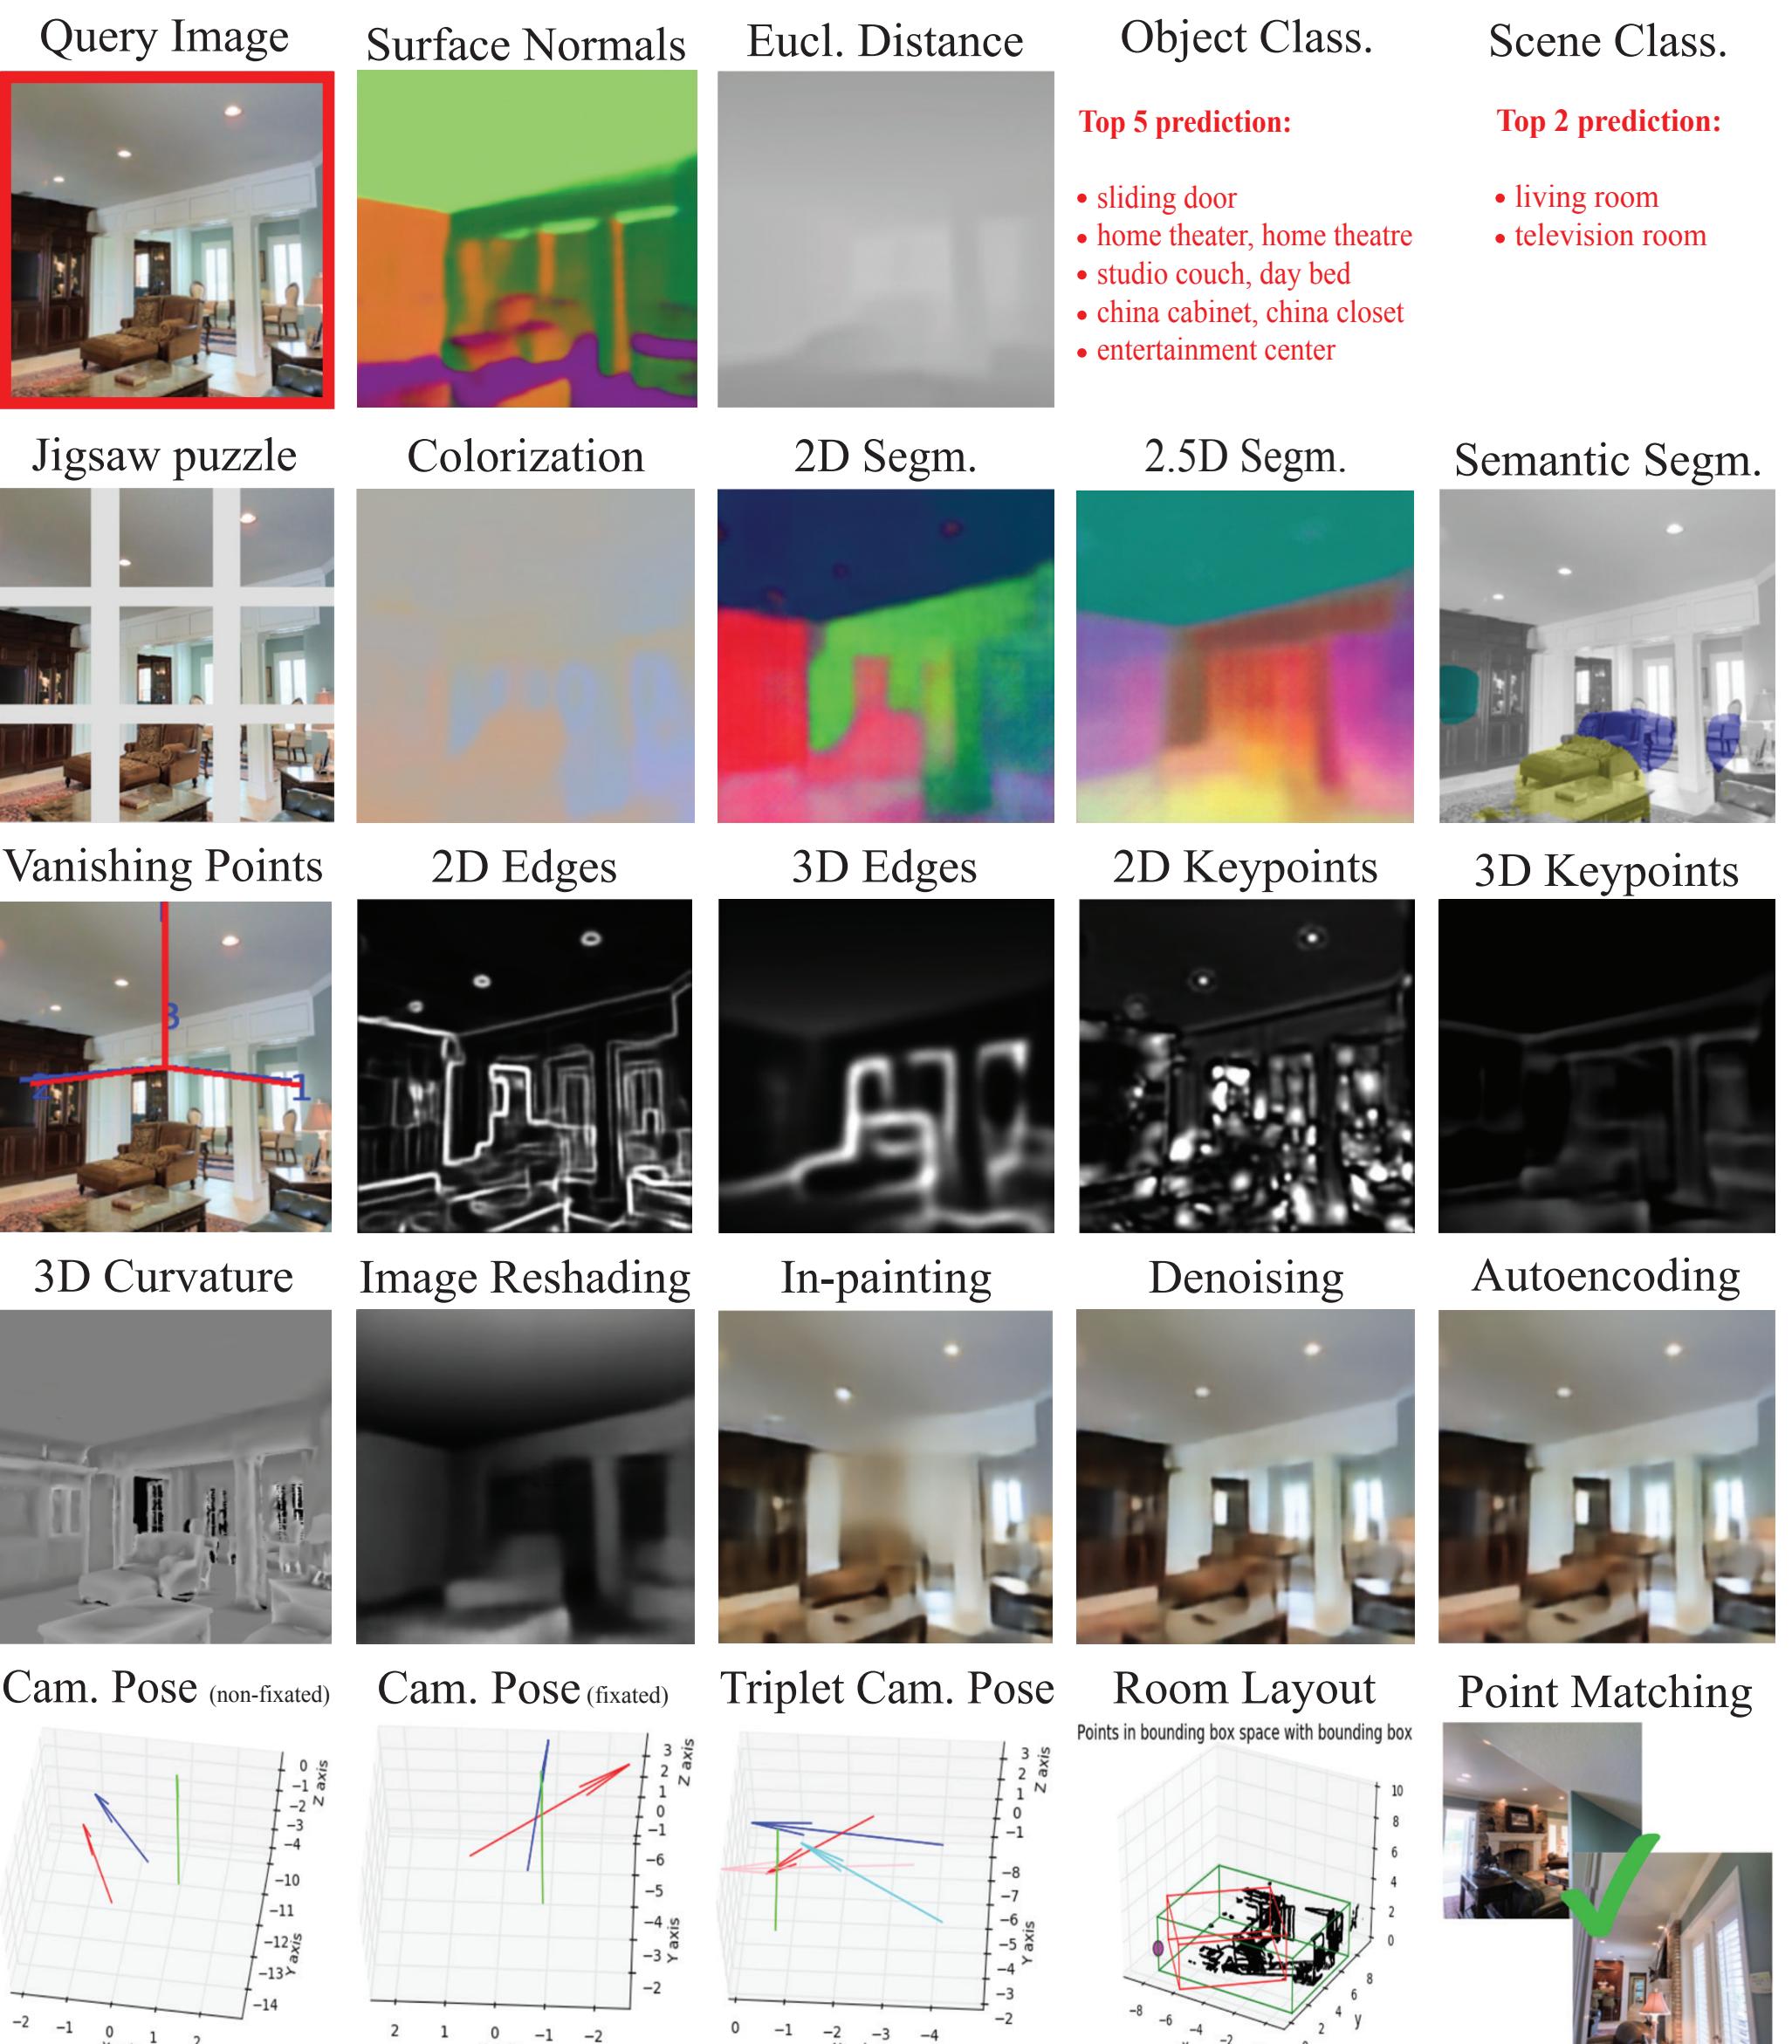


Taskonomy!

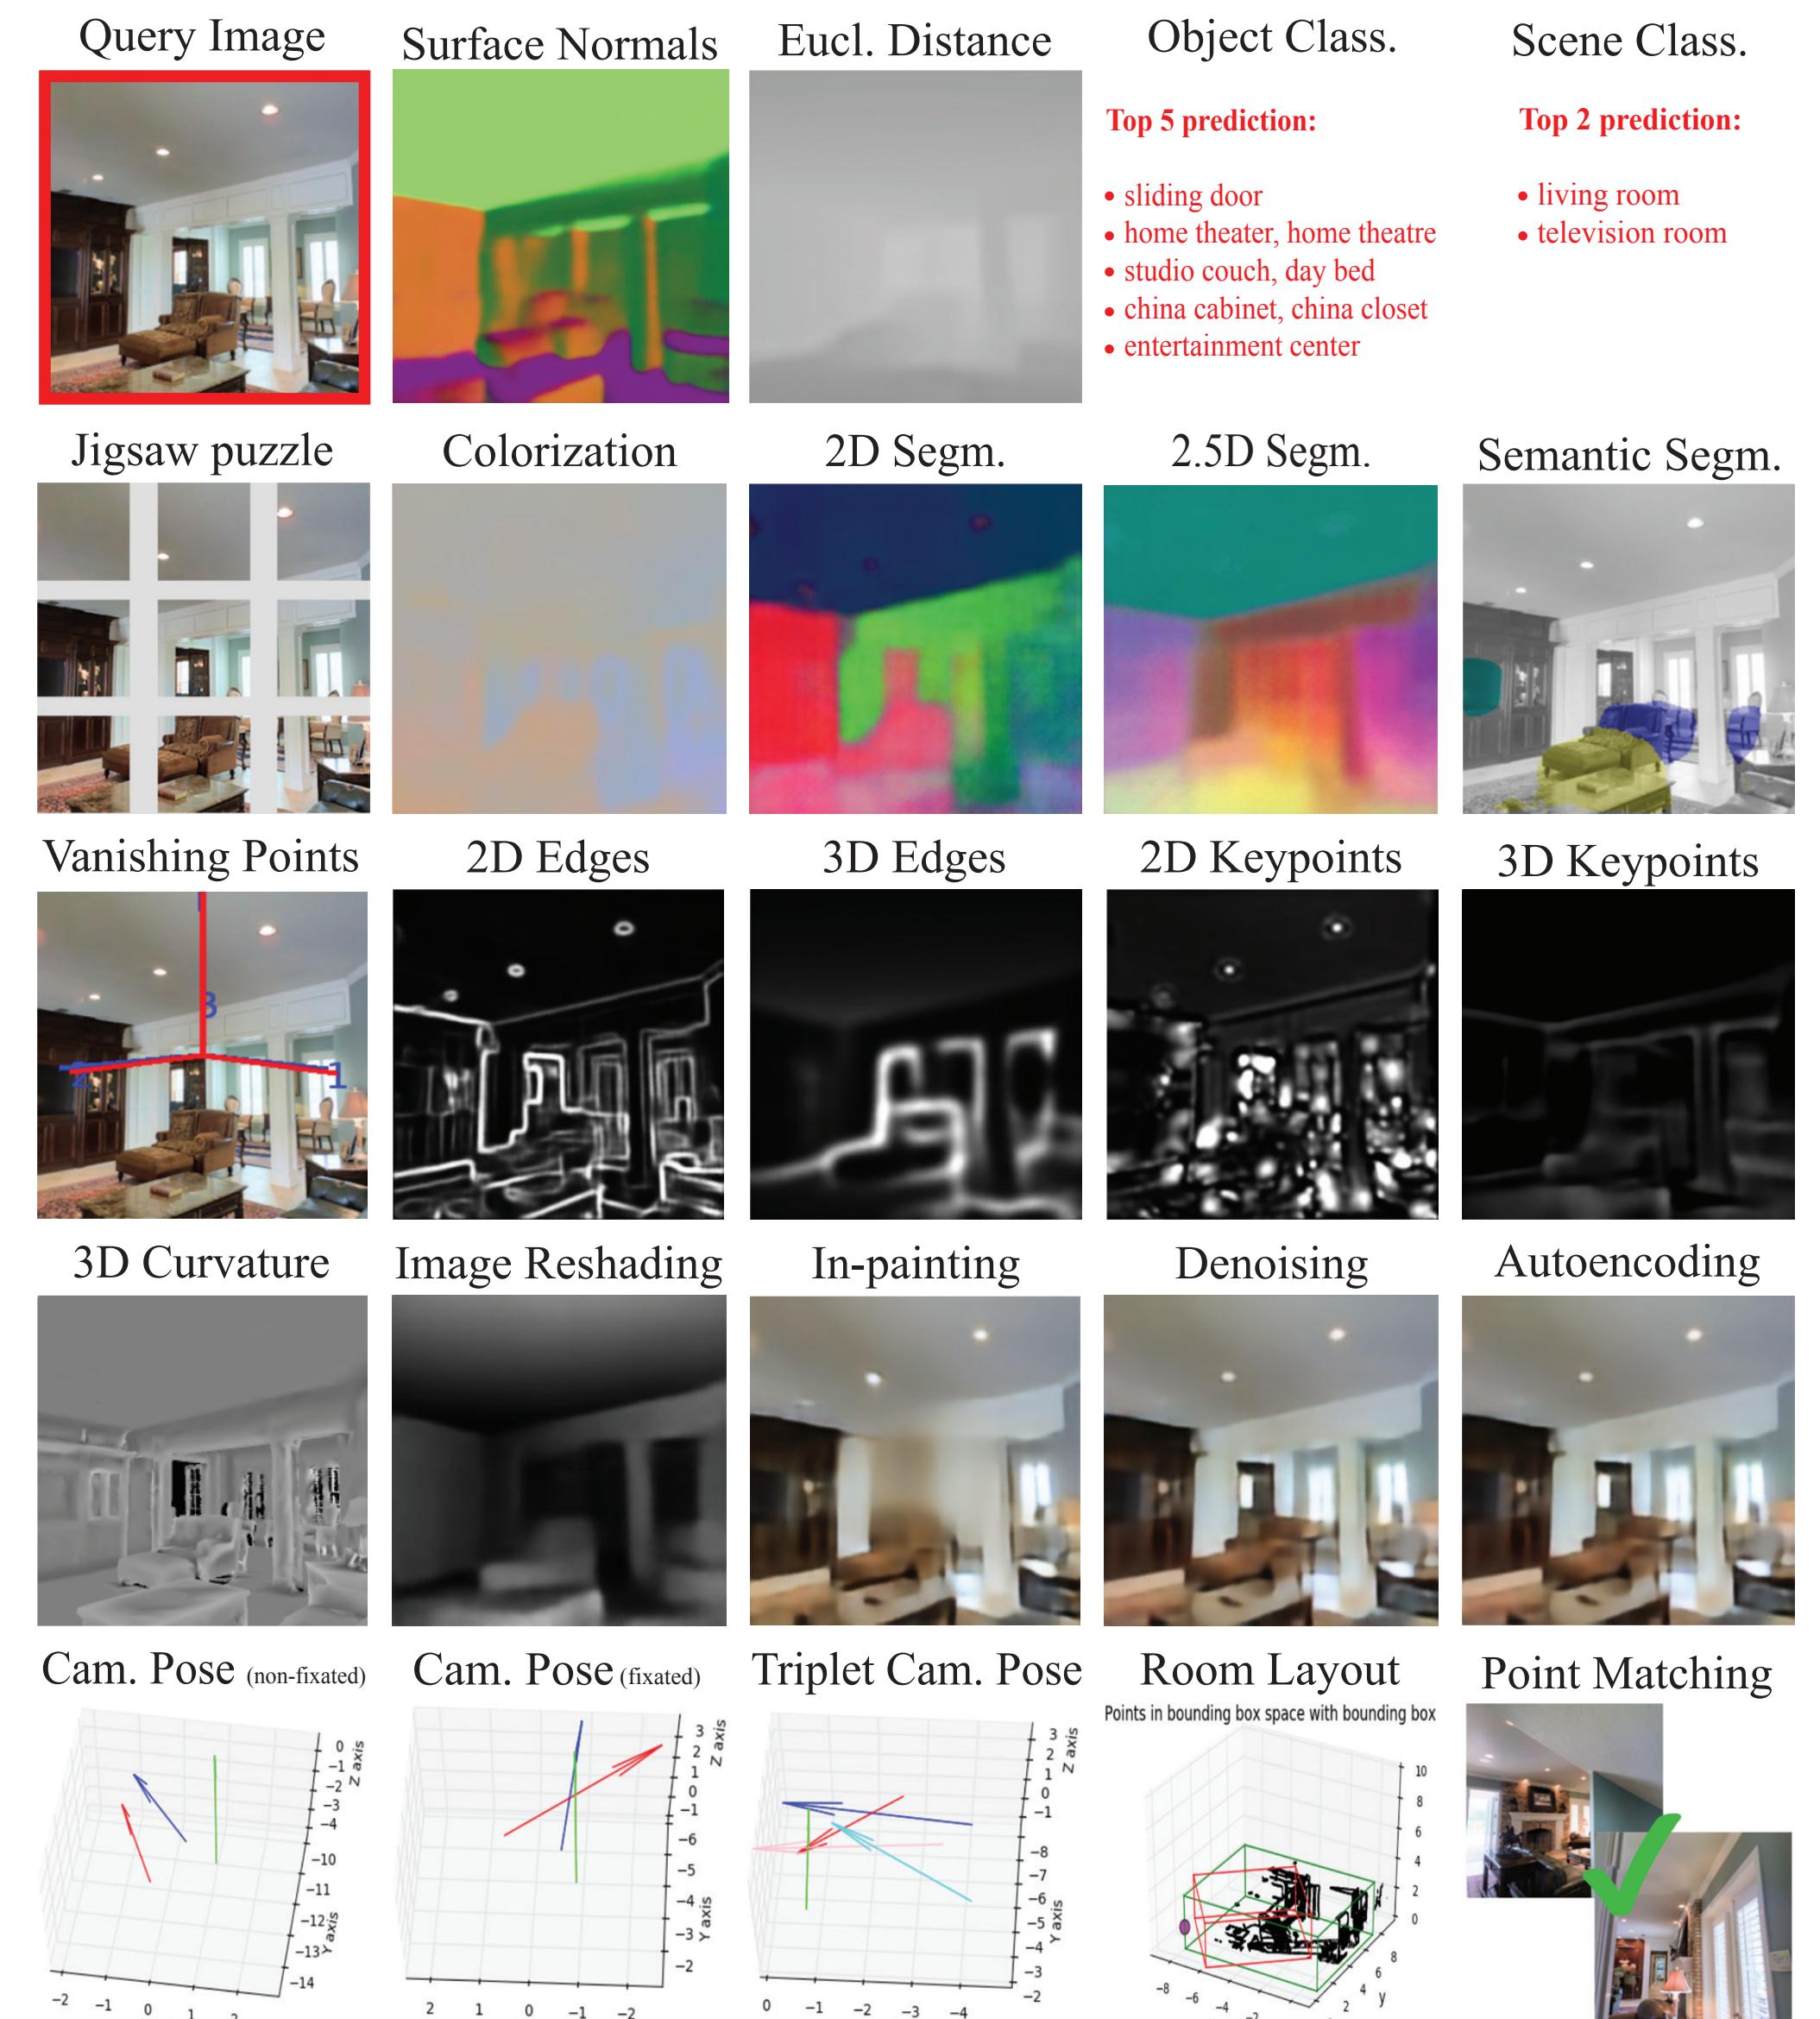
- A fully computational method for quantifying task relationships
- Extracting a structure out of them
- Unified model for transfer learning
- **task taxonomy \approx taskonomy**



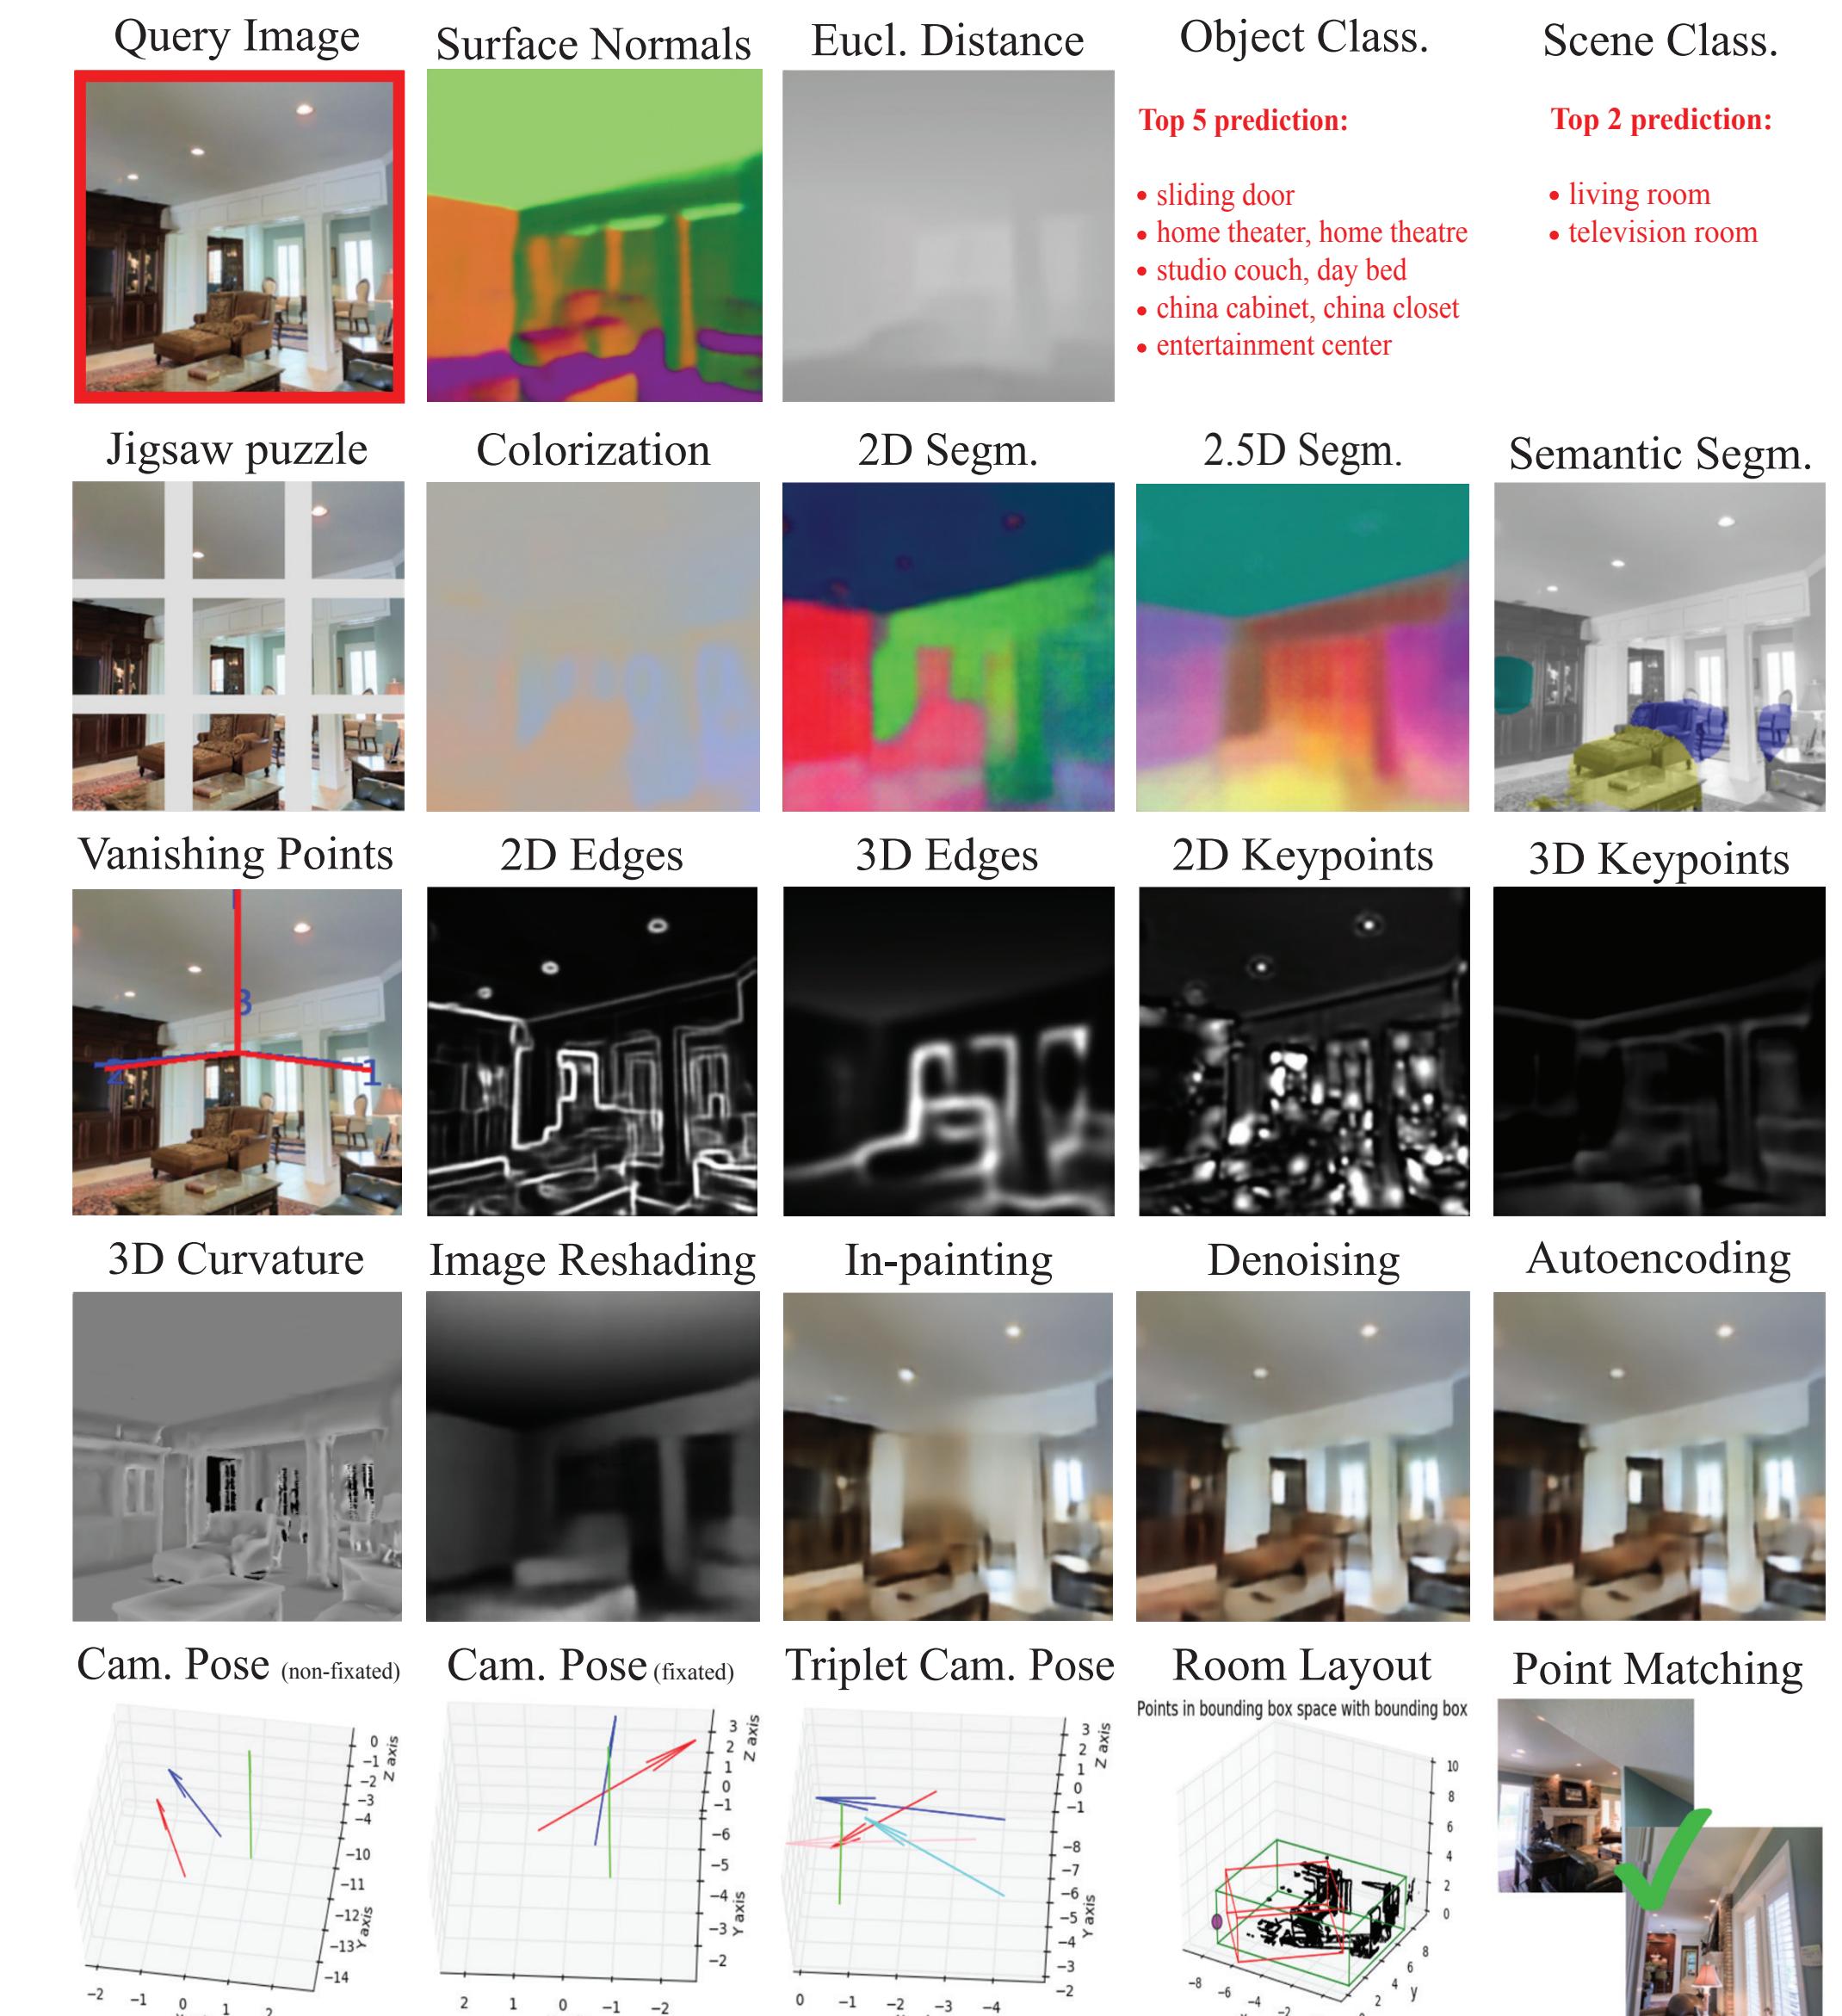
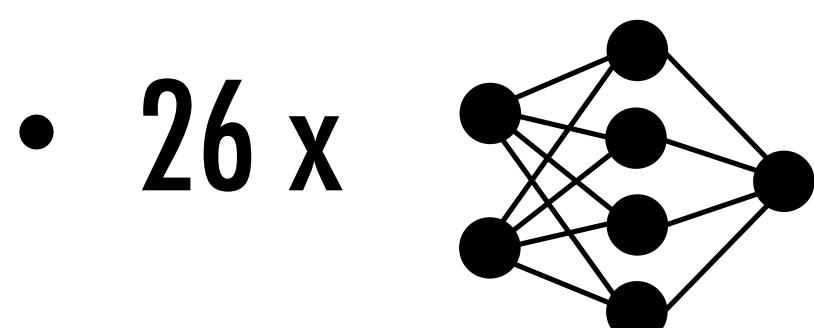
- **Task Bank**
- **26 Semantic, 2D, 3D, and tasks**



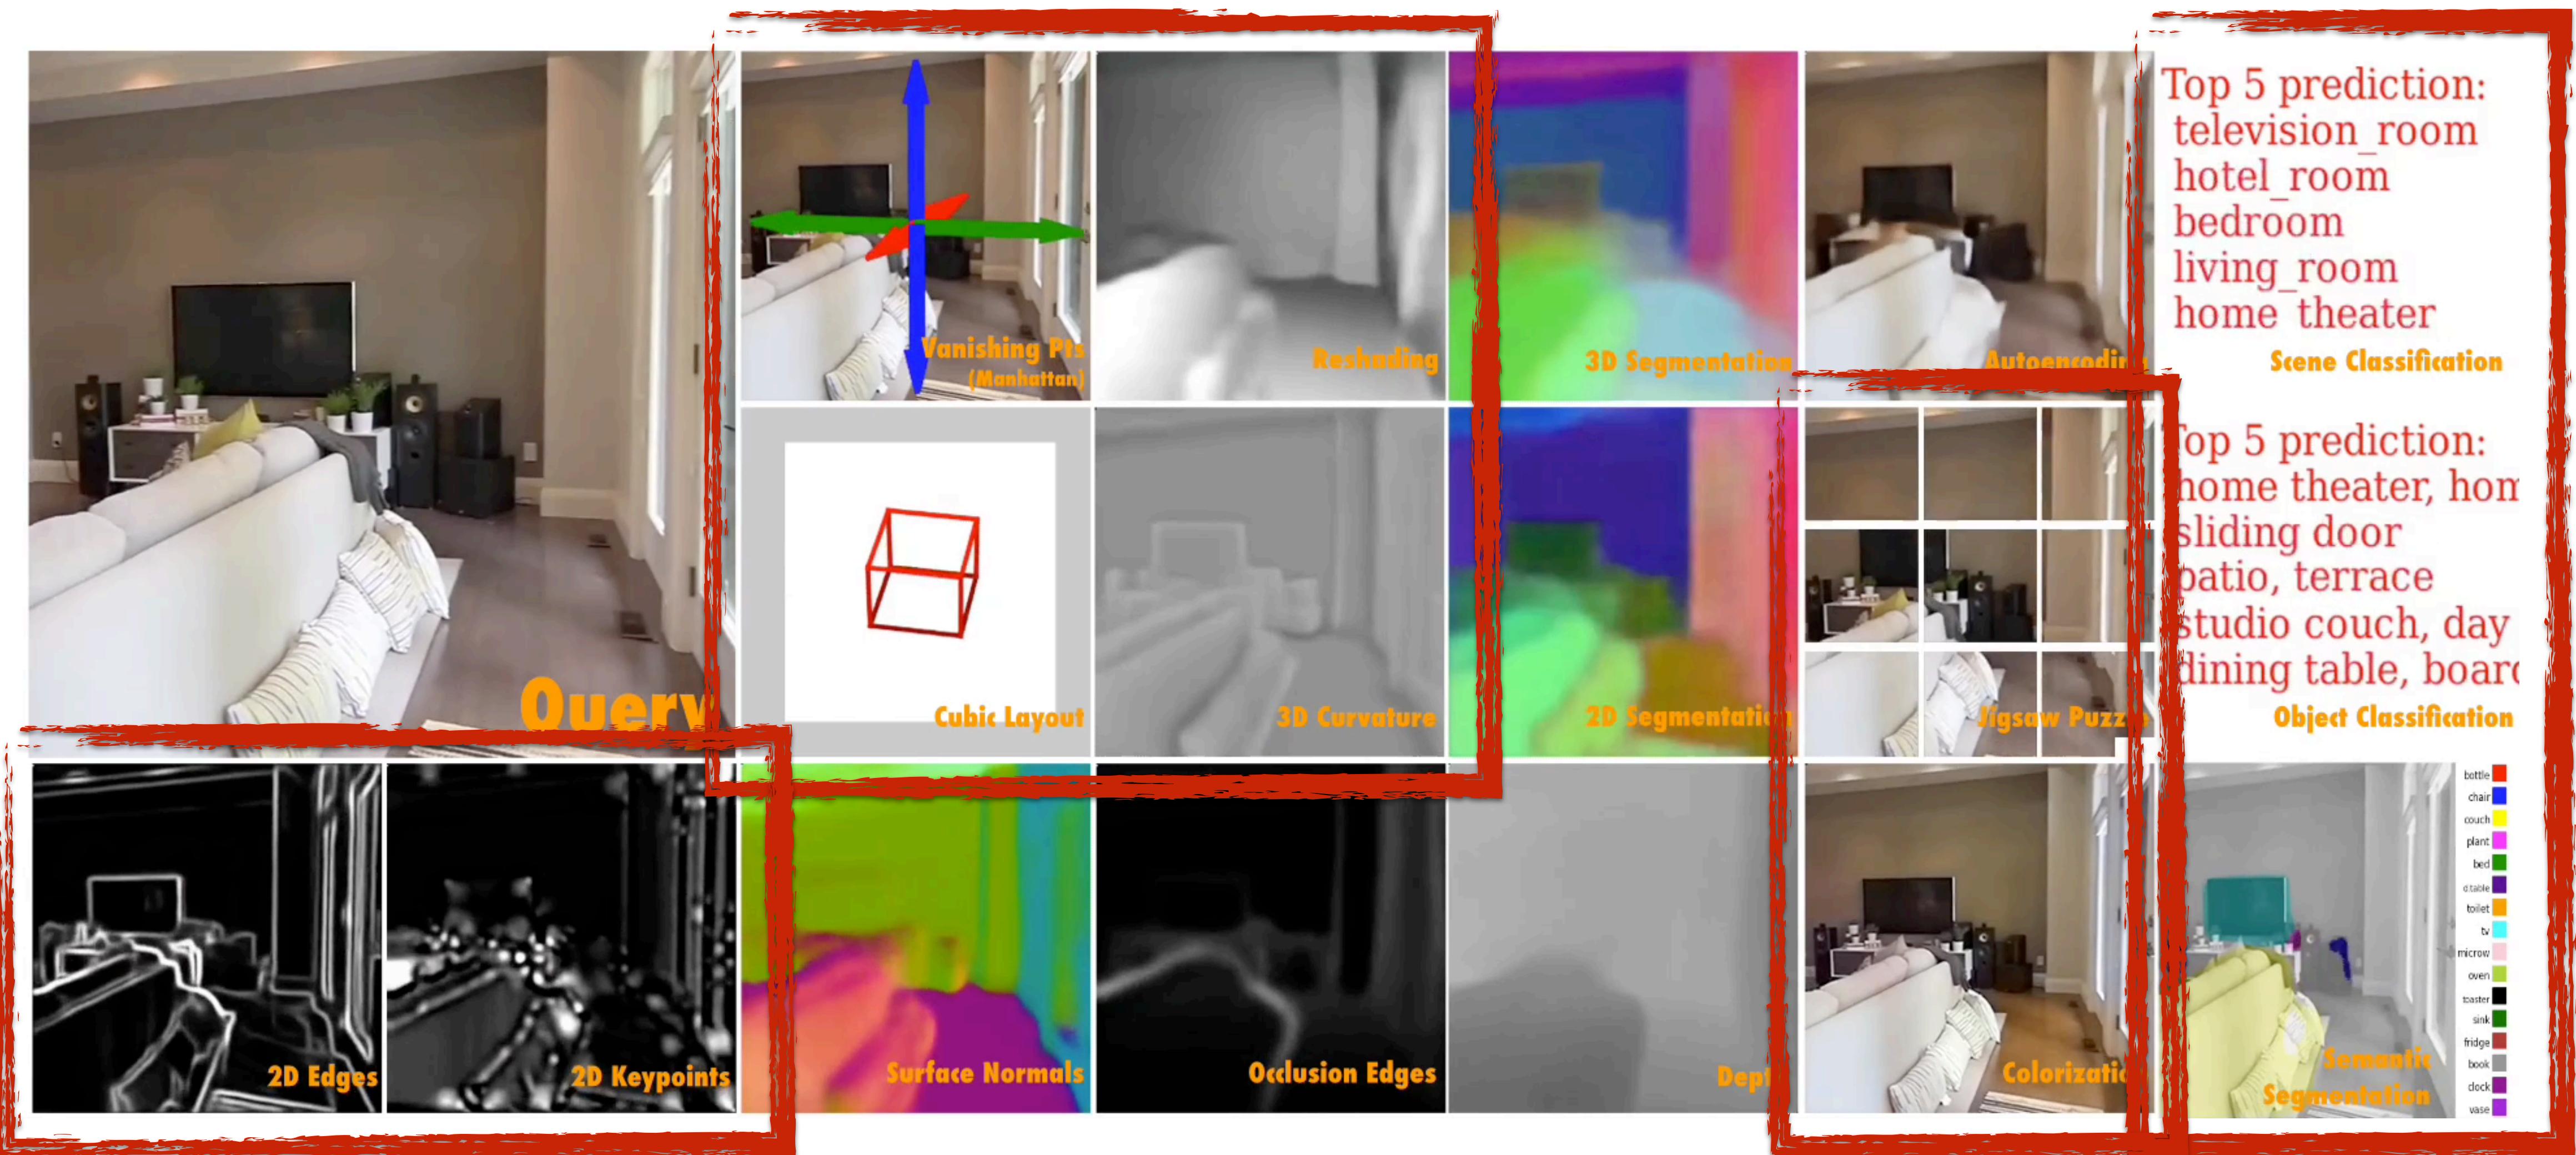
- **Task Bank**
 - 26 Semantic, 2D, 3D, and tasks
- **Dataset**
 - 4 million real images
 - Each image has the GT label for all tasks



- **Task Bank**
 - 26 Semantic, 2D, 3D, and tasks
- **Dataset**
 - 4 million real images
 - Each image has the GT label for all tasks
- **Task-Specific Networks**

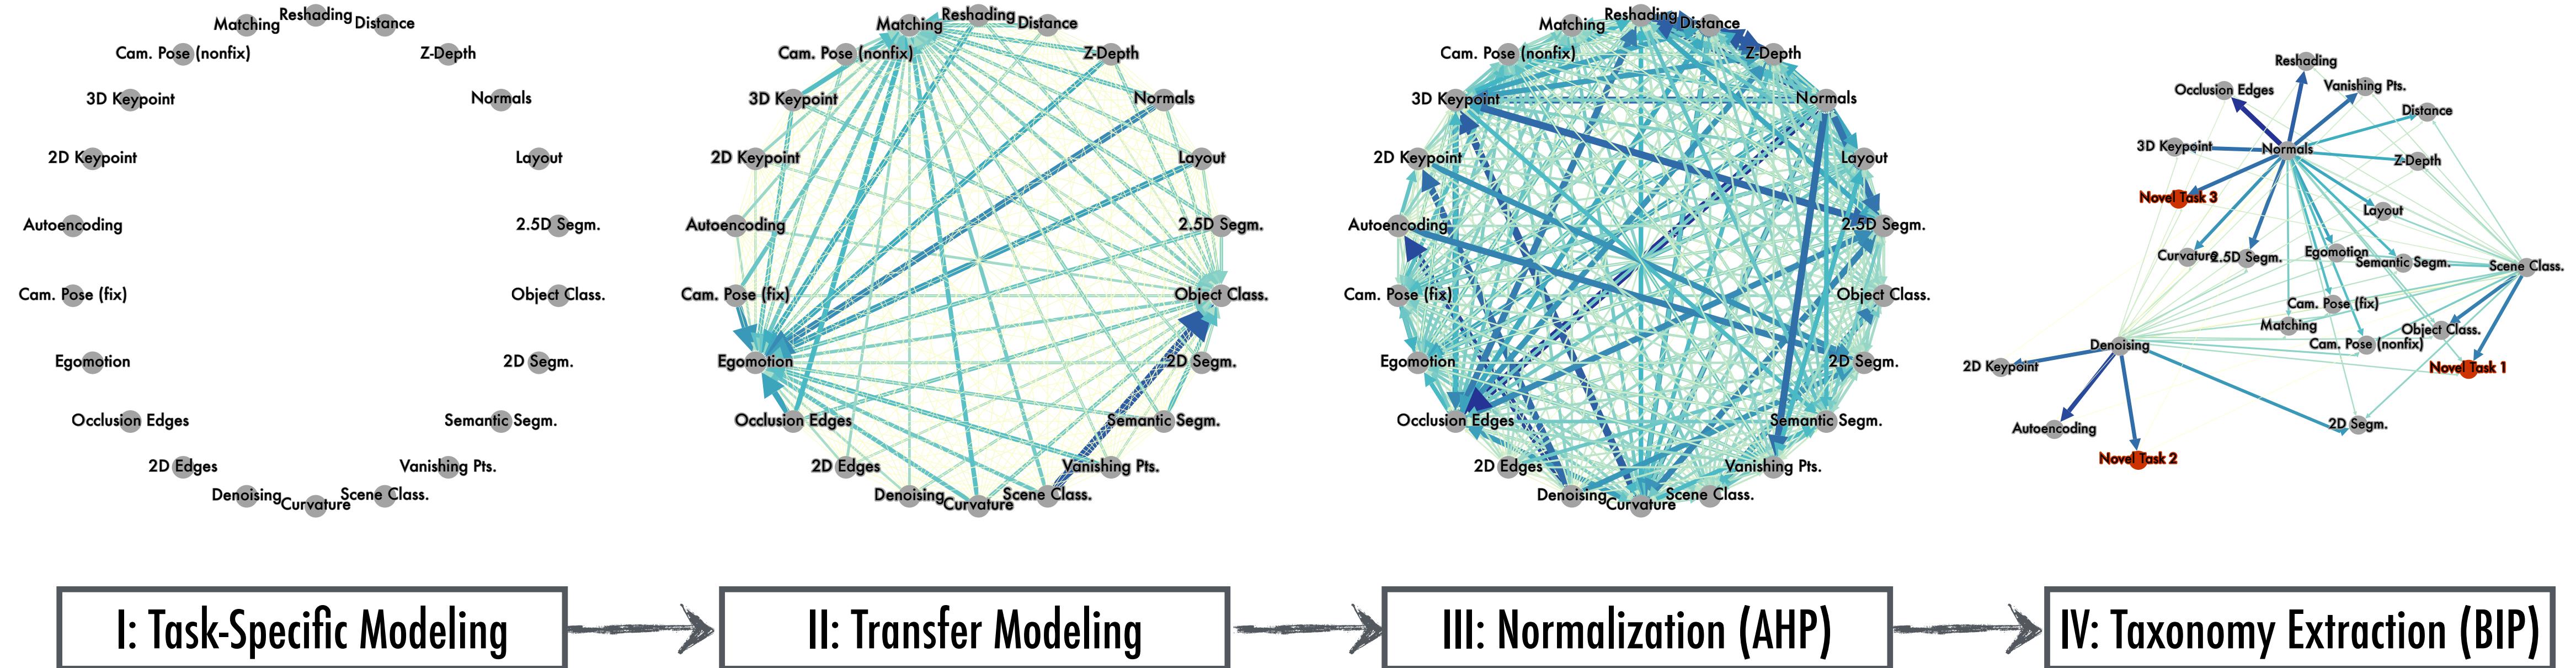


Task-Specific Networks

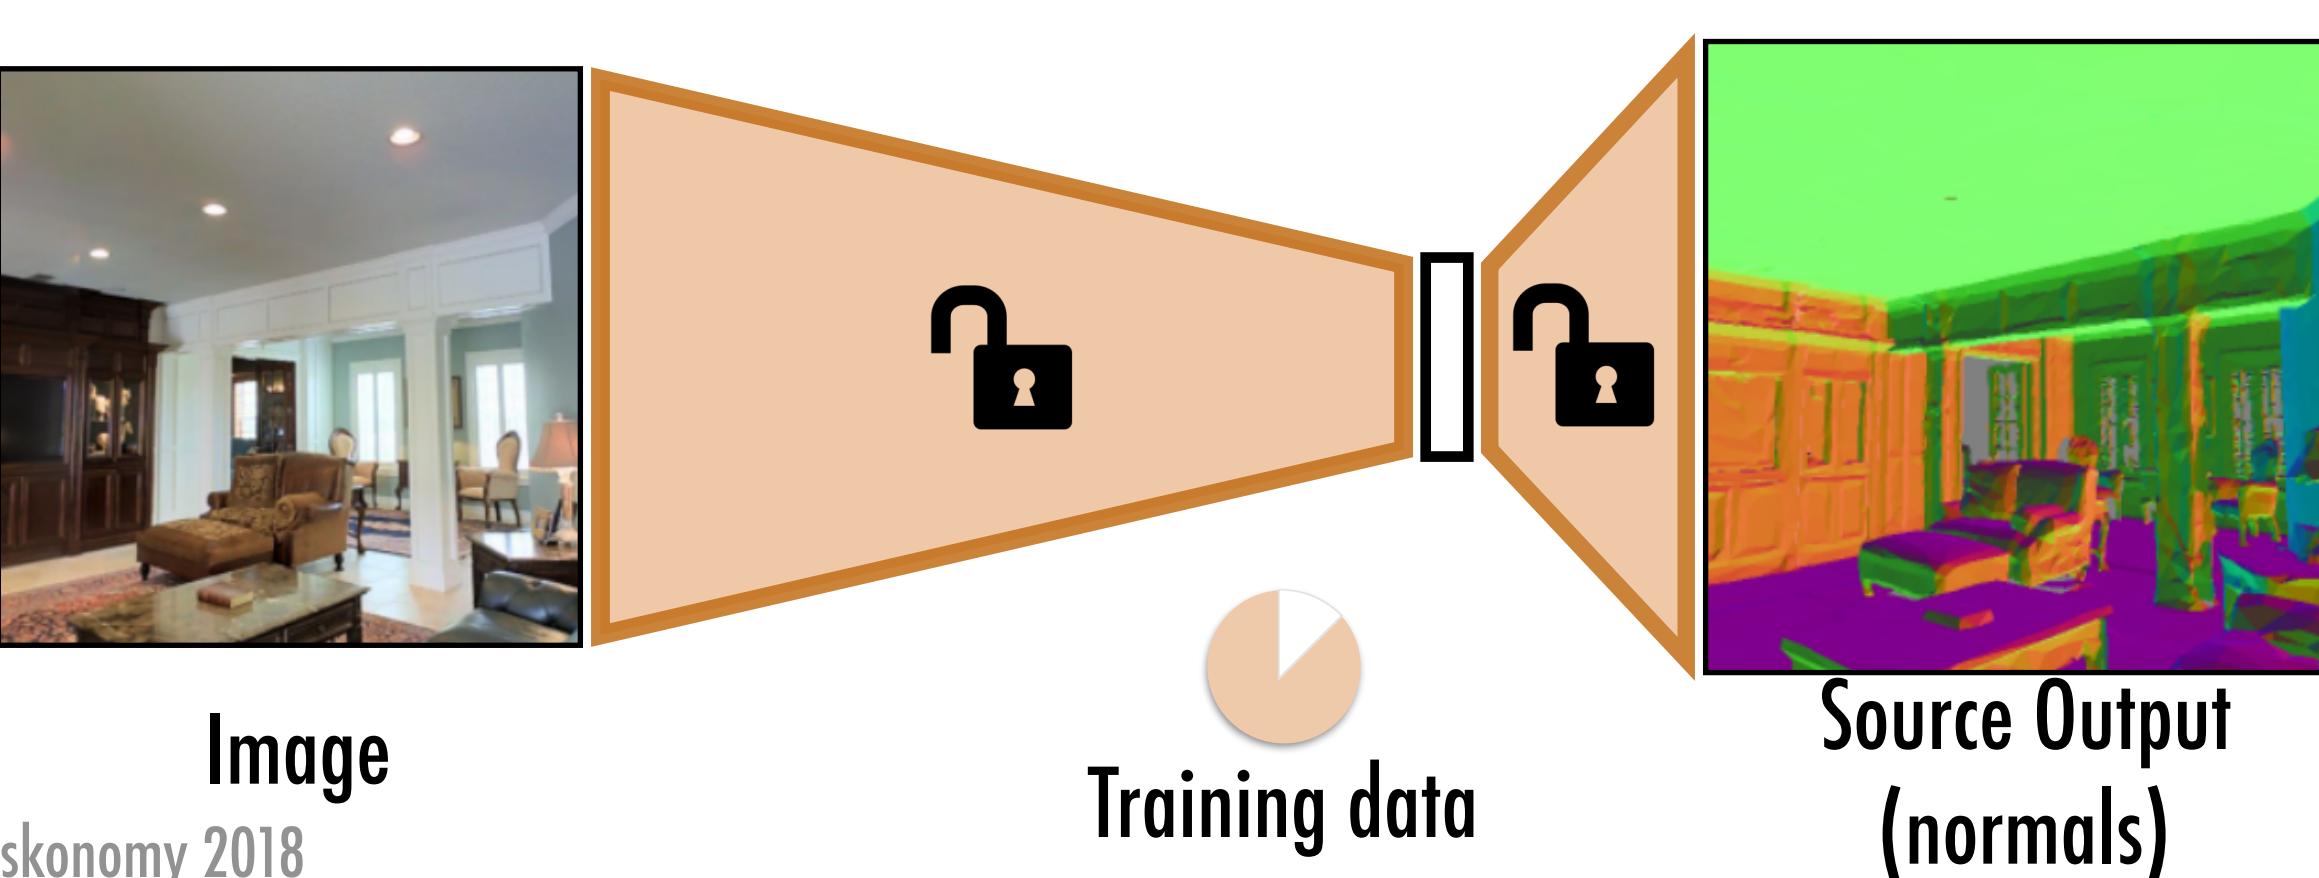
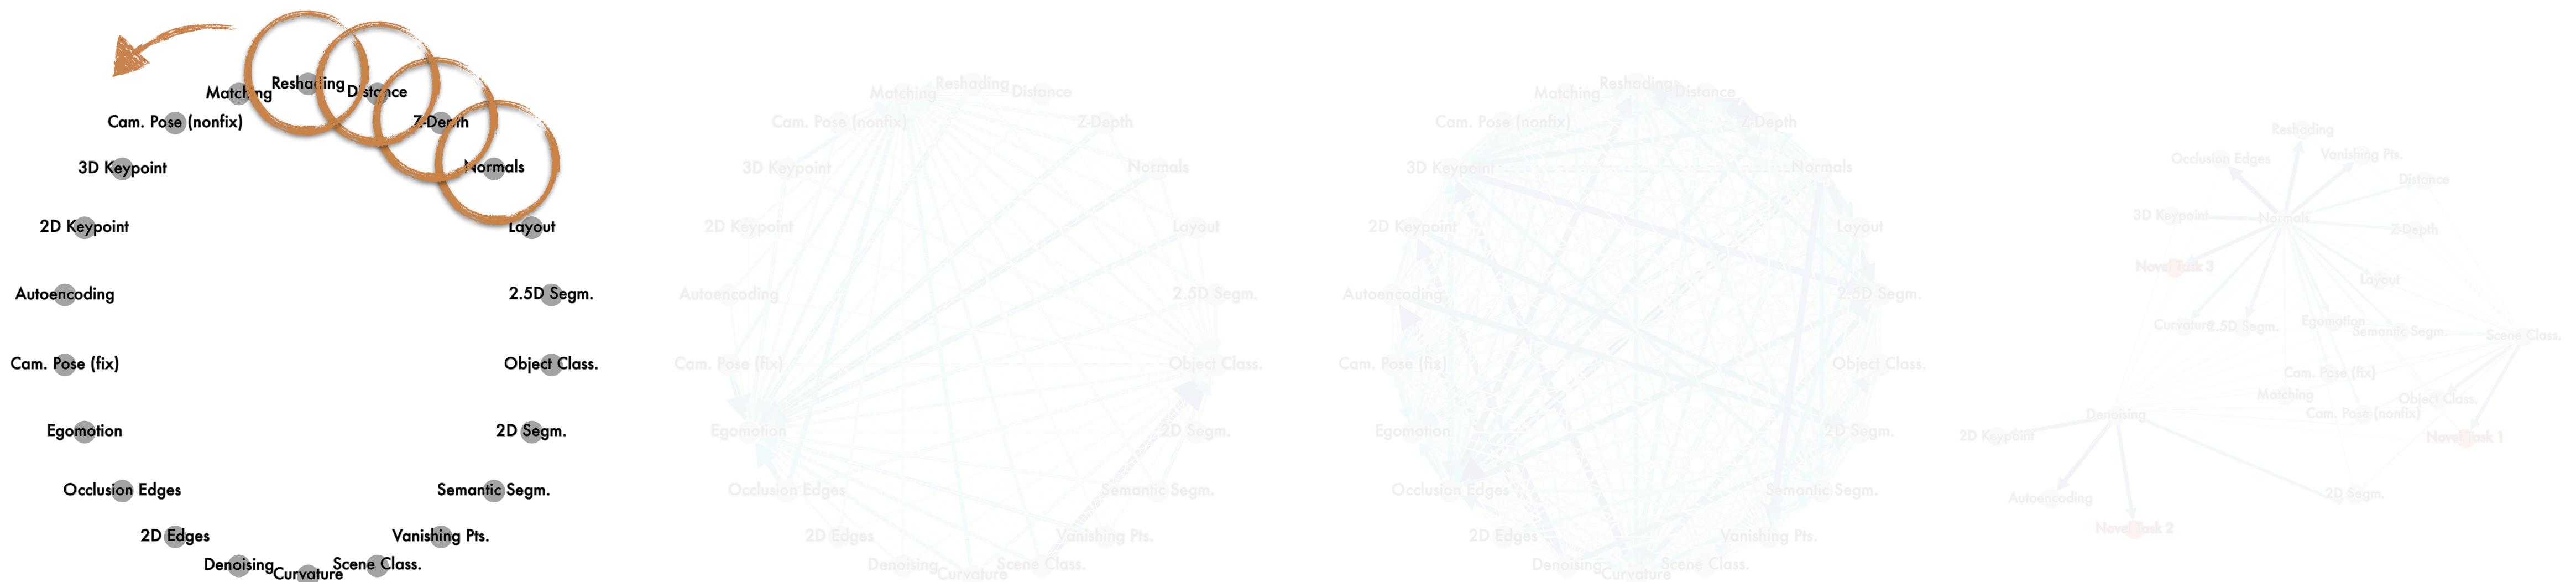


Frame-by-Frame results on a YouTube video. Live Demo: <https://taskonomy.vision/tasks>

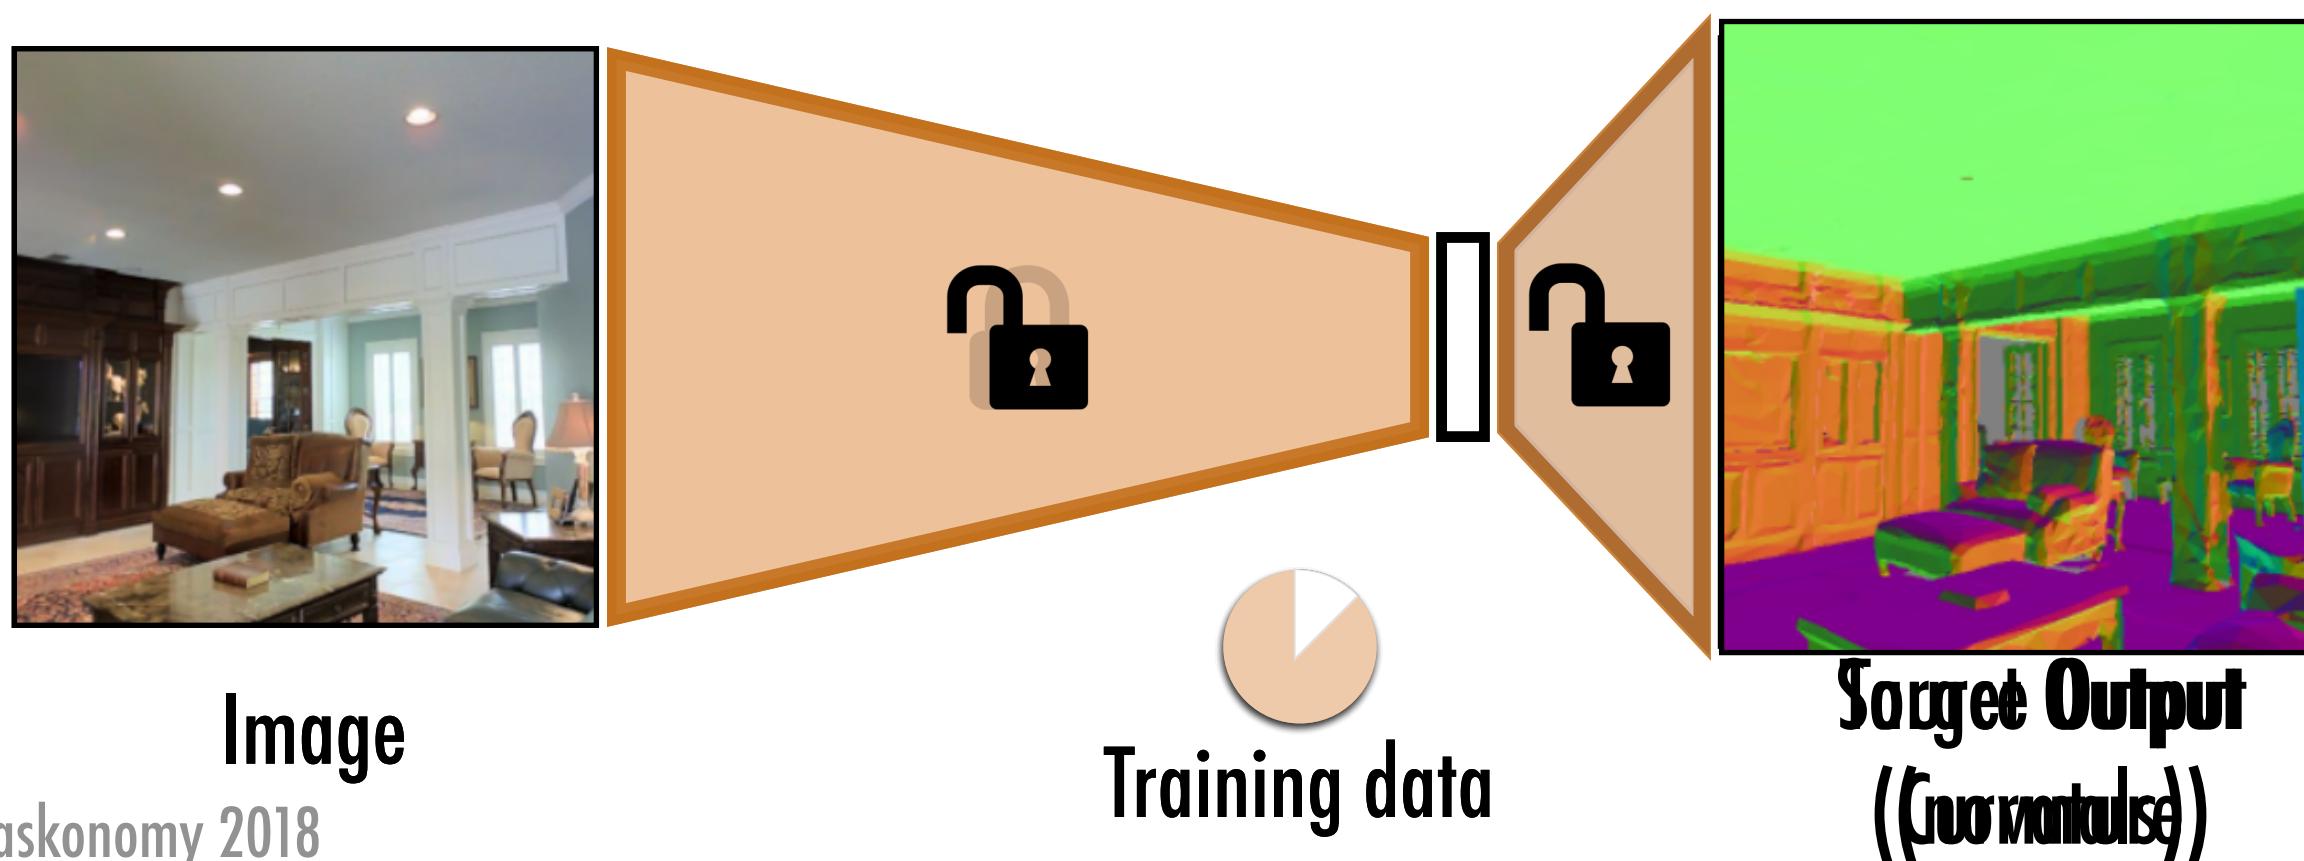
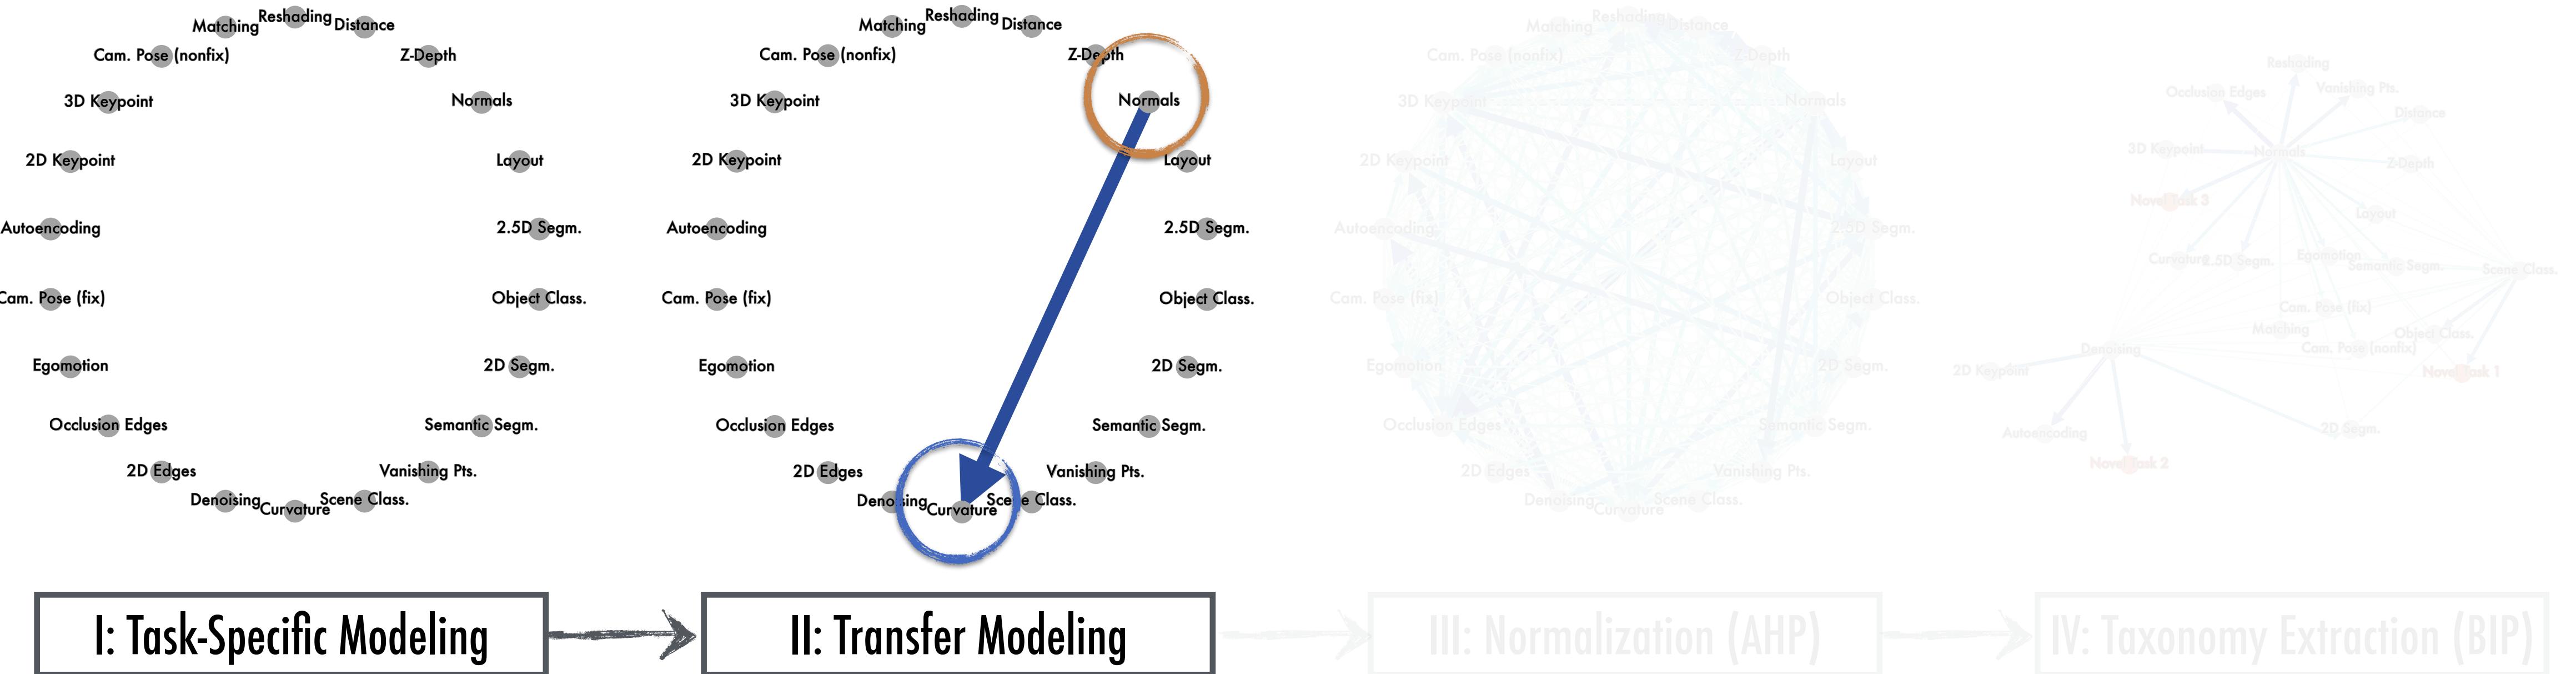
Modeling



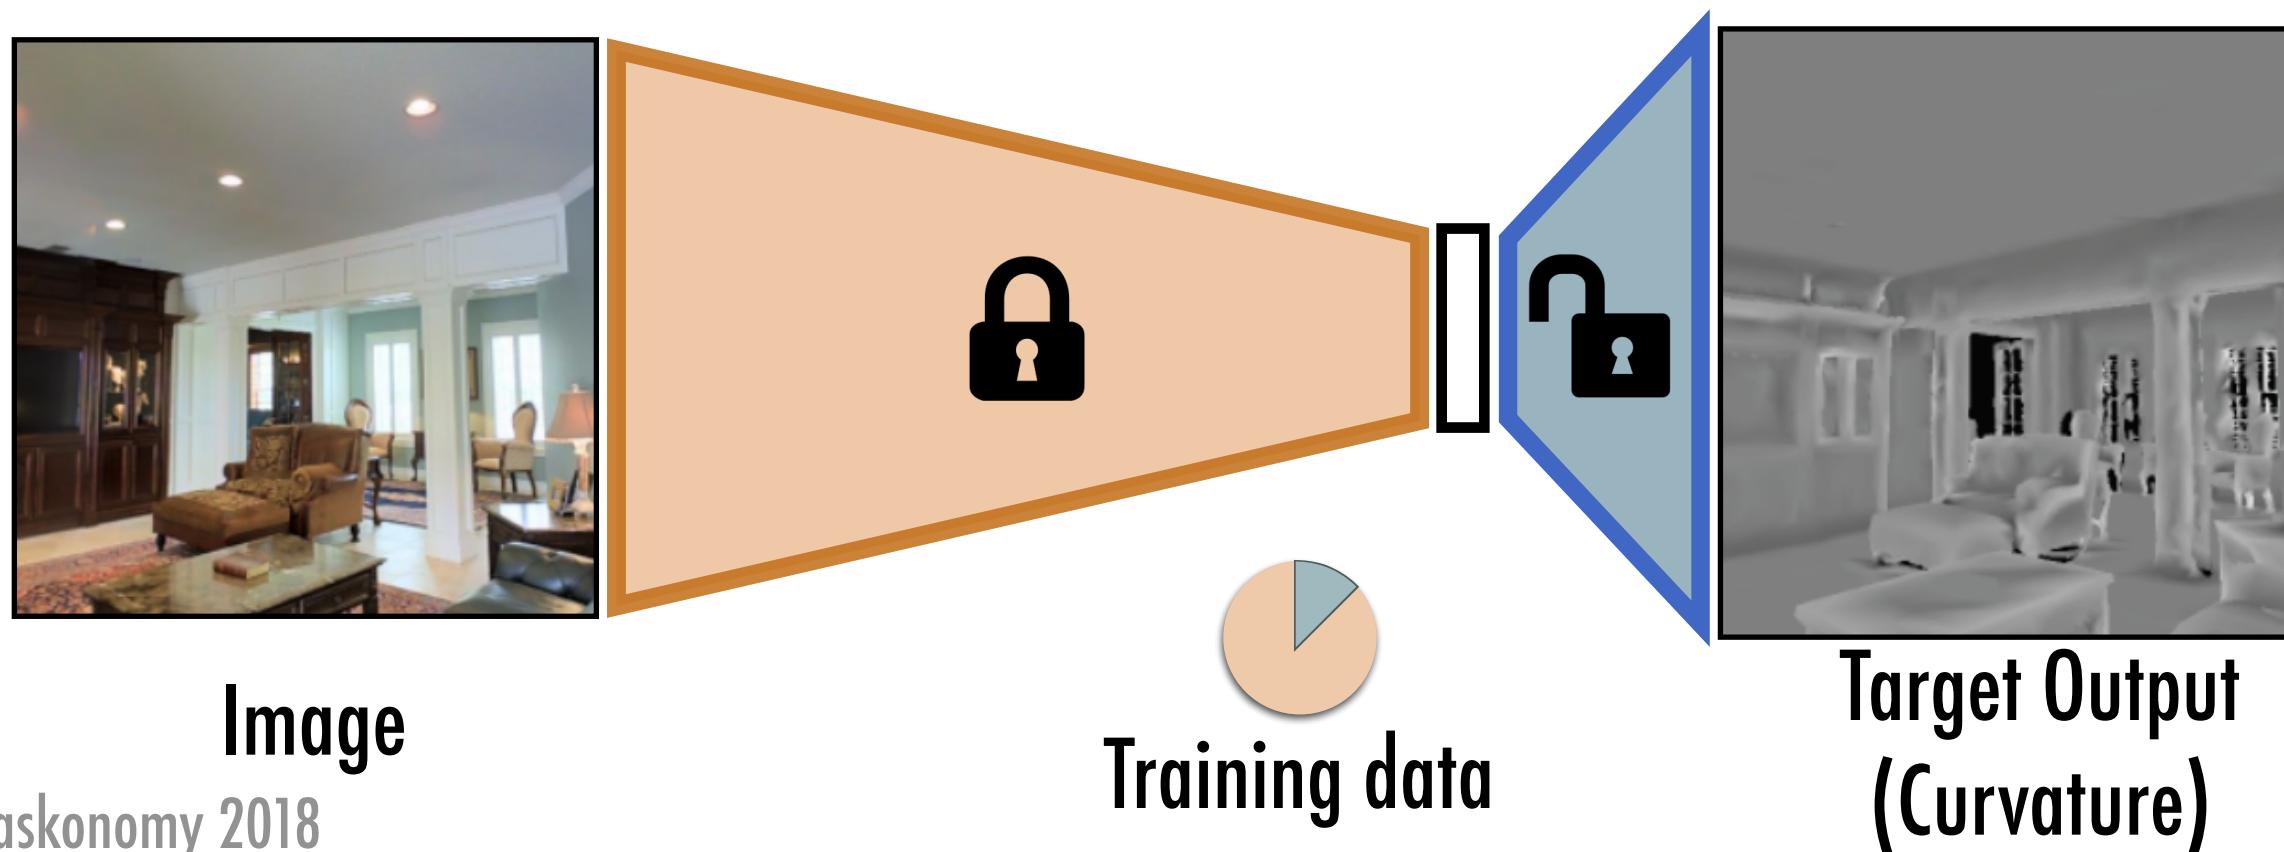
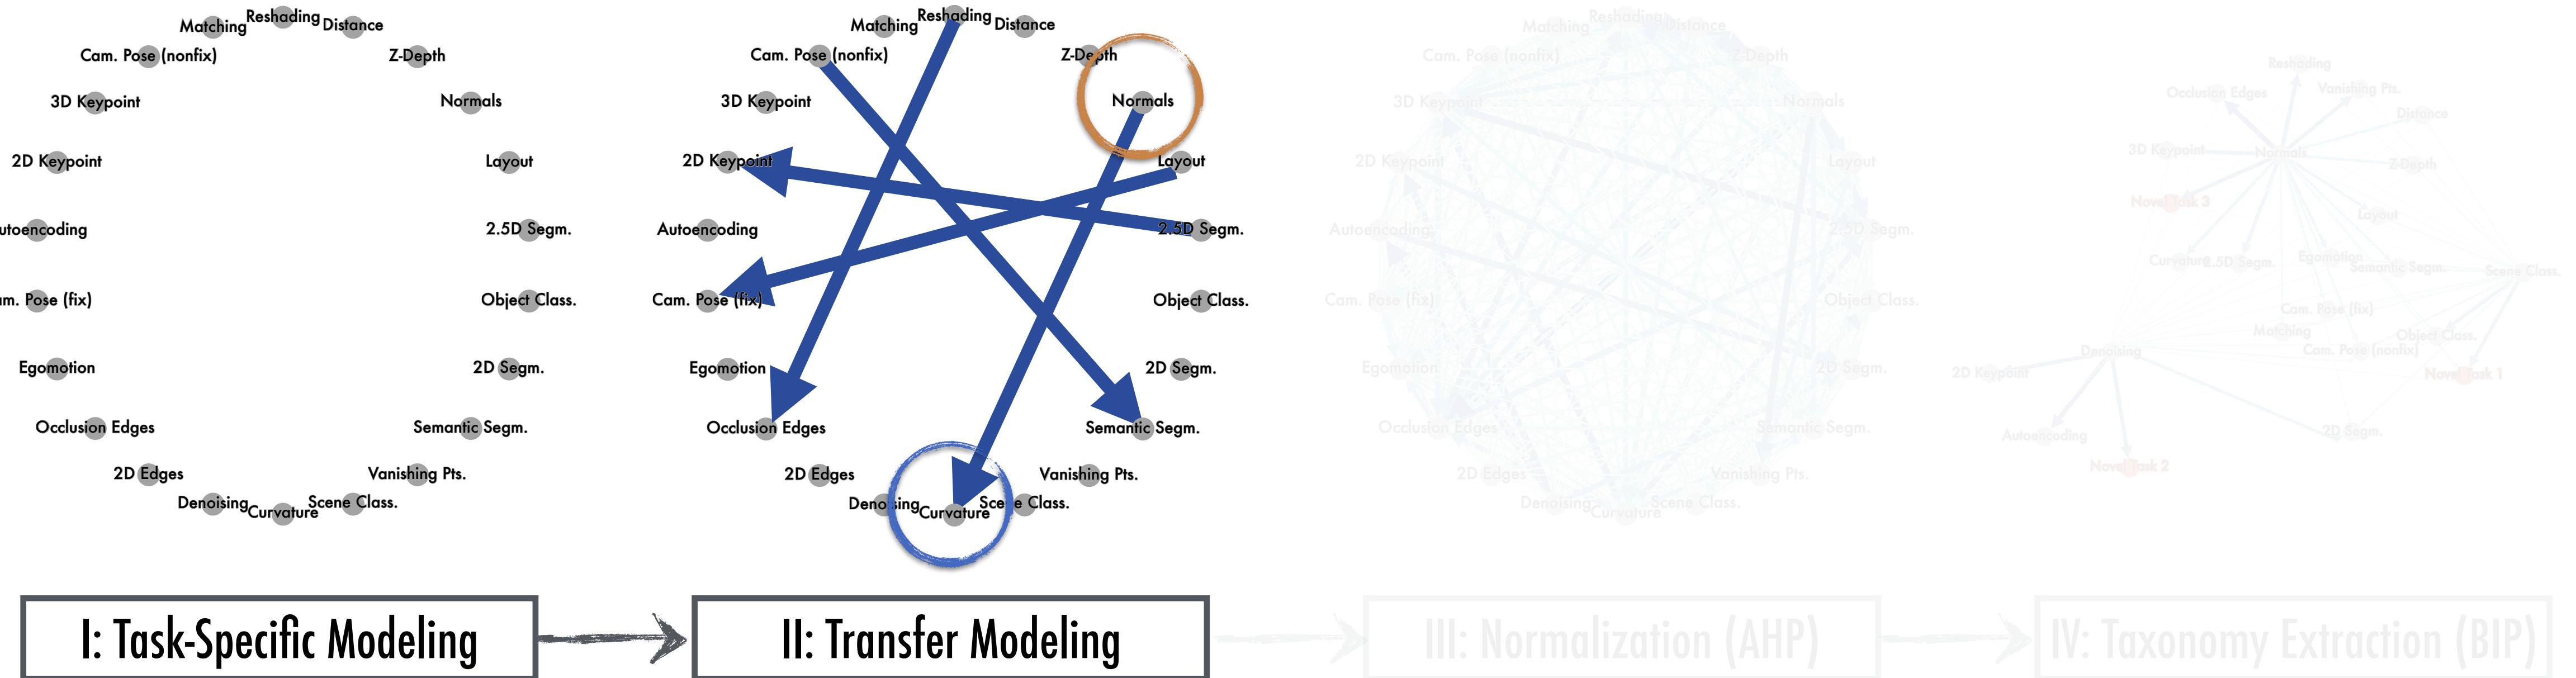
I: Task-Specific Modeling



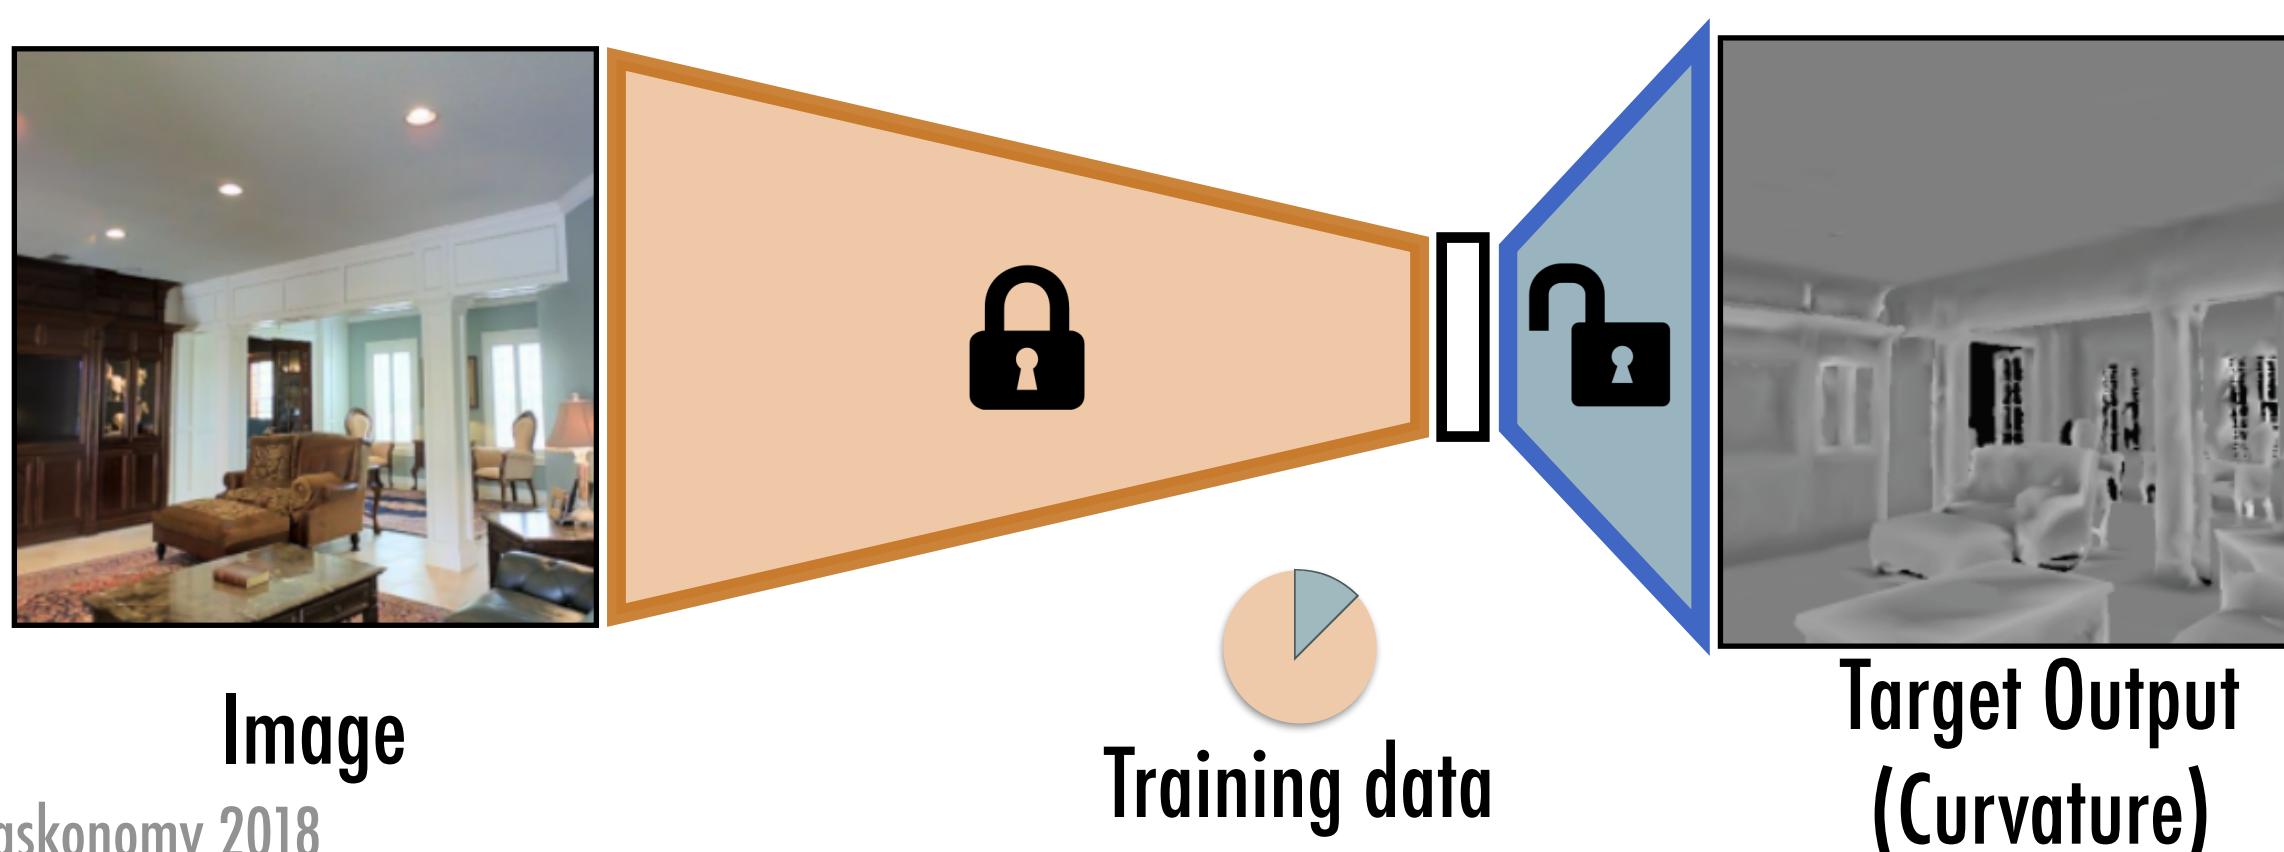
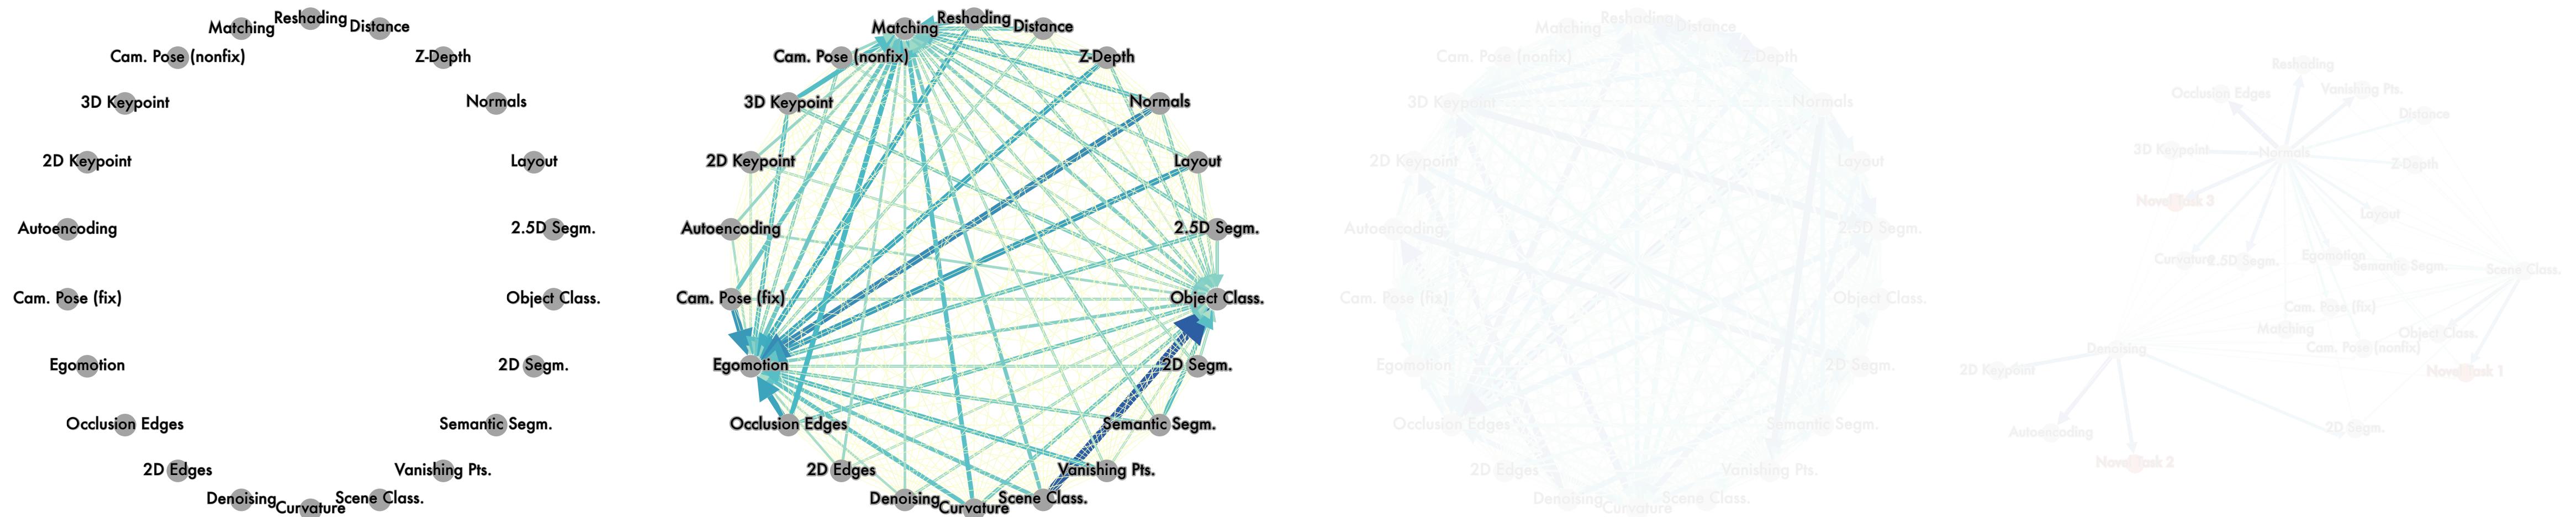
II: Transfer Modeling



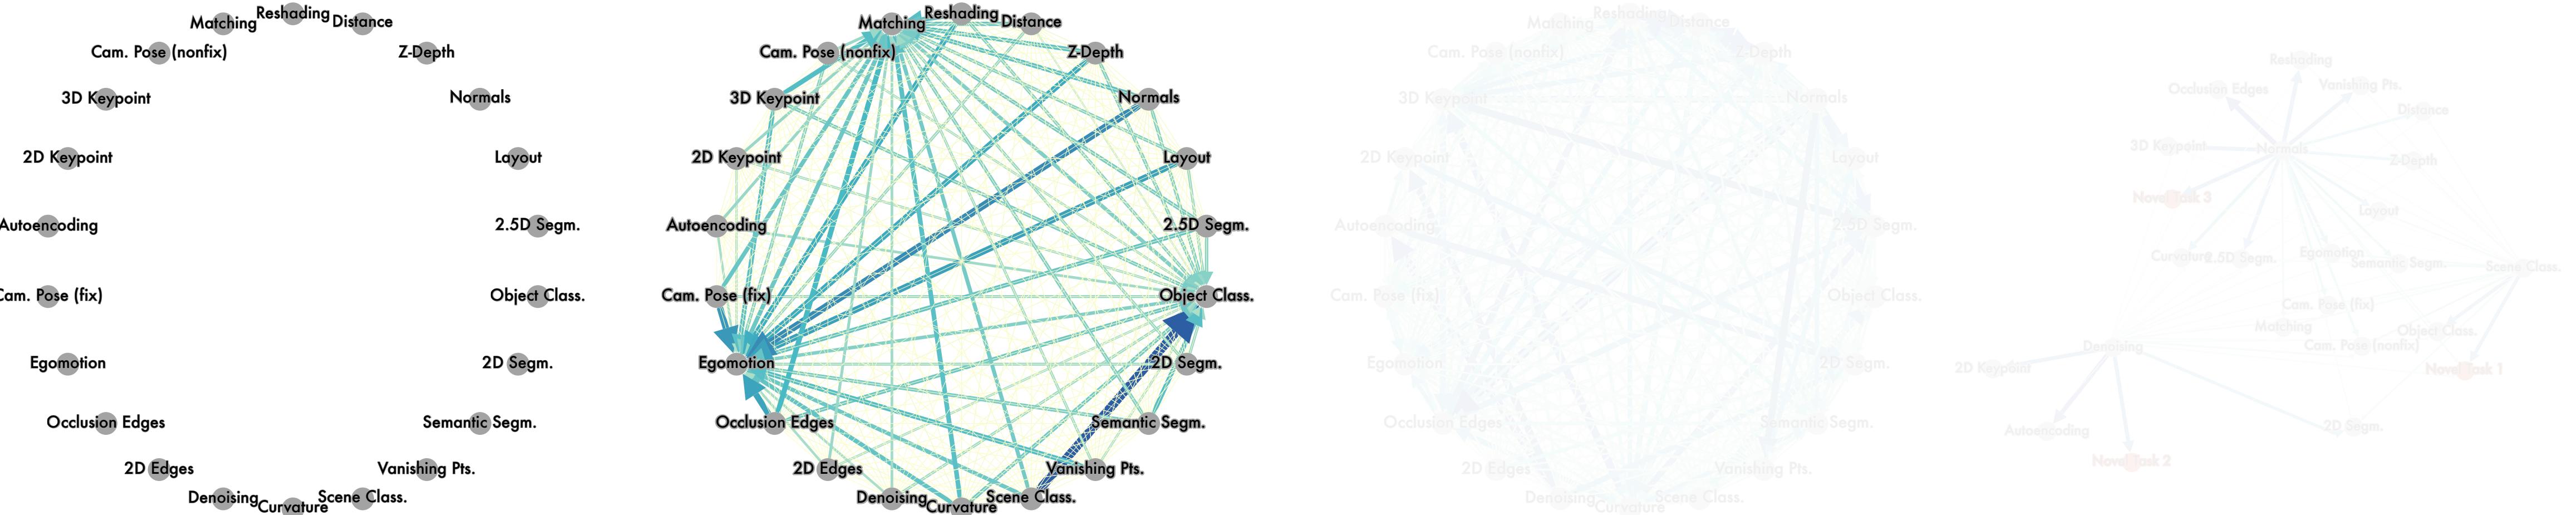
II: Transfer Modeling



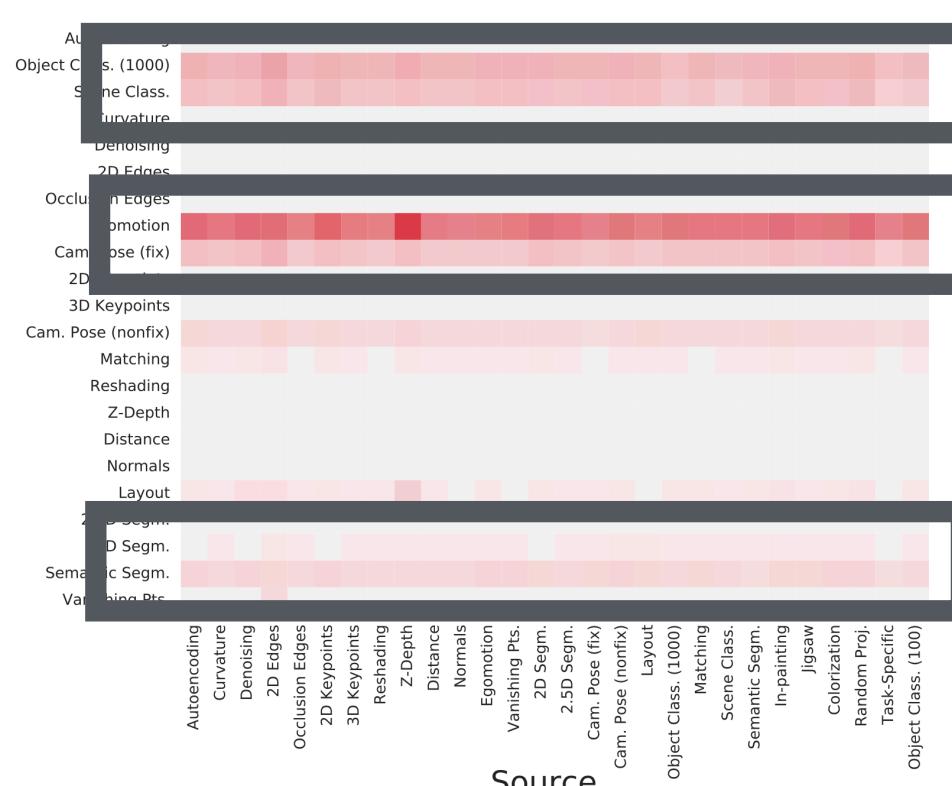
II: Transfer Modeling



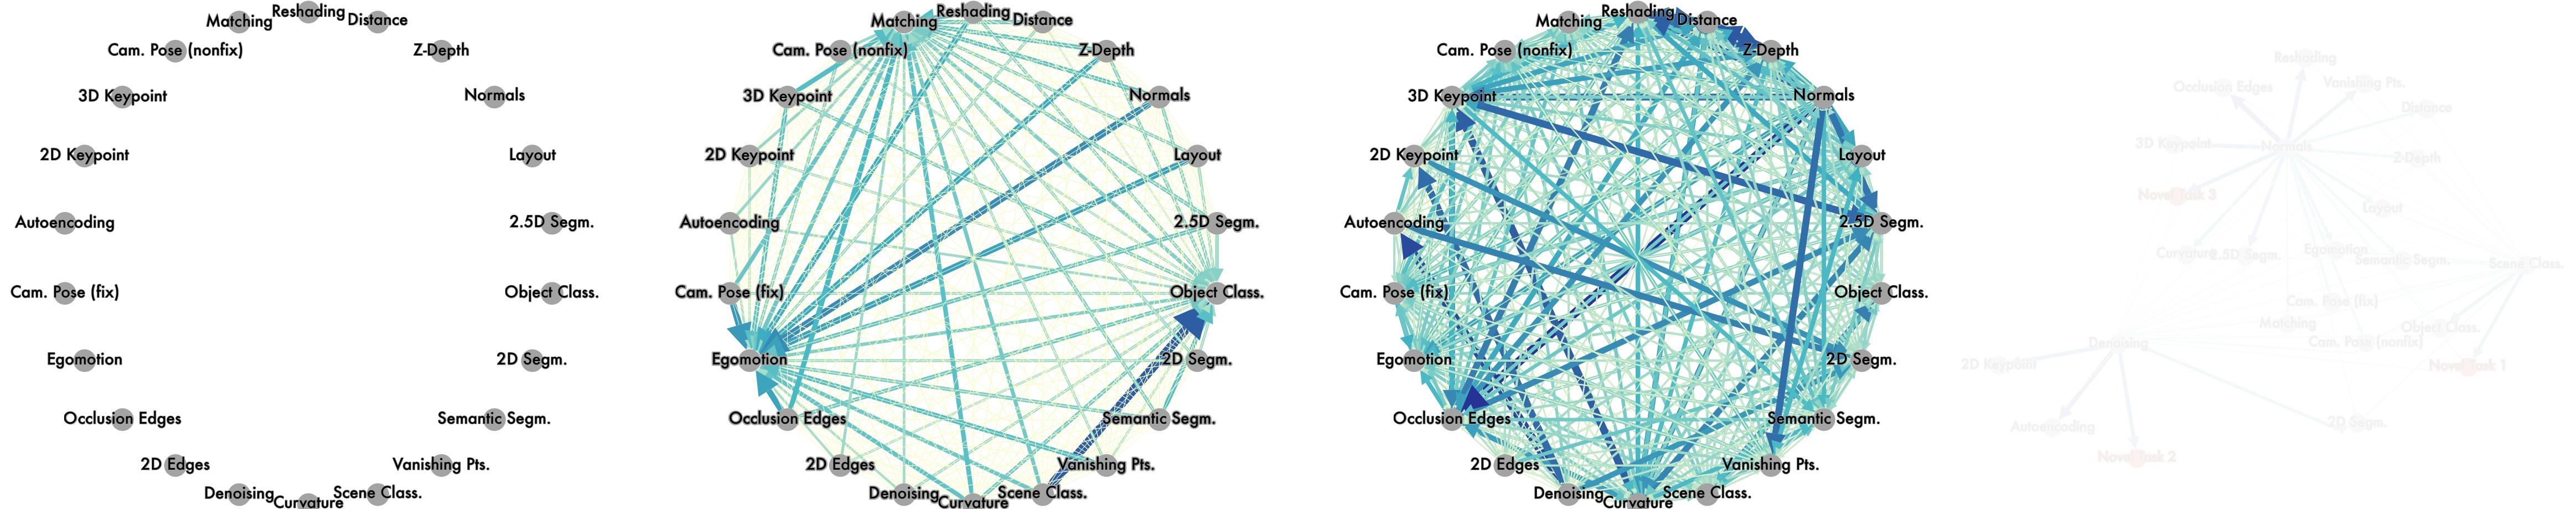
III: Normalization



Adjacency Matrix (pre-normalization)



III: Normalization



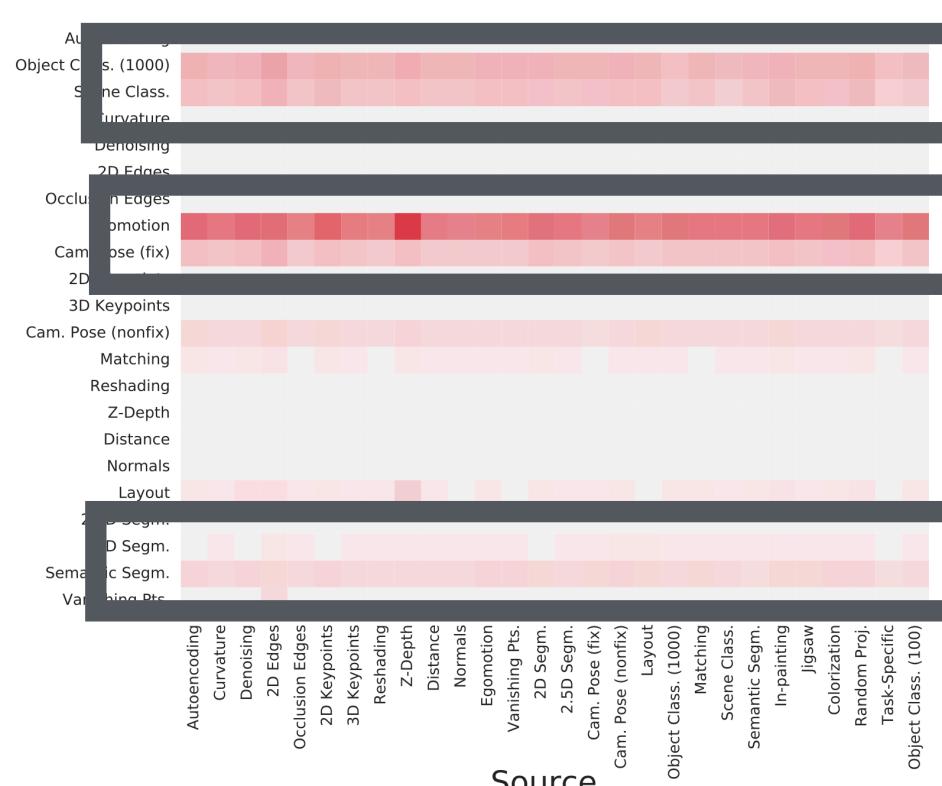
I: Task-Specific Modeling

II: Transfer Modeling

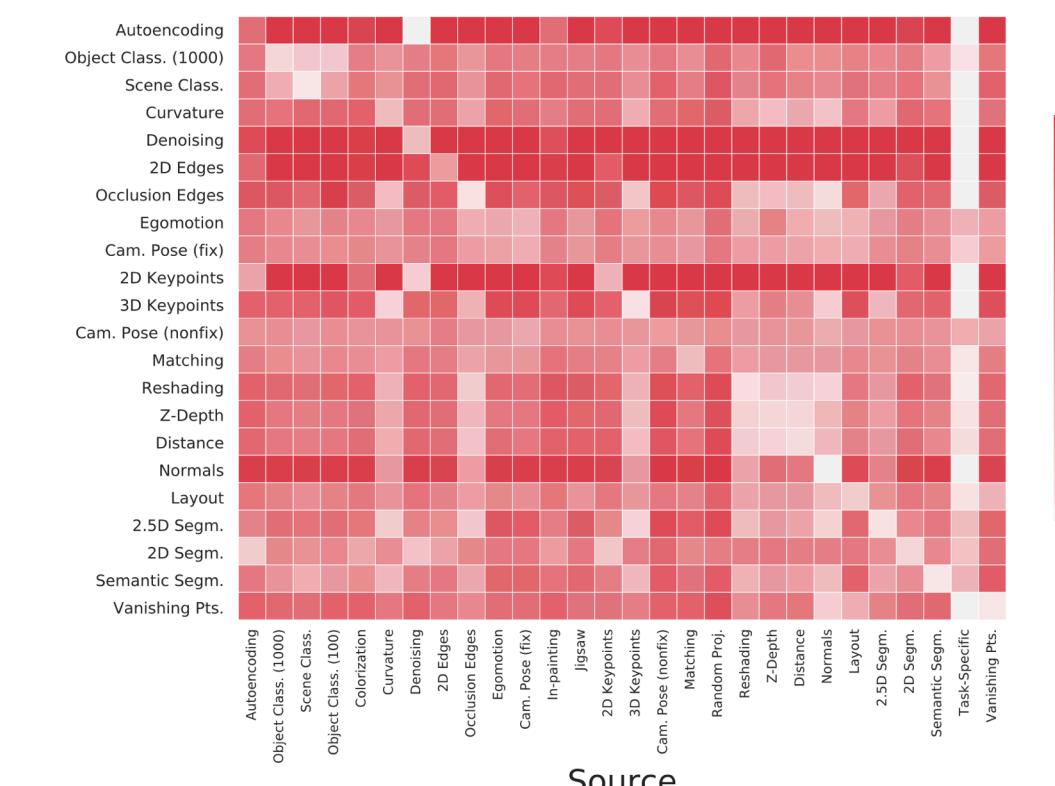
III: Normalization (AHP)

IV: Taxonomy Extraction (BIP)

Adjacency Matrix (pre-normalization)

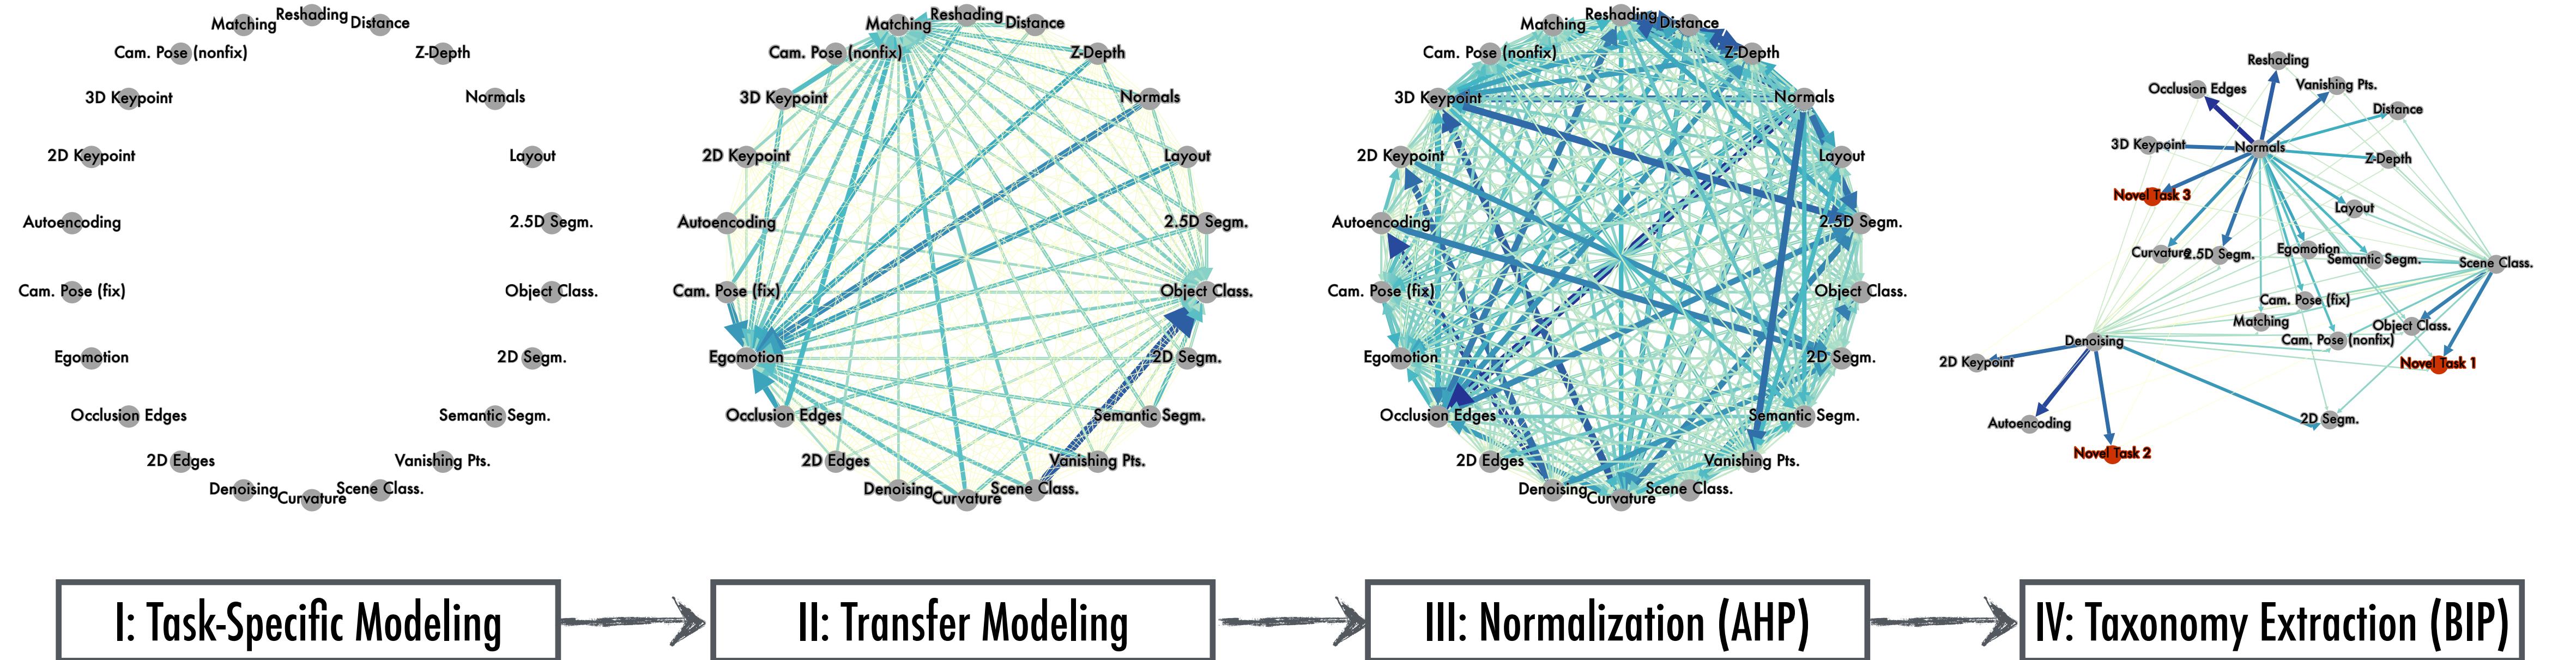


Adjacency Matrix (post-normalization)



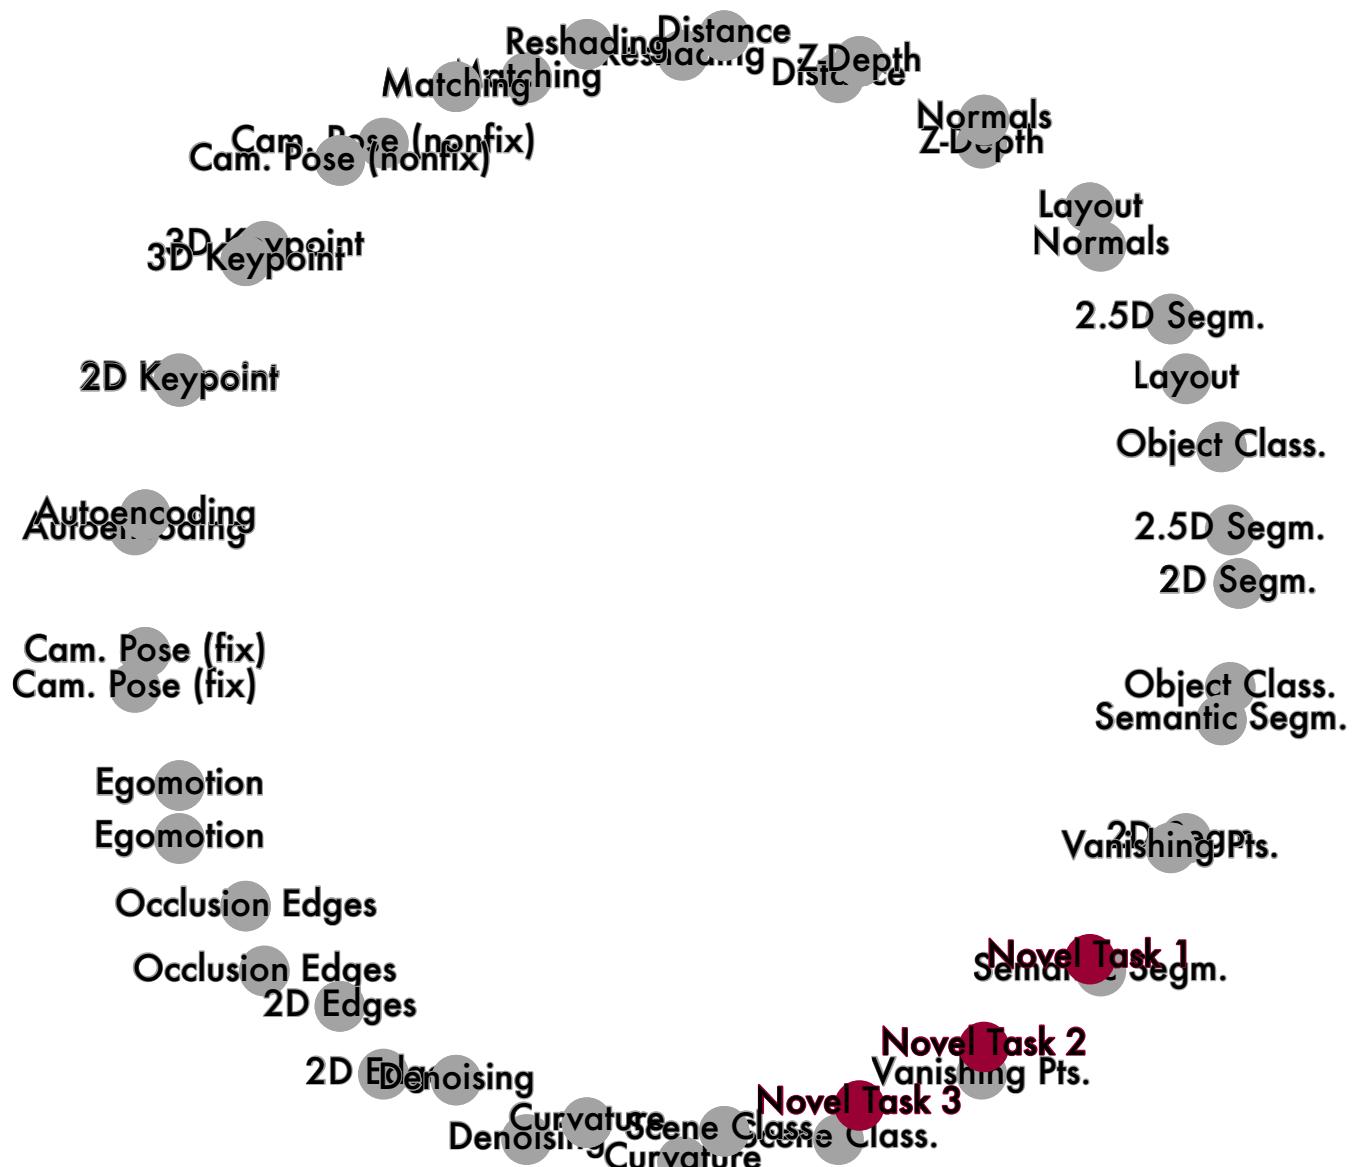
Ordinal Normalization -
Analytic Hierarchical Process.
(AHP)

IV: Taxonomy Extraction

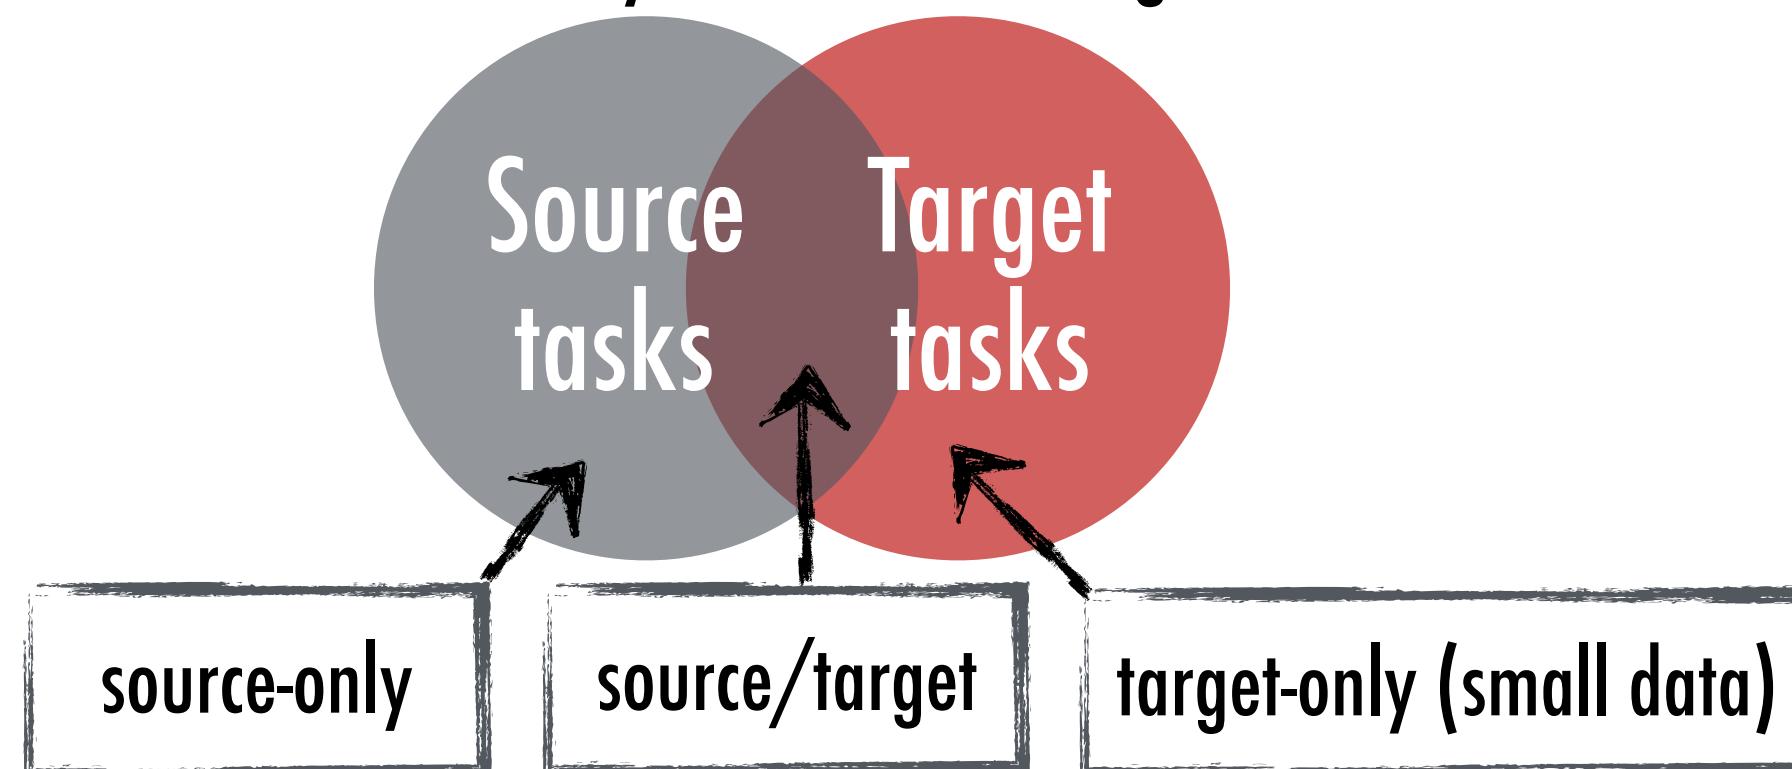


- **Taxonomical structure:**
 - Sparsified
 - What are best source tasks
 - What sources for each target
 - Out-of-dictionary tasks
 - Maximize performance while constrained by some budget

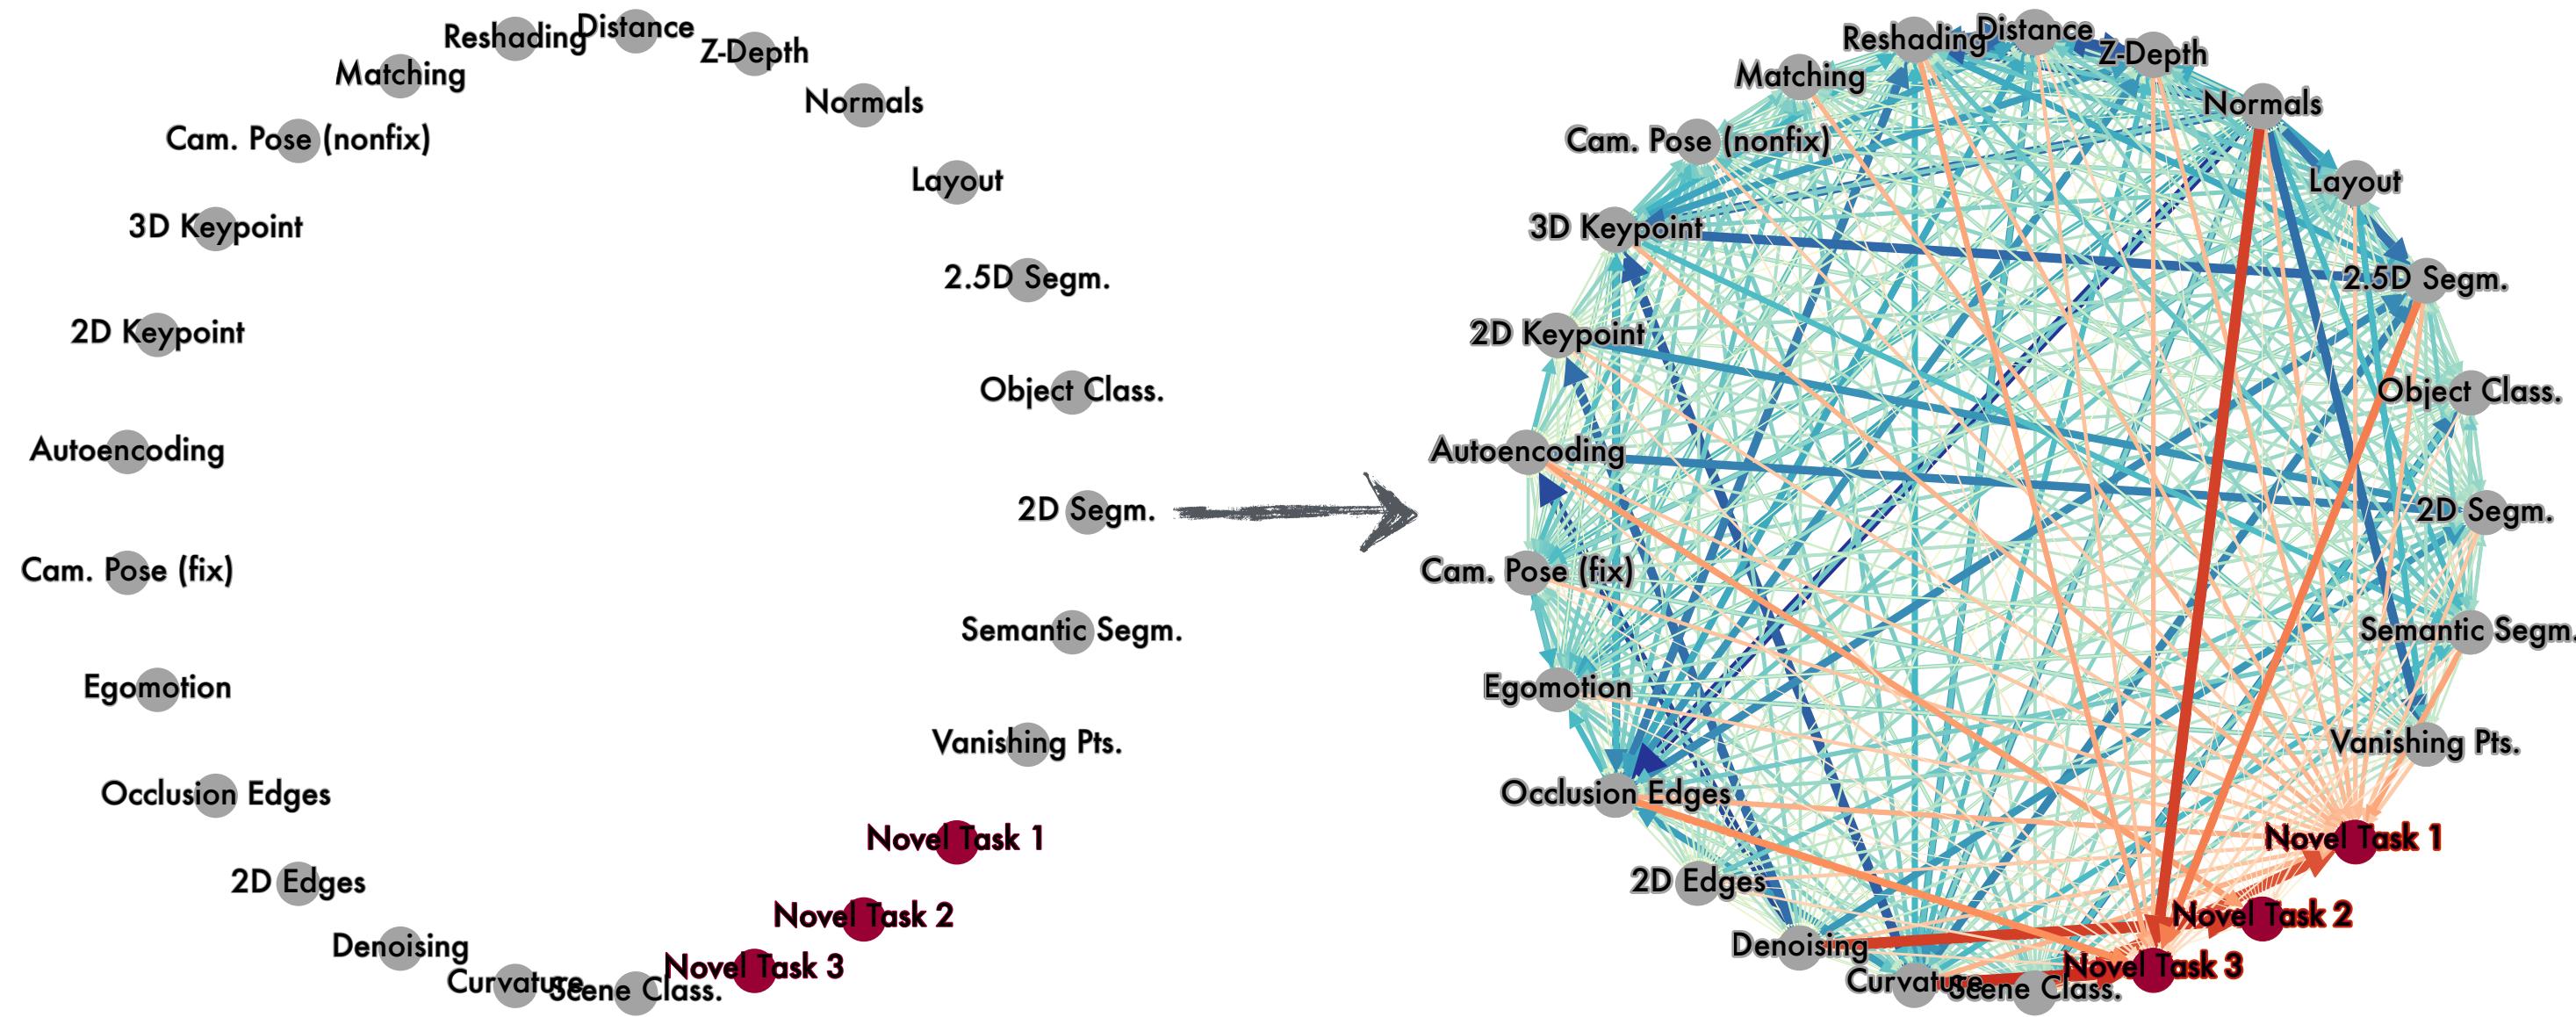
IV: Taxonomy Extraction



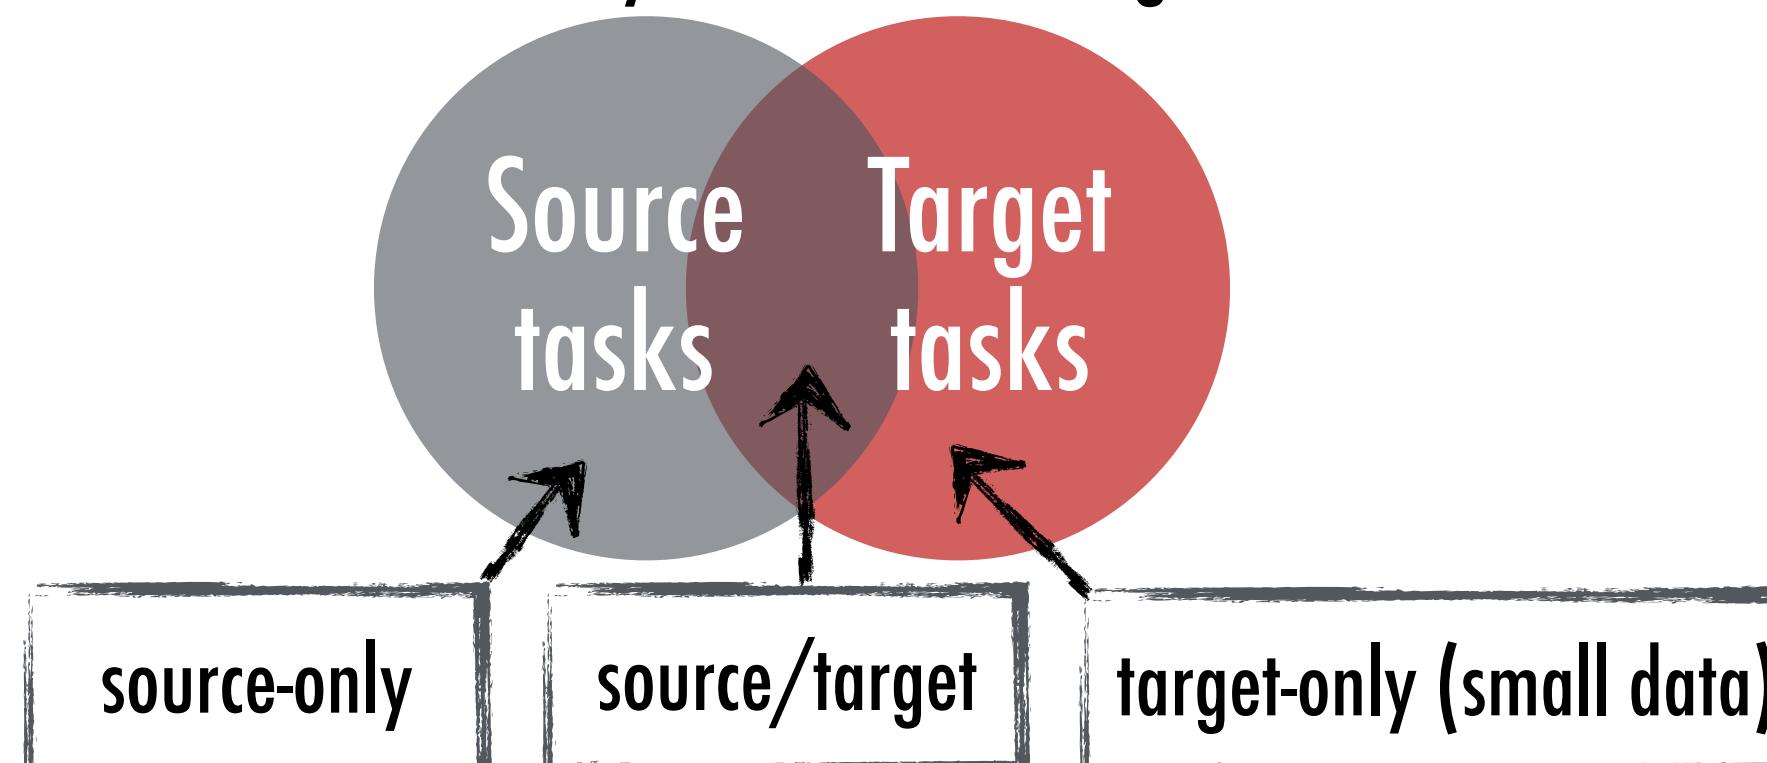
Dictionary= Sources \cup Targets



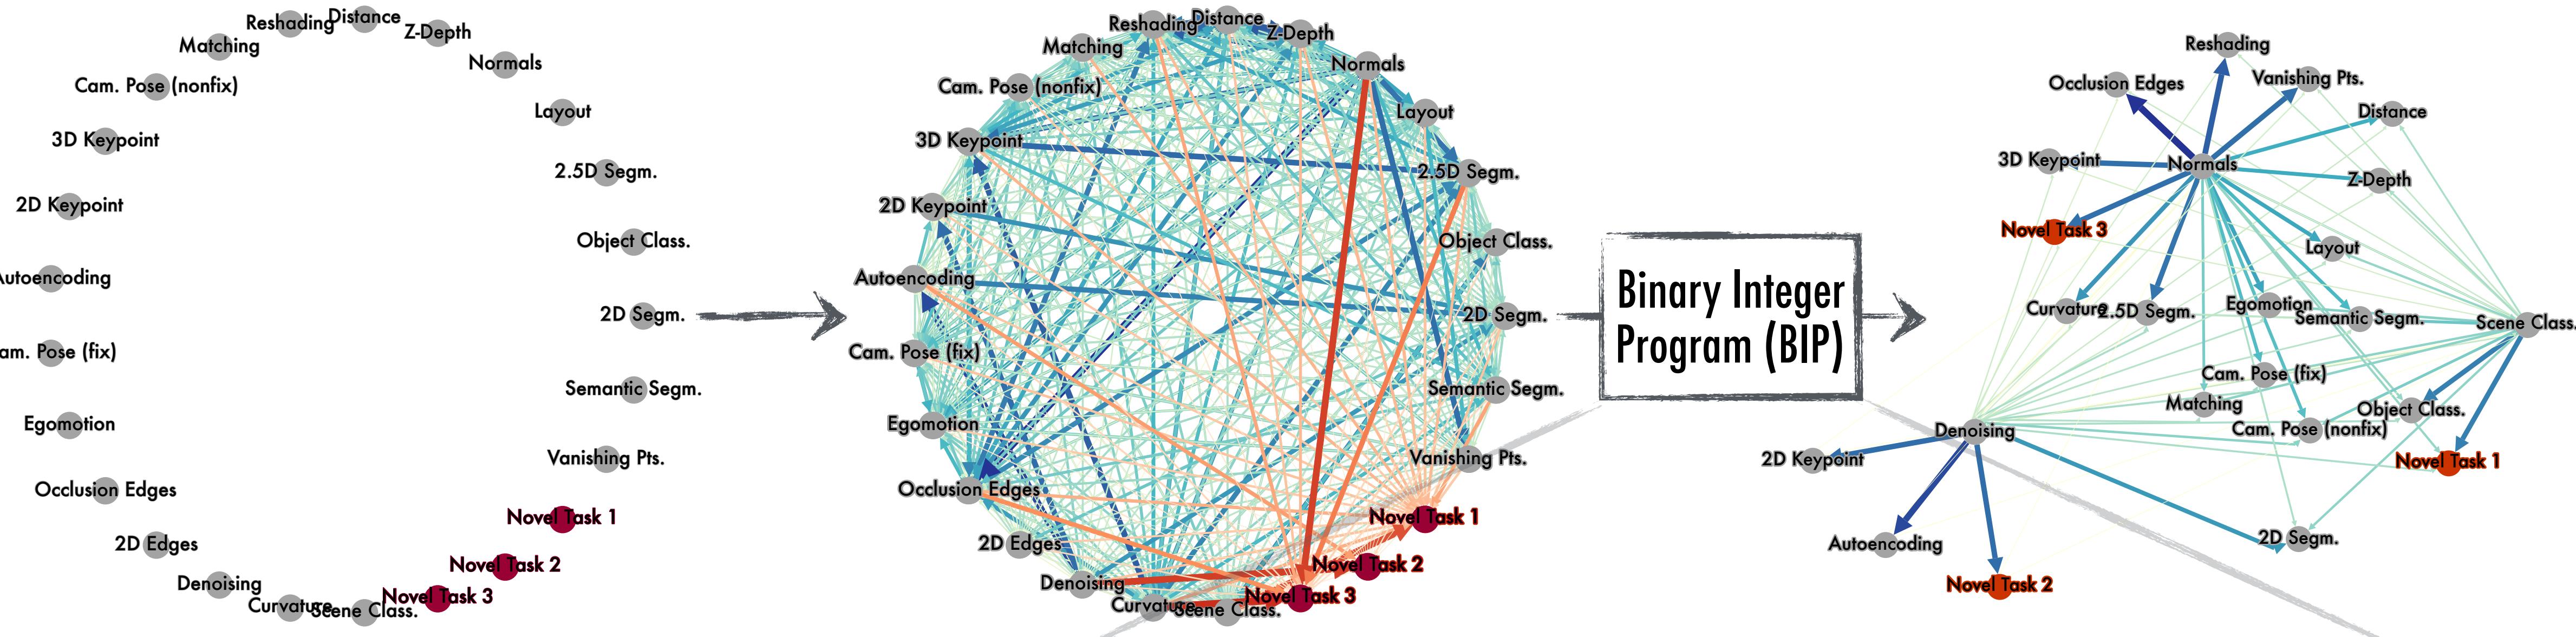
IV: Taxonomy Extraction



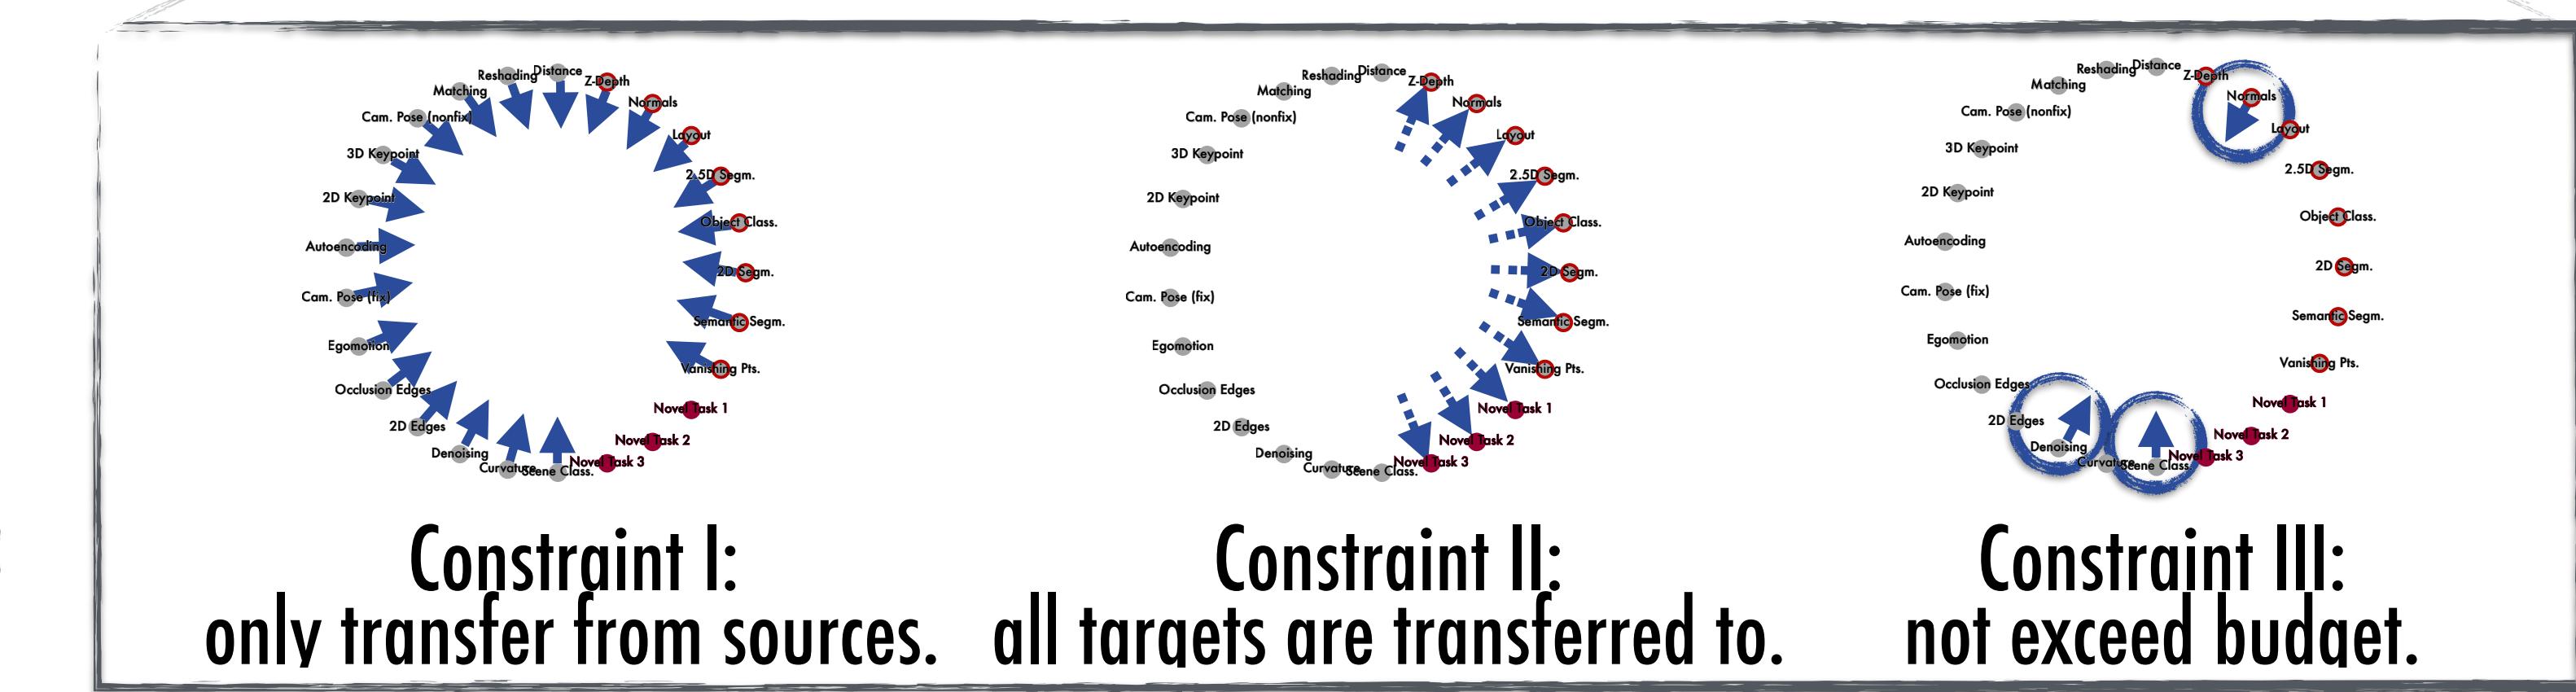
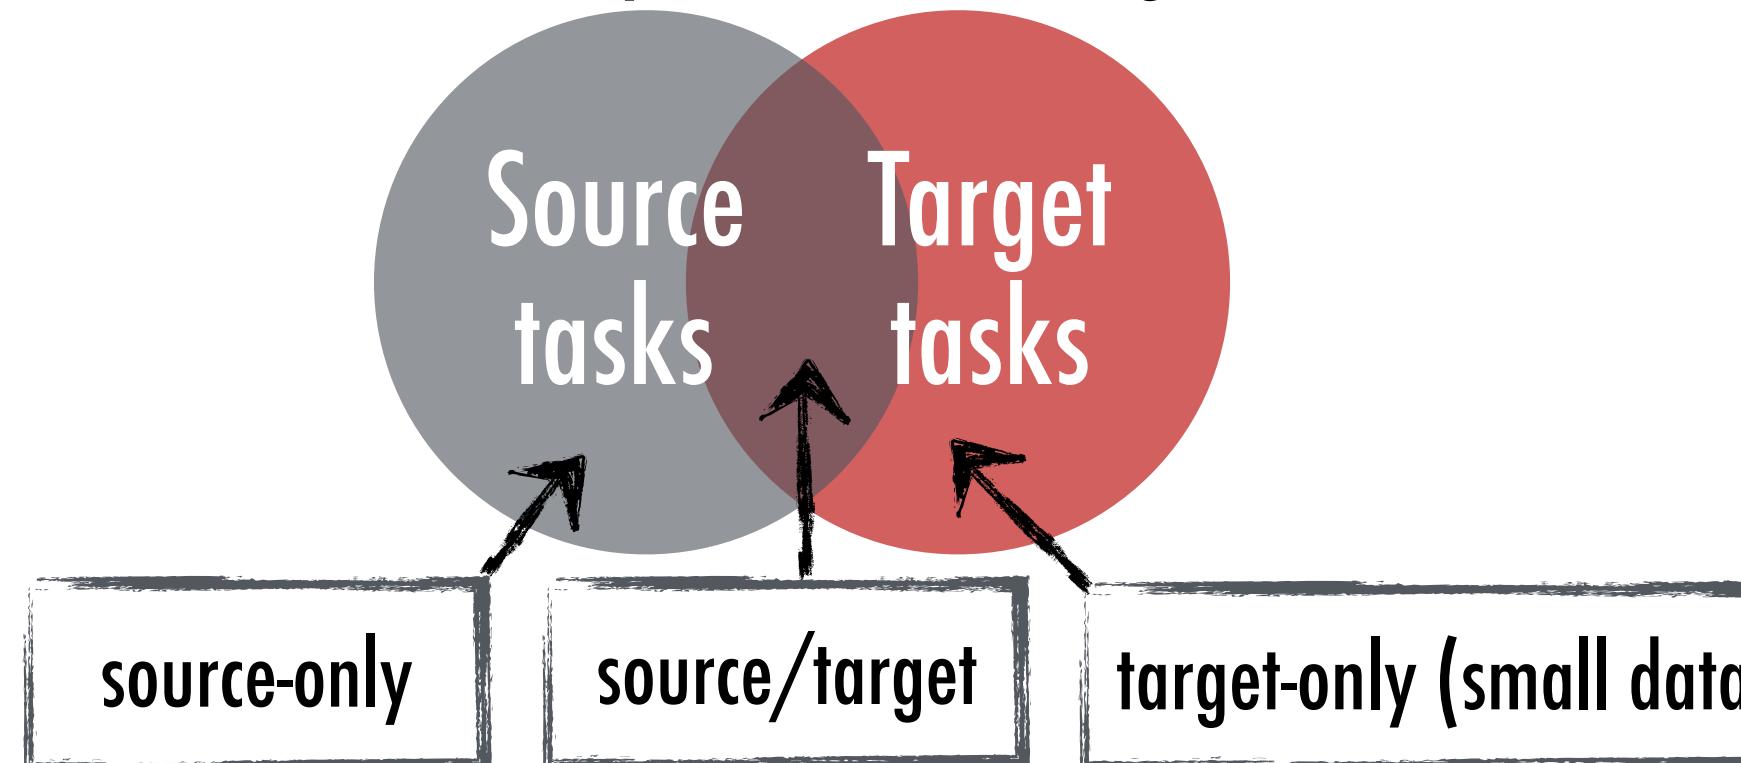
Dictionary= Sources \cup Targets



IV: Taxonomy Extraction

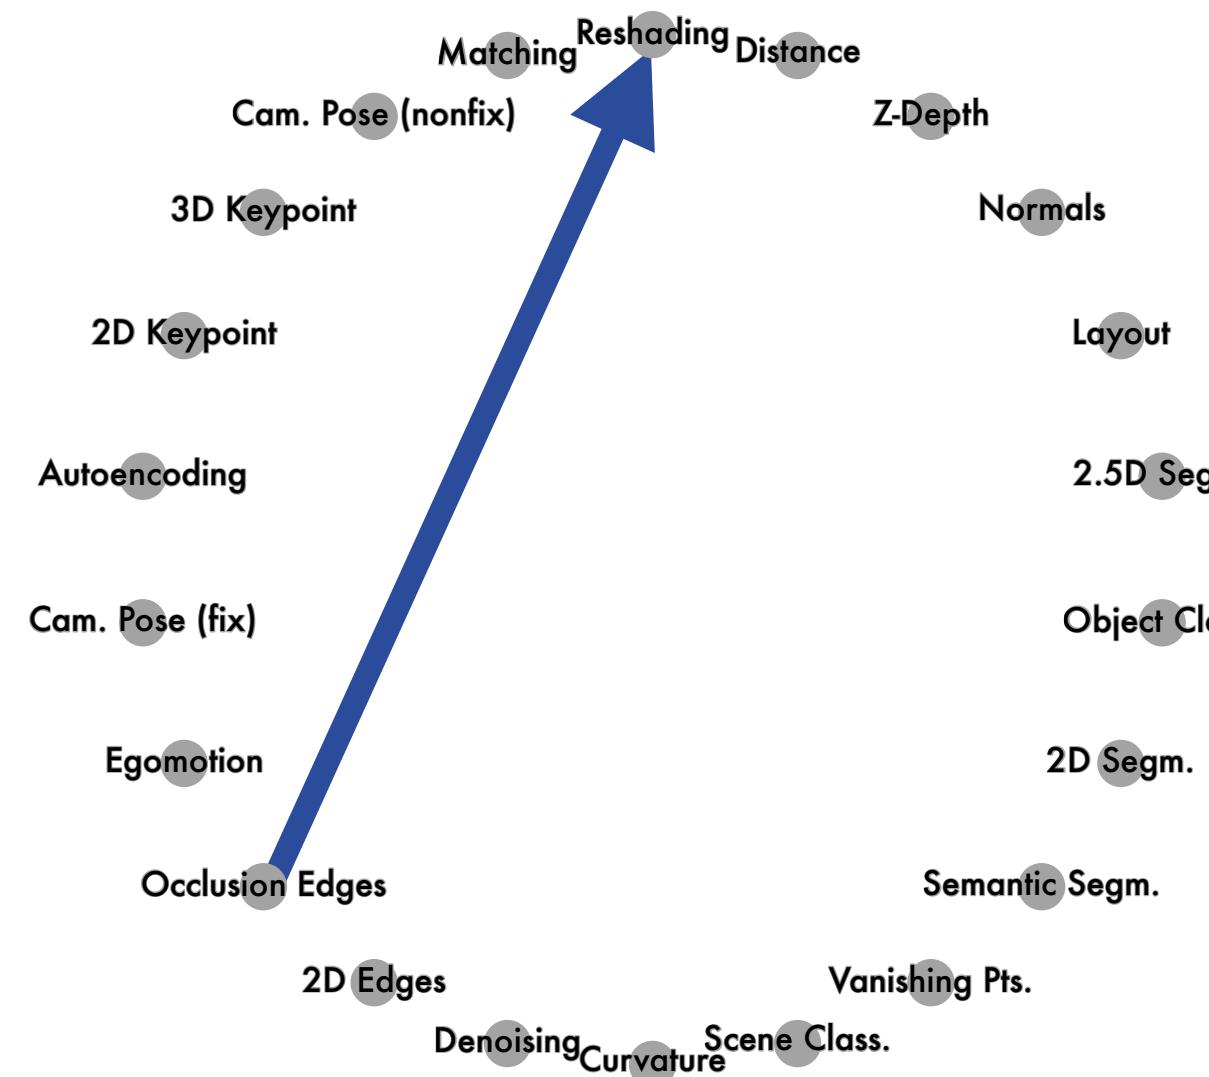


Dictionary= Sources \cup Targets

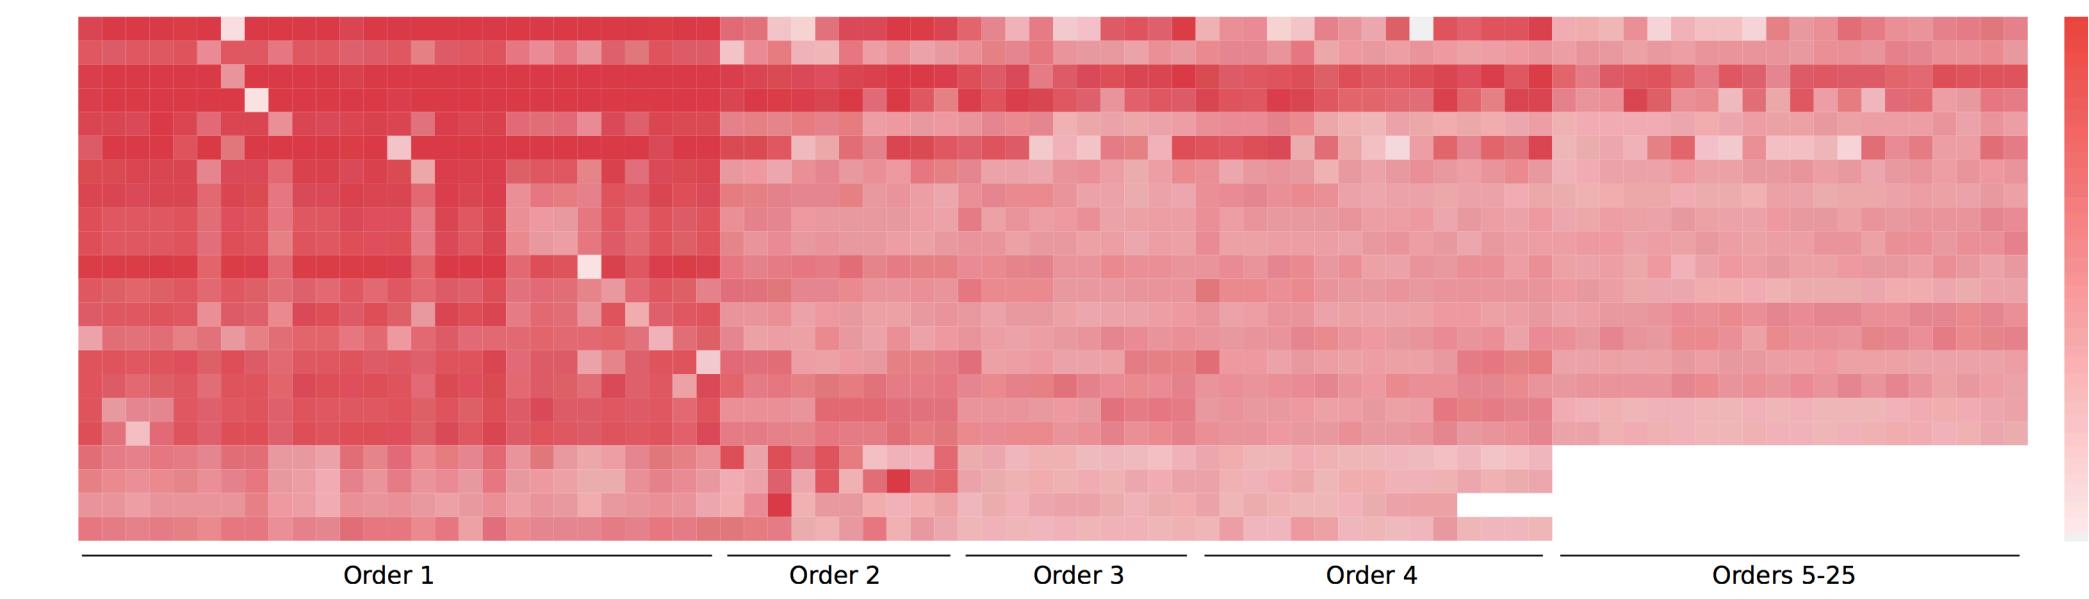
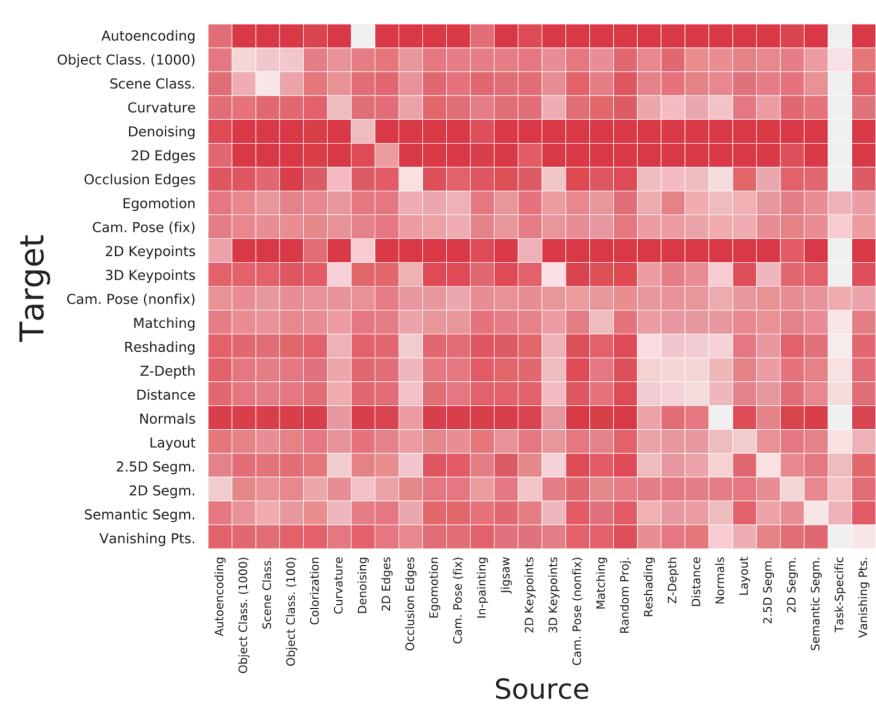
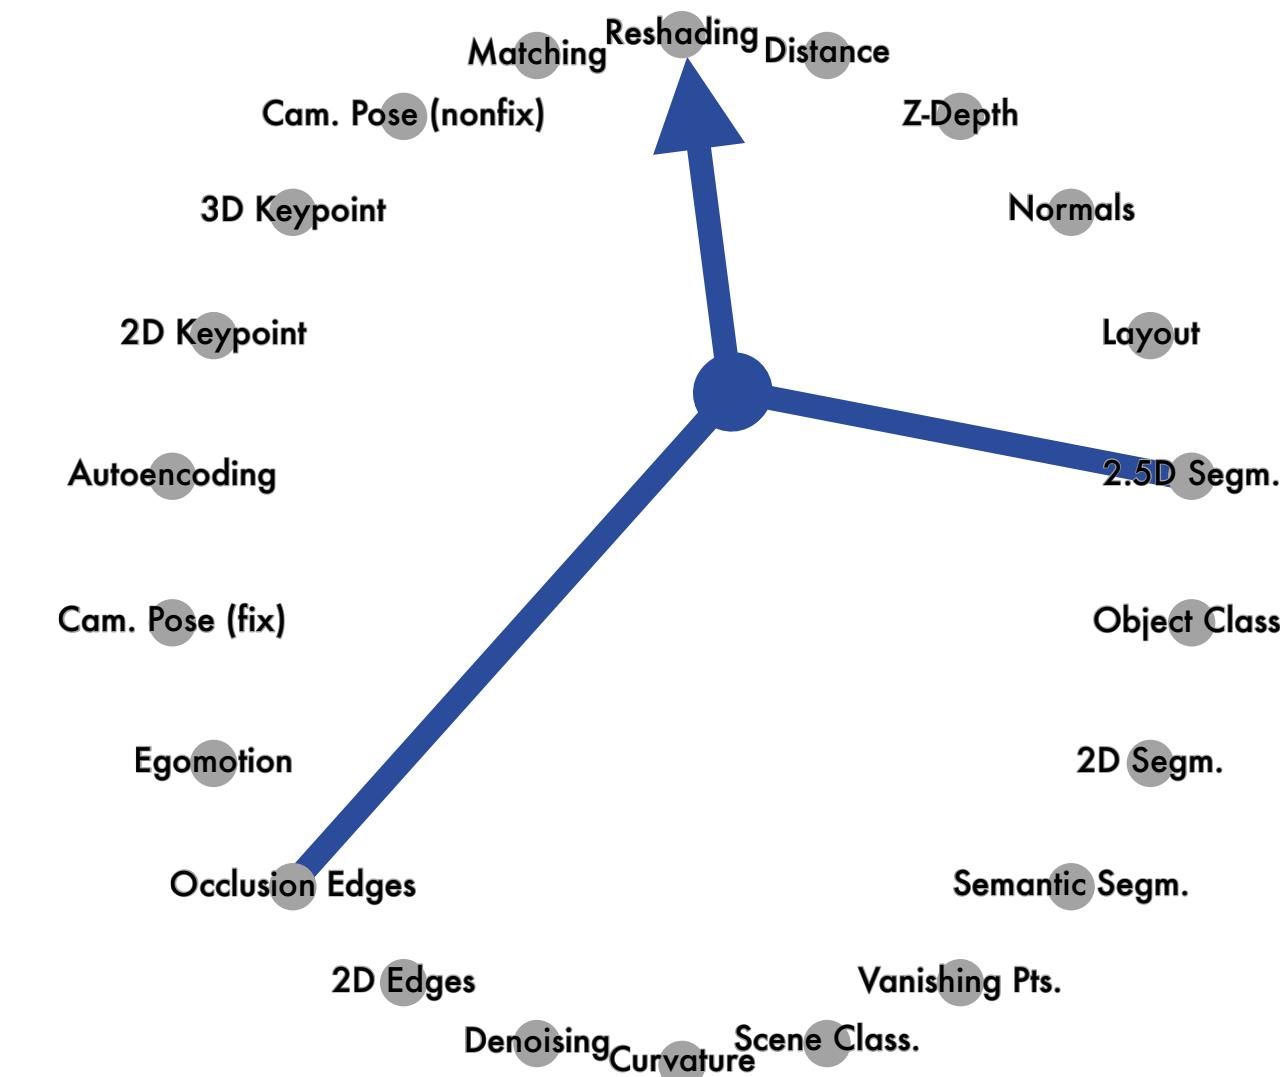


Higher Order Transfers

1st order transfer



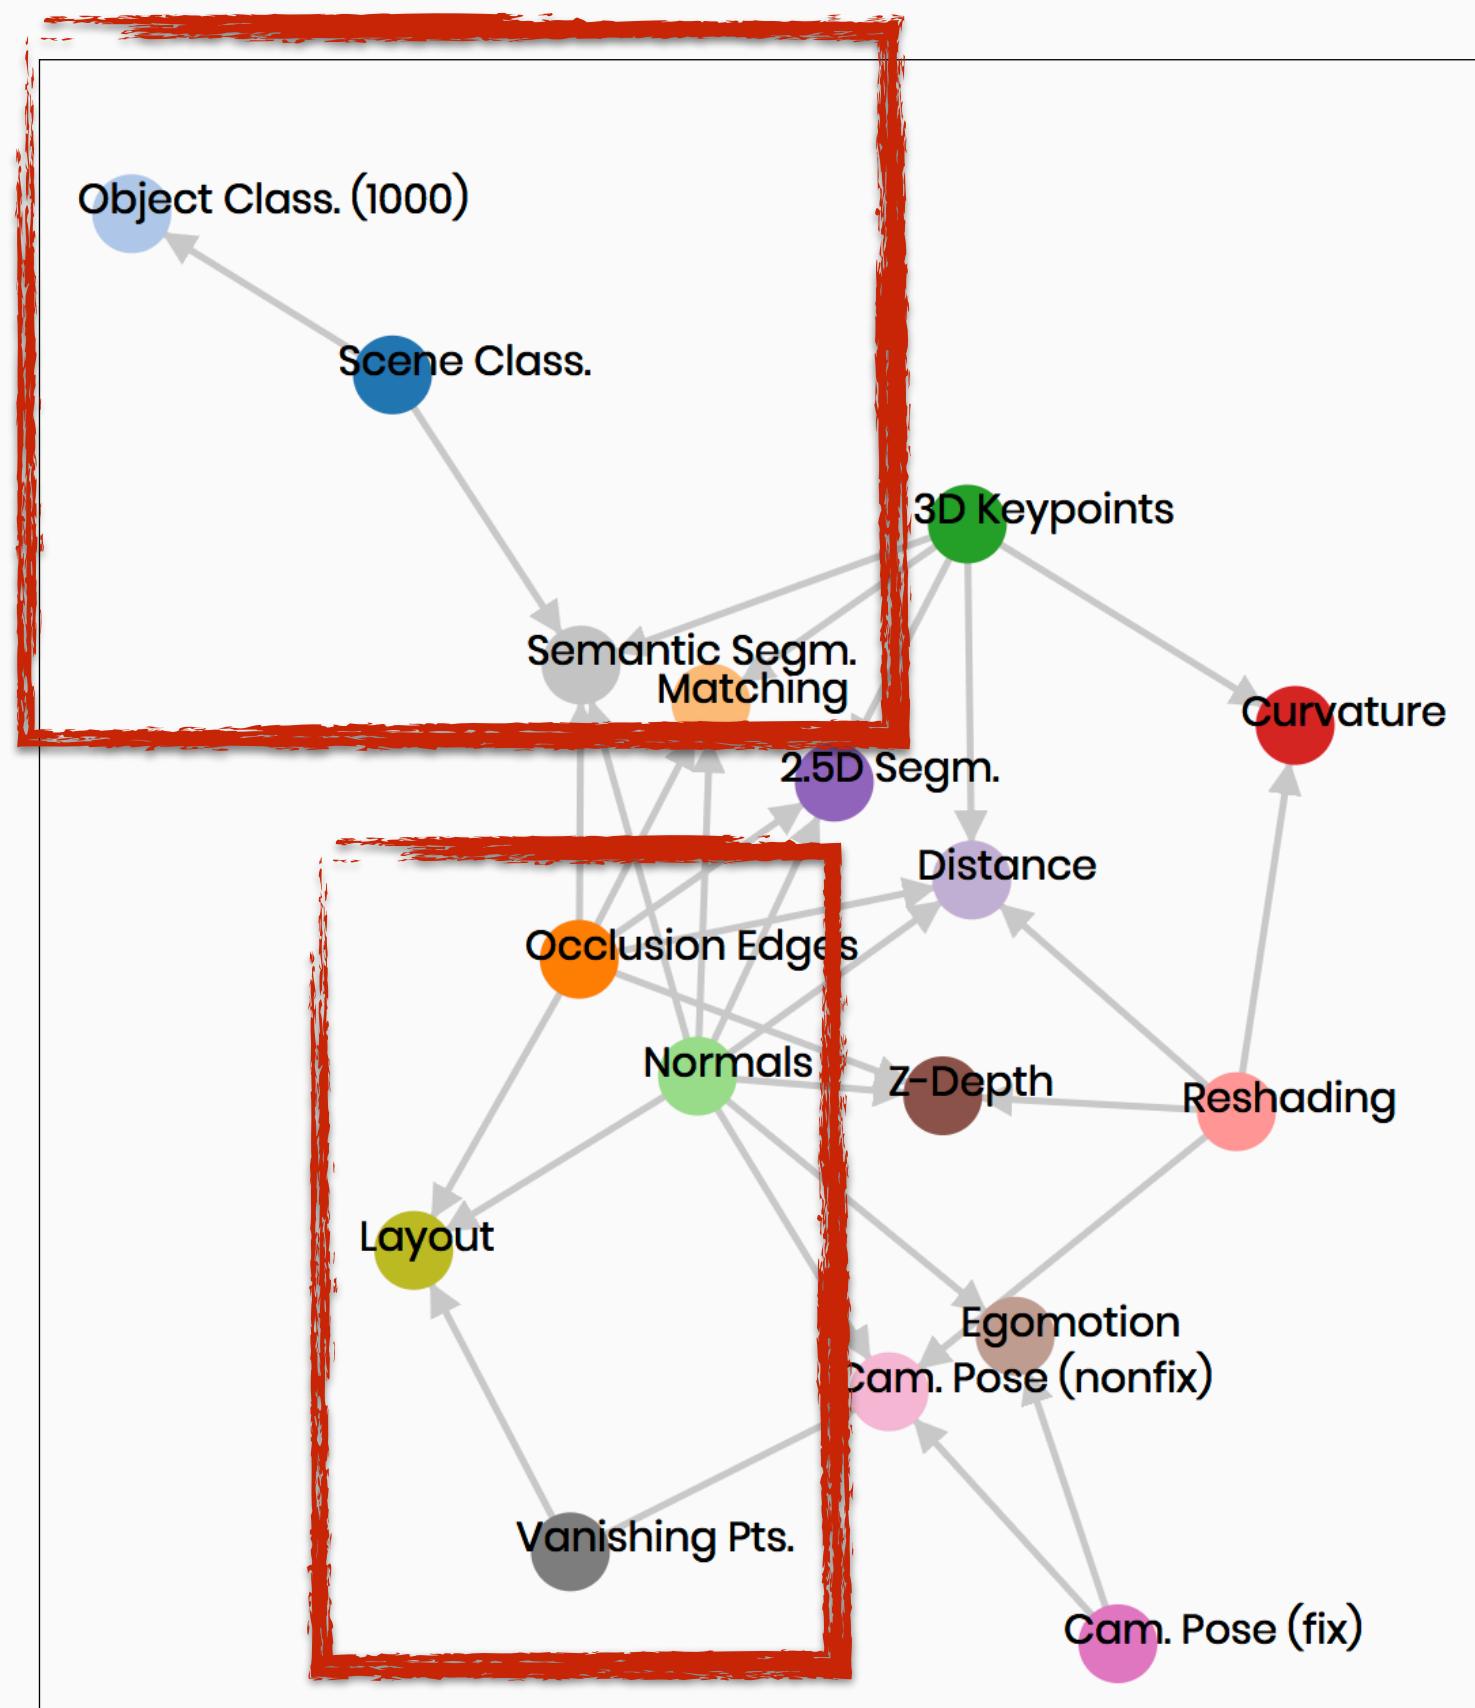
2nd order transfer



Experimental Results

- 26 Task-Specific Networks
- 3000 Transfer Networks
- 47,829 GPU hours 
- Transfers training data: 8x-120x less than task-specific

A Taxonomy

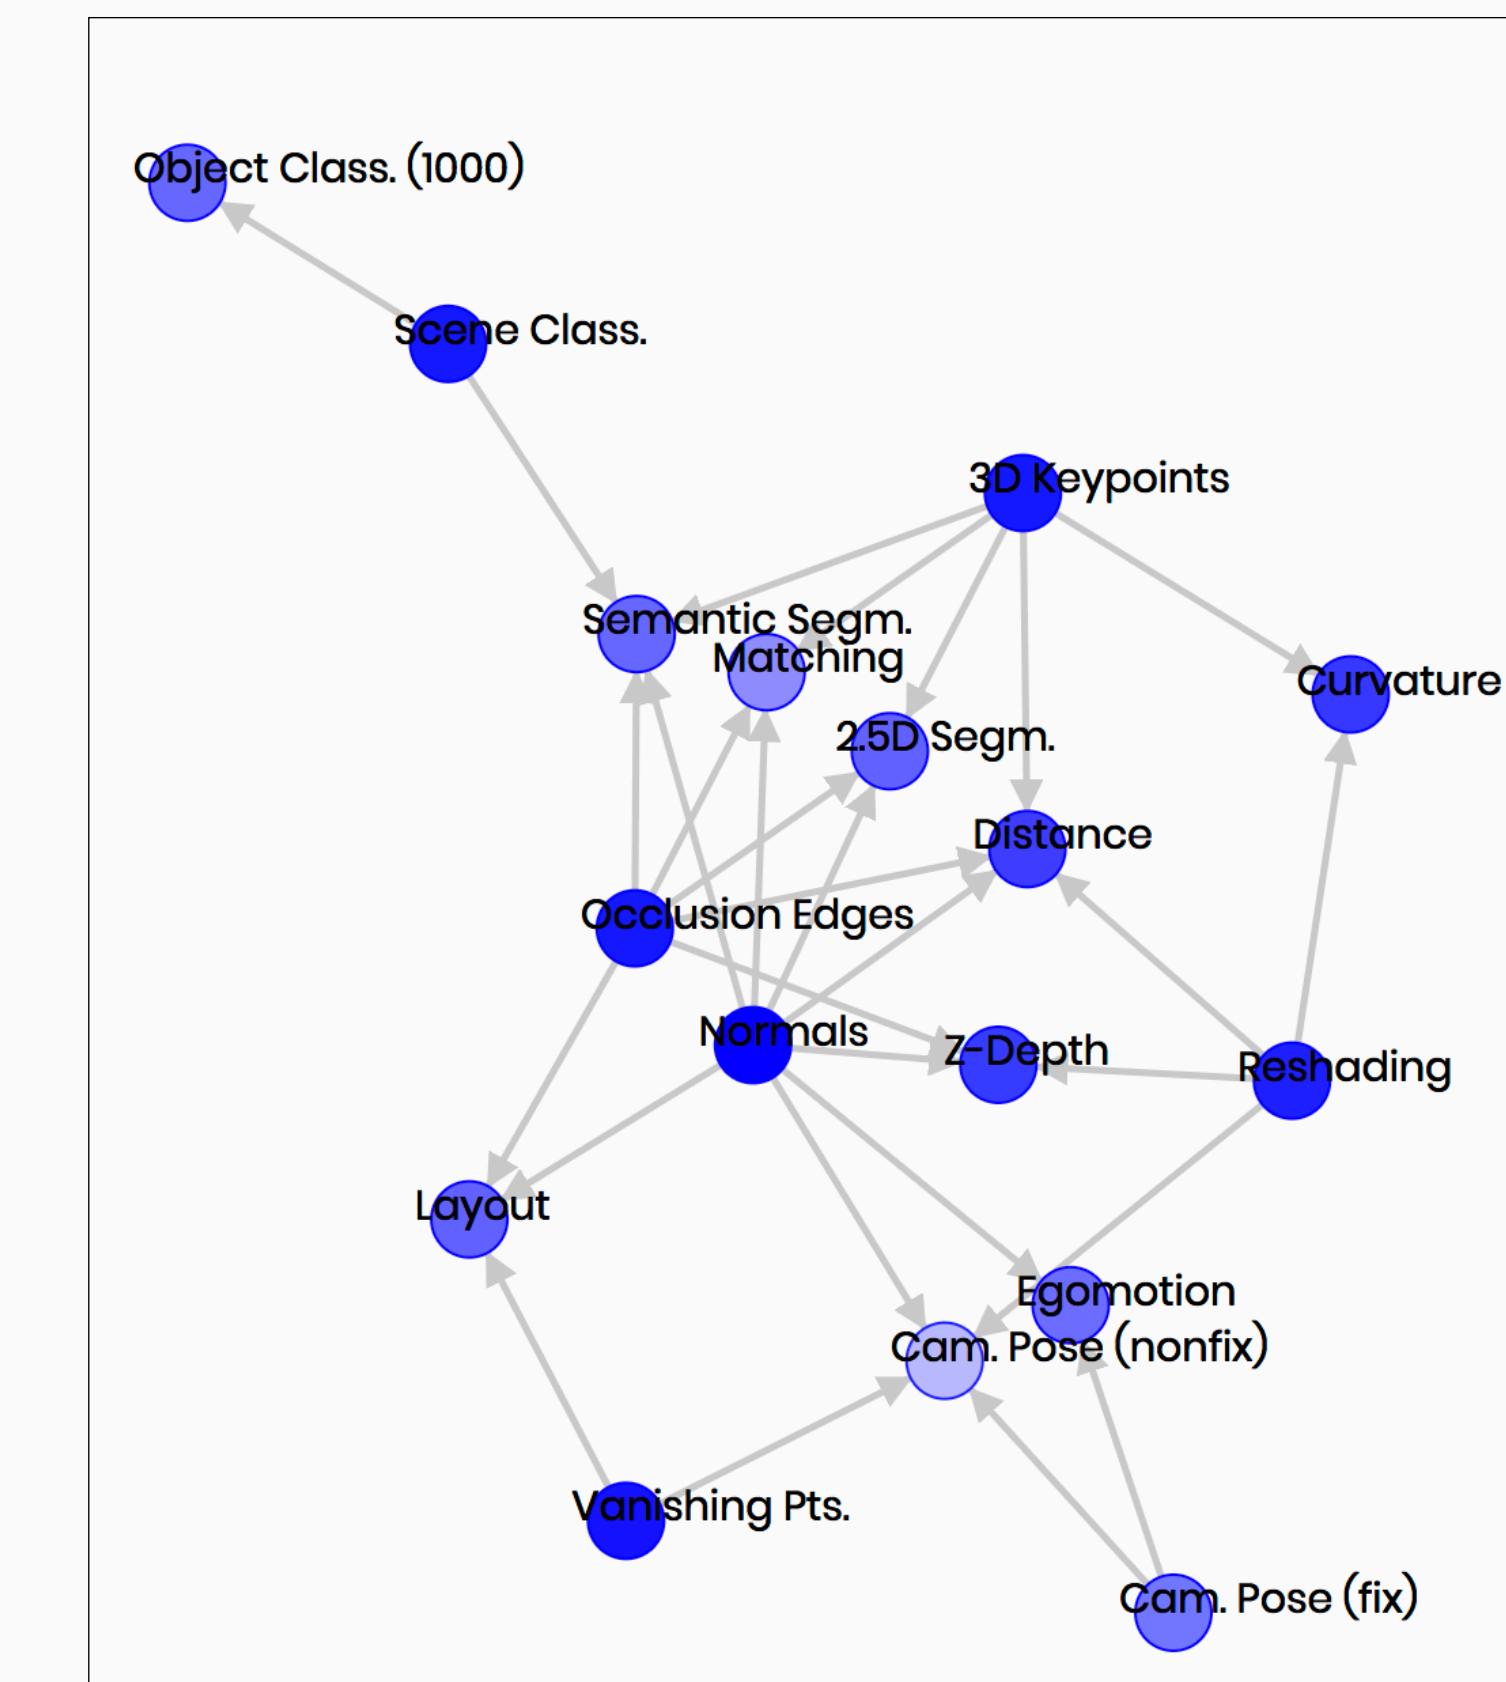
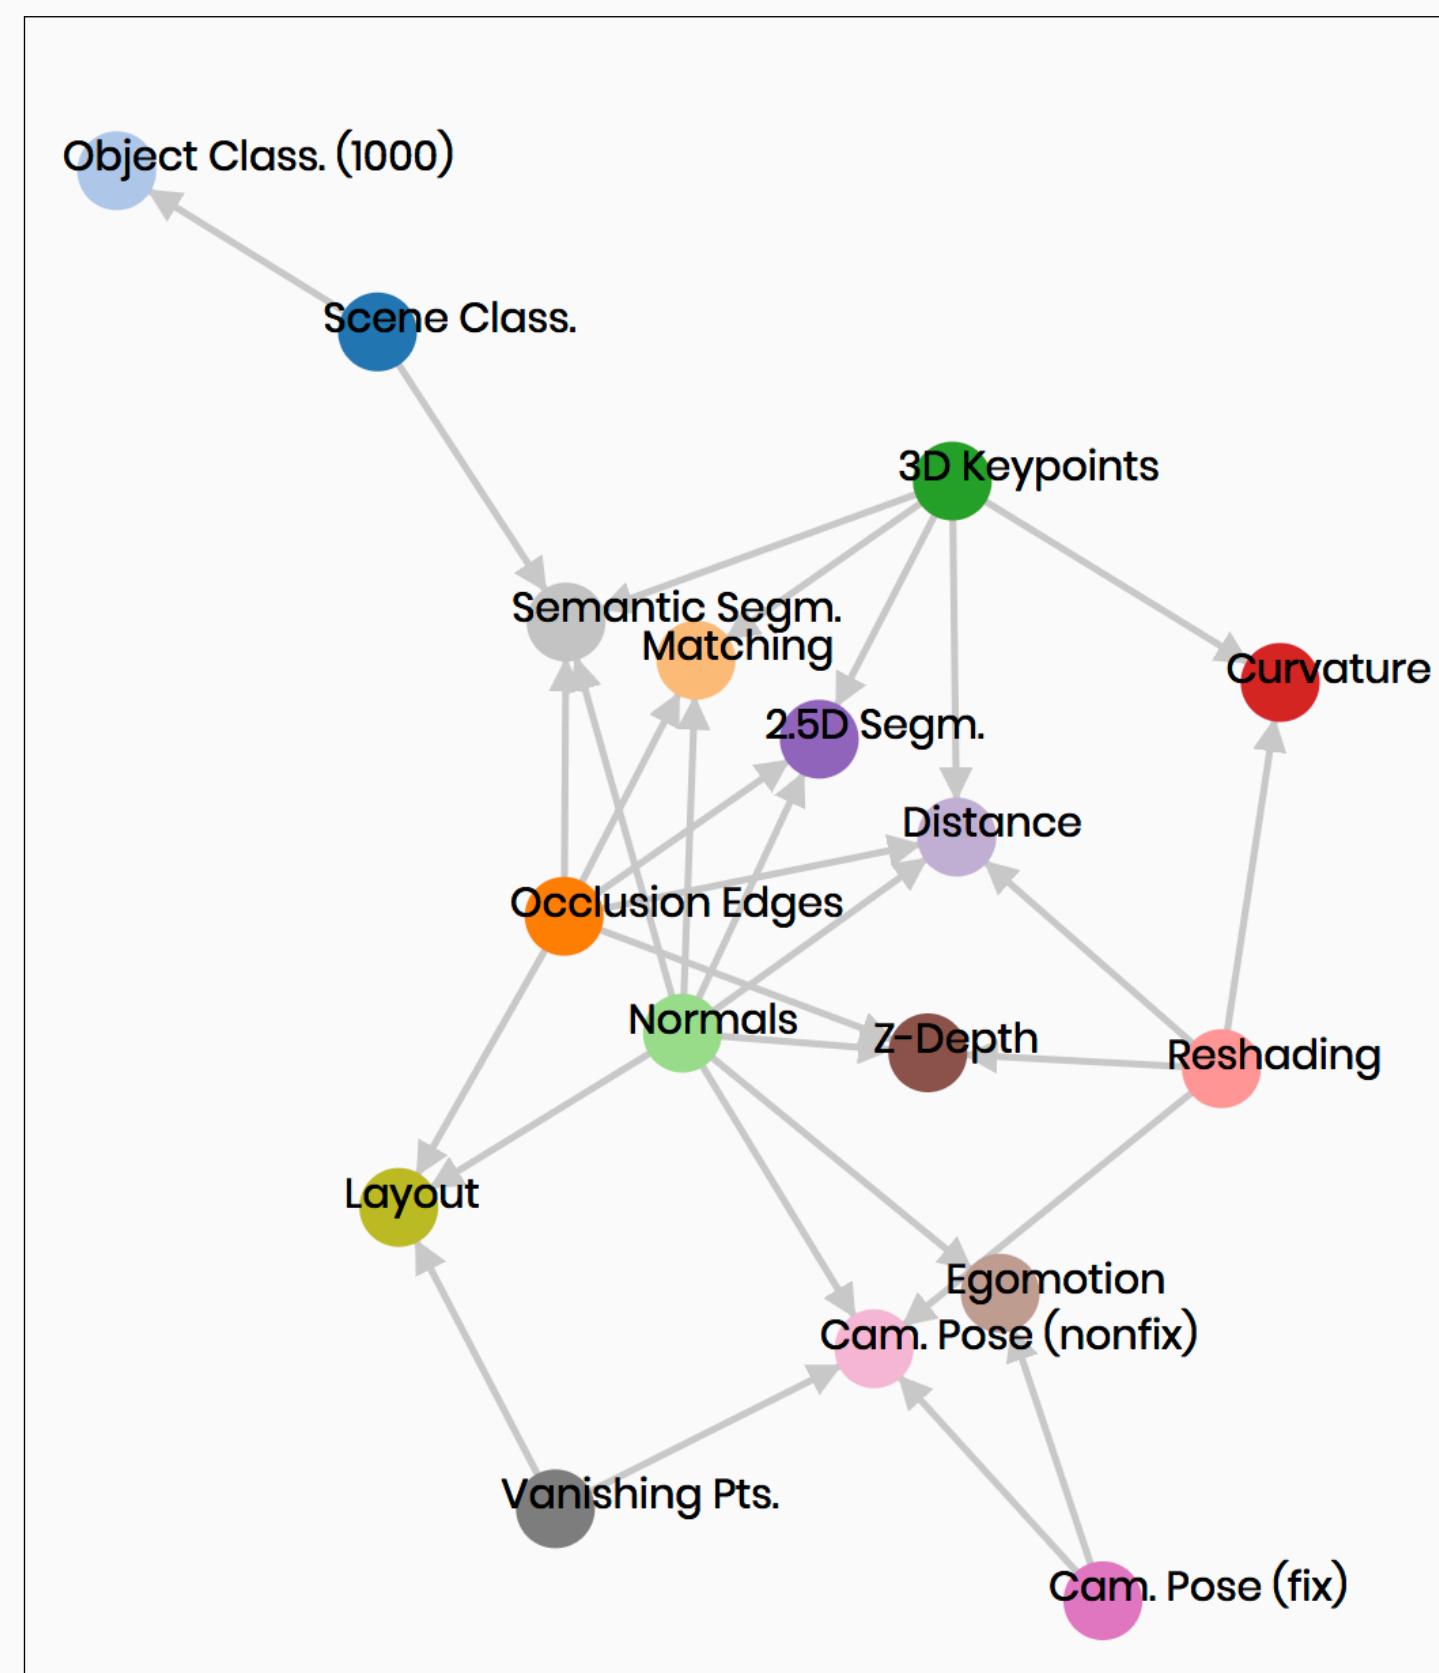


Target-only: Semantic Segmentation, Layout, Egomotion, 3D Curvature.

Source/Target: Depth, Occlusion Edges, Scene Classification, Normals, Camera Pose (non-fix), 3D Keypoint, Object Classification, Z-depth, Reshading, 3D Segmentation, Vanishing-Point, Camera pose (fix), Matching.

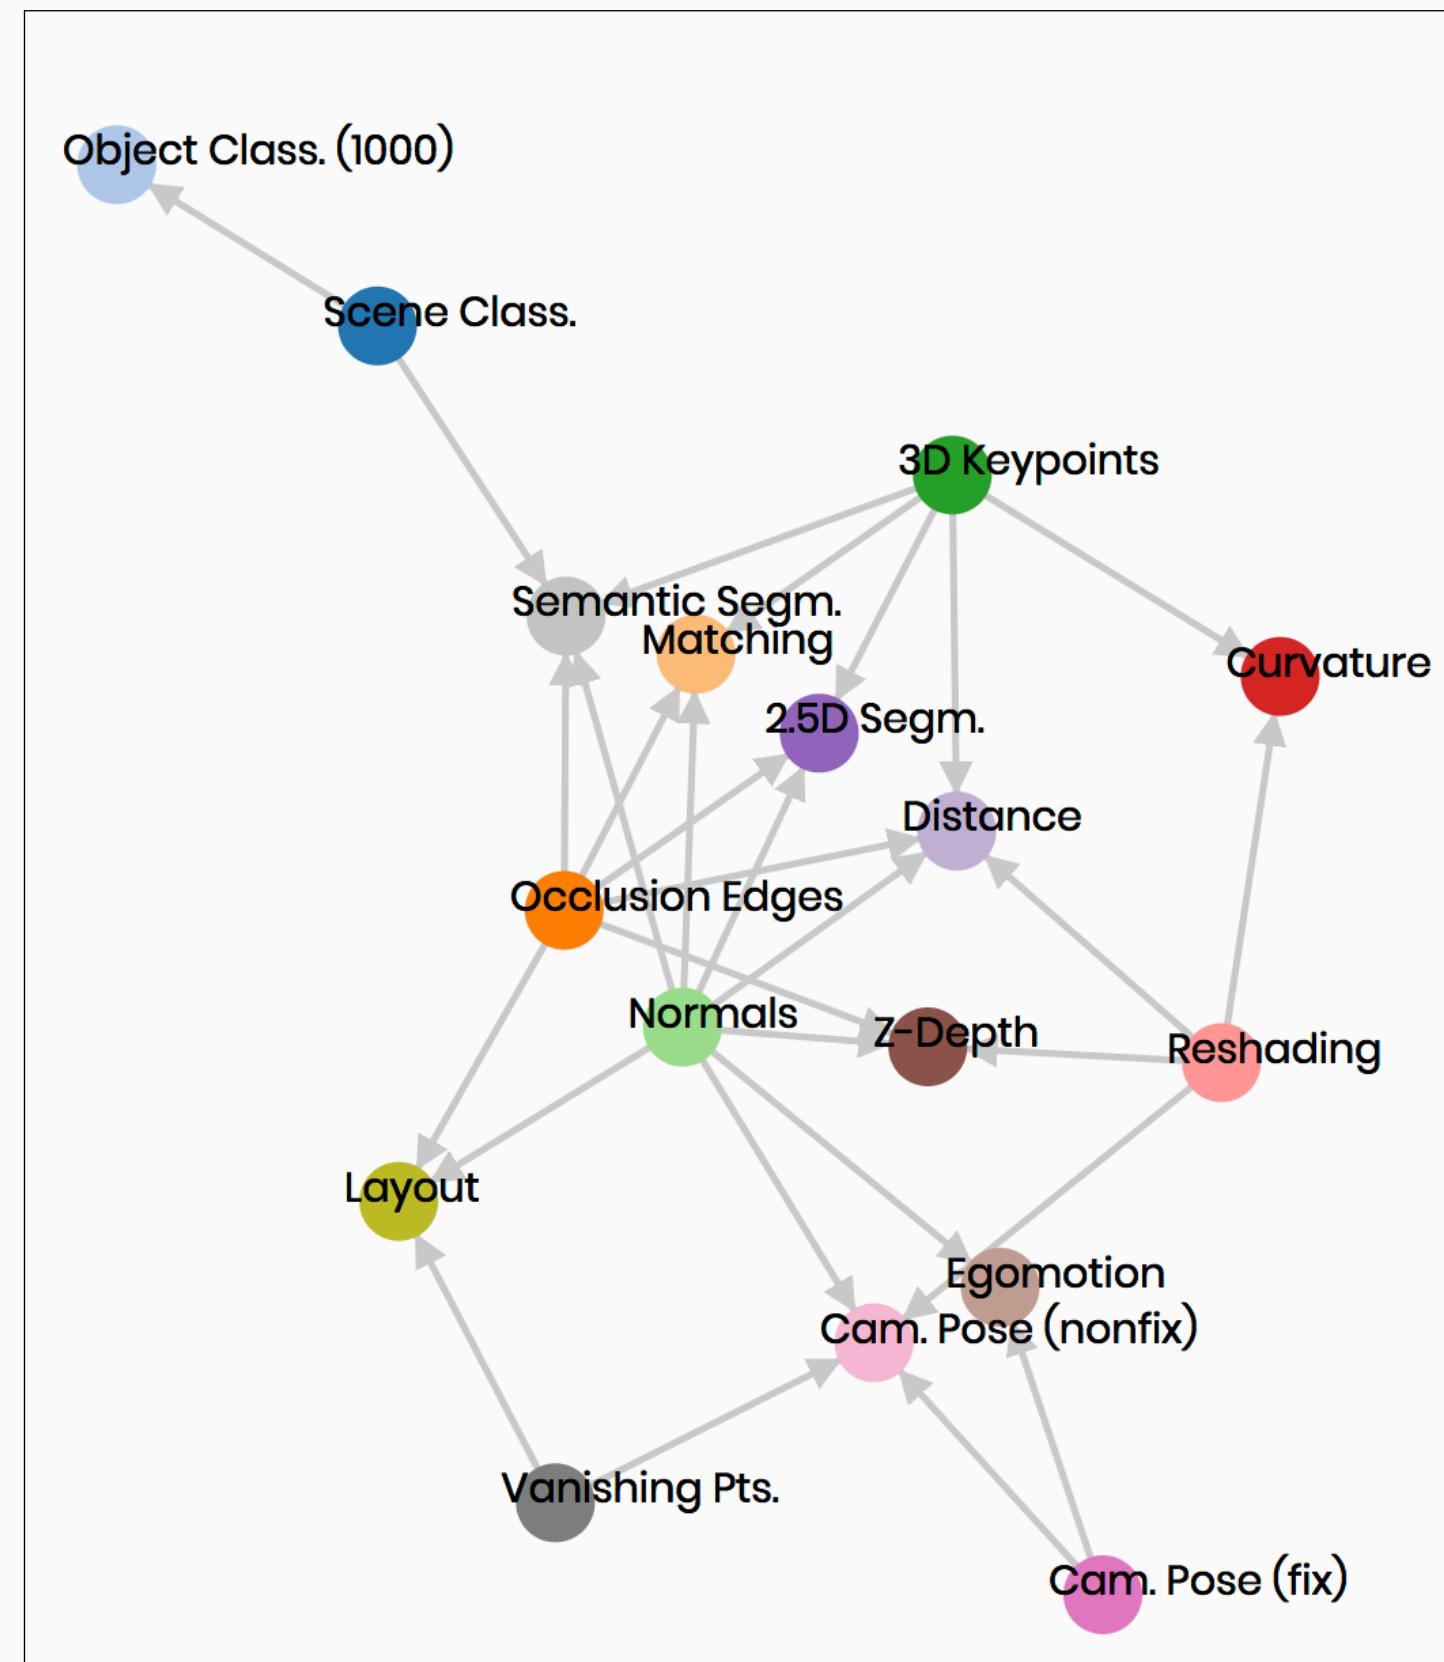
Source-only: Jigsaw, Colorization, Autoencoding, 2D edges, 2D edges, 2D keypoints, Denoising.

A Taxonomy

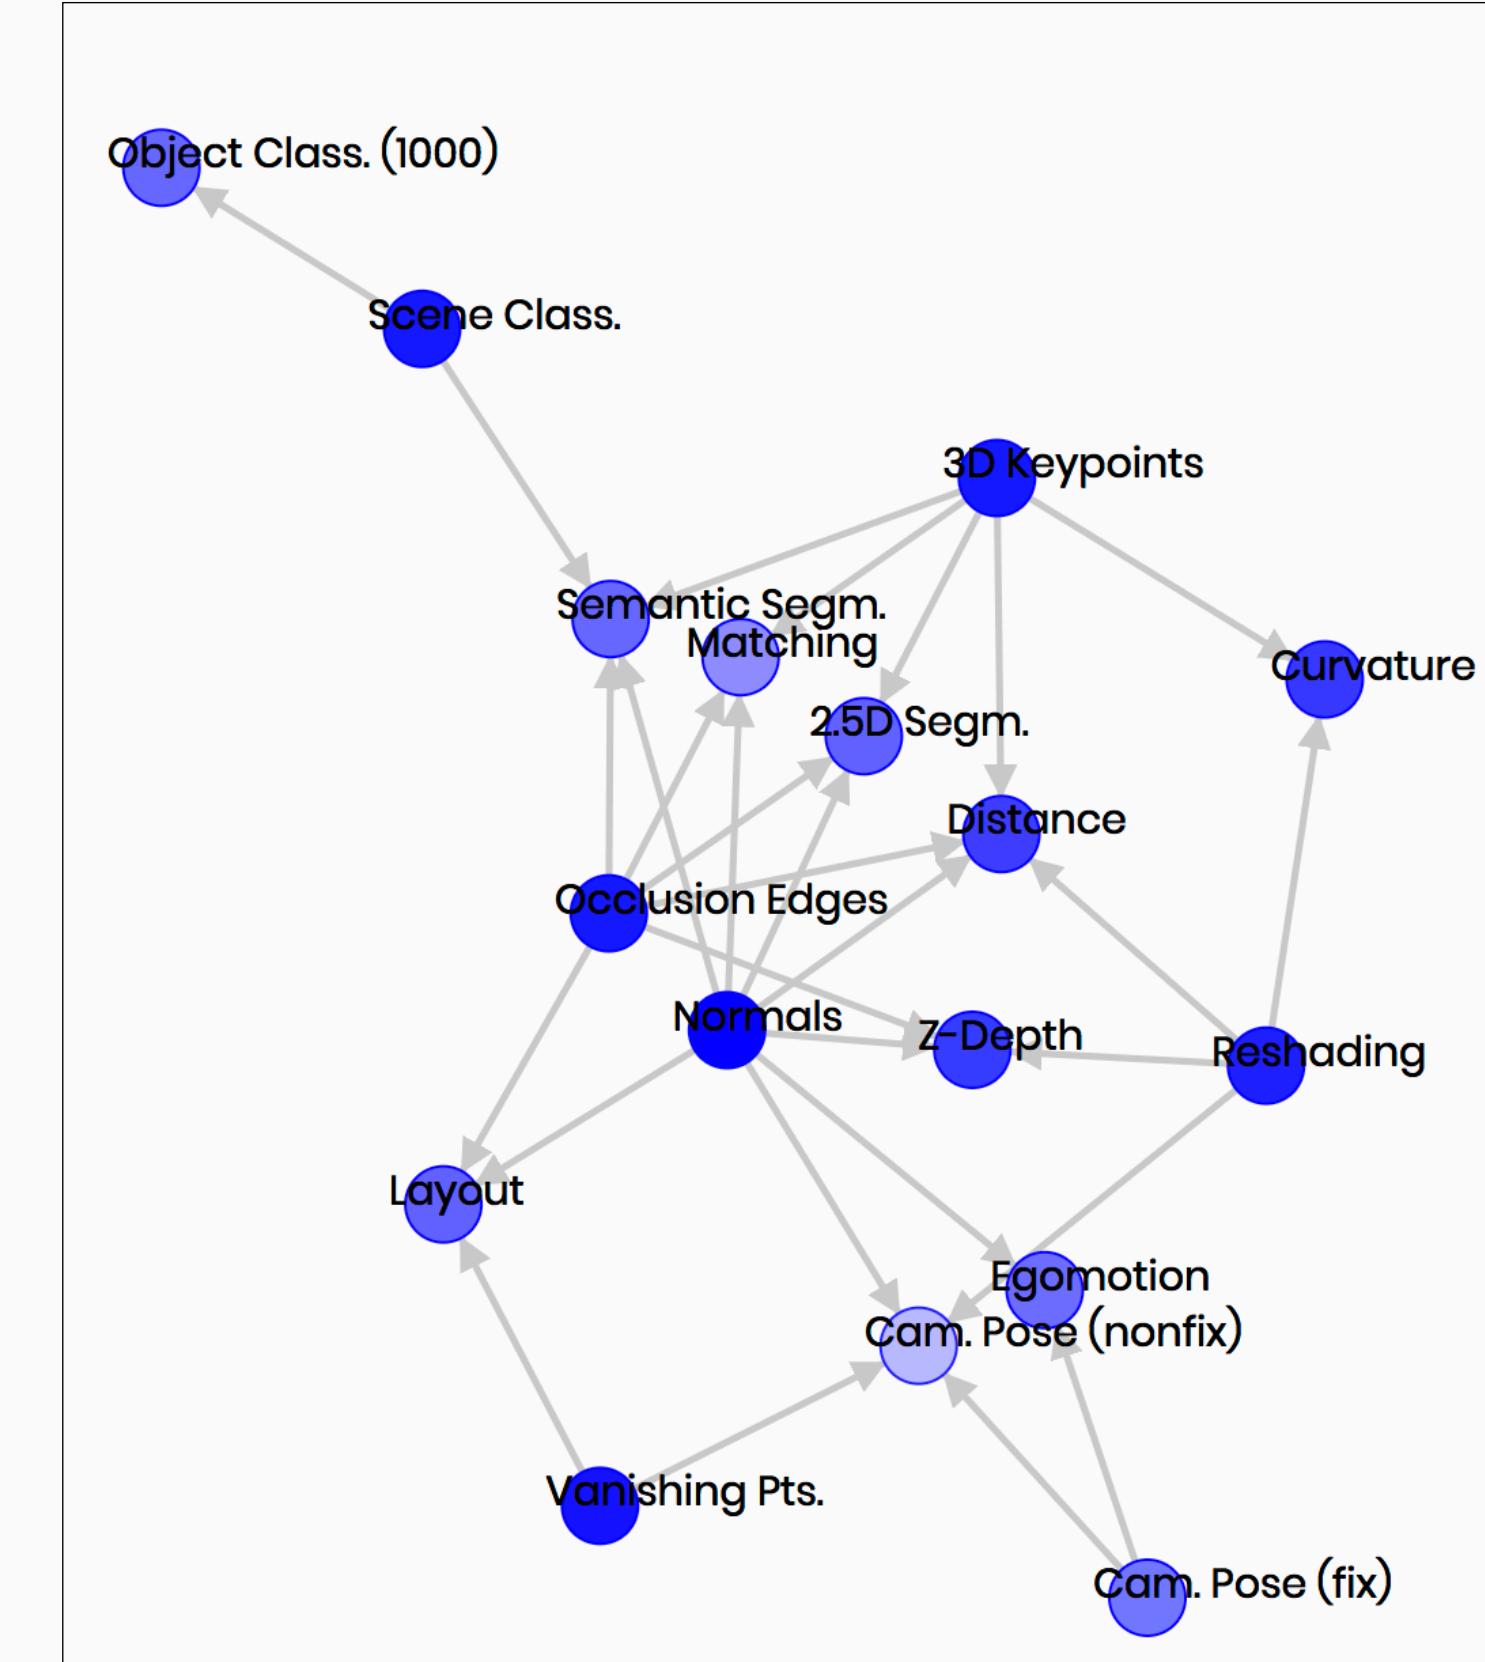


Gain: how much gained by transferring. (Win rate against a network trained from scratch using transfer data)

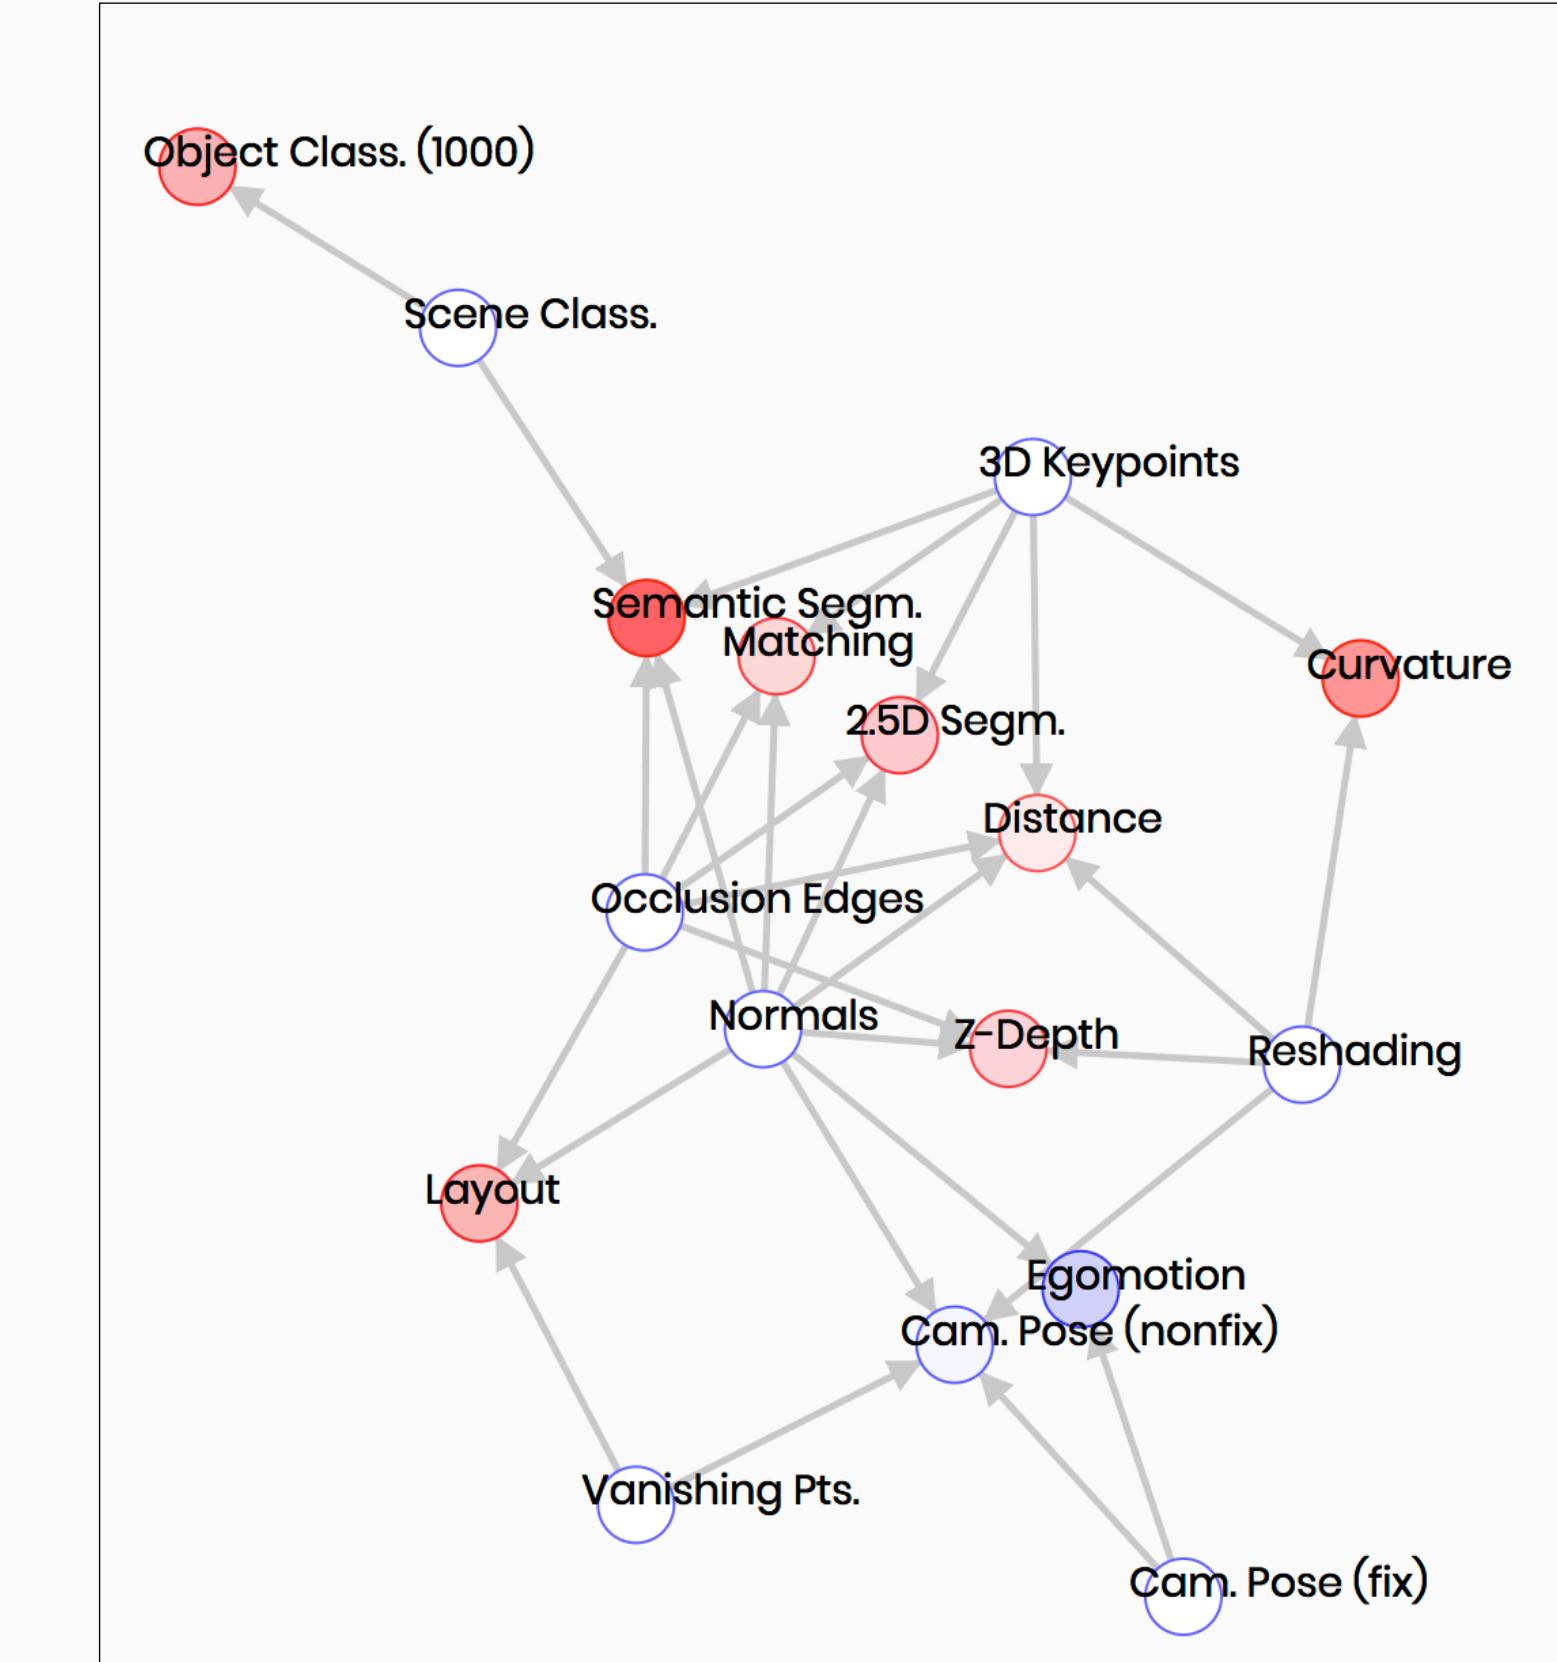
A Taxonomy



Gain



Quality

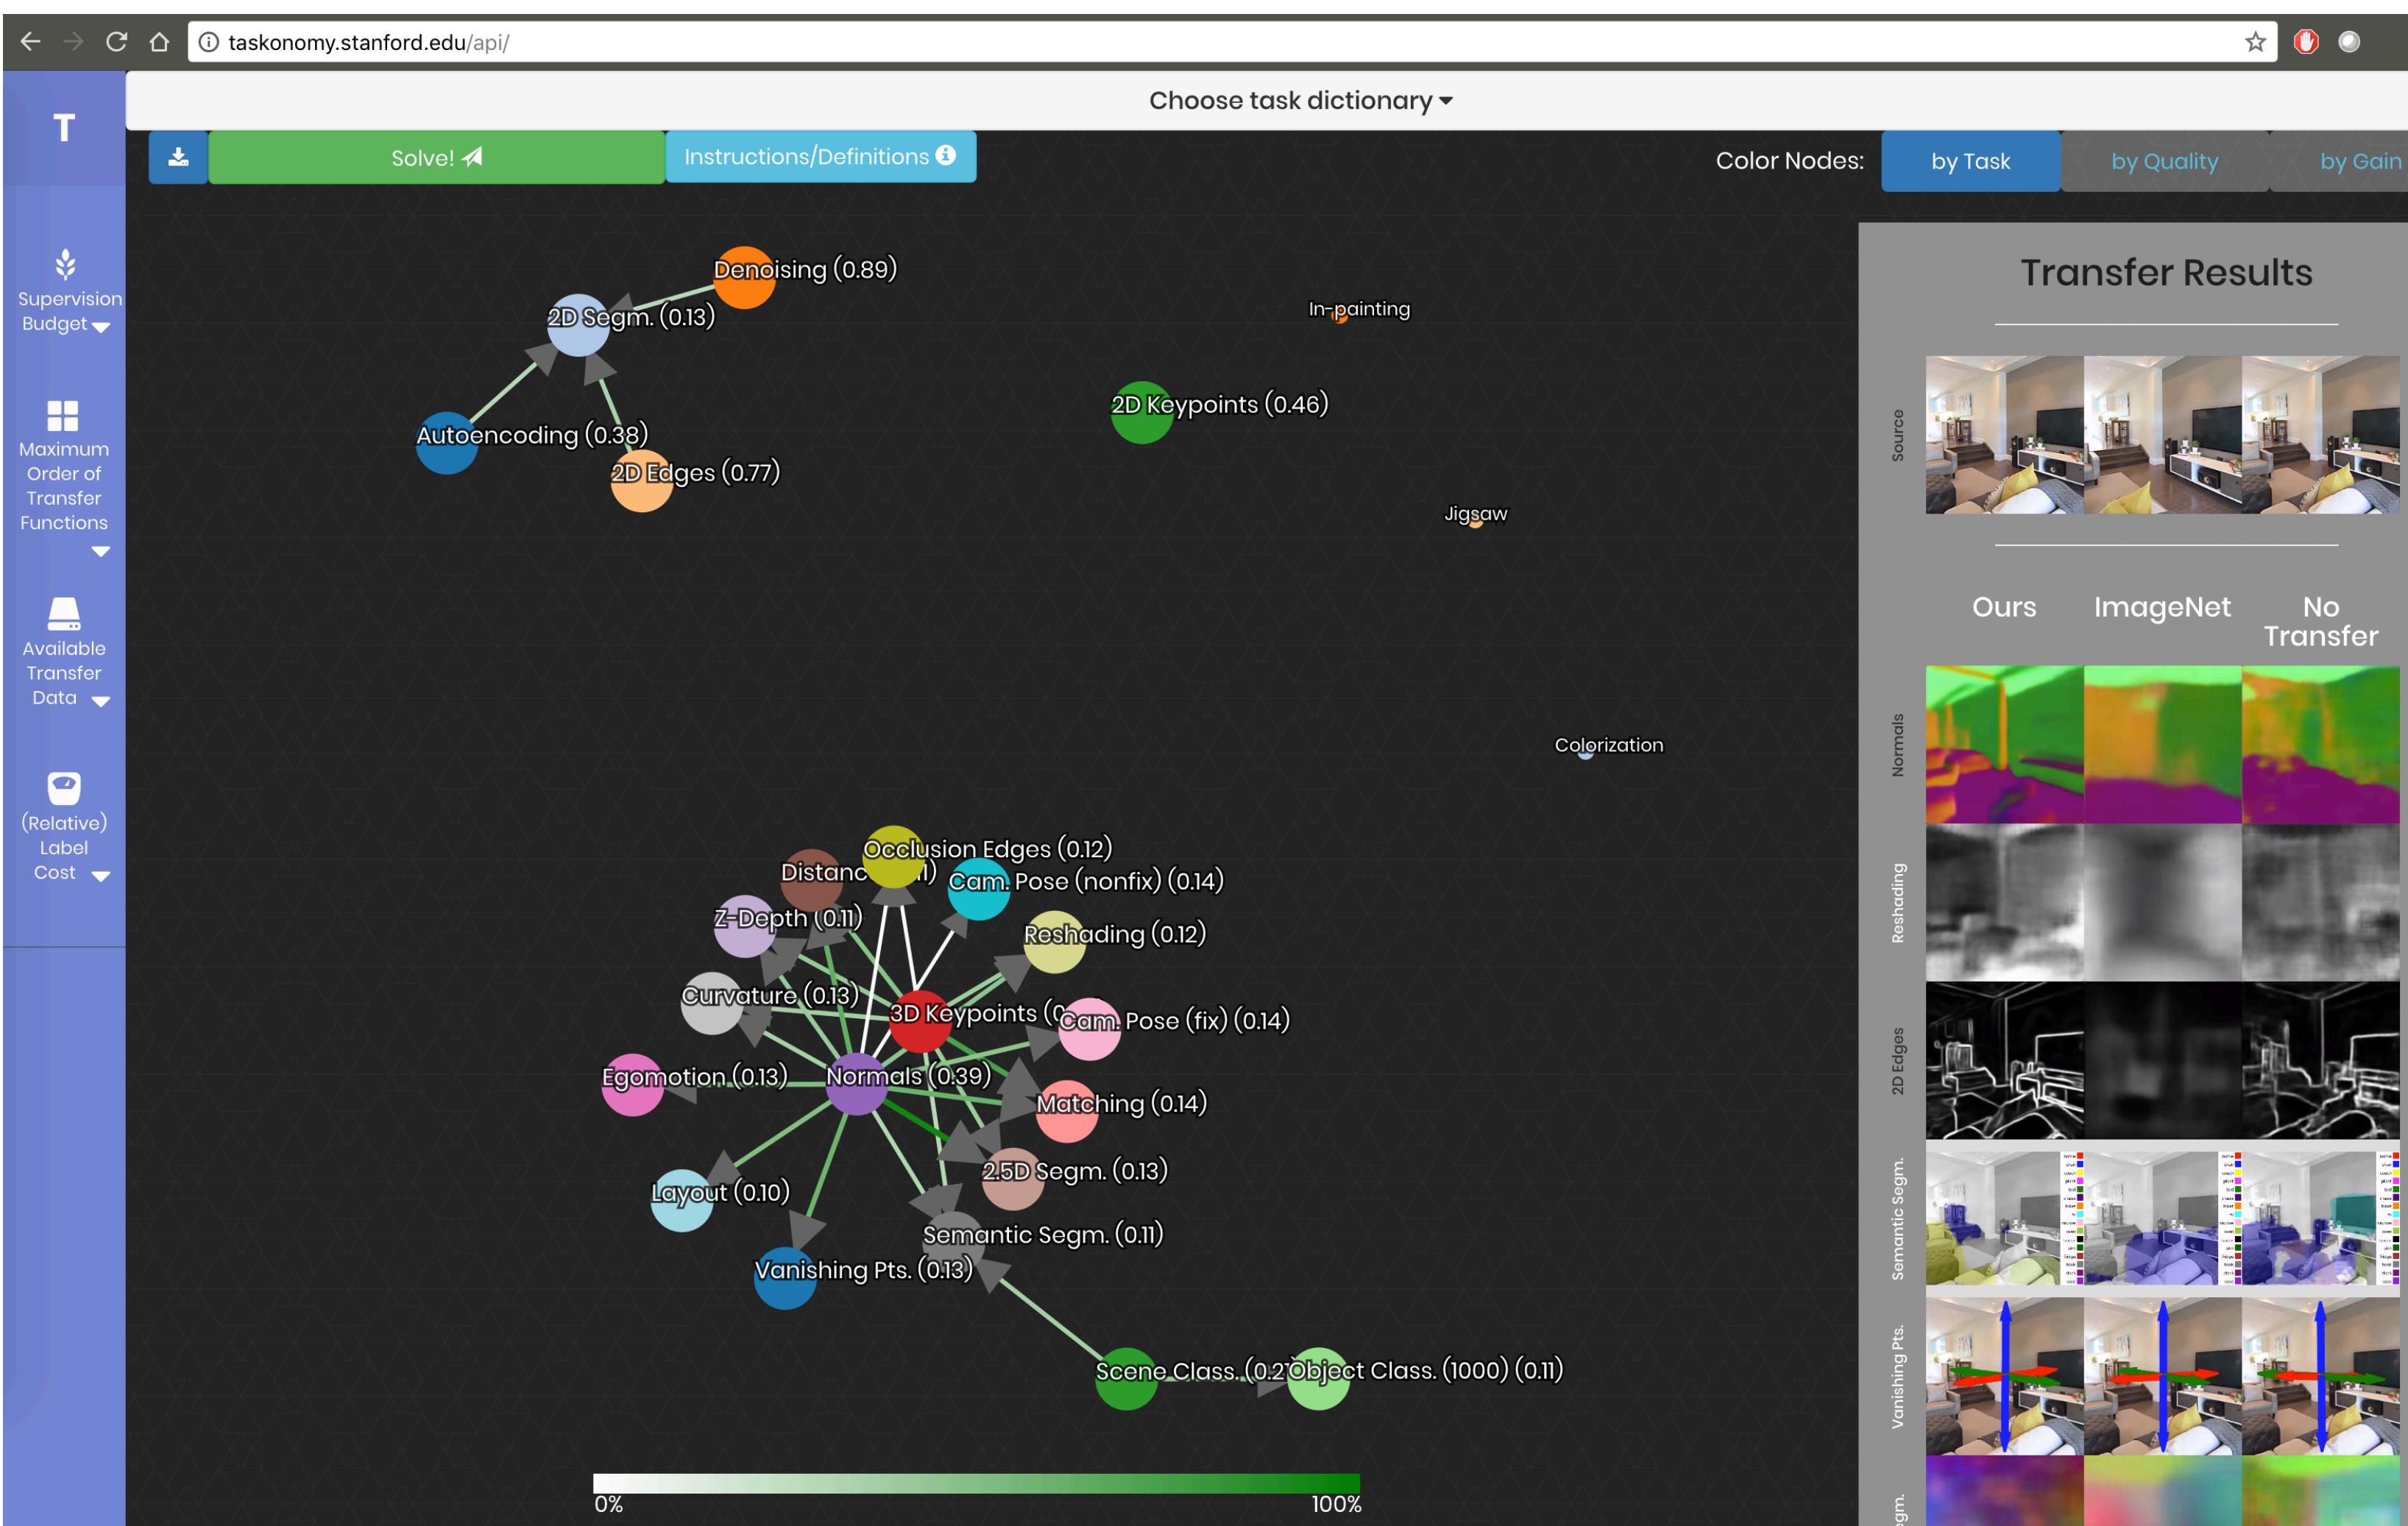


Gain: how much gained by transferring. (Win rate against a network trained from scratch using transfer data)

Quality: how close to gold-standard task-specific networks. (win rate against a fully supervised network)

Live Web API: <http://taskonomy.vision/api/>

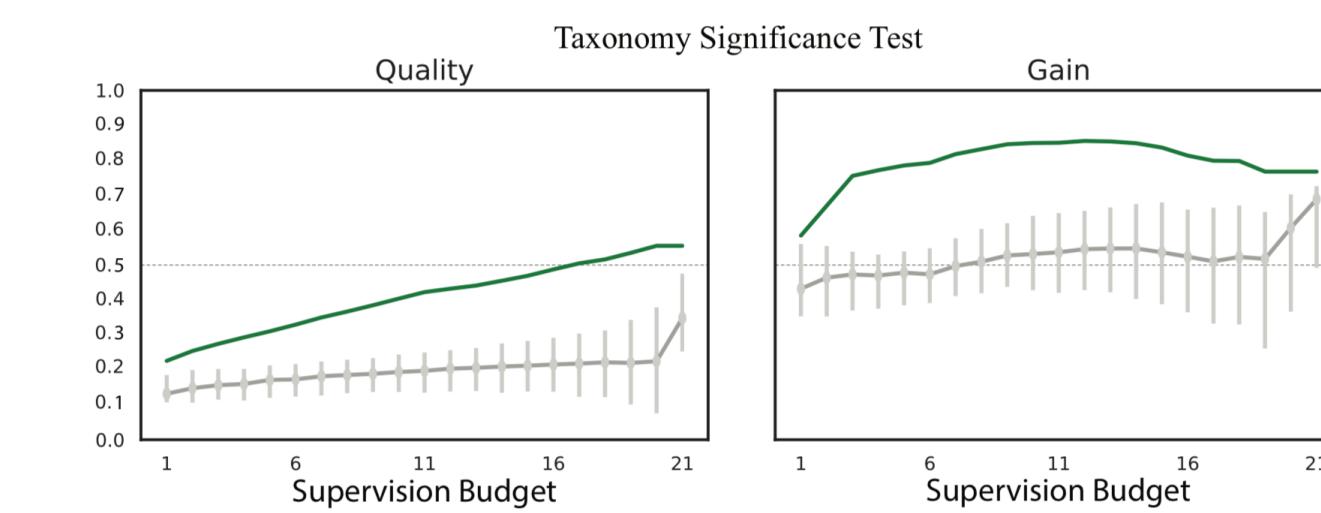
Run taskonomy with any arguments. See qualitative and quantitate results.



Experiments in paper/poster

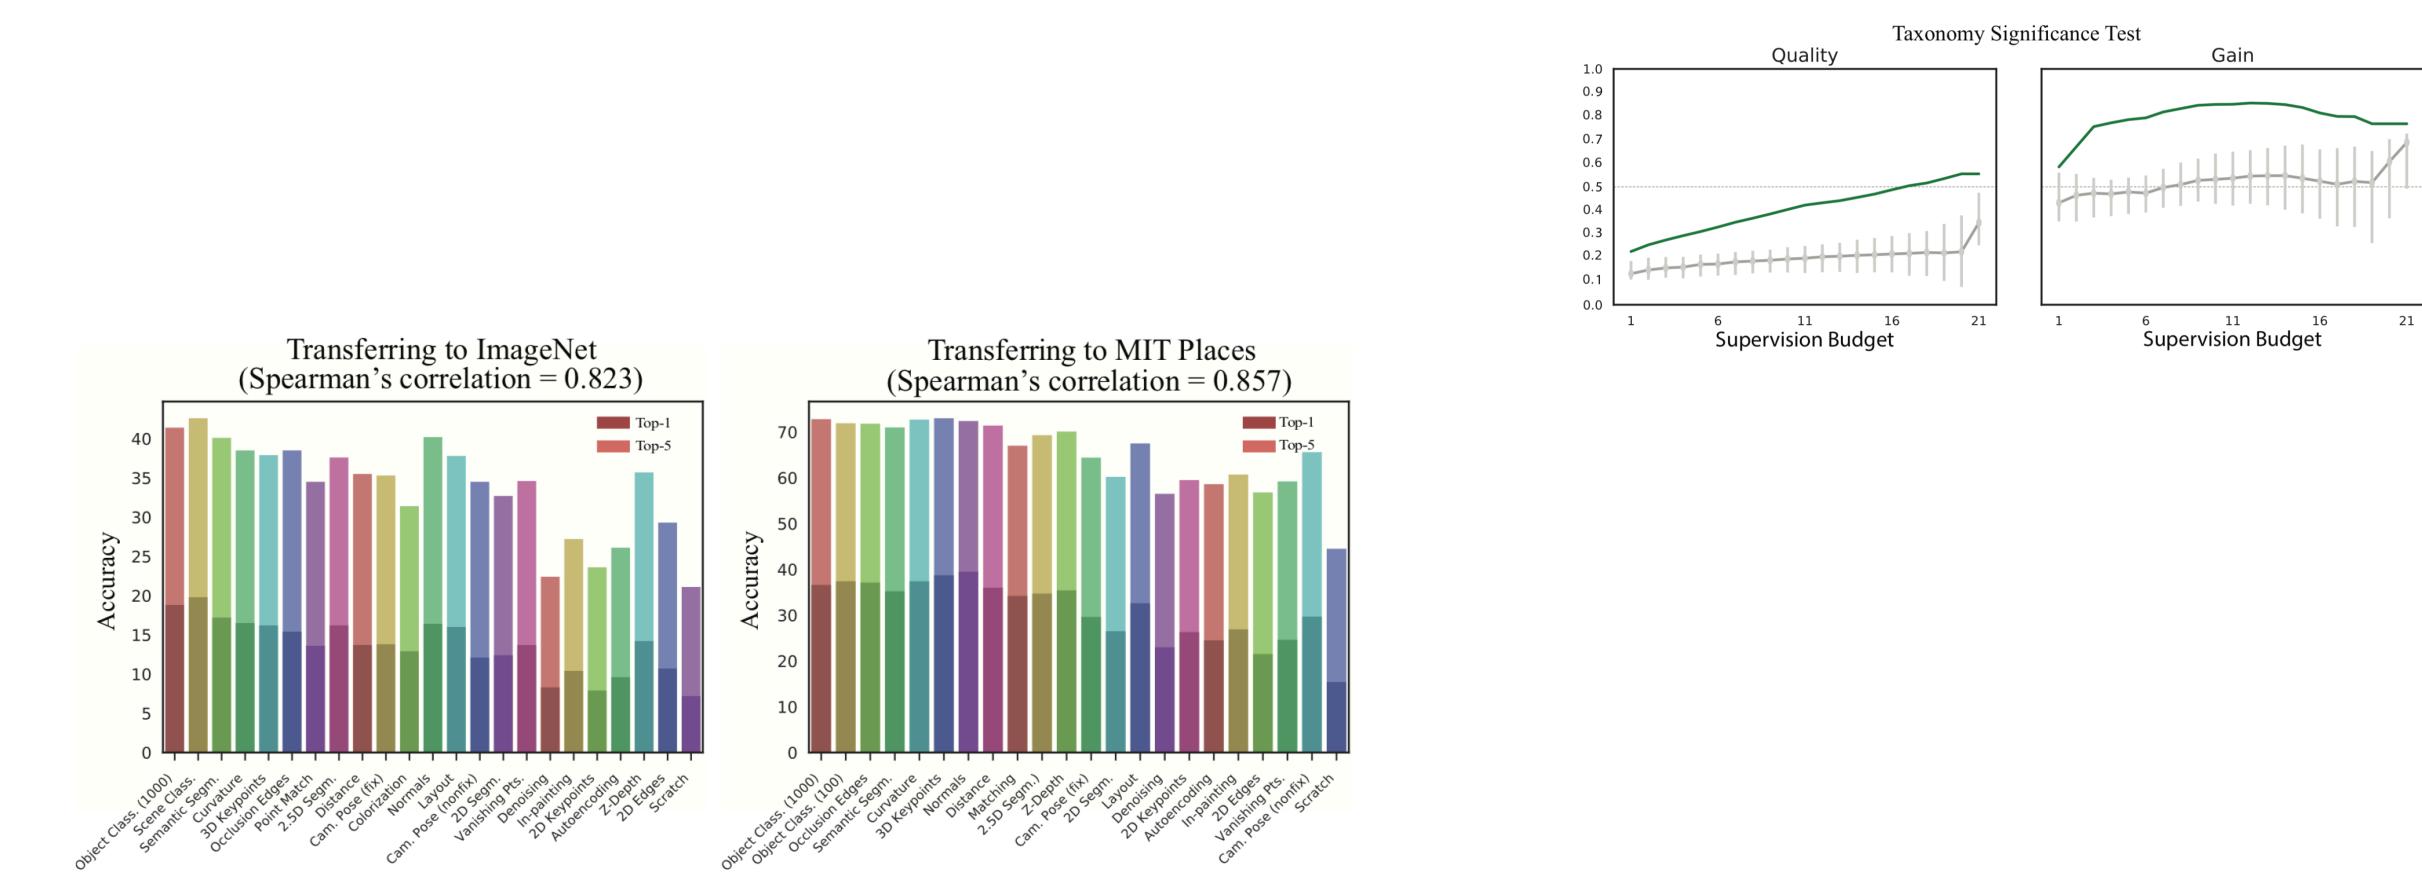
Experiments in paper/poster

- **Significance Test**



Experiments in paper/poster

- Significance Test
 - Trends on ImageNet and MIT Places



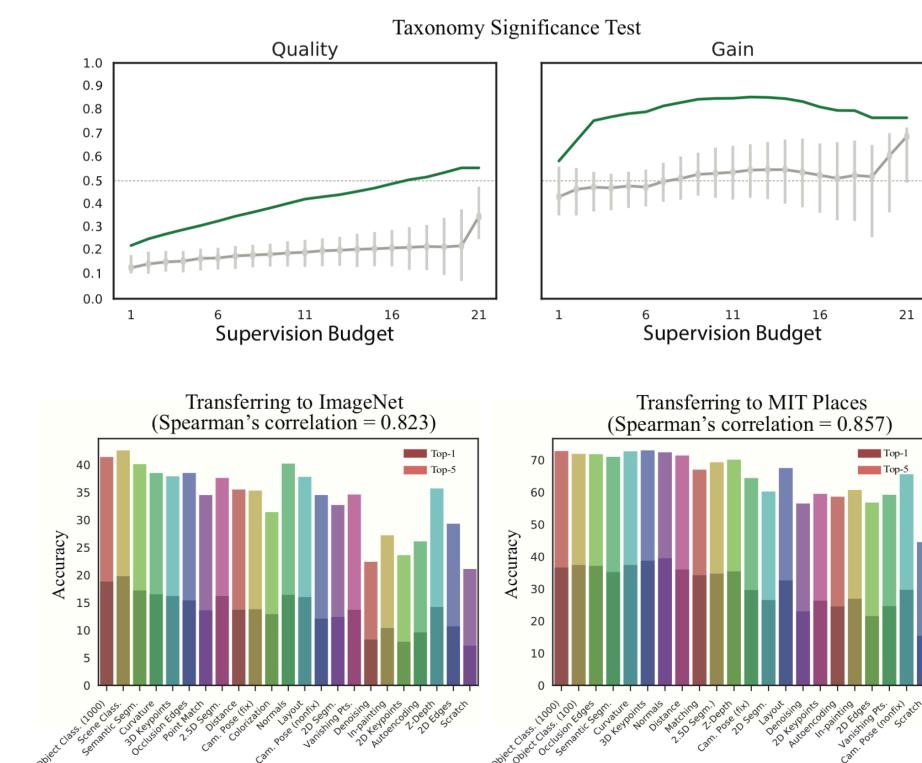
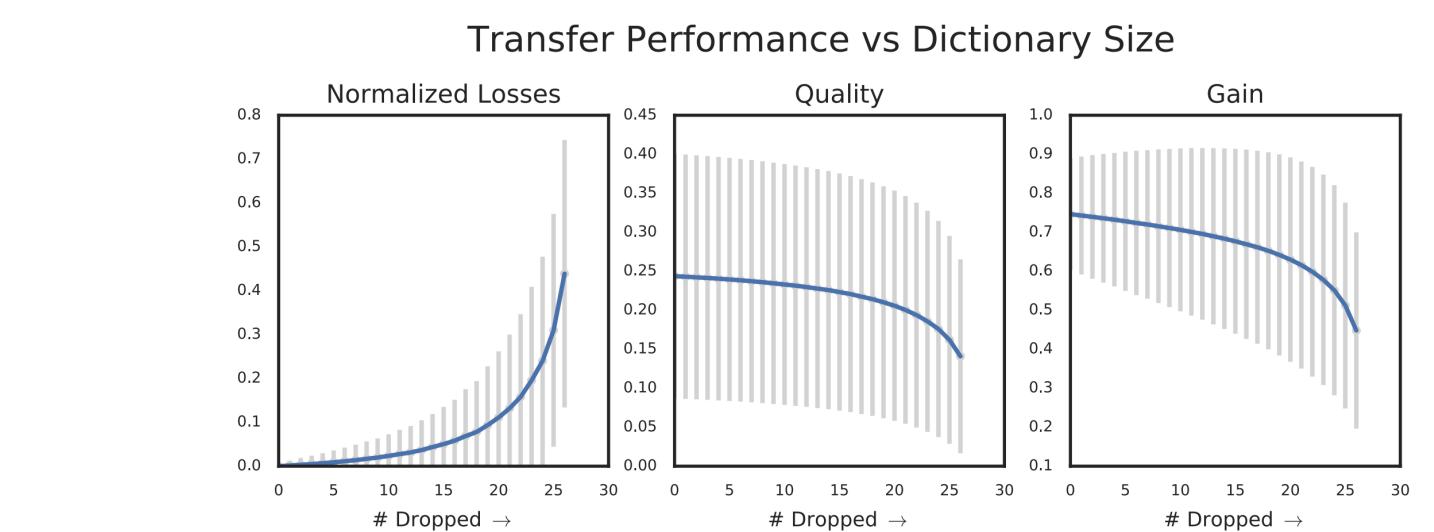
Experiments in paper/poster

- **Significance Test**
- **Trends on ImageNet and MIT Places**
- **Sensitivity to choice of dictionary**

Zamir et al., 2018
arXiv:1803.08941

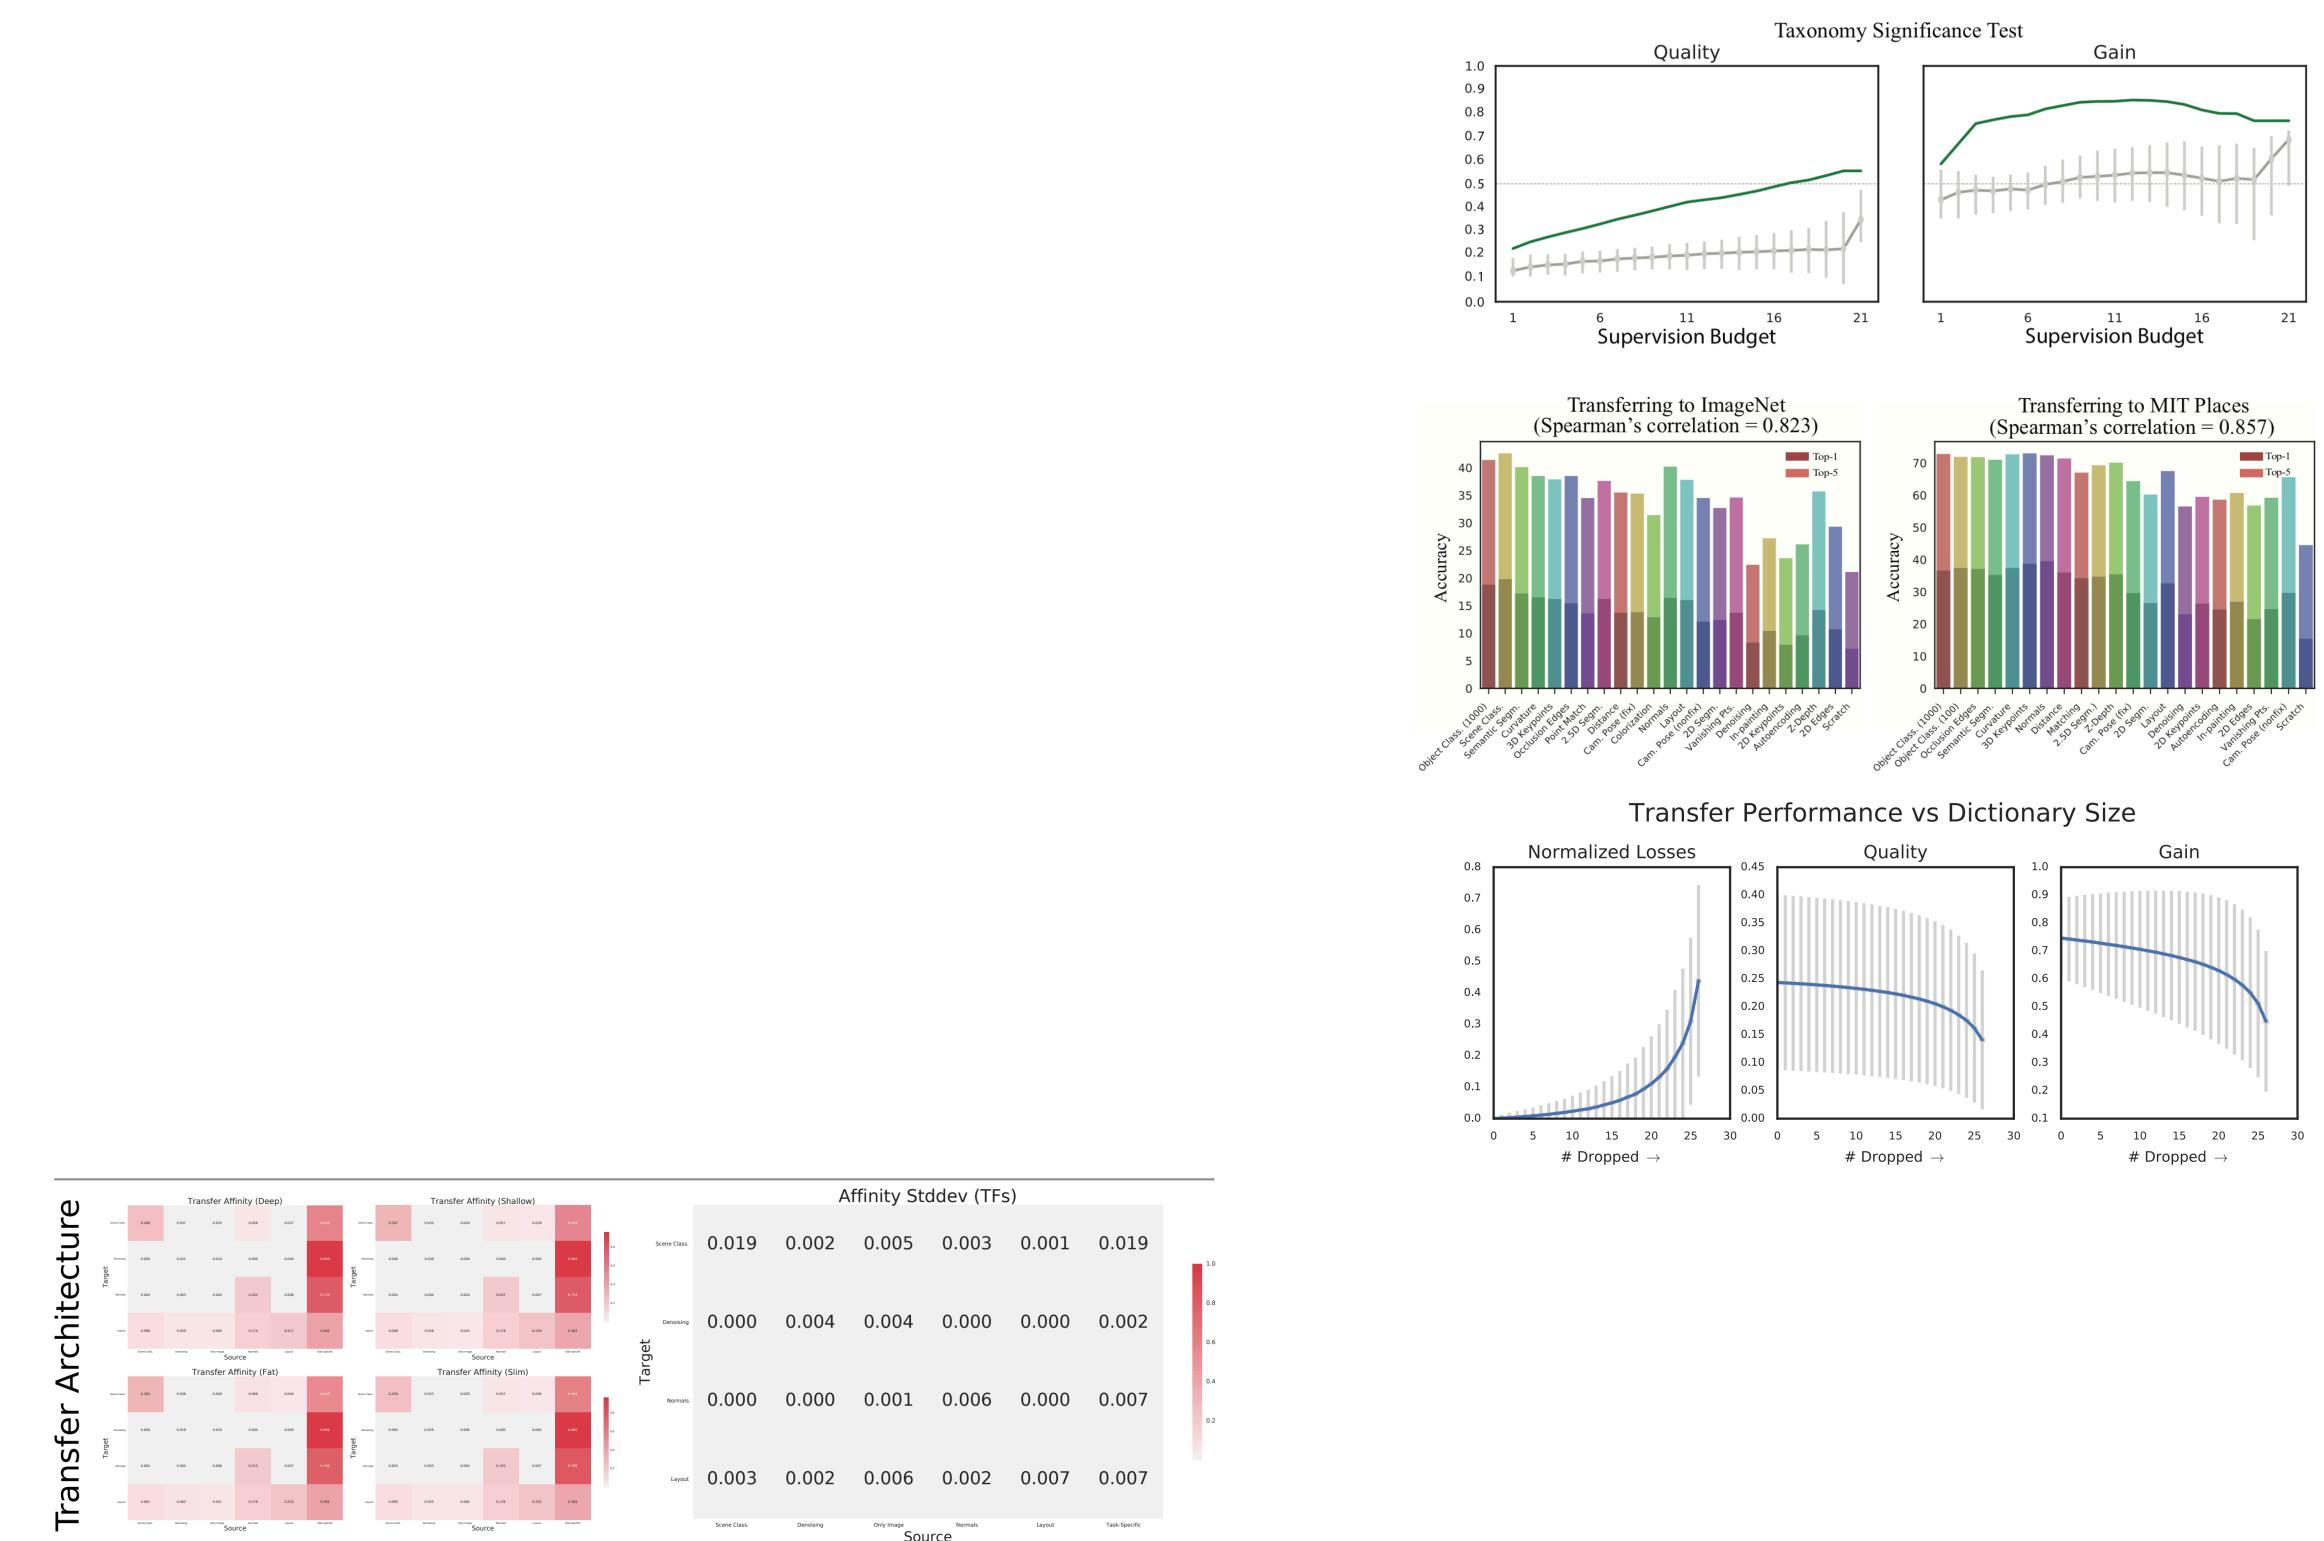
Taskonomy
CVPR 2018

Taskonomy
CVPR 2018



Experiments in paper/poster

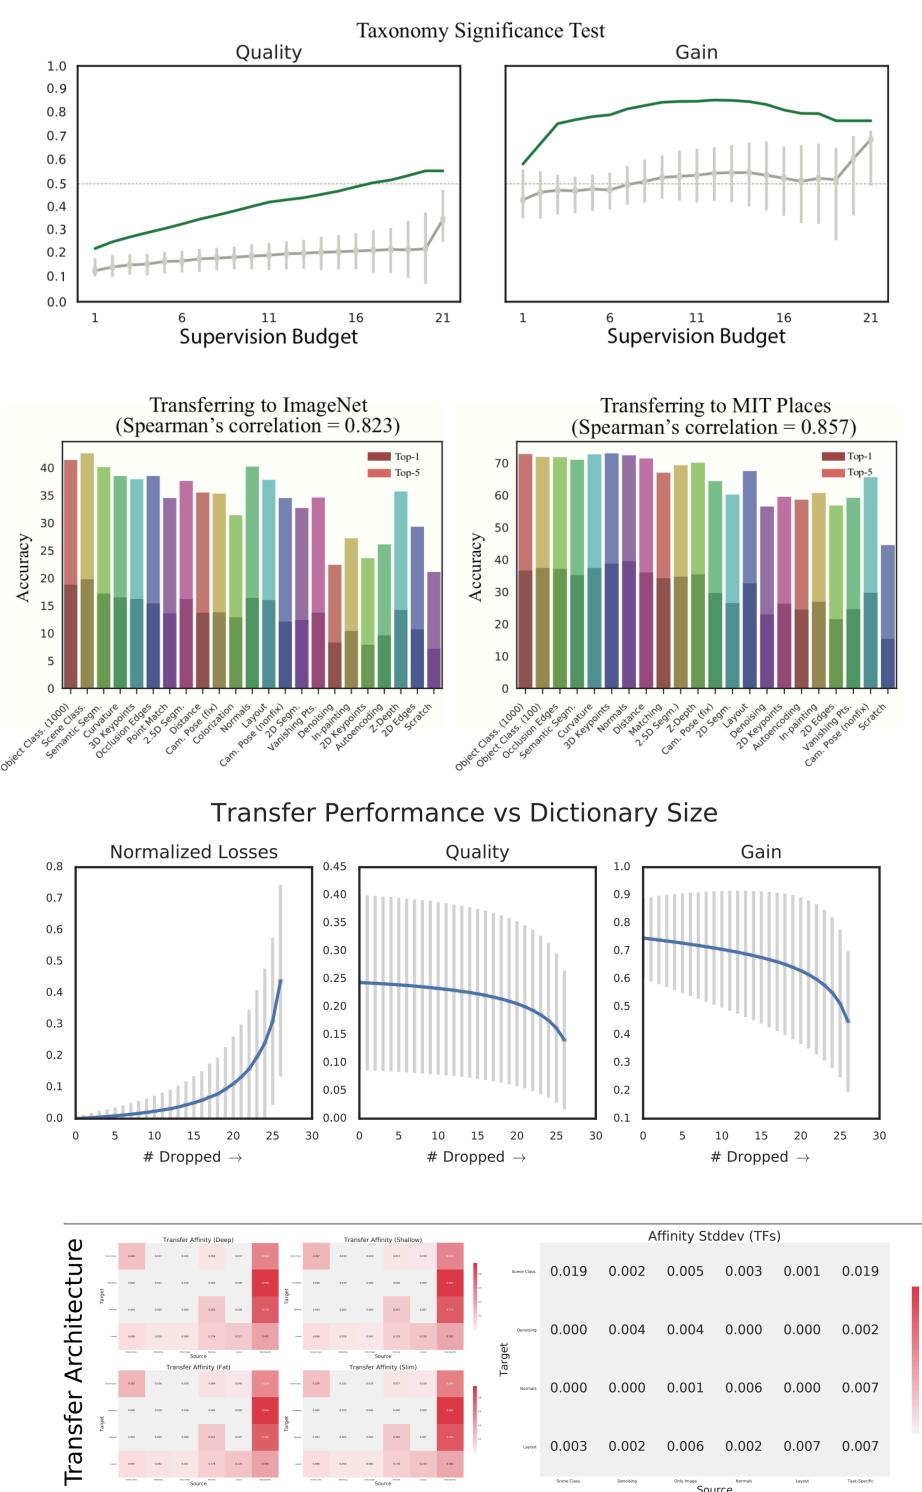
- **Significance Test**
- **Trends on ImageNet and MIT Places**
- **Sensitivity to choice of dictionary**
- **Architecture dependence**



Experiments in paper/poster

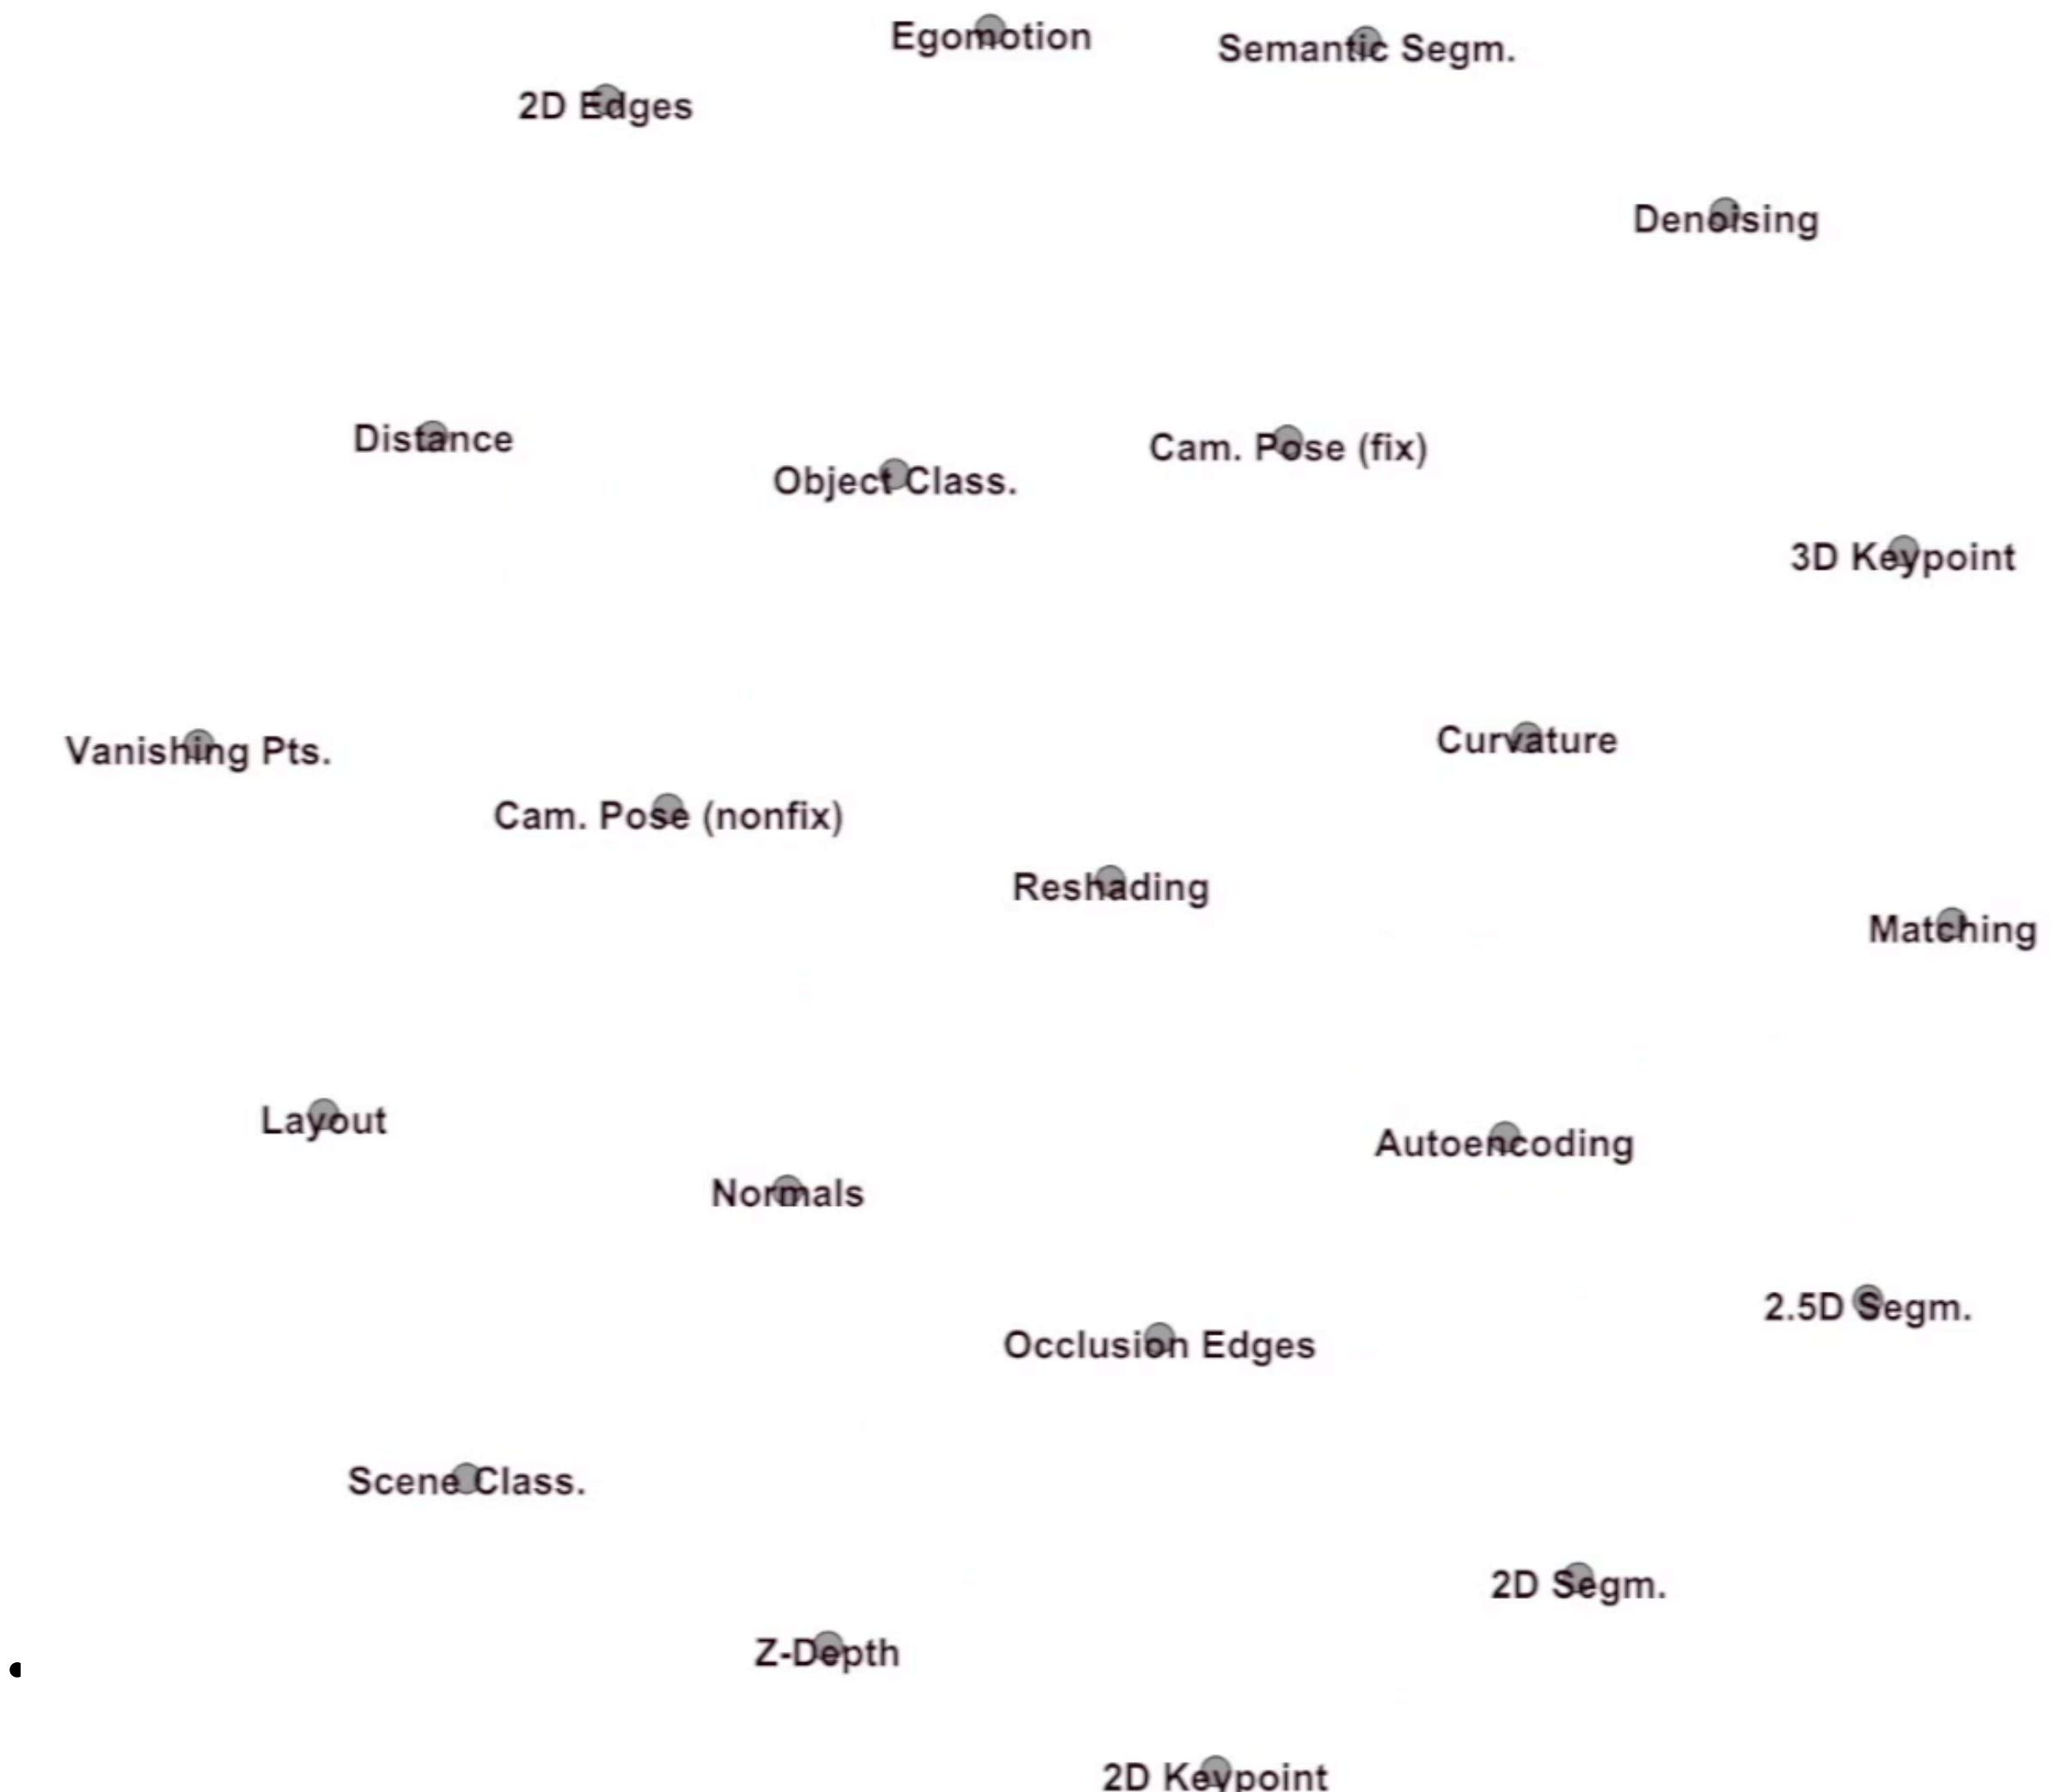
- **Significance Test**
- **Trends on ImageNet and MIT Places**
- **Sensitivity to choice of dictionary**
- **Architecture dependence**
- **Generalization to single out-of-dictionary tasks**
 - **Outperforming self-supervised, unsupervised, ImageNet-based baselines by large margin.**

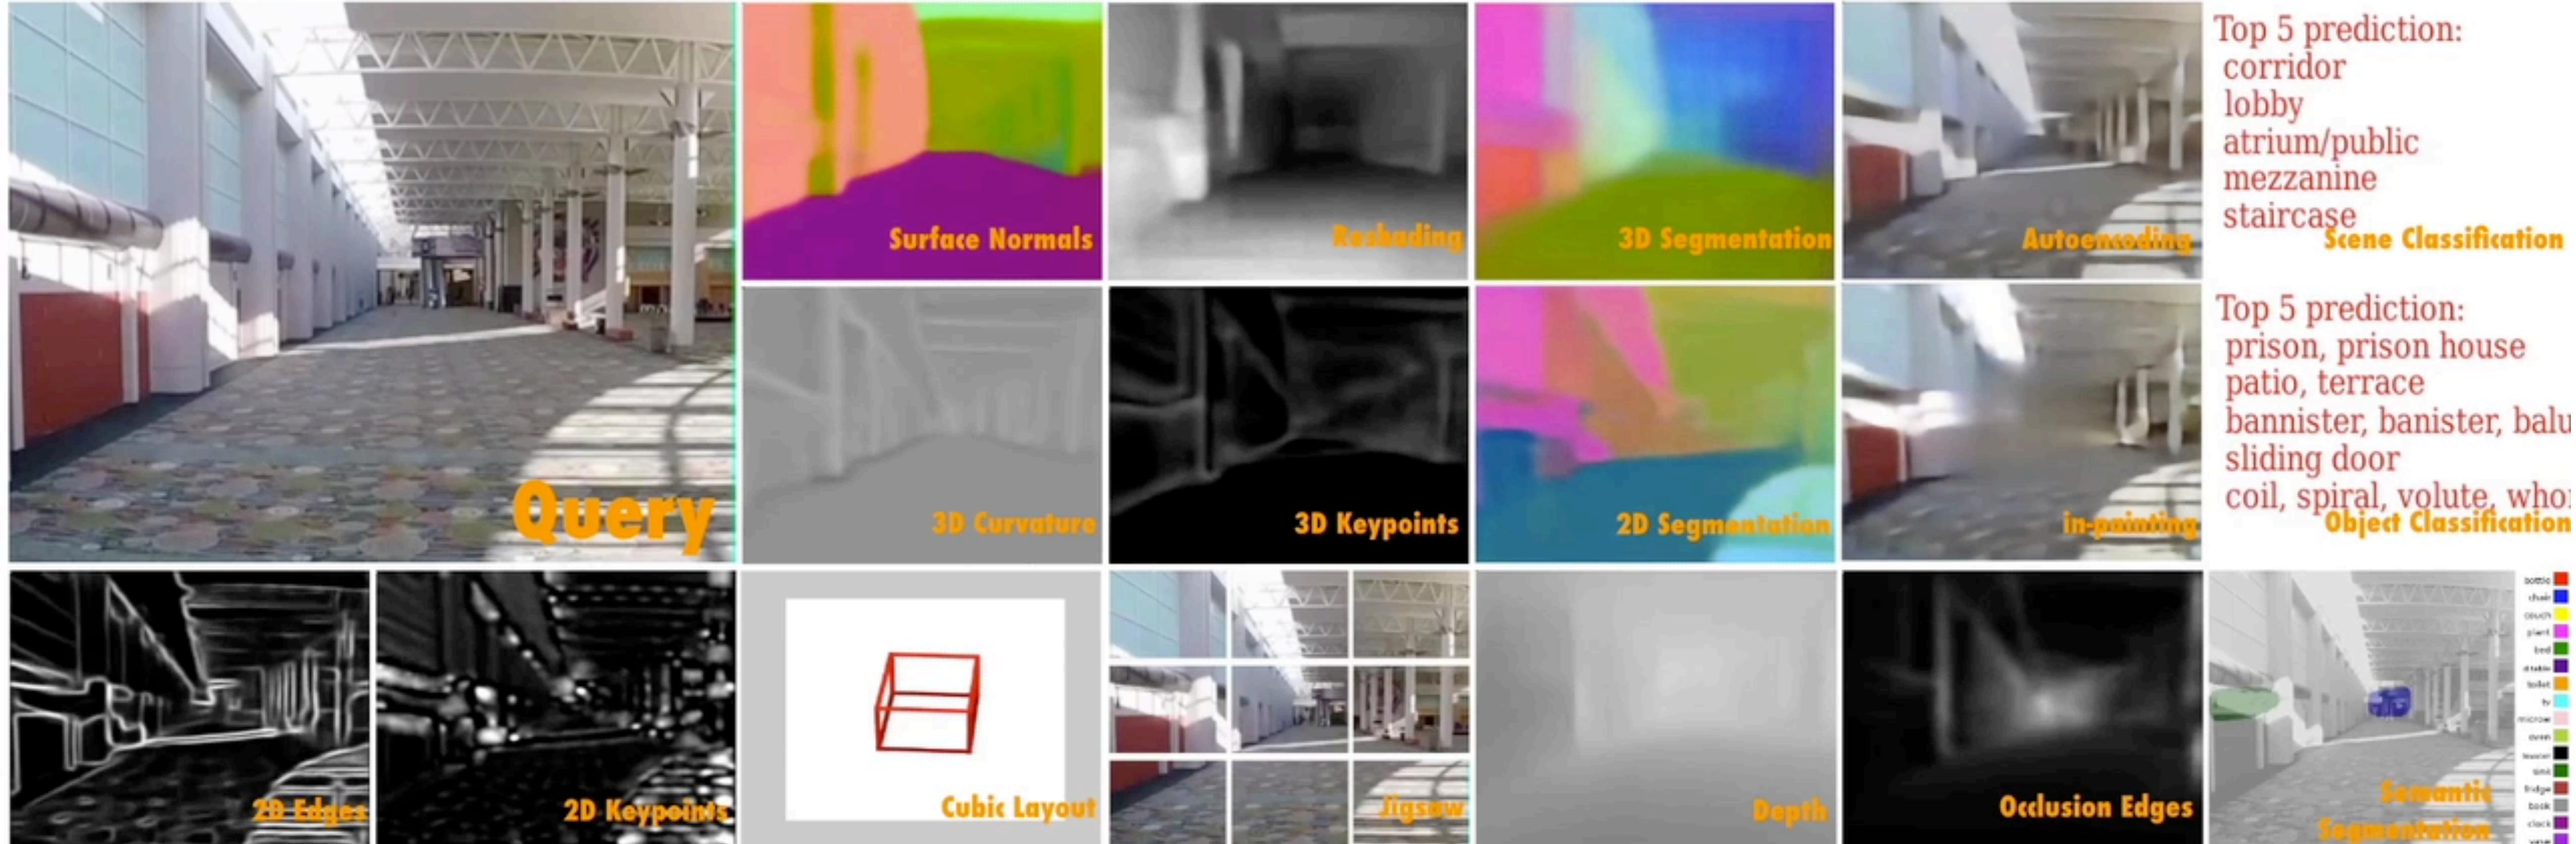
Task	scratch	ImageNet[51]	Wang.[96]	Agrawal.[1]	Zamit.[100]	Zhang.[103]	Noroozi.[68]	full sup.
Depth	88 .03	88 .04	93 .04	89 .03	88 .04	84 .03	86 .03	43 .02
Scene Cls.	80 3.30	52 2.76	83 3.56	74 3.15	74 3.17	71 3.09	75 3.19	15 2.23
Sem. Segm.	78 1.74	79 1.88	82 1.92	85 1.80	76 1.85	78 1.74	84 1.71	21 1.42
Object Cls.	79 4.08	54 3.57	82 4.27	76 3.99	75 3.98	76 4.00	76 3.97	34 3.26
Normals	97 .22	98 .30	98 .34	98 .28	98 .28	97 .23	97 .24	6 .12
2.5D Segm.	80 .21	93 .34	92 .34	89 .26	90 .29	84 .22	87 .24	40 .16
Occ. Edges	93 .16	96 .19	95 .18	93 .17	94 .18	93 .16	94 .17	42 .12
Curvature	88 .25	94 .28	89 .26	85 .25	88 .26	92 .26	88 .25	29 .21
Egomotion	79 8.60	78 8.58	83 9.26	77 8.41	76 8.34	74 8.15	71 7.94	59 7.32
Layout	80 .66	76 .66	85 .85	79 .65	77 .65	78 .62	70 .54	36 .37



Summary

- A striving step towards understanding the space of vision tasks.
- Treat tasks in concert, coming from a structured space, rather than isolated concepts.
- A fully computational framework.
- Transfer Learning. Generalist perception model.





Thank you!

<http://taskonomy.stanford.edu/>

Taskonomy: Disentangling Task Transfer Learning
 Zamir, Sax*, Shen*, Guibas, Malik, Savarese